

```
1: =====
2: F.A.R.F.A.N PIPELINE CODE AUDIT - BATCH 22
3: =====
4: Generated: 2025-12-07T06:17:28.392472
5: Files in this batch: 17
6: =====
7:
8:
9: =====
10: FILE: src/farfán_pipeline/utils/json_logger.py
11: =====
12:
13: """
14: Lightweight JSON Logging - Structured Event Logging
15: =====
16:
17: Provides structured JSON logging for the pipeline with:
18: - JSON formatter for LogRecord
19: - Helper for logging I/O events with envelope metadata
20: - No PII logging
21: - Correlation ID and event ID tracking
22:
23: Author: Policy Analytics Research Unit
24: Version: 1.0.0
25: License: Proprietary
26: """
27:
28: from __future__ import annotations
29:
30: import json
31: import logging
32: import time
33: from typing import Any
34: from farfan_pipeline.core.calibration.decorators import calibrated_method
35:
36: # Import will be available at runtime
37: try:
38:     from farfan_pipeline.utils.contract_io import ContractEnvelope
39: except ImportError:
40:     # Allow module to load for testing
41:     ContractEnvelope = None  # type: ignore
42:
43:
44: class JsonFormatter(logging.Formatter):
45:     """
46:     JSON formatter for structured logging.
47:
48:     Formats LogRecord as JSON with standard fields plus custom extras.
49:     """
50:
51:     @calibrated_method("farfan_core.utils.json_logger.JsonFormatter.format")
52:     def format(self, record: logging.LogRecord) -> str:
53:         """
54:             Format LogRecord as JSON string.
55:
56:             Args:
```

```
57:         record: LogRecord to format
58:
59:     Returns:
60:         JSON string representation
61:
62:     payload: dict[str, Any] = {
63:         "level": record.levelname,
64:         "logger": record.name,
65:         "message": record.getMessage(),
66:         "timestamp_utc": record.__dict__.get("timestamp_utc"),
67:         "event_id": record.__dict__.get("event_id"),
68:         "correlation_id": record.__dict__.get("correlation_id"),
69:         "policy_unit_id": record.__dict__.get("policy_unit_id"),
70:         "phase": record.__dict__.get("phase"),
71:         "latency_ms": record.__dict__.get("latency_ms"),
72:         "input_bytes": record.__dict__.get("input_bytes"),
73:         "output_bytes": record.__dict__.get("output_bytes"),
74:         "input_digest": record.__dict__.get("input_digest"),
75:         "output_digest": record.__dict__.get("output_digest"),
76:     }
77:     # Drop None values to keep JSON compact
78:     payload = {k: v for k, v in payload.items() if v is not None}
79:     return json.dumps(payload, separators=(",", ", "), ensure_ascii=False)
80:
81:
82: def get_json_logger(name: str = "farfan_core") -> logging.Logger:
83:     """
84:     Get or create a JSON logger.
85:
86:     Creates a logger with JSON formatting if not already configured.
87:
88:     Args:
89:         name: Logger name
90:
91:     Returns:
92:         Configured logger instance
93:
94:     Examples:
95:         >>> logger = get_json_logger("test")
96:         >>> logger.name
97:         'test'
98:         >>> logger.level
99:         20
100:    """
101:    logger = logging.getLogger(name)
102:    if not any(isinstance(h, logging.StreamHandler) for h in logger.handlers):
103:        h = logging.StreamHandler()
104:        h.setFormatter(JsonFormatter())
105:        logger.addHandler(h)
106:        logger.setLevel(logging.INFO)
107:        logger.propagate = False
108:    return logger
109:
110:
111: def log_io_event(
112:     logger: logging.Logger,
```

```
113:     *,
114:     phase: str,
115:     envelope_in: Any | None, # ContractEnvelope or None
116:     envelope_out: Any, # ContractEnvelope
117:     started_monotonic: float,
118: ) -> None:
119: """
120:     Log an I/O event with envelope metadata.
121:
122:     Args:
123:         logger: Logger instance
124:         phase: Phase name
125:         envelope_in: Input envelope (may be None)
126:         envelope_out: Output envelope
127:         started_monotonic: Monotonic start time
128:
129:     Examples:
130:         >>> import time
131:         >>> from farfan_core.utils.contract_io import ContractEnvelope
132:         >>> logger = get_json_logger("test")
133:         >>> out = ContractEnvelope.wrap(
134:             ... {"ok": True},
135:             ... policy_unit_id="PU_123",
136:             ... correlation_id="corr-1"
137:             ...
138:         ) # This will log JSON to stdout
139:         >>> log_io_event(
140:             ... logger,
141:             ... phase="normalize",
142:             ... envelope_in=None,
143:             ... envelope_out=out,
144:             ... started_monotonic=time.monotonic()
145:             ... ) # doctest: +SKIP
146: """
147: elapsed_ms = int((time.monotonic() - started_monotonic) * 1000)
148:
149: # Safely get payload sizes
150: input_bytes = None
151: if envelope_in is not None:
152:     try:
153:         payload = getattr(envelope_in, "payload", None)
154:         if payload is not None:
155:             input_bytes = len(json.dumps(payload, ensure_ascii=False))
156:     except (TypeError, AttributeError):
157:         pass
158:
159: output_bytes = None
160: try:
161:     output_bytes = len(json.dumps(envelope_out.payload, ensure_ascii=False))
162: except (TypeError, AttributeError):
163:     # If payload is missing or not serializable, skip logging output_bytes.
164:     # This is non-critical for logging; output_bytes will be None.
165:     pass
166:
167: logger.info(
168:     "phase_io",
```

```
169:         extra={
170:             "timestamp_utc": envelope_out.timestamp_utc,
171:             "event_id": envelope_out.event_id,
172:             "correlation_id": envelope_out.correlation_id,
173:             "policy_unit_id": envelope_out.policy_unit_id,
174:             "phase": phase,
175:             "latency_ms": elapsed_ms,
176:             "input_bytes": input_bytes,
177:             "output_bytes": output_bytes,
178:             "input_digest": getattr(envelope_in, "content_digest", None),
179:             "output_digest": envelope_out.content_digest,
180:         },
181:     )
182:
183:
184: if __name__ == "__main__":
185:     import doctest
186:     import time
187:
188:     # Run doctests
189:     print("Running doctests...")
190:     doctest.testmod(verbose=True)
191:
192:     # Integration tests
193:     print("\n" + "="*60)
194:     print("JSON Logger Integration Tests")
195:     print("="*60)
196:
197:     print("\n1. Testing JSON formatter:")
198:     logger = get_json_logger("demo")
199:     assert logger.level == logging.INFO
200:     assert len(logger.handlers) > 0
201:     assert isinstance(logger.handlers[0].formatter, JsonFormatter)
202:     print("  \u2192 Logger configured with JSON formatter")
203:
204:     print("\n2. Testing log output structure:")
205:     # Create a test record
206:     record = logging.LogRecord(
207:         name="test",
208:         level=logging.INFO,
209:         pathname="",
210:         lineno=0,
211:         msg="test message",
212:         args=(),
213:         exc_info=None,
214:     )
215:     record.event_id = "evt-123"
216:     record.correlation_id = "corr-456"
217:     record.latency_ms = 42
218:
219:     formatter = JsonFormatter()
220:     output = formatter.format(record)
221:     parsed = json.loads(output)
222:
223:     assert parsed["level"] == "INFO"
224:     assert parsed["message"] == "test message"
```

```
225:     assert parsed["event_id"] == "evt-123"
226:     assert parsed["correlation_id"] == "corr-456"
227:     assert parsed["latency_ms"] == 42
228:     print("  \u2192 223 JSON format includes all expected fields")
229:
230:     print("\n3. Testing I/O event logging:")
231:     # Only test if ContractEnvelope is available
232:     if ContractEnvelope is not None:
233:         from farfan_pipeline.utils.contract_io import ContractEnvelope
234:
235:         lg = get_json_logger("demo")
236:         out = ContractEnvelope.wrap(
237:             {"ok": True},
238:             policy_unit_id="PU_123",
239:             correlation_id="corr-1"
240:         )
241:
242:         # Capture the log output
243:         import io
244:         import sys
245:         old_stdout = sys.stdout
246:         sys.stdout = buffer = io.StringIO()
247:
248:         log_io_event(
249:             lg,
250:             phase="normalize",
251:             envelope_in=None,
252:             envelope_out=out,
253:             started_monotonic=time.monotonic()
254:         )
255:
256:         sys.stdout = old_stdout
257:         log_output = buffer.getvalue()
258:
259:         # Verify JSON output
260:         if log_output.strip():
261:             log_data = json.loads(log_output.strip())
262:             assert log_data["phase"] == "normalize"
263:             assert log_data["policy_unit_id"] == "PU_123"
264:             assert "latency_ms" in log_data
265:             print("  \u2192 223 I/O event logged with correct structure")
266:         else:
267:             print("  \u2192 223 I/O event logging executed (output suppressed)")
268:     else:
269:         print("  \u2192 230 Skipped (ContractEnvelope not available)")
270:
271:     print("\n" + "="*60)
272:     print("JSON logger doctest OK - All tests passed!")
273:     print("=". * 60)
274:
275:
276:
277: =====
278: FILE: src/farfan_pipeline/utils/metadata_loader.py
279: =====
280:
```

```
281: """
282: Metadata Loader with Supply-Chain Security
283: Implements fail-fast validation with version pinning, checksum verification, and schema validation
284: """
285:
286: import hashlib
287: import json
288: import logging
289: from pathlib import Path
290: from typing import Any
291:
292: import yaml
293:
294: from farfan_pipeline.utils.paths import proj_root
295: from farfan_pipeline.core.parameters import ParameterLoaderV2
296: from farfan_pipeline.core.calibration.decorators import calibrated_method
297:
298: try:
299:     import jsonschema
300:     JSONSCHEMA_AVAILABLE = True
301: except ImportError:
302:     JSONSCHEMA_AVAILABLE = False
303:     logging.warning("jsonschema not available - schema validation disabled")
304:
305: logger = logging.getLogger(__name__)
306:
307: class MetadataError(Exception):
308:     """Base exception for metadata errors"""
309:     pass
310:
311: class MetadataVersionError(MetadataError):
312:     """Version mismatch error"""
313:     def __init__(self, expected: str, actual: str, file_path: str) -> None:
314:         self.expected = expected
315:         self.actual = actual
316:         self.file_path = file_path
317:         super().__init__(
318:             f"Version mismatch in {file_path}: expected {expected}, got {actual}"
319:         )
320:
321: class MetadataIntegrityError(MetadataError):
322:     """Checksum/integrity violation error"""
323:     def __init__(self, file_path: str, expected_checksum: str | None = None, actual_checksum: str | None = None) -> None:
324:         self.file_path = file_path
325:         self.expected_checksum = expected_checksum
326:         self.actual_checksum = actual_checksum
327:         msg = f"Integrity violation in {file_path}"
328:         if expected_checksum and actual_checksum:
329:             msg += f": expected checksum {expected_checksum}, got {actual_checksum}"
330:         super().__init__(msg)
331:
332: class MetadataSchemaError(MetadataError):
333:     """Schema validation error"""
334:     def __init__(self, file_path: str, validation_errors: list) -> None:
335:         self.file_path = file_path
336:         self.validation_errors = validation_errors
```

```
337:         error_msgs = '\n'.join(f" - {err}" for err in validation_errors)
338:         super().__init__(
339:             f"Schema validation failed for {file_path}:\n{error_msgs}"
340:         )
341:
342: class MetadataMissingKeyError(MetadataError):
343:     """Required key missing in metadata"""
344:     def __init__(self, file_path: str, missing_key: str, context: str = "") -> None:
345:         self.file_path = file_path
346:         self.missing_key = missing_key
347:         self.context = context
348:         msg = f"Required key '{missing_key}' missing in {file_path}"
349:         if context:
350:             msg += f" ({context})"
351:         super().__init__(msg)
352:
353: class MetadataLoader:
354:     """
355:         Unified metadata loader with strict validation
356:
357:     Features:
358:     - Version pinning with semantic versioning
359:     - SHA-256 checksum verification
360:     - JSON Schema validation
361:     - Fail-fast on any violation
362:     - Structured logging of all errors
363:     """
364:
365:     def __init__(self, workspace_root: Path | None = None) -> None:
366:         self.workspace_root = Path(workspace_root) if workspace_root else proj_root()
367:         self.schemas_dir = self.workspace_root / "schemas"
368:
369:         # Loaded schemas cache
370:         self._schema_cache: dict[str, dict] = {}
371:
372:     def load_and_validate_metadata(
373:         self,
374:         path: Path,
375:         schema_ref: str | None = None,
376:         required_version: str | None = None,
377:         expected_checksum: str | None = None,
378:         checksum_algorithm: str = "sha256"
379:     ) -> dict[str, Any]:
380:         """
381:             Load and validate metadata file with all safeguards
382:
383:         Args:
384:             path: Path to metadata file (JSON or YAML)
385:             schema_ref: Schema file name (e.g., "rubric.schema.json")
386:             required_version: Required version string (e.g., "2.0.0")
387:             expected_checksum: Expected SHA-256 checksum (hex)
388:             checksum_algorithm: Hash algorithm ("sha256", "md5")
389:
390:         Returns:
391:             Validated metadata dictionary
```

```
392:
393:     Raises:
394:         MetadataVersionError: Version mismatch
395:         MetadataIntegrityError: Checksum mismatch
396:         MetadataSchemaError: Schema validation failure
397:         MetadataMissingKeyError: Required key missing
398:     """
399:
400:     # 1. Load file
401:     metadata = self._load_file(path)
402:
403:     # 2. Version check
404:     if required_version:
405:         actual_version = metadata.get("version")
406:         if not actual_version:
407:             raise MetadataMissingKeyError(str(path), "version", "version field required")
408:
409:         if actual_version != required_version:
410:             self._log_error(
411:                 rule_id="VERSION_MISMATCH",
412:                 file_path=str(path),
413:                 expected=required_version,
414:                 actual=actual_version
415:             )
416:             raise MetadataVersionError(required_version, actual_version, str(path))
417:
418:             logger.info(f"\u234\u234 Version validated: {path.name} v{actual_version}")
419:
420:     # 3. Checksum verification
421:     if expected_checksum:
422:         actual_checksum = self._calculate_checksum(metadata, checksum_algorithm)
423:
424:         if actual_checksum != expected_checksum:
425:             self._log_error(
426:                 rule_id="CHECKSUM_MISMATCH",
427:                 file_path=str(path),
428:                 expected=expected_checksum,
429:                 actual=actual_checksum
430:             )
431:             raise MetadataIntegrityError(str(path), expected_checksum, actual_checksum)
432:
433:             logger.info(f"\u234\u234 Checksum validated: {path.name} ({checksum_algorithm})")
434:
435:     # 4. Schema validation
436:     if schema_ref and JSONSCHEMA_AVAILABLE:
437:         schema = self._load_schema(schema_ref)
438:         errors = self._validate_schema(metadata, schema)
439:
440:         if errors:
441:             self._log_error(
442:                 rule_id="SCHEMA_VALIDATION_FAILED",
443:                 file_path=str(path),
444:                 errors=errors
445:             )
446:             raise MetadataSchemaError(str(path), errors)
447:
```

```
448:         logger.info(f"â\234\223 Schema validated: {path.name}")
449:
450:     return metadata
451:
452:     @calibrated_method("farfan_core.utils.metadata_loader.MetadataLoader._load_file")
453:     def _load_file(self, path: Path) -> dict[str, Any]:
454:         """Load JSON or YAML file"""
455:         if not path.exists():
456:             raise FileNotFoundError(f"Metadata file not found: {path}")
457:
458:         try:
459:             with open(path, encoding='utf-8') as f:
460:                 content = f.read()
461:
462:                 if path.suffix in ['.json']:
463:                     return json.loads(content)
464:                 elif path.suffix in ['.yaml', '.yml']:
465:                     return yaml.safe_load(content)
466:                 else:
467:                     raise ValueError(f"Unsupported file type: {path.suffix}")
468:
469:             except (json.JSONDecodeError, yaml.YAMLError) as e:
470:                 raise MetadataError(f"Failed to parse {path}: {e}")
471:
472:     @calibrated_method("farfan_core.utils.metadata_loader.MetadataLoader._calculate_checksum")
473:     def _calculate_checksum(self, metadata: dict[str, Any], algorithm: str = "sha256") -> str:
474:         """
475:             Calculate reproducible checksum of metadata
476:
477:             Normalization:
478:             - JSON serialization with sorted keys
479:             - UTF-8 encoding
480:             - No whitespace variations
481:         """
482:             normalized = json.dumps(metadata, sort_keys=True, separators=(',', ':'))
483:
484:             if algorithm == "sha256":
485:                 return hashlib.sha256(normalized.encode('utf-8')).hexdigest()
486:             elif algorithm == "md5":
487:                 return hashlib.md5(normalized.encode('utf-8')).hexdigest()
488:             else:
489:                 raise ValueError(f"Unsupported algorithm: {algorithm}")
490:
491:     @calibrated_method("farfan_core.utils.metadata_loader.MetadataLoader._load_schema")
492:     def _load_schema(self, schema_ref: str) -> dict[str, Any]:
493:         """Load JSON Schema from schemas directory"""
494:         if schema_ref in self._schema_cache:
495:             return self._schema_cache[schema_ref]
496:
497:             schema_path = self.schemas_dir / schema_ref
498:
499:             if not schema_path.exists():
500:                 raise FileNotFoundError(f"Schema not found: {schema_path}")
501:
502:             with open(schema_path, encoding='utf-8') as f:
503:                 schema = json.load(f)
```

```
504:         self._schema_cache[schema_ref] = schema
505:     return schema
506:
507:
508:     @calibrated_method("farfan_core.utils.metadata_loader.MetadataLoader._validate_schema")
509:     def _validate_schema(self, metadata: dict[str, Any], schema: dict[str, Any]) -> list:
510:         """Validate metadata against JSON Schema"""
511:         if not JSONSCHEMA_AVAILABLE:
512:             logger.warning("jsonschema not available - skipping schema validation")
513:             return []
514:
515:         # Import check for type checker
516:         import jsonschema as js
517:         validator = js.Draft7Validator(schema)
518:         errors = []
519:
520:         for error in validator.iter_errors(metadata):
521:             error_path = '.'.join(str(p) for p in error.path) if error.path else 'root'
522:             errors.append(f"{error_path}: {error.message}")
523:
524:         return errors
525:
526:     @calibrated_method("farfan_core.utils.metadata_loader.MetadataLoader._log_error")
527:     def _log_error(self, rule_id: str, file_path: str, **kwargs) -> None:
528:         """Structured error logging"""
529:         from datetime import datetime, timezone
530:
531:         log_entry = {
532:             "timestamp": datetime.now(timezone.utc).isoformat(),
533:             "level": "ERROR",
534:             "rule_id": rule_id,
535:             "file_path": file_path,
536:             **kwargs
537:         }
538:
539:         logger.error(json.dumps(log_entry, indent=2))
540:
541: # REMOVED: load_cuestionario() - LEGACY FUNCTION
542: # Questionnaire monolith must ONLY be loaded via factory.load_questionnaire_monolith()
543: # This enforces architectural requirement: Single I/O boundary in factory.py
544: # See: src/farfan_core/core/orchestrator/factory.py::load_questionnaire_monolith()
545:
546: def load_execution_mapping(
547:     path: Path | None = None,
548:     required_version: str = "2.0"
549: ) -> dict[str, Any]:
550:
551:     """
552:     Load and validate execution_mapping.yaml
553:
554:     Args:
555:         path: Path to execution mapping (default: execution_mapping.yaml)
556:         required_version: Required version
557:
558:     Returns:
559:         Validated execution mapping
560:     """
561:
```

```
560:     if path is None:
561:         path = proj_root() / "execution_mapping.yaml"
562:
563:     loader = MetadataLoader()
564:     return loader.load_and_validate_metadata(
565:         path=path,
566:         schema_ref="execution_mapping.schema.json",
567:         required_version=required_version
568:     )
569:
570: def load_rubric_scoring(
571:     path: Path | None = None,
572:     required_version: str = "2.0"
573: ) -> dict[str, Any]:
574:     """
575:         Load and validate rubric_scoring.json
576:
577:     Args:
578:         path: Path to rubric scoring (default: rubric_scoring.json)
579:         required_version: Required version
580:
581:     Returns:
582:         Validated rubric scoring configuration
583:     """
584:     if path is None:
585:         path = proj_root() / "rubric_scoring.json"
586:
587:     loader = MetadataLoader()
588:     return loader.load_and_validate_metadata(
589:         path=path,
590:         schema_ref="rubric.schema.json",
591:         required_version=required_version
592:     )
593:
594:
595:
596: =====
597: FILE: src/farfan_pipeline/utils/method_config_loader.py
598: =====
599:
600: """
601: Method Configuration Loader for Canonical JSON Specification.
602:
603: Provides unified access to method parameters from the canonical
604: parameterization specification.
605: """
606: import ast
607: import json
608: from pathlib import Path
609: from typing import Any
610: from farfan_pipeline.core.parameters import ParameterLoaderV2
611: from farfan_pipeline.core.calibration.decorators import calibrated_method
612:
613:
614: class MethodConfigLoader:
615:     """
```

```
616:     Loads and provides access to method parameters from canonical JSON.
617:
618:     Usage:
619:         loader = MethodConfigLoader("CANONICAL_METHOD_PARAMETERIZATION_SPEC.json")
620:         threshold = loader.get_method_parameter(
621:             "CAUSAL.BMI.infer_mech_v1",
622:             "kl_divergence_threshold"
623:         )
624:
625:     Note:
626:         The loader expects the JSON spec to follow the canonical schema with
627:         keys: specification_metadata, methods, and epistemic_validation_summary.
628: """
629:
630:     def __init__(self, spec_path: str | Path) -> None:
631:         self.spec_path = Path(spec_path)
632:         with open(self.spec_path) as f:
633:             self.spec = json.load(f)
634:
635:         # Validate schema before use
636:         self.validate_spec_schema()
637:
638:         # Build index for fast lookup
639:         self._method_index = {
640:             method["canonical_id"]: method
641:             for method in self.spec["methods"]
642:         }
643:
644:     @calibrated_method("farfan_core.utils.method_config_loader.MethodConfigLoader.validate_spec_schema")
645:     def validate_spec_schema(self) -> None:
646:         """
647:             Validate JSON spec matches expected schema.
648:
649:             Raises:
650:                 ValueError: If spec is missing required keys
651: """
652:     required_keys = {"specification_metadata", "methods"}
653:     # Note: epistemic_validation_summary is optional in some versions
654:     if not required_keys.issubset(self.spec.keys()):
655:         missing = required_keys - set(self.spec.keys())
656:         raise ValueError(f"Spec missing required keys: {missing}")
657:
658:     def get_method_parameter(
659:         self,
660:         canonical_id: str,
661:         param_name: str,
662:         override: Any = None
663:     ) -> Any:
664:         """
665:             Get parameter value for a method.
666:
667:             Args:
668:                 canonical_id: Canonical method ID (e.g., "CAUSAL.BMI.infer_mech_v1")
669:                 param_name: Parameter name
670:                 override: Optional override value (takes precedence over default)
671:
```

```
672:     Returns:
673:         Parameter value (default or override)
674:
675:     Raises:
676:         KeyError: If method or parameter not found
677:     """
678:     if canonical_id not in self._method_index:
679:         raise KeyError(f"Method {canonical_id} not found in canonical spec")
680:
681:     method = self._method_index[canonical_id]
682:
683:     for param in method["parameters"]:
684:         if param["name"] == param_name:
685:             return override if override is not None else param["default"]
686:
687:     raise KeyError(f"Parameter {param_name} not found for method {canonical_id}")
688:
689: @calibrated_method("farfan_core.utils.method_config_loader.MethodConfigLoader.get_method_description")
690: def get_method_description(self, canonical_id: str) -> str:
691:     """Get method description."""
692:     return self._method_index[canonical_id]["description"]
693:
694: @calibrated_method("farfan_core.utils.method_config_loader.MethodConfigLoader.get_parameter_spec")
695: def get_parameter_spec(self, canonical_id: str, param_name: str) -> dict:
696:     """Get full parameter specification including allowed values."""
697:     method = self._method_index[canonical_id]
698:     for param in method["parameters"]:
699:         if param["name"] == param_name:
700:             return param
701:     raise KeyError(f"Parameter {param_name} not found")
702:
703: def validate_parameter_value(
704:     self,
705:     canonical_id: str,
706:     param_name: str,
707:     value: Any
708: ) -> bool:
709:     """
710:     Validate parameter value against allowed_values specification.
711:
712:     Returns:
713:         True if valid, raises ValueError if invalid
714:     """
715:     param_spec = self.get_parameter_spec(canonical_id, param_name)
716:     allowed = param_spec["allowed_values"]
717:
718:     if allowed["kind"] == "range":
719:         spec = allowed["spec"]
720:         min_val, max_val = self._parse_range(spec)
721:         if not (min_val <= value <= max_val):
722:             raise ValueError(f"{param_name}={value} out of range {spec}")
723:
724:     elif allowed["kind"] == "set":
725:         spec = allowed["spec"]
726:         valid_values = self._parse_set(spec)
727:         if value not in valid_values:
```

```
728:             raise ValueError(f"{{param_name}}={{value}} not in allowed set {{spec}}")
729:
730:     return True
731:
732:     @calibrated_method("farfan_core.utils.method_config_loader.MethodConfigLoader._parse_range")
733:     def _parse_range(self, spec: str) -> tuple[float, float]:
734:         """
735:             Parse range specification like '[ParameterLoaderV2.get("farfan_core.utils.method_config_loader.MethodConfigLoader._parse_range", "auto_param_L135_41", 0.0), ParameterLoaderV2.get("farfan_core.utils.method_config_loader.MethodConfigLoader._parse_range", "auto_param_L135_46", 1.0)], inclusive' or '[100, 10000], integer'.
736:
737:             Args:
738:                 spec: Range specification string with format "[min, max], modifiers"
739:                     Modifiers can include: inclusive, exclusive, integer
740:
741:             Returns:
742:                 Tuple of (min_val, max_val) as floats
743:
744:             Raises:
745:                 ValueError: If spec format is invalid
746:
747:             Note:
748:                 The inclusive/exclusive and integer modifiers are parsed but not
749:                 currently enforced in validation. This maintains compatibility with
750:                 the current spec while allowing future enhancement.
751:         """
752:         try:
753:             # Extract bracketed part before any modifiers
754:             bracket_part = spec.split("[]")[0] + "]"
755:             parts = bracket_part.replace("[", "").replace("]", "").split(",")
756:             min_val = float(parts[0].strip())
757:             max_val = float(parts[1].strip())
758:             return min_val, max_val
759:         except (IndexError, ValueError) as e:
760:             raise ValueError(f"Invalid range spec: {spec}") from e
761:
762:     @calibrated_method("farfan_core.utils.method_config_loader.MethodConfigLoader._parse_set")
763:     def _parse_set(self, spec: str | list) -> set:
764:         """
765:             Parse set specification safely.
766:
767:             Args:
768:                 spec: Either a list or a string representation of a Python literal
769:
770:             Returns:
771:                 Set of allowed values
772:
773:             Raises:
774:                 ValueError: If spec cannot be parsed safely
775:
776:             Note:
777:                 Uses ast.literal_eval() for safe parsing of string specs.
778:                 Only Python literals (strings, numbers, tuples, lists, dicts,
779:                 booleans, None) are supported - no arbitrary code execution.
780:         """
781:         if isinstance(spec, list):
```

```
782:         return set(spec)
783:     try:
784:         # Use ast.literal_eval for safer parsing
785:         return set(ast.literal_eval(spec))
786:     except (ValueError, SyntaxError) as e:
787:         raise ValueError(f"Invalid set spec: {spec}") from e
788:
789:
790:
791: =====
792: FILE: src/farfan_pipeline/utils/paths.py
793: =====
794:
795: """
796: Portable, secure, and deterministic path utilities for SAAAAAA.
797:
798: This module provides cross-platform path operations that ensure:
799: - Portability across Linux, macOS, and Windows
800: - Security through path traversal protection
801: - Determinism via normalized paths
802: - Controlled write locations (never in source tree)
803:
804: All path operations in the repository MUST use these utilities instead of:
805: - Direct __file__ usage for resource access
806: - sys.path manipulation
807: - Hardcoded absolute paths
808: - os.path functions (use pathlib.Path instead)
809: """
810:
811: from __future__ import annotations
812:
813: import os
814: import unicodedata
815: from pathlib import Path
816: from typing import Final
817: from farfan_pipeline.core.calibration.decorators import calibrated_method
818:
819:
820: # Custom exception types for path errors
821: class PathError(Exception):
822:     """Base exception for path-related errors."""
823:     pass
824:
825:
826: class PathTraversalError(PathError):
827:     """Raised when a path attempts to escape workspace boundaries."""
828:     pass
829:
830:
831: class PathNotFoundError(PathError):
832:     """Raised when a required path does not exist."""
833:     pass
834:
835:
836: class PathOutsideWorkspaceError(PathError):
837:     """Raised when a path is outside the allowed workspace."""
```

```
838:     pass
839:
840:
841: class UnnormalizedPathError(PathError):
842:     """Raised when a path is not properly normalized."""
843:     pass
844:
845:
846: # Project root detection - computed once at module load
847: def _detect_project_root() -> Path:
848:     """
849:         Detect the project root directory using filesystem markers.
850:
851:         This function uses a multi-strategy approach to locate the project root:
852:
853:             1. Primary strategy: Search for pyproject.toml
854:                 - Walks up the directory tree from this file's location
855:                 - Returns the first directory containing pyproject.toml
856:
857:             2. Secondary strategy: Search for src/farfan_core layout
858:                 - Looks for directories with both src/farfan_core and setup.py
859:                 - This supports older project structures
860:
861:             3. Fallback strategy: Relative path calculation
862:                 - If no markers found, assumes standard layout (src/farfan_core/utils)
863:                 - Returns path 3 levels up from this file
864:
865:         The function is called once at module load time, and the result is
866:         cached in the PROJECT_ROOT constant.
867:
868:     Returns:
869:         Path: Absolute path to the project root directory
870:
871:     Raises:
872:         No exceptions raised; always returns a path (uses fallback if needed)
873:
874:     Note:
875:         This function is intended for internal use. External code should use
876:         the PROJECT_ROOT constant instead of calling this directly.
877:     """
878:     # Start from this file's location
879:     current = Path(__file__).resolve().parent
880:
881:     # Walk up to find pyproject.toml
882:     for parent in [current] + list(current.parents):
883:         if (parent / "pyproject.toml").exists():
884:             return parent
885:         if (parent / "src" / "farfan_core").exists() and (parent / "setup.py").exists():
886:             return parent
887:
888:     # Fallback: if we can't find it, assume we're in src/farfan_core/utils
889:     # and go up 3 levels
890:     return current.parent.parent.parent
891:
892:
893: # Global constants for common directories
```

```
894: PROJECT_ROOT: Final[Path] = _detect_project_root()
895: SRC_DIR: Final[Path] = PROJECT_ROOT / "src"
896: DATA_DIR: Final[Path] = PROJECT_ROOT / "data"
897: TESTS_DIR: Final[Path] = PROJECT_ROOT / "tests"
898:
899:
900: def proj_root() -> Path:
901:     """
902:         Get the project root directory.
903:
904:         Returns:
905:             Absolute path to the project root (where pyproject.toml lives)
906: """
907:     return PROJECT_ROOT
908:
909:
910: def src_dir() -> Path:
911:     """Get the src directory path."""
912:     return SRC_DIR
913:
914:
915: def data_dir() -> Path:
916:     """
917:         Get the data directory path.
918:         Creates it if it doesn't exist.
919: """
920:     DATA_DIR.mkdir(parents=True, exist_ok=True)
921:     return DATA_DIR
922:
923:
924: def tmp_dir() -> Path:
925:     """
926:         Get a project-specific temporary directory.
927:
928:         Uses PROJECT_ROOT/tmp to keep temporary files within the workspace
929:         and avoid polluting system temp directories.
930:
931:         Returns:
932:             Path to tmp directory (created if needed)
933: """
934:     tmp = PROJECT_ROOT / "tmp"
935:     tmp.mkdir(parents=True, exist_ok=True)
936:     return tmp
937:
938:
939: def build_dir() -> Path:
940:     """
941:         Get the build directory for generated artifacts.
942:
943:         Returns:
944:             Path to build directory (created if needed)
945: """
946:     build = PROJECT_ROOT / "build"
947:     build.mkdir(parents=True, exist_ok=True)
948:     return build
949:
```

```
950:
951: def cache_dir() -> Path:
952:     """
953:         Get the cache directory.
954:
955:     Returns:
956:         Path to cache directory (created if needed)
957:     """
958:     cache = build_dir() / "cache"
959:     cache.mkdir(parents=True, exist_ok=True)
960:     return cache
961:
962:
963: def reports_dir() -> Path:
964:     """
965:         Get the reports directory for generated reports.
966:
967:     Returns:
968:         Path to reports directory (created if needed)
969:     """
970:     reports = build_dir() / "reports"
971:     reports.mkdir(parents=True, exist_ok=True)
972:     return reports
973:
974:
975: def is_within(base: Path, child: Path) -> bool:
976:     """
977:         Check if child path is within base directory (no traversal outside).
978:
979:     Args:
980:         base: Base directory that should contain child
981:         child: Path to check
982:
983:     Returns:
984:         True if child is within base, False otherwise
985:
986:     Example:
987:         >>> project_root = Path("project_root")
988:         >>> is_within(project_root, project_root / "src" / "file.py")
989:         True
990:         >>> other_root = Path("other_project")
991:         >>> is_within(project_root, other_root / "file.py")
992:         False
993:     """
994:     try:
995:         base_resolved = base.resolve()
996:         child_resolved = child.resolve()
997:
998:         # Check if child is relative to base
999:         child_resolved.relative_to(base_resolved)
1000:     return True
1001: except (ValueError, RuntimeError):
1002:     return False
1003:
1004:
1005: def safe_join(base: Path, *parts: str) -> Path:
```

```
1006: """
1007:     Safely join path components, preventing traversal outside base.
1008:
1009:     This prevents directory traversal attacks using ".." components.
1010:
1011:     Args:
1012:         base: Base directory
1013:         *parts: Path components to join
1014:
1015:     Returns:
1016:         Resolved path within base
1017:
1018:     Raises:
1019:         PathTraversalError: If the resulting path would be outside base
1020:
1021:     Example:
1022:         >>> project_root = Path("project_root")
1023:         >>> safe_join(project_root, "src", "file.py")
1024:         project_root/src/file.py
1025:         >>> safe_join(project_root, "..", "other")  # raises
1026:         PathTraversalError
1027: """
1028: result = base.joinpath(*parts).resolve()
1029:
1030: if not is_within(base, result):
1031:     raise PathTraversalError(
1032:         f"Path traversal detected: '{result}' is outside base '{base}'. "
1033:         f"Use paths within the workspace."
1034:     )
1035:
1036: return result
1037:
1038:
1039: def normalize_unicode(path: Path, form: str = "NFC") -> Path:
1040: """
1041:     Normalize Unicode in path for cross-platform consistency.
1042:
1043:     Different filesystems handle Unicode differently:
1044:     - macOS (HFS+) uses NFD normalization
1045:     - Linux typically uses NFC
1046:     - Windows uses UTF-16
1047:
1048:     Args:
1049:         path: Path to normalize
1050:         form: Unicode normalization form ("NFC", "NFD", "NFKC", "NFKD")
1051:             Default "NFC" for maximum compatibility
1052:
1053:     Returns:
1054:         Path with normalized Unicode
1055: """
1056: normalized_str = unicodedata.normalize(form, str(path))
1057: return Path(normalized_str)
1058:
1059:
1060: def normalize_case(path: Path) -> Path:
1061: """
```

```
1062:     Normalize path case for case-insensitive filesystems.
1063:
1064:     On case-insensitive filesystems (Windows, macOS default), this ensures
1065:     consistent casing. On case-sensitive systems (Linux), returns unchanged.
1066:
1067:     Args:
1068:         path: Path to normalize
1069:
1070:     Returns:
1071:         Path with normalized case
1072:     """
1073:     # Check if filesystem is case-sensitive
1074:     # This is a heuristic - we check if we can create files differing only in case
1075:     if path.exists():
1076:         # Use actual case from filesystem
1077:         try:
1078:             # On Windows/macOS this will resolve to actual case
1079:             return path.resolve()
1080:         except Exception:
1081:             pass
1082:
1083:     return path
1084:
1085:
1086: def resources(package: str, *path_parts: str) -> Path:
1087:     """
1088:     Access packaged resource files in a portable way.
1089:
1090:     This uses importlib.resources (Python 3.9+) to access resources that
1091:     are included in the installed package, whether from source or wheel.
1092:
1093:     Args:
1094:         package: Package name (e.g., "farfan_core.core")
1095:         *path_parts: Path components within the package
1096:
1097:     Returns:
1098:         Path to the resource
1099:
1100:     Raises:
1101:         PathNotFoundError: If resource doesn't exist
1102:
1103:     Example:
1104:         >>> resources("farfan_core.core", "config", "default.yaml")
1105:         Path('/path/to/farfan_core/core/config/default.yaml')
1106:     """
1107:     try:
1108:         # Python 3.9+ way
1109:         from importlib.resources import files
1110:
1111:         pkg_path = files(package)
1112:         for part in path_parts:
1113:             pkg_path = pkg_path.joinpath(part)
1114:
1115:         # Convert to Path - files() returns Traversable
1116:         if hasattr(pkg_path, '__fspath__'):
1117:             return Path(pkg_path)
```

```
1118:     else:
1119:         # Fallback for Traversable that doesn't support __fspath__
1120:         # Read the resource and return a path to it
1121:         raise PathNotFoundError(
1122:             f"Resource '{'.'.join(path_parts)}' in package '{package}' "
1123:             f"is not accessible as a filesystem path. "
1124:             f"Consider using importlib.resources.read_text() or read_binary() instead."
1125:         )
1126:     except (ImportError, ModuleNotFoundError, FileNotFoundError, TypeError) as e:
1127:         raise PathNotFoundError(
1128:             f"Resource '{'.'.join(path_parts)}' not found in package '{package}'. "
1129:             f"Ensure it's declared in pyproject.toml [tool.setuptools.package-data]. "
1130:             f"Error: {e}"
1131:         ) from e
1132:
1133:
1134: def validate_read_path(path: Path) -> None:
1135:     """
1136:     Validate a path before reading from it.
1137:
1138:     Args:
1139:         path: Path to validate
1140:
1141:     Raises:
1142:         PathNotFoundError: If path doesn't exist
1143:         PermissionError: If path is not readable
1144:     """
1145:     if not path.exists():
1146:         raise PathNotFoundError(f"Path does not exist: '{path}'")
1147:
1148:     if not os.access(path, os.R_OK):
1149:         raise PermissionError(f"Path is not readable: '{path}'")
1150:
1151:
1152: def validate_write_path(path: Path, allow_source_tree: bool = False) -> None:
1153:     """
1154:     Validate a path before writing to it.
1155:
1156:     By default, prohibits writing to the source tree to prevent
1157:     accidental modification of versioned code.
1158:
1159:     Args:
1160:         path: Path to validate
1161:         allow_source_tree: If True, allow writing to source tree
1162:                         (for special cases like code generation)
1163:
1164:     Raises:
1165:         PathOutsideWorkspaceError: If path is outside workspace
1166:         PermissionError: If parent directory is not writable
1167:         ValueError: If trying to write to source tree when not allowed
1168:     """
1169:     # Ensure it's within the workspace
1170:     if not is_within(PROJECT_ROOT, path):
1171:         raise PathOutsideWorkspaceError(
1172:             f"Cannot write to '{path}' - outside workspace '{PROJECT_ROOT}'"
1173:         )
```

```
1174:  
1175:     # Prohibit writing to source tree unless explicitly allowed  
1176:     if not allow_source_tree and is_within(SRC_DIR, path):  
1177:         raise ValueError(  
1178:             f"Cannot write to source tree: '{path}'. "  
1179:             f"Write to build/, cache/, or reports/ instead. "  
1180:             f"If you need to write to source (e.g., code generation), "  
1181:             f"set allow_source_tree=True."  
1182:     )  
1183:  
1184:     # Ensure parent directory exists and is writable  
1185:     parent = path.parent  
1186:     if parent.exists() and not os.access(parent, os.W_OK):  
1187:         raise PermissionError(f"Parent directory is not writable: '{parent}'")  
1188:  
1189:  
1190: # Environment variable accessors (typed and safe)  
1191:  
1192: def get_env_path(key: str, default: Path | None = None) -> Path | None:  
1193:     """  
1194:     Get a path from environment variable.  
1195:  
1196:     Args:  
1197:         key: Environment variable name  
1198:         default: Default value if not set  
1199:  
1200:     Returns:  
1201:         Path from environment or default  
1202:     """  
1203:     value = os.getenv(key)  
1204:     if value is None:  
1205:         return default  
1206:     return Path(value).resolve()  
1207:  
1208:  
1209: def get_workdir() -> Path:  
1210:     """  
1211:     Get the working directory from FLUX_WORKDIR env var or default to project root.  
1212:     """  
1213:     return get_env_path("FLUX_WORKDIR", PROJECT_ROOT) or PROJECT_ROOT  
1214:  
1215:  
1216: def get_tmpdir() -> Path:  
1217:     """  
1218:     Get the temporary directory from FLUX_TMPDIR env var or default to project tmp.  
1219:     """  
1220:     result = get_env_path("FLUX_TMPDIR", tmp_dir()) or tmp_dir()  
1221:     result.mkdir(parents=True, exist_ok=True)  
1222:     return result  
1223:  
1224:  
1225: def get_reports_dir() -> Path:  
1226:     """  
1227:     Get the reports directory from FLUX_REPORTS env var or default to build/reports.  
1228:     """  
1229:     result = get_env_path("FLUX_REPORTS", reports_dir()) or reports_dir()
```

```
1230:     result.mkdir(parents=True, exist_ok=True)
1231:     return result
1232:
1233:
1234: __all__ = [
1235:     # Exceptions
1236:     "PathError",
1237:     "PathTraversalError",
1238:     "PathNotFoundError",
1239:     "PathOutsideWorkspaceError",
1240:     "UnnormalizedPathError",
1241:     # Constants
1242:     "PROJECT_ROOT",
1243:     "SRC_DIR",
1244:     "DATA_DIR",
1245:     "TESTS_DIR",
1246:     # Directory accessors
1247:     "proj_root",
1248:     "src_dir",
1249:     "data_dir",
1250:     "tmp_dir",
1251:     "build_dir",
1252:     "cache_dir",
1253:     "reports_dir",
1254:     # Path operations
1255:     "is_within",
1256:     "safe_join",
1257:     "normalize_unicode",
1258:     "normalize_case",
1259:     "resources",
1260:     # Validation
1261:     "validate_read_path",
1262:     "validate_write_path",
1263:     # Environment
1264:     "get_env_path",
1265:     "get_workdir",
1266:     "get_tmpdir",
1267:     "get_reports_dir",
1268: ]
1269:
1270:
1271:
1272: =====
1273: FILE: src/farfan_pipeline/utils/proof_generator.py
1274: =====
1275:
1276: """Cryptographic proof generation for pipeline execution verification.
1277:
1278: This module generates cryptographic proof files that allow non-engineers to verify
1279: that a pipeline execution was successful and complete. It produces:
1280: - proof.json: Contains execution metadata, phase/question counts, and SHA-256 hashes
1281: - proof.hash: SHA-256 hash of the proof.json file for verification
1282:
1283: The proof is ONLY generated when ALL success conditions are met:
1284: - All phases report success=True
1285: - No abort is active
```

```
1286: - Non-empty artifacts exist (JSON/MD/logs)
1287: """
1288:
1289: import hashlib
1290: import json
1291: from dataclasses import dataclass, field
1292: from pathlib import Path
1293: from typing import Any
1294: from farfan_pipeline.core.calibration.decorators import calibrated_method
1295:
1296:
1297: @dataclass
1298: class ProofData:
1299:     """Container for proof generation data.
1300:
1301:     All fields must be populated from real execution data.
1302:     No values should be invented or hardcoded.
1303:     """
1304:     run_id: str
1305:     timestamp_utc: str
1306:     phases_total: int
1307:     phases_success: int
1308:     questions_total: int
1309:     questions_answered: int
1310:     evidence_records: int
1311:     monolith_hash: str
1312:     questionnaire_hash: str
1313:     catalog_hash: str
1314:     method_map_hash: str
1315:     code_signature: dict[str, str]
1316:
1317:     # Optional fields for additional verification
1318:     input_pdf_hash: str | None = None
1319:     artifacts_manifest: dict[str, str] = field(default_factory=dict)
1320:     execution_metadata: dict[str, Any] = field(default_factory=dict)
1321:
1322:     # Calibration metadata for traceability
1323:     calibration_version: str | None = None
1324:     calibration_hash: str | None = None
1325:
1326:
1327: def compute_file_hash(file_path: Path) -> str:
1328:     """Compute SHA-256 hash of a file.
1329:
1330:     Args:
1331:         file_path: Path to the file to hash
1332:
1333:     Returns:
1334:         Hex string of SHA-256 hash
1335:
1336:     Raises:
1337:         FileNotFoundError: If file doesn't exist
1338:     """
1339:     if not file_path.exists():
1340:         raise FileNotFoundError(f"Cannot hash missing file: {file_path}")
1341:
```

```
1342:     sha256 = hashlib.sha256()
1343:     with open(file_path, 'rb') as f:
1344:         # Read in chunks for large files
1345:         for chunk in iter(lambda: f.read(65536), b''):
1346:             sha256.update(chunk)
1347:     return sha256.hexdigest()
1348:
1349:
1350: def compute_dict_hash(data: dict[str, Any]) -> str:
1351:     """Compute SHA-256 hash of a dictionary.
1352:
1353:     The dictionary is serialized with sort_keys=True and ensure_ascii=True
1354:     to ensure deterministic hashing.
1355:
1356:     Args:
1357:         data: Dictionary to hash
1358:
1359:     Returns:
1360:         Hex string of SHA-256 hash
1361:     """
1362:     json_str = json.dumps(data, sort_keys=True, ensure_ascii=True, separators=(',', ':'))
1363:     return hashlib.sha256(json_str.encode('utf-8')).hexdigest()
1364:
1365:
1366: def compute_code_signatures(src_root: Path) -> dict[str, str]:
1367:     """Compute SHA-256 hashes of core orchestrator files.
1368:
1369:     Args:
1370:         src_root: Root path to the src/farfan_core directory
1371:
1372:     Returns:
1373:         Dictionary mapping filename to SHA-256 hash
1374:
1375:     Raises:
1376:         FileNotFoundError: If any required file is missing
1377:     """
1378:     core_files = {
1379:         'core.py': src_root / 'core' / 'orchestrator' / 'core.py',
1380:         'executors.py': src_root / 'core' / 'orchestrator' / 'executors.py',
1381:         'factory.py': src_root / 'core' / 'orchestrator' / 'factory.py',
1382:     }
1383:
1384:     signatures = {}
1385:     for name, path in core_files.items():
1386:         if not path.exists():
1387:             raise FileNotFoundError(f"Required core file missing: {path}")
1388:         signatures[name] = compute_file_hash(path)
1389:
1390:     return signatures
1391:
1392:
1393: def verify_success_conditions(
1394:     phase_results: list[Any],
1395:     abort_active: bool,
1396:     output_dir: Path,
1397: ) -> tuple[bool, list[str]]:
```

```

1398:     """Verify that all success conditions are met before generating proof.
1399:
1400:     Args:
1401:         phase_results: List of PhaseResult objects from orchestrator
1402:         abort_active: Whether an abort signal is active
1403:         output_dir: Directory where artifacts should exist
1404:
1405:     Returns:
1406:         Tuple of (success: bool, errors: list[str])
1407:     """
1408:     errors = []
1409:
1410:     # Check all phases succeeded
1411:     if not phase_results:
1412:         errors.append("No phase results available")
1413:         return False, errors
1414:
1415:     failed_phases = [
1416:         i for i, result in enumerate(phase_results)
1417:         if not result.success
1418:     ]
1419:     if failed_phases:
1420:         errors.append(f"Phases failed: {failed_phases}")
1421:
1422:     # Check no abort
1423:     if abort_active:
1424:         errors.append("Abort signal is active")
1425:
1426:     # Check for artifacts (at minimum, directory should exist and have content)
1427:     if not output_dir.exists():
1428:         errors.append(f"Output directory does not exist: {output_dir}")
1429:     else:
1430:         # Check for at least some artifacts
1431:         artifacts = list(output_dir.rglob('*.*json')) + list(output_dir.rglob('*.*md'))
1432:         if not artifacts:
1433:             errors.append(f"No artifacts (JSON/MD) found in {output_dir}")
1434:
1435:     return len(errors) == 0, errors
1436:
1437:
1438: def generate_proof(
1439:     proof_data: ProofData,
1440:     output_dir: Path,
1441: ) -> tuple[Path, Path]:
1442:     """Generate proof.json and proof.hash files.
1443:
1444:     Args:
1445:         proof_data: Proof data to serialize
1446:         output_dir: Directory where proof files will be written
1447:
1448:     Returns:
1449:         Tuple of (proof.json path, proof.hash path)
1450:
1451:     Raises:
1452:         ValueError: If proof_data is incomplete
1453:     """

```

```
1454:     # Validate required fields are not empty
1455:     required_fields = [
1456:         'run_id', 'timestamp_utc', 'monolith_hash', 'questionnaire_hash',
1457:         'catalog_hash', 'code_signature'
1458:     ]
1459:     for field_name in required_fields:
1460:         value = getattr(proof_data, field_name)
1461:         if not value:
1462:             raise ValueError(f"Required field '{field_name}' is empty")
1463:
1464:     # Ensure output directory exists
1465:     output_dir.mkdir(parents=True, exist_ok=True)
1466:
1467:     # Build proof dictionary
1468:     proof_dict = {
1469:         'run_id': proof_data.run_id,
1470:         'timestamp_utc': proof_data.timestamp_utc,
1471:         'phases_total': proof_data.phases_total,
1472:         'phases_success': proof_data.phases_success,
1473:         'questions_total': proof_data.questions_total,
1474:         'questions_answered': proof_data.questions_answered,
1475:         'evidence_records': proof_data.evidence_records,
1476:         'monolith_hash': proof_data.monolith_hash,
1477:         'questionnaire_hash': proof_data.questionnaire_hash,
1478:         'catalog_hash': proof_data.catalog_hash,
1479:         'method_map_hash': proof_data.method_map_hash,
1480:         'code_signature': proof_data.code_signature,
1481:     }
1482:
1483:     # Add optional fields if present
1484:     if proof_data.input_pdf_hash:
1485:         proof_dict['input_pdf_hash'] = proof_data.input_pdf_hash
1486:     if proof_data.artifacts_manifest:
1487:         proof_dict['artifacts_manifest'] = proof_data.artifacts_manifest
1488:     if proof_data.execution_metadata:
1489:         proof_dict['execution_metadata'] = proof_data.execution_metadata
1490:
1491:     # Add calibration metadata for traceability
1492:     if proof_data.calibration_version:
1493:         proof_dict['calibration_version'] = proof_data.calibration_version
1494:     if proof_data.calibration_hash:
1495:         proof_dict['calibration_hash'] = proof_data.calibration_hash
1496:
1497:     # Write proof.json with deterministic serialization
1498:     proof_json_path = output_dir / 'proof.json'
1499:     with open(proof_json_path, 'w', encoding='utf-8') as f:
1500:         json.dump(
1501:             proof_dict,
1502:             f,
1503:             sort_keys=True,
1504:             ensure_ascii=True,
1505:             separators=(',', ':'),
1506:         )
1507:
1508:     # Compute hash of proof.json (using compact serialization for hash)
1509:     proof_hash = compute_dict_hash(proof_dict)
```

```
1510:  
1511:     # Write proof.hash  
1512:     proof_hash_path = output_dir / 'proof.hash'  
1513:     with open(proof_hash_path, 'w', encoding='utf-8') as f:  
1514:         f.write(proof_hash)  
1515:  
1516:     return proof_json_path, proof_hash_path  
1517:  
1518:  
1519: def collect_artifacts_manifest(output_dir: Path) -> dict[str, str]:  
1520:     """Collect hashes of all artifacts in output directory.  
1521:  
1522:     Args:  
1523:         output_dir: Directory containing artifacts  
1524:  
1525:     Returns:  
1526:         Dictionary mapping relative path to SHA-256 hash  
1527:     """  
1528:     manifest = {}  
1529:  
1530:     # Find all artifacts (JSON, MD, logs)  
1531:     patterns = ['*.json', '*.md', '*.log', '*.txt']  
1532:     for pattern in patterns:  
1533:         for artifact_path in output_dir.rglob(pattern):  
1534:             # Skip proof files themselves  
1535:             if artifact_path.name in ('proof.json', 'proof.hash'):  
1536:                 continue  
1537:  
1538:             try:  
1539:                 rel_path = artifact_path.relative_to(output_dir)  
1540:                 manifest[str(rel_path)] = compute_file_hash(artifact_path)  
1541:             except (OSErr, PermissionError, ValueError):  
1542:                 # Skip files we can't read or hash  
1543:                 pass  
1544:  
1545:     return manifest  
1546:  
1547:  
1548: def verify_proof(proof_json_path: Path, proof_hash_path: Path) -> tuple[bool, str]:  
1549:     """Verify that a proof.json file matches its proof.hash.  
1550:  
1551:     This allows anyone to verify that the proof hasn't been tampered with.  
1552:  
1553:     Args:  
1554:         proof_json_path: Path to proof.json  
1555:         proof_hash_path: Path to proof.hash  
1556:  
1557:     Returns:  
1558:         Tuple of (valid: bool, message: str)  
1559:     """  
1560:     try:  
1561:         # Read proof.json  
1562:         with open(proof_json_path, encoding='utf-8') as f:  
1563:             proof_dict = json.load(f)  
1564:  
1565:         # Read proof.hash
```

```
1566:         with open(proof_hash_path, encoding='utf-8') as f:
1567:             stored_hash = f.read().strip()
1568:
1569:             # Recompute hash
1570:             computed_hash = compute_dict_hash(proof_dict)
1571:
1572:             # Compare
1573:             if computed_hash == stored_hash:
1574:                 return True, "\u234\u205 Proof verified: hash matches"
1575:             else:
1576:                 return False, f"\u235\u214 Proof verification failed: hash mismatch\n  Expected: {stored_hash}\n  Got: {computed_hash}"
1577:
1578:     except Exception as e:
1579:         return False, f"\u235\u214 Proof verification error: {e}"
1580:
1581:
1582:
1583: =====
1584: FILE: src/farfan_pipeline/utils/qmcm_hooks.py
1585: =====
1586:
1587: """
1588: QMCM (Quality Method Call Monitoring) hooks for ReportAssemblyProducer
1589:
1590: Records method calls for registry tracking and quality assurance.
1591: Does NOT summarize or analyze - only records method invocations.
1592: """
1593:
1594: import json
1595: import logging
1596: from datetime import datetime
1597: from functools import wraps
1598: from pathlib import Path
1599: from typing import Any
1600: from farfan_pipeline.core.parameters import ParameterLoaderV2
1601: from farfan_pipeline.core.calibration.decorators import calibrated_method
1602:
1603: logger = logging.getLogger(__name__)
1604:
1605: class QMCMRecorder:
1606:     """
1607:     Records method calls for quality monitoring
1608:
1609:     Tracks:
1610:         - Method invocations
1611:         - Call frequency
1612:         - Input/output types
1613:         - Execution status
1614:
1615:     Does NOT track:
1616:         - Actual data content (no summarization leakage)
1617:         - User-specific information
1618:     """
1619:
1620:     def __init__(self, recording_path: Path | None = None) -> None:
1621:         """Initialize QMCM recorder"""
```

```
1622:         self.recording_path = recording_path or Path(".qmcm_recording.json")
1623:         self.calls: list[dict[str, Any]] = []
1624:         self.enabled = True
1625:
1626:     def record_call(
1627:         self,
1628:         method_name: str,
1629:         input_types: dict[str, str],
1630:         output_type: str,
1631:         execution_status: str = "success",
1632:         execution_time_ms: float = ParameterLoaderV2.get("farfan_core.utils.qmcm_hooks.QMCMRecorder.__init__", "auto_param_L45_39", 0.0),
1633:         monolith_hash: str | None = None
1634:     ) -> None:
1635:         """
1636:             Record a method call
1637:
1638:         Args:
1639:             method_name: Name of the method called
1640:             input_types: Dictionary mapping parameter names to type names
1641:             output_type: Type name of the return value
1642:             execution_status: 'success', 'error', or 'skipped'
1643:             execution_time_ms: Execution time in milliseconds
1644:             monolith_hash: SHA-256 hash of questionnaire_monolith.json (recommended)
1645:
1646:             ARCHITECTURAL NOTE: Including monolith_hash ties each method call
1647:             to the specific questionnaire version, enabling reproducibility.
1648:             Use factory.compute_monolith_hash() to generate this value.
1649:         """
1650:
1651:         if not self.enabled:
1652:             return
1653:
1654:         call_record = {
1655:             "timestamp": datetime.now().isoformat(),
1656:             "method_name": method_name,
1657:             "input_types": input_types,
1658:             "output_type": output_type,
1659:             "execution_status": execution_status,
1660:             "execution_time_ms": round(execution_time_ms, 2)
1661:         }
1662:
1663:         # Include monolith_hash if provided
1664:         if monolith_hash is not None:
1665:             call_record["monolith_hash"] = monolith_hash
1666:
1667:         self.calls.append(call_record)
1668:         logger.debug(f"QMCM recorded: {method_name}")
1669: @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.get_statistics")
1670: def get_statistics(self) -> dict[str, Any]:
1671:     """
1672:         Get recording statistics
1673:
1674:         Returns summary of method call patterns without data content
1675:     """
1676:
1677:         if not self.calls:
1678:             return {
```

```
1678:         "total_calls": 0,
1679:         "unique_methods": 0,
1680:         "method_frequency": {},
1681:         "success_rate": ParameterLoaderV2.get("farfan_core.utils.qmcm_hooks.QMCMRecorder.get_statistics", "auto_param_L94_32", 0.0),
1682:         "most_called_method": None
1683:     }
1684:
1685:     method_counts = {}
1686:     success_count = 0
1687:
1688:     for call in self.calls:
1689:         method_name = call["method_name"]
1690:         method_counts[method_name] = method_counts.get(method_name, 0) + 1
1691:
1692:         if call["execution_status"] == "success":
1693:             success_count += 1
1694:
1695:     most_called = None
1696:     if method_counts:
1697:         most_called = max(method_counts.items(), key=lambda x: x[1])[0]
1698:
1699:     return {
1700:         "total_calls": len(self.calls),
1701:         "unique_methods": len(method_counts),
1702:         "method_frequency": method_counts,
1703:         "success_rate": success_count / len(self.calls) if self.calls else ParameterLoaderV2.get("farfan_core.utils.qmcm_hooks.QMCMRecorder.get_statistics",
1704: "auto_param_L116_79", 0.0),
1705:         "most_called_method": most_called
1706:     }
1707:
1708:     @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.save_recording")
1709:     def save_recording(self) -> None:
1710:         """Save recording to disk"""
1711:         recording_data = {
1712:             "recording_metadata": {
1713:                 "generated_at": datetime.now().isoformat(),
1714:                 "total_calls": len(self.calls)
1715:             },
1716:             "statistics": self.get_statistics(),
1717:             "calls": self.calls
1718:         }
1719:
1720:         with open(self.recording_path, 'w') as f:
1721:             json.dump(recording_data, f, indent=2)
1722:
1723:         logger.info(f"QMCM recording saved: {self.recording_path}")
1724:
1725:     @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.load_recording")
1726:     def load_recording(self) -> None:
1727:         """Load recording from disk"""
1728:         if not self.recording_path.exists():
1729:             logger.warning(f"No recording found: {self.recording_path}")
1730:
1731:         with open(self.recording_path) as f:
1732:             recording_data = json.load(f)
```

```
1733:         self.calls = recording_data.get("calls", [])
1734:         logger.info(f"QMCM recording loaded: {len(self.calls)} calls")
1735:
1736:
1737:     @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.clear_recording")
1738:     def clear_recording(self) -> None:
1739:         """Clear all recorded calls"""
1740:         self.calls = []
1741:         logger.info("QMCM recording cleared")
1742:
1743:     @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.enable")
1744:     def enable(self) -> None:
1745:         """Enable recording"""
1746:         self.enabled = True
1747:
1748:     @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.disable")
1749:     def disable(self) -> None:
1750:         """Disable recording"""
1751:         self.enabled = False
1752:
1753: # Global recorder instance
1754: _global_recorder: QMCMRecorder | None = None
1755:
1756: def get_global_recorder() -> QMCMRecorder:
1757:     """Get or create global QMCM recorder"""
1758:     global _global_recorder
1759:     if _global_recorder is None:
1760:         _global_recorder = QMCMRecorder()
1761:     return _global_recorder
1762:
1763: def qmcm_record(method=None, *, monolith_hash: str | None = None):
1764:     """
1765:     Decorator to record method calls in QMCM
1766:
1767:     Usage:
1768:         @qmcm_record
1769:         @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.my_method")
1770:         def my_method(self, arg1: str, arg2: int) -> dict:
1771:             return {"result": "data"}
1772:
1773:         # With monolith hash (recommended for questionnaire-dependent methods)
1774:         @qmcm_record(monolith_hash=compute_monolith_hash(questionnaire))
1775:         @calibrated_method("farfan_core.utils.qmcm_hooks.QMCMRecorder.my_questionnaire_method")
1776:         def my_questionnaire_method(self, question_id: str) -> dict:
1777:             return {"result": "data"}
1778:     """
1779:     def decorator(func):
1780:         @wraps(func)
1781:         def wrapper(*args, **kwargs):
1782:             recorder = get_global_recorder()
1783:
1784:             import time
1785:             start_time = time.time()
1786:
1787:             try:
1788:                 result = func(*args, **kwargs)
```

```
1789:         execution_time_ms = (time.time() - start_time) * 1000
1790:
1791:     # Record the call
1792:     input_types = {}
1793:     if args:
1794:         # Skip self argument
1795:         for i, arg in enumerate(args[1:], 1):
1796:             input_types[f"arg{i}"] = type(arg).__name__
1797:     for key, value in kwargs.items():
1798:         input_types[key] = type(value).__name__
1799:
1800:     output_type = type(result).__name__
1801:
1802:     recorder.record_call(
1803:         method_name=func.__name__,
1804:         input_types=input_types,
1805:         output_type=output_type,
1806:         execution_status="success",
1807:         execution_time_ms=execution_time_ms,
1808:         monolith_hash=monolith_hash
1809:     )
1810:
1811:     return result
1812:
1813: except Exception:
1814:     execution_time_ms = (time.time() - start_time) * 1000
1815:
1816:     recorder.record_call(
1817:         method_name=func.__name__,
1818:         input_types={},
1819:         output_type="error",
1820:         execution_status="error",
1821:         execution_time_ms=execution_time_ms,
1822:         monolith_hash=monolith_hash
1823:     )
1824:
1825:     raise
1826:
1827:     return wrapper
1828:
1829: # Handle both @qmcm_record and @qmcm_record() usage
1830: if method is None:
1831:     return decorator
1832: else:
1833:     return decorator(method)
1834:
1835: # Export public API
1836: __all__ = [
1837:     'QMCMRecorder',
1838:     'get_global_recorder',
1839:     'qmcm_record'
1840: ]
1841:
1842:
1843:
1844: =====
```

```
1845: FILE: src/farfan_pipeline/utils/runtime_error_fixes.py
1846: =====
1847:
1848: """
1849: Runtime Error Fixes for Policy Analysis
1850:
1851: This module contains fixes for three critical runtime errors:
1852: 1. 'bool' object is not iterable - Functions returning bool instead of list
1853: 2. 'str' object has no attribute 'text' - String passed where spacy object expected
1854: 3. can't multiply sequence by non-int of type 'float' - List multiplication by float
1855:
1856: These fixes are applied defensively to prevent crashes in production.
1857: """
1858:
1859: from typing import TYPE_CHECKING, Any
1860: from farfan_pipeline.core.calibration.decorators import calibrated_method
1861:
1862: if TYPE_CHECKING:
1863:     import numpy as np
1864:     NumpyArray = np.ndarray
1865: else:
1866:     NumpyArray = Any # type: ignore[misc]
1867:
1868: try:
1869:     HAS_NUMPY = True
1870: except ImportError:
1871:     HAS_NUMPY = False
1872:
1873: def ensure_list_return(value: Any) -> list[Any]:
1874: """
1875: Ensure a value is a list, converting bool/None to empty list.
1876:
1877: Fixes: 'bool' object is not iterable
1878:
1879: Args:
1880:     value: Value that should be a list
1881:
1882: Returns:
1883:     Empty list if value is False/None/bool, otherwise the value as-is
1884: """
1885: if isinstance(value, bool) or value is None:
1886:     return []
1887: if isinstance(value, list):
1888:     return value
1889: # If it's iterable but not a list, convert it
1890: try:
1891:     return list(value)
1892: except (TypeError, ValueError):
1893:     return []
1894:
1895: def safe_text_extract(obj: Any) -> str:
1896: """
1897: Safely extract text from object that might be str or have .text attribute.
1898:
1899: Fixes: 'str' object has no attribute 'text'
1900:
```

```
1901:     Args:
1902:         obj: Object that is either str or has .text attribute (e.g., spacy Doc/Span)
1903:
1904:     Returns:
1905:         Extracted text string
1906:         """
1907:     # If it's already a string, return it
1908:     if isinstance(obj, str):
1909:         return obj
1910:
1911:     # If it has a .text attribute, extract it
1912:     if hasattr(obj, 'text'):
1913:         text_value = obj.text
1914:         if isinstance(text_value, str):
1915:             return text_value
1916:
1917:     # Fallback: convert to string
1918:     return str(obj)
1919:
1920: def safe_weighted_multiply(items: list[float] | NumpyArray, weight: float) -> list[float] | NumpyArray:
1921:     """
1922:     Safely multiply a list or array by a weight.
1923:
1924:     Fixes: can't multiply sequence by non-int of type 'float'
1925:
1926:     Args:
1927:         items: List or array of numbers
1928:         weight: Weight to multiply by
1929:
1930:     Returns:
1931:         New list/array with each element multiplied by weight
1932:         """
1933:     # If it's a numpy array, use numpy multiplication
1934:     if HAS_NUMPY and hasattr(items, '__array_interface__'):
1935:         import numpy as np # Import here for runtime use
1936:         if isinstance(items, np.ndarray):
1937:             return items * weight
1938:
1939:     # If it's a list, use list comprehension
1940:     if isinstance(items, list):
1941:         return [item * weight for item in items]
1942:
1943:     # If it's something else iterable, convert and multiply
1944:     try:
1945:         return [item * weight for item in items]
1946:     except (TypeError, ValueError):
1947:         # If multiplication fails, return empty list
1948:         return []
1949:
1950: def safe_list_iteration(value: Any) -> list[Any]:
1951:     """
1952:     Ensure a value can be safely iterated over.
1953:
1954:     Converts bool, None, or non-iterables to empty list.
1955:     Handles the common error of trying to iterate over bool.
1956:
```

```
1957:     Args:
1958:         value: Value to iterate over
1959:
1960:     Returns:
1961:         Iterable list
1962:         """
1963:     # Reject booleans explicitly
1964:     if isinstance(value, bool):
1965:         return []
1966:
1967:     # Handle None
1968:     if value is None:
1969:         return []
1970:
1971:     # If it's already a list, return it
1972:     if isinstance(value, list):
1973:         return value
1974:
1975:     # If it's a string, don't iterate over characters - return as single item
1976:     if isinstance(value, str):
1977:         return [value]
1978:
1979:     # Try to convert to list
1980:     try:
1981:         return list(value)
1982:     except (TypeError, ValueError):
1983:         return []
1984:
1985:
1986:
1987: =====
1988: FILE: src/farfan_pipeline/utils/schema_monitor.py
1989: =====
1990:
1991: """
1992: SCHEMA DRIFT MONITORING - Watch Production Payloads
1993: =====
1994:
1995: Sample payloads in staging/prod and validate shapes.
1996: Emit metrics on key presence/type.
1997: Page when new keys appear or required keys vanish.
1998:
1999: Catches upstream changes (or LLM output drift) instantly.
2000: """
2001:
2002: from __future__ import annotations
2003:
2004: import json
2005: import logging
2006: import random
2007: from collections import Counter, defaultdict
2008: from dataclasses import dataclass, field
2009: from datetime import datetime
2010: from typing import TYPE_CHECKING, Any, TypedDict
2011: from farfan_pipeline.core.parameters import ParameterLoaderV2
2012: from farfan_pipeline.core.calibration.decorators import calibrated_method
```

```
2013:
2014: if TYPE_CHECKING:
2015:     from collections.abc import Mapping
2016:     from pathlib import Path
2017:
2018: logger = logging.getLogger(__name__)
2019:
2020: # =====
2021: # SCHEMA SHAPE TRACKING
2022: # =====
2023:
2024: class SchemaShape(TypedDict):
2025:     """Shape of a data payload."""
2026:
2027:     keys: set[str]
2028:     types: dict[str, str]
2029:     sample_values: dict[str, Any]
2030:     timestamp: str
2031:
2032: @dataclass
2033: class SchemaStats:
2034:     """Statistics about schema shape over time."""
2035:
2036:     key_frequency: Counter[str] = field(default_factory=Counter)
2037:     type_by_key: dict[str, Counter[str]] = field(default_factory=lambda: defaultdict(Counter))
2038:     new_keys: set[str] = field(default_factory=set)
2039:     missing_keys: set[str] = field(default_factory=set)
2040:     total_samples: int = 0
2041:     last_updated: datetime | None = None
2042:
2043: class SchemaDriftDetector:
2044:     """
2045:         Detects schema drift by sampling payloads and tracking shape changes.
2046:
2047:     Usage:
2048:         detector = SchemaDriftDetector(sample_rate=0.05)
2049:
2050:         # In your API/pipeline
2051:         if detector.should_sample():
2052:             detector.record_payload(data, source="api_input")
2053:
2054:         # Check for drift
2055:         alerts = detector.get_alerts()
2056:
2057:
2058:     def __init__(
2059:         self,
2060:         *,
2061:         sample_rate: float = 0.05,
2062:         baseline_path: Path | None = None,
2063:         alert_threshold: float = 0.1,
2064:     ) -> None:
2065:         """
2066:             Initialize schema drift detector.
2067:
2068:     Args:
```

```
2069:         sample_rate: Percentage of payloads to sample (0.01 = 1%, 0.05 = 5%)
2070:         baseline_path: Path to baseline schema file
2071:         alert_threshold: Threshold for drift alert (% of samples with drift)
2072:         """
2073:         self.sample_rate = sample_rate
2074:         self.baseline_path = baseline_path
2075:         self.alert_threshold = alert_threshold
2076:
2077:         # Tracking state
2078:         self.stats_by_source: dict[str, SchemaStats] = defaultdict(SchemaStats)
2079:         self.baseline_schema: dict[str, SchemaShape] = {}
2080:
2081:         # Load baseline if provided
2082:         if baseline_path and baseline_path.exists():
2083:             self._load_baseline()
2084:
2085:     @calibrated_method("farfan_core.utils.schema_monitor.SchemaDriftDetector.should_sample")
2086:     def should_sample(self) -> bool:
2087:         """Decide whether to sample this payload (probabilistic)."""
2088:         return random.random() < self.sample_rate
2089:
2090:     def record_payload(
2091:         self,
2092:         payload: Mapping[str, Any],
2093:         *,
2094:         source: str,
2095:         timestamp: datetime | None = None,
2096:     ) -> None:
2097:         """
2098:             Record a payload for schema tracking.
2099:
2100:         Args:
2101:             payload: Data payload to analyze
2102:             source: Source identifier (e.g., "api_input", "document_loader")
2103:             timestamp: Optional timestamp, defaults to now
2104:         """
2105:         if timestamp is None:
2106:             timestamp = datetime.utcnow()
2107:
2108:         stats = self.stats_by_source[source]
2109:
2110:         # Extract shape
2111:         keys = set(payload.keys())
2112:         types = {k: type(v).__name__ for k, v in payload.items()}
2113:
2114:         # Update statistics
2115:         stats.total_samples += 1
2116:         stats.last_updated = timestamp
2117:
2118:         for key in keys:
2119:             stats.key_frequency[key] += 1
2120:             stats.type_by_key[key][types[key]] += 1
2121:
2122:         # Detect new keys (compared to baseline)
2123:         if source in self.baseline_schema:
2124:             baseline_keys = self.baseline_schema[source]["keys"]
```

```

2125:         new_keys = keys - baseline_keys
2126:         if new_keys:
2127:             stats.new_keys.update(new_keys)
2128:             logger.warning(
2129:                 f"SCHEMA_DRIFT[source={source}]: New keys detected: {new_keys}"
2130:             )
2131:
2132:         missing_keys = baseline_keys - keys
2133:         if missing_keys:
2134:             stats.missing_keys.update(missing_keys)
2135:             logger.warning(
2136:                 f"SCHEMA_DRIFT[source={source}]: Missing keys detected: {missing_keys}"
2137:             )
2138:
2139:     @calibrated_method("farfan_core.utils.schema_monitor.SchemaDriftDetector.get_alerts")
2140:     def get_alerts(self, *, source: str | None = None) -> list[dict[str, Any]]:
2141:         """
2142:             Get schema drift alerts.
2143:
2144:             Args:
2145:                 source: Optional source filter
2146:
2147:             Returns:
2148:                 List of alert dicts
2149:             """
2150:         alerts: list[dict[str, Any]] = []
2151:
2152:         sources = [source] if source else list(self.stats_by_source.keys())
2153:
2154:         for src in sources:
2155:             stats = self.stats_by_source[src]
2156:
2157:             if stats.new_keys:
2158:                 alerts.append({
2159:                     "level": "WARNING",
2160:                     "source": src,
2161:                     "type": "NEW_KEYS",
2162:                     "keys": list(stats.new_keys),
2163:                     "timestamp": stats.last_updated.isoformat() if stats.last_updated else None,
2164:                 })
2165:
2166:             if stats.missing_keys:
2167:                 alerts.append({
2168:                     "level": "CRITICAL",
2169:                     "source": src,
2170:                     "type": "MISSING_KEYS",
2171:                     "keys": list(stats.missing_keys),
2172:                     "timestamp": stats.last_updated.isoformat() if stats.last_updated else None,
2173:                 })
2174:
2175:             # Check for type inconsistencies
2176:             for key, type_counts in stats.type_by_key.items():
2177:                 if len(type_counts) > 1:
2178:                     # Multiple types seen for same key
2179:                     dominant_type = type_counts.most_common(1)[0][0]
2180:                     other_types = [t for t in type_counts if t != dominant_type]

```

```
2181:             alerts.append({
2182:                 "level": "WARNING",
2183:                 "source": src,
2184:                 "type": "TYPE_INCONSISTENCY",
2185:                 "key": key,
2186:                 "expected_type": dominant_type,
2187:                 "observed_types": other_types,
2188:                 "timestamp": stats.last_updated.isoformat() if stats.last_updated else None,
2189:             })
2190:
2191:     return alerts
2192:
2193:
2194: @calibrated_method("farfan_core.utils.schema_monitor.SchemaDriftDetector.save_baseline")
2195: def save_baseline(self, output_path: Path) -> None:
2196:     """
2197:     Save current schema shapes as baseline.
2198:
2199:     Args:
2200:         output_path: Path to save baseline JSON
2201:     """
2202:     baseline: dict[str, dict[str, Any]] = {}
2203:
2204:     for source, stats in self.stats_by_source.items():
2205:         # Get most common keys (present in >50% of samples)
2206:         threshold = stats.total_samples * ParameterLoaderV2.get("farfan_core.utils.schema_monitor.SchemaDriftDetector.save_baseline", "auto_param_L215_4
6", 0.5)
2207:         common_keys = [
2208:             key for key, count in stats.key_frequency.items()
2209:             if count >= threshold
2210:         ]
2211:
2212:         # Get dominant type for each key
2213:         types = [
2214:             key: type_counts.most_common(1)[0][0]
2215:             for key, type_counts in stats.type_by_key.items()
2216:         ]
2217:
2218:         baseline[source] = {
2219:             "keys": list(common_keys),
2220:             "types": types,
2221:             "timestamp": datetime.utcnow().isoformat(),
2222:         }
2223:
2224:     output_path.write_text(json.dumps(baseline, indent=2))
2225:     logger.info(f"Saved schema baseline to {output_path}")
2226:
2227: @calibrated_method("farfan_core.utils.schema_monitor.SchemaDriftDetector._load_baseline")
2228: def _load_baseline(self) -> None:
2229:     """Load baseline schema from file."""
2230:     if not self.baseline_path:
2231:         return
2232:
2233:     try:
2234:         data = json.loads(self.baseline_path.read_text())
2235:
```

```
2236:         for source, shape_data in data.items():
2237:             self.baseline_schema[source] = {
2238:                 "keys": set(shape_data["keys"]),
2239:                 "types": shape_data["types"],
2240:                 "sample_values": {},
2241:                 "timestamp": shape_data["timestamp"],
2242:             }
2243:
2244:             logger.info(f"Loaded schema baseline from {self.baseline_path}")
2245:         except Exception as e:
2246:             logger.error(f"Failed to load baseline: {e}")
2247:
2248:     @calibrated_method("farfan_core.utils.schema_monitor.SchemaDriftDetector.get_metrics")
2249:     def get_metrics(self, *, source: str | None = None) -> dict[str, Any]:
2250:         """
2251:             Get monitoring metrics.
2252:
2253:             Args:
2254:                 source: Optional source filter
2255:
2256:             Returns:
2257:                 Dict of metrics
2258:         """
2259:         if source:
2260:             stats = self.stats_by_source.get(source)
2261:             if not stats:
2262:                 return {}
2263:
2264:             return {
2265:                 "source": source,
2266:                 "total_samples": stats.total_samples,
2267:                 "unique_keys": len(stats.key_frequency),
2268:                 "new_keys_count": len(stats.new_keys),
2269:                 "missing_keys_count": len(stats.missing_keys),
2270:                 "type_inconsistencies": sum(
2271:                     1 for counts in stats.type_by_key.values()
2272:                         if len(counts) > 1
2273:                 ),
2274:             },
2275:         }
2276:
2277:         # Aggregate across all sources
2278:         return {
2279:             "sources": list(self.stats_by_source.keys()),
2280:             "total_samples": sum(s.total_samples for s in self.stats_by_source.values()),
2281:             "sources_with_drift": len([
2282:                 s for s in self.stats_by_source.values()
2283:                     if s.new_keys or s.missing_keys
2284:             ]),
2285:         }
2286:     # =====
2287:     # PAYLOAD VALIDATOR
2288:     # =====
2289:
2290:     class PayloadValidator:
2291:         """
```

```
2292:     Validate payloads against expected schema.
2293:
2294:     Usage:
2295:         validator = PayloadValidator(schema_path=Path("schemas/api_input.json"))
2296:
2297:         try:
2298:             validator.validate(data, source="api_endpoint")
2299:         except ValueError as e:
2300:             logger.error(f"Validation failed: {e}")
2301: """
2302:
2303: def __init__(self, *, schema_path: Path | None = None) -> None:
2304: """
2305:     Initialize payload validator.
2306:
2307:     Args:
2308:         schema_path: Path to schema definition JSON
2309: """
2310:     self.schema_path = schema_path
2311:     self.schemas: dict[str, dict[str, Any]] = {}
2312:
2313:     if schema_path and schema_path.exists():
2314:         self._load_schemas()
2315:
2316:     def validate(
2317:         self,
2318:         payload: Mapping[str, Any],
2319:         *,
2320:         source: str,
2321:         strict: bool = True,
2322:     ) -> None:
2323: """
2324:     Validate payload against schema.
2325:
2326:     Args:
2327:         payload: Data payload to validate
2328:         source: Source identifier
2329:         strict: If True, raise on missing keys; if False, only warn
2330:
2331:     Raises:
2332:         ValueError: If validation fails in strict mode
2333:         TypeError: If value types don't match schema
2334: """
2335:     if source not in self.schemas:
2336:         logger.warning(f"No schema defined for source '{source}'")
2337:         return
2338:
2339:     schema = self.schemas[source]
2340:     required_keys = set(schema.get("required_keys", []))
2341:     expected_types = schema.get("types", {})
2342:
2343:     # Check required keys
2344:     payload_keys = set(payload.keys())
2345:     missing = required_keys - payload_keys
2346:
2347:     if missing:
```

```
2348:             msg = f"VALIDATION_ERROR[source={source}]: Missing required keys: {missing}"
2349:             if strict:
2350:                 raise ValueError(msg)
2351:             else:
2352:                 logger.warning(msg)
2353:
2354:             # Check types
2355:             for key, expected_type in expected_types.items():
2356:                 if key in payload:
2357:                     actual_type = type(payload[key]).__name__
2358:                     if actual_type != expected_type:
2359:                         msg = (
2360:                             f"VALIDATION_ERROR[source={source}, key={key}]: "
2361:                             f"Expected type {expected_type}, got {actual_type}"
2362:                         )
2363:                     if strict:
2364:                         raise TypeError(msg)
2365:                     else:
2366:                         logger.warning(msg)
2367:
2368:             @calibrated_method("farfan_core.utils.schema_monitor.PayloadValidator._load_schemas")
2369:             def _load_schemas(self) -> None:
2370:                 """Load schema definitions from file."""
2371:                 if not self.schema_path:
2372:                     return
2373:
2374:                 try:
2375:                     self.schemas = json.loads(self.schema_path.read_text())
2376:                     logger.info(f"Loaded schemas from {self.schema_path}")
2377:                 except Exception as e:
2378:                     logger.error(f"Failed to load schemas: {e}")
2379:
2380: # =====
2381: # GLOBAL INSTANCE (optional convenience)
2382: # =====
2383:
2384: # Singleton detector for application-wide use
2385: _global_detector: SchemaDriftDetector | None = None
2386:
2387: def get_detector() -> SchemaDriftDetector:
2388:     """Get or create global schema drift detector."""
2389:     global _global_detector
2390:     if _global_detector is None:
2391:         _global_detector = SchemaDriftDetector(sample_rate=ParameterLoaderV2.get("farfan_core.utils.schema_monitor.PayloadValidator._load_schemas", "auto_param_L400_59"), 0.05)
2392:     return _global_detector
2393:
2394:
2395:
2396: =====
2397: FILE: src/farfan_pipeline/utils/seed_factory.py
2398: =====
2399:
2400: """
2401: Deterministic Seed Factory
2402: Generates reproducible seeds for all stochastic operations
```

```
2403: """
2404:
2405: import hashlib
2406: import hmac
2407: import random
2408: from typing import Any
2409: from farfan_pipeline.core.calibration.decorators import calibrated_method
2410:
2411: try:
2412:     import numpy as np
2413:     NUMPY_AVAILABLE = True
2414: except ImportError:
2415:     NUMPY_AVAILABLE = False
2416:     np = None # type: ignore
2417:
2418: class SeedFactory:
2419:     """
2420:         Factory for generating deterministic seeds
2421:
2422:     Ensures:
2423:         - Reproducibility: Same inputs \u2192 same seed
2424:         - Uniqueness: Different contexts \u2192 different seeds
2425:         - Cryptographic quality: HMAC-SHA256 derivation
2426:     """
2427:
2428:     # Fixed salt for seed derivation (should be configured per deployment)
2429:     DEFAULT_SALT = b"PDM_EVALUATOR_V2_DETERMINISTIC_SEED_2025"
2430:
2431:     def __init__(self, fixed_salt: bytes | None = None) -> None:
2432:         self.salt = fixed_salt or self.DEFAULT_SALT
2433:
2434:     def create_deterministic_seed(
2435:         self,
2436:         correlation_id: str,
2437:         file_checksums: dict[str, str] | None = None,
2438:         context: dict[str, Any] | None = None
2439:     ) -> int:
2440:         """
2441:             Generate deterministic seed from correlation ID and context
2442:
2443:         Args:
2444:             correlation_id: Unique workflow instance identifier
2445:             file_checksums: Dict of {filename: sha256_checksum}
2446:             context: Additional context (question_id, policy_area, etc.)
2447:
2448:         Returns:
2449:             32-bit integer seed (0 to 2^32-1)
2450:
2451:         Example:
2452:             >>> factory = SeedFactory()
2453:             >>> seed1 = factory.create_deterministic_seed("run-001", {"data.json": "abc123"})
2454:             >>> seed2 = factory.create_deterministic_seed("run-001", {"data.json": "abc123"})
2455:             >>> assert seed1 == seed2 # Reproducible
2456:         """
2457:
2458:         # Build deterministic input string
```

```
2459:     components = [correlation_id]
2460:
2461:     # Add file checksums (sorted for determinism)
2462:     if file_checksums:
2463:         sorted_checksums = sorted(file_checksums.items())
2464:         for filename, checksum in sorted_checksums:
2465:             components.append(f"{filename}:{checksum}")
2466:
2467:     # Add context (sorted for determinism)
2468:     if context:
2469:         sorted_context = sorted(context.items())
2470:         for key, value in sorted_context:
2471:             components.append(f"{key}={value}")
2472:
2473:     # Combine with deterministic separator
2474:     seed_input = "|".join(components).encode('utf-8')
2475:
2476:     # HMAC-SHA256 for cryptographic quality
2477:     seed_hmac = hmac.new(
2478:         key=self.salt,
2479:         msg=seed_input,
2480:         digestmod=hashlib.sha256
2481:     )
2482:
2483:     # Convert to 32-bit integer
2484:     seed_bytes = seed_hmac.digest()[:4]  # First 4 bytes
2485:     seed_int = int.from_bytes(seed_bytes, byteorder='big')
2486:
2487:     return seed_int
2488:
2489: @calibrated_method("farfan_core.utils.seed_factory.SeedFactory.configure_global_random_state")
2490: def configure_global_random_state(self, seed: int) -> None:
2491:     """
2492:         Configure all random number generators with seed
2493:
2494:         Sets:
2495:             - Python random module
2496:             - NumPy random state
2497:             - (Add torch, tensorflow if needed)
2498:
2499:         Args:
2500:             seed: Deterministic seed
2501:         """
2502:
2503:         # Python random module
2504:         random.seed(seed)
2505:
2506:         # NumPy
2507:         if NUMPY_AVAILABLE and np is not None:
2508:             np.random.seed(seed)
2509:
2510:         # TODO: Add torch.manual_seed(seed) if PyTorch is used
2511:         # TODO: Add tf.random.set_seed(seed) if TensorFlow is used
2512:
2513: class DeterministicContext:
2514:     """
```

```
2515:     Context manager for deterministic execution
2516:
2517:     Usage:
2518:         with DeterministicContext(correlation_id="run-001") as seed:
2519:             # All random operations are deterministic
2520:             result = some_stochastic_function()
2521:
2522:
2523:     def __init__(
2524:         self,
2525:         correlation_id: str,
2526:         file_checksums: dict[str, str] | None = None,
2527:         context: dict[str, Any] | None = None,
2528:         fixed_salt: bytes | None = None
2529:     ) -> None:
2530:         self.correlation_id = correlation_id
2531:         self.file_checksums = file_checksums
2532:         self.context = context
2533:         self.factory = SeedFactory(fixed_salt=fixed_salt)
2534:
2535:         self.seed: int | None = None
2536:         self.previous_random_state = None
2537:         self.previous_numpy_state = None
2538:
2539:     def __enter__(self) -> int:
2540:         """Enter deterministic context"""
2541:
2542:         # Generate deterministic seed
2543:         self.seed = self.factory.create_deterministic_seed(
2544:             correlation_id=self.correlation_id,
2545:             file_checksums=self.file_checksums,
2546:             context=self.context
2547:         )
2548:
2549:         # Save current random states
2550:         self.previous_random_state = random.getstate()
2551:         if NUMPY_AVAILABLE and np is not None:
2552:             self.previous_numpy_state = np.random.get_state()
2553:
2554:         # Configure with deterministic seed
2555:         self.factory.configure_global_random_state(self.seed)
2556:
2557:         return self.seed
2558:
2559:     def __exit__(self, exc_type, exc_val, exc_tb):
2560:         """Exit deterministic context and restore previous state"""
2561:
2562:         # Restore previous random states
2563:         if self.previous_random_state:
2564:             random.setstate(self.previous_random_state)
2565:
2566:         if NUMPY_AVAILABLE and np is not None and self.previous_numpy_state:
2567:             np.random.set_state(self.previous_numpy_state)
2568:
2569:         return False
2570:
```

```
2571: def create_deterministic_seed(
2572:     correlation_id: str,
2573:     file_checksums: dict[str, str] | None = None,
2574:     **context_kwargs
2575: ) -> int:
2576:     """
2577:         Convenience function for creating deterministic seed
2578:
2579:     Args:
2580:         correlation_id: Unique workflow instance ID
2581:         file_checksums: Dict of file checksums
2582:         **context_kwargs: Additional context as keyword arguments
2583:
2584:     Returns:
2585:         Deterministic 32-bit integer seed
2586:
2587:     Example:
2588:         >>> seed = create_deterministic_seed(
2589:             ...      "run-001",
2590:             ...      question_id="P1-D1-Q001",
2591:             ...      policy_area="P1"
2592:             ...
2593:         """
2594:         factory = SeedFactory()
2595:         return factory.create_deterministic_seed(
2596:             correlation_id=correlation_id,
2597:             file_checksums=file_checksums,
2598:             context=context_kwargs if context_kwargs else None
2599:         )
2600:
2601:
2602:
2603: =====
2604: FILE: src/farfan_pipeline/utils/signature_validator.py
2605: =====
2606:
2607: """
2608: Signature Validation and Interface Governance System
2609: =====
2610:
2611: Implements automated signature consistency auditing, runtime validation,
2612: and interface governance to prevent function signature mismatches.
2613:
2614: Based on the Strategic Mitigation Plan for addressing interface inconsistencies.
2615:
2616: Author: Signature Governance Team
2617: Version: 1.0.0
2618: """
2619:
2620: import ast
2621: import functools
2622: import hashlib
2623: import inspect
2624: import json
2625: import logging
2626: from collections.abc import Callable
```

```
2627: from dataclasses import asdict, dataclass, field
2628: from datetime import datetime
2629: from pathlib import Path
2630: from typing import Any, TypeVar, get_type_hints
2631: from farfan_pipeline.core.calibration.decorators import calibrated_method
2632:
2633: logger = logging.getLogger(__name__)
2634:
2635: # Type variable for decorated functions
2636: F = TypeVar('F', bound=Callable[..., Any])
2637:
2638: # =====
2639: # SIGNATURE METADATA STORAGE
2640: # =====
2641:
2642: @dataclass
2643: class FunctionSignature:
2644:     """Stores metadata about a function's signature"""
2645:     module: str
2646:     class_name: str | None
2647:     function_name: str
2648:     parameters: list[str]
2649:     parameter_types: dict[str, str]
2650:     return_type: str
2651:     signature_hash: str
2652:     timestamp: str = field(default_factory=lambda: datetime.now().isoformat())
2653:
2654:     @calibrated_method("farfan_core.utils.signature_validator.FunctionSignature.to_dict")
2655:     def to_dict(self) -> dict[str, Any]:
2656:         return asdict(self)
2657:
2658: class SignatureRegistry:
2659:     """
2660:         Maintains a registry of function signatures with version tracking
2661:         Implements signature snapshotting as described in the mitigation plan
2662:     """
2663:
2664:     def __init__(self, registry_path: Path = Path("data/signature_registry.json")) -> None:
2665:         self.registry_path = registry_path
2666:         self.signatures: dict[str, FunctionSignature] = {}
2667:         self.load()
2668:
2669:     @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.compute_signature_hash")
2670:     def compute_signature_hash(self, func: Callable) -> str:
2671:         """Compute a hash of the function's signature"""
2672:         sig = inspect.signature(func)
2673:         sig_str = str(sig)
2674:         return hashlib.sha256(sig_str.encode()).hexdigest()[:16]
2675:
2676:     @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.register_function")
2677:     def register_function(self, func: Callable) -> FunctionSignature:
2678:         """Register a function's signature"""
2679:         sig = inspect.signature(func)
2680:
2681:         # Extract parameter information
2682:         parameters = list(sig.parameters.keys())
```

```
2683:  
2684:     # Get type hints if available  
2685:     try:  
2686:         type_hints = get_type_hints(func)  
2687:         parameter_types = {  
2688:             name: str(type_hints.get(name, 'Any'))  
2689:             for name in parameters  
2690:         }  
2691:         return_type = str(type_hints.get('return', 'Any'))  
2692:     except (TypeError, AttributeError, NameError) as e:  
2693:         # get_type_hints can fail for various reasons:  
2694:         # - TypeError: if func is not a callable  
2695:         # - AttributeError: if func doesn't have required attributes  
2696:         # - NameError: if type hints reference undefined names  
2697:         logger.debug(f"Could not extract type hints for {func.__name__}: {e}")  
2698:         parameter_types = dict.fromkeys(parameters, 'Any')  
2699:         return_type = 'Any'  
2700:  
2701:     # Get module and class information  
2702:     module = func.__module__ if hasattr(func, '__module__') else 'unknown'  
2703:     class_name = None  
2704:     if hasattr(func, '__qualname__') and '.' in func.__qualname__:  
2705:         class_name = func.__qualname__.rsplit('.', 1)[0]  
2706:  
2707:     signature_hash = self.compute_signature_hash(func)  
2708:  
2709:     func_sig = FunctionSignature(  
2710:         module=module,  
2711:         class_name=class_name,  
2712:         function_name=func.__name__,  
2713:         parameters=parameters,  
2714:         parameter_types=parameter_types,  
2715:         return_type=return_type,  
2716:         signature_hash=signature_hash  
2717:     )  
2718:  
2719:     # Store in registry  
2720:     key = self._get_function_key(module, class_name, func.__name__)  
2721:     self.signatures[key] = func_sig  
2722:  
2723:     return func_sig  
2724:  
2725: @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry._get_function_key")  
2726: def _get_function_key(self, module: str, class_name: str | None, func_name: str) -> str:  
2727:     """Generate a unique key for a function"""  
2728:     if class_name:  
2729:         return f"{module}.{class_name}.{func_name}"  
2730:     return f"{module}.{func_name}"  
2731:  
2732: @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.get_signature")  
2733: def get_signature(self, module: str, class_name: str | None, func_name: str) -> FunctionSignature | None:  
2734:     """Retrieve a stored signature"""  
2735:     key = self._get_function_key(module, class_name, func_name)  
2736:     return self.signatures.get(key)  
2737:  
2738: @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.has_signature_changed")
```

```
2739:     def has_signature_changed(self, func: Callable) -> tuple[bool, FunctionSignature | None, FunctionSignature | None]:
2740:         """Check if a function's signature has changed from the registered version"""
2741:         module = func.__module__ if hasattr(func, '__module__') else 'unknown'
2742:         class_name = None
2743:         if hasattr(func, '__qualname__') and '.' in func.__qualname__:
2744:             class_name = func.__qualname__.rsplit('.', 1)[0]
2745:
2746:         old_sig = self.get_signature(module, class_name, func.__name__)
2747:         if old_sig is None:
2748:             return False, None, None # No previous signature
2749:
2750:         new_sig = self.register_function(func)
2751:         changed = old_sig.signature_hash != new_sig.signature_hash
2752:
2753:         return changed, old_sig, new_sig
2754:
2755:     @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.save")
2756:     def save(self) -> None:
2757:         """Save registry to disk"""
2758:         self.registry_path.parent.mkdir(parents=True, exist_ok=True)
2759:
2760:         registry_data = {
2761:             key: sig.to_dict()
2762:             for key, sig in self.signatures.items()
2763:         }
2764:
2765:         with open(self.registry_path, 'w') as f:
2766:             json.dump(registry_data, f, indent=2)
2767:
2768:         logger.info(f"Saved {len(self.signatures)} signatures to {self.registry_path}")
2769:
2770:     @calibrated_method("farfan_core.utils.signature_validator.SignatureRegistry.load")
2771:     def load(self) -> None:
2772:         """Load registry from disk"""
2773:         if not self.registry_path.exists():
2774:             logger.info(f"No existing registry found at {self.registry_path}")
2775:             return
2776:
2777:         try:
2778:             with open(self.registry_path) as f:
2779:                 registry_data = json.load(f)
2780:
2781:                 self.signatures = {
2782:                     key: FunctionSignature(**data)
2783:                         for key, data in registry_data.items()
2784:                 }
2785:
2786:                 logger.info(f"Loaded {len(self.signatures)} signatures from {self.registry_path}")
2787:             except Exception as e:
2788:                 logger.error(f"Failed to load registry: {e}")
2789:
2790: # Global registry instance
2791: _signature_registry = SignatureRegistry()
2792:
2793: # =====
2794: # RUNTIME VALIDATION DECORATOR
```

```
2795: # =====
2796:
2797: def validate_signature(enforce: bool = True, track: bool = True):
2798:     """
2799:         Decorator to validate function calls against expected signatures at runtime
2800:
2801:     Args:
2802:         enforce: If True, raise TypeError on signature violations
2803:         track: If True, register signature in the global registry
2804:
2805:     Example:
2806:         @validate_signature(enforce=True)
2807:         def my_function(arg1: str, arg2: int) -> bool:
2808:             return True
2809:     """
2810:     def decorator(func: F) -> F:
2811:         # Register function signature if tracking is enabled
2812:         if track:
2813:             _signature_registry.register_function(func)
2814:
2815:         # Get the function signature
2816:         sig = inspect.signature(func)
2817:
2818:         @functools.wraps(func)
2819:         def wrapper(*args, **kwargs):
2820:             # Bind arguments to signature
2821:             try:
2822:                 bound_args = sig.bind(*args, **kwargs)
2823:                 bound_args.apply_defaults()
2824:             except TypeError as e:
2825:                 error_msg = (
2826:                     f"Signature mismatch in {func.__module__}.{func.__qualname__}: {e}\n"
2827:                     f"Expected signature: {sig}\n"
2828:                     f"Called with args: {args}, kwargs: {kwargs}"
2829:                 )
2830:                 logger.error(error_msg)
2831:
2832:                 if enforce:
2833:                     raise TypeError(error_msg) from e
2834:                 else:
2835:                     logger.warning(f"Signature validation failed but enforcement is disabled: {e}")
2836:
2837:             # Call the original function
2838:             return func(*args, **kwargs)
2839:
2840:         return wrapper # type: ignore
2841:
2842:     return decorator
2843:
2844: def validate_call_signature(func: Callable, *args, **kwargs) -> bool:
2845:     """
2846:         Validate that a function call matches the expected signature without actually calling it
2847:
2848:     Args:
2849:         func: Function to validate
2850:         *args: Positional arguments
```

```
2851:         **kwargs: Keyword arguments
2852:
2853:     Returns:
2854:         True if signature is valid, False otherwise
2855:
2856:     try:
2857:         sig = inspect.signature(func)
2858:         sig.bind(*args, **kwargs)
2859:         return True
2860:     except TypeError:
2861:         return False
2862:
2863: # =====
2864: # STATIC SIGNATURE AUDITOR
2865: # =====
2866:
2867: @dataclass
2868: class SignatureMismatch:
2869:     """Represents a detected signature mismatch"""
2870:     caller_module: str
2871:     caller_function: str
2872:     caller_line: int
2873:     callee_module: str
2874:     callee_class: str | None
2875:     callee_function: str
2876:     expected_signature: str
2877:     actual_call: str
2878:     severity: str # 'high', 'medium', 'low'
2879:     description: str
2880:
2881: class SignatureAuditor:
2882:     """
2883:         Static introspection tool to cross-validate function definitions against call sites
2884:         Implements automated signature consistency audit from the mitigation plan
2885:     """
2886:
2887:     def __init__(self) -> None:
2888:         self.mismatches: list[SignatureMismatch] = []
2889:         self.call_graph: dict[str, list[tuple[str, int, list[str], dict[str, str]]]] = {}
2890:
2891:     @calibrated_method("farfan_core.utils.signature_validator.SignatureAuditor.audit_module")
2892:     def audit_module(self, module_path: Path) -> list[SignatureMismatch]:
2893:         """
2894:             Audit a Python module for signature mismatches
2895:
2896:             Args:
2897:                 module_path: Path to the Python module
2898:
2899:             Returns:
2900:                 List of detected signature mismatches
2901:         """
2902:         logger.info(f"Auditing module: {module_path}")
2903:
2904:         # Skip test files, virtual environments, and build directories
2905:         exclude_patterns = ['test', 'venv', '.venv', '__pycache__', '.git', 'build', 'dist']
2906:         if any(module_path.match(f'*/{pattern}/*') or module_path.match(f'*/{pattern}'))
```

```
2907:         for pattern in exclude_patterns):
2908:             logger.debug(f"Skipping excluded path: {module_path}")
2909:             return []
2910:
2911:     try:
2912:         with open(module_path, encoding='utf-8') as f:
2913:             source_code = f.read()
2914:
2915:         tree = ast.parse(source_code, filename=str(module_path))
2916:
2917:         # Extract function definitions
2918:         function_defs = self._extract_function_definitions(tree, module_path.stem)
2919:
2920:         # Extract function calls
2921:         function_calls = self._extract_function_calls(tree, module_path.stem)
2922:
2923:         # Cross-validate
2924:         mismatches = self._cross_validate(function_defs, function_calls)
2925:
2926:         self.mismatches.extend(mismatches)
2927:
2928:     return mismatches
2929:
2930: except Exception as e:
2931:     logger.error(f"Failed to audit {module_path}: {e}")
2932:     return []
2933:
2934: @calibrated_method("farfan_core.utils.signature_validator.SignatureAuditor._extract_function_definitions")
2935: def _extract_function_definitions(self, tree: ast.AST, module_name: str) -> dict[str, ast.FunctionDef]:
2936:     """Extract all function definitions from AST"""
2937:     functions = {}
2938:
2939:     for node in ast.walk(tree):
2940:         if isinstance(node, ast.FunctionDef):
2941:             # Generate full qualified name
2942:             full_name = f"{module_name}.{node.name}"
2943:             functions[full_name] = node
2944:
2945:     return functions
2946:
2947: @calibrated_method("farfan_core.utils.signature_validator.SignatureAuditor._extract_function_calls")
2948: def _extract_function_calls(self, tree: ast.AST, module_name: str) -> list[tuple[str, int, ast.Call]]:
2949:     """Extract all function calls from AST"""
2950:     calls = []
2951:
2952:     for node in ast.walk(tree):
2953:         if isinstance(node, ast.Call):
2954:             # Try to get the function name
2955:             func_name = None
2956:             if isinstance(node.func, ast.Name):
2957:                 func_name = node.func.id
2958:             elif isinstance(node.func, ast.Attribute):
2959:                 func_name = node.func.attr
2960:
2961:             if func_name:
2962:                 calls.append((func_name, node.lineno, node))
```

```
2963:  
2964:         return calls  
2965:  
2966:     def _cross_validate(  
2967:         self,  
2968:         function_defs: dict[str, ast.FunctionDef],  
2969:         function_calls: list[tuple[str, int, ast.Call]]  
2970:     ) -> list[SignatureMismatch]:  
2971:         """Cross-validate function calls against definitions"""  
2972:         mismatches = []  
2973:  
2974:         # This is a simplified implementation  
2975:         # A full implementation would need more sophisticated analysis  
2976:  
2977:         return mismatches  
2978:  
2979:     @calibrated_method("farfan_core.utils.signature_validator.SignatureAuditor.export_report")  
2980:     def export_report(self, output_path: Path) -> None:  
2981:         """Export audit report to JSON"""  
2982:         output_path.parent.mkdir(parents=True, exist_ok=True)  
2983:  
2984:         report = {  
2985:             "audit_timestamp": datetime.now().isoformat(),  
2986:             "total_mismatches": len(self.mismatches),  
2987:             "mismatches": [asdict(m) for m in self.mismatches]  
2988:         }  
2989:  
2990:         with open(output_path, 'w') as f:  
2991:             json.dump(report, f, indent=2)  
2992:  
2993:         logger.info(f"Exported audit report to {output_path}")  
2994:  
2995: # =====  
2996: # COMPATIBILITY LAYER  
2997: # =====  
2998:  
2999: def create_adapter(  
3000:     func: Callable,  
3001:     old_params: list[str],  
3002:     new_params: list[str],  
3003:     param_mapping: dict[str, str] | None = None  
3004: ) -> Callable:  
3005:     """  
3006:     Create a backward-compatible adapter for a function with changed signature  
3007:  
3008:     Args:  
3009:         func: The new function with updated signature  
3010:         old_params: List of old parameter names  
3011:         new_params: List of new parameter names  
3012:         param_mapping: Optional mapping from old to new parameter names  
3013:  
3014:     Returns:  
3015:         Adapter function that accepts old signature and calls new function  
3016:     """  
3017:     param_mapping = param_mapping or {}  
3018:
```

```
3019:     @functools.wraps(func)
3020:     def adapter(*args, **kwargs):
3021:         # Remap old parameter names to new ones
3022:         new_kwargs = {}
3023:         for old_key, value in kwargs.items():
3024:             new_key = param_mapping.get(old_key, old_key)
3025:             new_kwargs[new_key] = value
3026:
3027:         return func(*args, **new_kwargs)
3028:
3029:     return adapter
3030:
3031: # =====
3032: # MODULE INITIALIZATION
3033: # =====
3034:
3035: def initialize_signature_registry(project_root: Path) -> None:
3036:     """
3037:     Initialize signature registry by scanning all Python files in the project
3038:
3039:     Args:
3040:         project_root: Root directory of the project
3041:     """
3042:     logger.info(f"Initializing signature registry for project: {project_root}")
3043:
3044:     python_files = list(project_root.glob("*/*.py"))
3045:     logger.info(f"Found {len(python_files)} Python files")
3046:
3047:     # This would require dynamic import which is complex
3048:     # For now, we rely on decorators to register functions
3049:
3050:     _signature_registry.save()
3051:
3052: def audit_project_signatures(project_root: Path, output_path: Path | None = None) -> list[SignatureMismatch]:
3053:     """
3054:     Audit all Python files in a project for signature mismatches
3055:
3056:     Args:
3057:         project_root: Root directory of the project
3058:         output_path: Optional path to save audit report
3059:
3060:     Returns:
3061:         List of detected signature mismatches
3062:     """
3063:     auditor = SignatureAuditor()
3064:
3065:     python_files = list(project_root.glob("*/*.py"))
3066:     logger.info(f"Auditing {len(python_files)} Python files")
3067:
3068:     # Define patterns to exclude
3069:     exclude_patterns = ['test', 'venv', '.venv', '__pycache__', '.git', 'build', 'dist']
3070:
3071:     all_mismatches = []
3072:     for py_file in python_files:
3073:         # Skip excluded patterns
3074:         if any(py_file.match(f'*/{pattern}/*') or py_file.match(f'*/{pattern}'))
```

```
3075:         for pattern in exclude_patterns):
3076:             continue
3077:
3078:             mismatches = auditor.audit_module(py_file)
3079:             all_mismatches.extend(mismatches)
3080:
3081:     if output_path:
3082:         auditor.export_report(output_path)
3083:
3084:     logger.info(f"Audit complete: {len(all_mismatches)} mismatches detected")
3085:
3086:     return all_mismatches
3087:
3088: # =====
3089: # CLI INTERFACE
3090: # =====
3091:
3092: # Note: Main entry point removed to maintain I/O boundary separation.
3093: # For CLI usage, see examples/ directory or create a dedicated CLI script.
3094:
3095:
3096:
3097: =====
3098: FILE: src/farfan_pipeline/utils/spc_adapter.py
3099: =====
3100:
3101: """SPC to Orchestrator Adapter (Shim).
3102:
3103: This module is a shim for backward compatibility. The canonical implementation
3104: has been moved to 'farfan_core.utils.cpp_adapter' to align with the Canon Policy Package (CPP)
3105: terminology.
3106:
3107: Please use 'farfan_core.utils.cpp_adapter.CPPAdapter' instead.
3108: """
3109:
3110: from __future__ import annotations
3111:
3112: import warnings
3113: from farfan_pipeline.utils.cpp_adapter import (
3114:     CPPAdapter as SPCAdapter,
3115:     CPPAdapterError as SPCAdapterError,
3116:     adapt_cpp_to_orchestrator as adapt_spc_to_orchestrator
3117: )
3118: from farfan_pipeline.core.calibration.decorators import calibrated_method
3119:
3120: # Issue deprecation warning when module is imported
3121: warnings.warn(
3122:     "farfan_core.utils.spc_adapter is deprecated. Use farfan_core.utils.cpp_adapter instead.",
3123:     DeprecationWarning,
3124:     stacklevel=2
3125: )
3126:
3127: __all__ = [
3128:     'SPCAdapter',
3129:     'SPCAdapterError',
3130:     'adapt_spc_to_orchestrator',
```

```
3131: ]
3132:
3133:
3134:
3135: =====
3136: FILE: src/farfan_pipeline/utils/validation/__init__.py
3137: =====
3138:
3139: """Validation module for pre-execution checks and preconditions."""
3140:
3141: from farfan_pipeline.utils.validation.aggregation_models import (
3142:     AggregationWeights,
3143:     AreaAggregationConfig,
3144:     ClusterAggregationConfig,
3145:     DimensionAggregationConfig,
3146:     MacroAggregationConfig,
3147:     validate_dimension_config,
3148:     validate_weights,
3149: )
3150: from farfan_pipeline.utils.validation.architecture_validator import (
3151:     ArchitectureValidationResult,
3152:     validate_architecture,
3153:     write_validation_report,
3154: )
3155: from farfan_pipeline.utils.validation.golden_rule import GoldenRuleValidator, GoldenRuleViolation
3156: from farfan_pipeline.utils.validation.schema_validator import (
3157:     MonolithIntegrityReport,
3158:     MonolithSchemaValidator,
3159:     SchemaInitializationError,
3160:     validate_monolith_schema,
3161: )
3162:
3163: __all__ = [
3164:     "ArchitectureValidationResult",
3165:     "GoldenRuleValidator",
3166:     "GoldenRuleViolation",
3167:     "validate_architecture",
3168:     "write_validation_report",
3169:     "AggregationWeights",
3170:     "DimensionAggregationConfig",
3171:     "AreaAggregationConfig",
3172:     "ClusterAggregationConfig",
3173:     "MacroAggregationConfig",
3174:     "validate_weights",
3175:     "validate_dimension_config",
3176:     "MonolithSchemaValidator",
3177:     "MonolithIntegrityReport",
3178:     "SchemaInitializationError",
3179:     "validate_monolith_schema",
3180: ]
3181:
3182:
3183:
3184: =====
3185: FILE: src/farfan_pipeline/utils/validation/aggregation_models.py
3186: =====
```

```
3187:  
3188: """  
3189: Pydantic models for aggregation weight validation.  
3190:  
3191: This module provides strict type-safe validation for aggregation weights,  
3192: ensuring zero-tolerance for invalid values at ingestion time.  
3193: """  
3194:  
3195: from pydantic import BaseModel, ConfigDict, Field, field_validator, model_validator  
3196: from typing_extensions import Self  
3197: from farfan_pipeline.core.parameters import ParameterLoaderV2  
3198: from farfan_pipeline.core.calibration.decorators import calibrated_method  
3199:  
3200:  
3201: class AggregationWeights(BaseModel):  
3202:     """  
3203:         Validation model for aggregation weights.  
3204:  
3205:         Enforces:  
3206:             - All weights must be non-negative (>= 0)  
3207:             - All weights must be <= 1.0  
3208:             - Weights must sum to 1.0 (within tolerance)  
3209:     """  
3210:  
3211:     model_config = ConfigDict(frozen=True, extra='forbid')  
3212:  
3213:     weights: list[float] = Field(..., min_length=1, description="List of aggregation weights")  
3214:     tolerance: float = Field(default=1e-6, ge=0, description="Tolerance for sum validation")  
3215:  
3216:     @field_validator('weights')  
3217:     @classmethod  
3218:     def validate_non_negative(cls, v: list[float]) -> list[float]:  
3219:         """Ensure all weights are non-negative."""  
3220:         for i, weight in enumerate(v):  
3221:             if weight < 0:  
3222:                 raise ValueError(  
3223:                     f"Invalid aggregation weight at index {i}: {weight}. "  
3224:                     f"All weights must be non-negative (>= 0)."  
3225:                 )  
3226:             if weight > 1.0:  
3227:                 raise ValueError(  
3228:                     f"Invalid aggregation weight at index {i}: {weight}. "  
3229:                     f"All weights must be <= 1.0."  
3230:                 )  
3231:         return v  
3232:  
3233:     @model_validator(mode='after')  
3234:     @calibrated_method("farfan_core.utils.validation.aggregation_models.AggregationWeights.validate_sum")  
3235:     def validate_sum(self) -> Self:  
3236:         """Ensure weights sum to ParameterLoaderV2.get("farfan_core.utils.validation.aggregation_models.AggregationWeights.validate_sum", "auto_param_L48_33  
", 1.0) within tolerance."""  
3237:         weight_sum = sum(self.weights)  
3238:         expected_sum = ParameterLoaderV2.get("farfan_core.utils.validation.aggregation_models.AggregationWeights.validate_sum", "auto_param_L52_79", 1.0)  
3239:         diff = abs(weight_sum - expected_sum)  
3240:         if diff > self.tolerance:  
3241:             raise ValueError(
```

```
3242:             f"Weight sum validation failed: sum={weight_sum:.6f}, expected={expected_sum}. "
3243:             f"Difference {diff:.6f} exceeds tolerance {self.tolerance:.6f}.)"
3244:         )
3245:     return self
3246:
3247: class DimensionAggregationConfig(BaseModel):
3248:     """Configuration for dimension-level aggregation."""
3249:
3250:     model_config = ConfigDict(frozen=True, extra='forbid')
3251:
3252:     dimension_id: str = Field(..., pattern=r'^DIM\d{2}$')
3253:     area_id: str = Field(..., pattern=r'^PA\d{2}$')
3254:     weights: AggregationWeights | None = None
3255:     expected_question_count: int = Field(default=5, ge=1, le=10)
3256:     group_by_keys: list[str] = Field(default=['dimension', 'policy_area'], min_length=1)
3257:
3258:
3259: class AreaAggregationConfig(BaseModel):
3260:     """Configuration for area-level aggregation."""
3261:
3262:     model_config = ConfigDict(frozen=True, extra='forbid')
3263:
3264:     area_id: str = Field(..., pattern=r'^PA\d{2}$')
3265:     expected_dimension_count: int = Field(default=6, ge=1, le=10)
3266:     weights: AggregationWeights | None = None
3267:     group_by_keys: list[str] = Field(default=['area_id'], min_length=1)
3268:
3269:
3270: class ClusterAggregationConfig(BaseModel):
3271:     """Configuration for cluster-level aggregation."""
3272:
3273:     model_config = ConfigDict(frozen=True, extra='forbid')
3274:
3275:     cluster_id: str = Field(..., pattern=r'^CL\d{2}$')
3276:     policy_area_ids: list[str] = Field(..., min_length=1)
3277:     weights: AggregationWeights | None = None
3278:     group_by_keys: list[str] = Field(default=['cluster_id'], min_length=1)
3279:
3280:     @field_validator('policy_area_ids')
3281:     @classmethod
3282:     def validate_policy_areas(cls, v: list[str]) -> list[str]:
3283:         """Ensure all policy area IDs follow the correct pattern."""
3284:         for pa_id in v:
3285:             if len(pa_id) < 3 or not pa_id.startswith('PA') or not pa_id[2:].isdigit():
3286:                 raise ValueError(f"Invalid policy area ID: {pa_id}. Expected format: PA##")
3287:         return v
3288:
3289: class MacroAggregationConfig(BaseModel):
3290:     """Configuration for macro-level aggregation."""
3291:
3292:     model_config = ConfigDict(frozen=True, extra='forbid')
3293:
3294:     cluster_ids: list[str] = Field(..., min_length=1)
3295:     weights: AggregationWeights | None = None
3296:
3297:     @field_validator('cluster_ids')
```

```
3298:     @classmethod
3299:     def validate_clusters(cls, v: list[str]) -> list[str]:
3300:         """Ensure all cluster IDs follow the correct pattern."""
3301:         for cl_id in v:
3302:             if len(cl_id) < 3 or not cl_id.startswith('CL') or not cl_id[2:].isdigit():
3303:                 raise ValueError(f"Invalid cluster ID: {cl_id}. Expected format: CL##")
3304:         return v
3305:
3306:     def validate_weights(weights: list[float], tolerance: float = 1e-6) -> AggregationWeights:
3307:         """
3308:             Convenience function to validate a list of weights.
3309:
3310:             Args:
3311:                 weights: List of weights to validate
3312:                 tolerance: Tolerance for sum validation
3313:
3314:             Returns:
3315:                 Validated AggregationWeights instance
3316:
3317:             Raises:
3318:                 ValueError: If validation fails
3319:         """
3320:         return AggregationWeights(weights=weights, tolerance=tolerance)
3321:
3322:     def validate_dimension_config(
3323:         dimension_id: str,
3324:         area_id: str,
3325:         weights: list[float] | None = None,
3326:         expected_question_count: int = 5
3327:     ) -> DimensionAggregationConfig:
3328:         """
3329:             Validate dimension aggregation configuration.
3330:
3331:             Args:
3332:                 dimension_id: Dimension ID (e.g., "DIM01")
3333:                 area_id: Area ID (e.g., "PA01")
3334:                 weights: Optional list of weights
3335:                 expected_question_count: Expected number of questions
3336:
3337:             Returns:
3338:                 Validated configuration
3339:
3340:             Raises:
3341:                 ValueError: If validation fails
3342:         """
3343:         weight_model = None
3344:         if weights is not None:
3345:             weight_model = validate_weights(weights)
3346:
3347:         return DimensionAggregationConfig(
3348:             dimension_id=dimension_id,
3349:             area_id=area_id,
3350:             weights=weight_model,
3351:             expected_question_count=expected_question_count
3352:         )
3353:
```

```
3354:  
3355:  
3356: =====  
3357: FILE: src/farfan_pipeline/utils/validation/architecture_validator.py  
3358: =====  
3359:  
3360: """Architecture validation utilities for the municipal policy analysis system.  
3361:  
3362: This module provides helpers to enforce that the municipal policy  
3363: analysis architecture blueprint references real, implemented methods.  
3364: It parses the ``policy_analysis_architecture.json`` specification,  
3365: compares every referenced method against the inventoried codebase and  
3366: produces coverage reports per analytical dimension.  
3367: """  
3368:  
3369: from __future__ import annotations  
3370:  
3371: import ast  
3372: import json  
3373: import re  
3374: from collections.abc import Iterable, Mapping  
3375: from dataclasses import dataclass, field  
3376: from pathlib import Path  
3377: from farfan_pipeline.core.calibration.decorators import calibrated_method  
3378:  
3379: # Regular expression used to capture fully-qualified method references such as  
3380: # ``ClassName.method_name``.  
3381: METHOD_PATTERN = re.compile(r"^[A-Za-z_][A-Za-z0-9_]*\.[A-Za-z_][A-Za-z0-9_]*$")  
3382:  
3383: _EXTERNAL_REFERENCE = object()  
3384:  
3385: ALIAS_MAP: dict[str, object] = {  
3386:     # Performance analyzer exposes the functionality through a private helper.  
3387:     "PerformanceAnalyzer.analyze_loss_function": "PerformanceAnalyzer._calculate_loss_functions",  
3388:     # The municipal plan analyzer generates recommendations with a private helper.  
3389:     "PDET Municipal Plan Analyzer.generate_recommendations": "PDET Municipal Plan Analyzer._generate_recommendations",  
3390:     # Advanced DAG validation leverages TeoriaCambio utilities internally.  
3391:     "AdvancedDAGValidator.validacion_completa": "TeoriaCambio.validacion_completa",  
3392:     "AdvancedDAGValidator._validar_orden_causal": "TeoriaCambio._validar_orden_causal",  
3393:     "AdvancedDAGValidator._encontrar_caminos_completos": "TeoriaCambio._encontrar_caminos_completos",  
3394:     # External dependency references (networkx graphs).  
3395:     "nx.DiGraph": _EXTERNAL_REFERENCE,  
3396: }  
3397:  
3398: # Root directory of the repository (two levels above this file).  
3399: ROOT_DIR = Path(__file__).resolve().parent.parent  
3400:  
3401: @dataclass(frozen=True)  
3402: class ArchitectureValidationResult:  
3403:     """Container with the outcome of the architecture validation process."""  
3404:  
3405:     resolved_methods: set[str]  
3406:     missing_methods: Mapping[str, Mapping[str, list[str]]]  
3407:     coverage: float  
3408:     total_spec_methods: int  
3409:     total_available_methods: int
```

```
3410:     per_dimension: Mapping[str, Mapping[str, list[str]]]
3411:     global_methods: tuple[str, ...] = field(default_factory=tuple)
3412:
3413:     @calibrated_method("farfan_core.utils.validation.architecture_validator.ArchitectureValidationResult.to_dict")
3414:     def to_dict(self) -> dict[str, object]:
3415:         """Serialise the validation result into a JSON-compatible dict."""
3416:
3417:         return {
3418:             "coverage": self.coverage,
3419:             "total_spec_methods": self.total_spec_methods,
3420:             "total_available_methods": self.total_available_methods,
3421:             "resolved_methods": sorted(self.resolved_methods),
3422:             "missing_methods": {
3423:                 dimension: dict(question_map.items())
3424:                 for dimension, question_map in self.missing_methods.items()
3425:             },
3426:             "per_dimension": {
3427:                 dimension: dict(question_map.items())
3428:                 for dimension, question_map in self.per_dimension.items()
3429:             },
3430:             "global_methods": list(self.global_methods),
3431:         }
3432:
3433:     def load_architecture_spec(path: Path) -> dict[str, object]:
3434:         """Load the JSON architecture specification from `path`."""
3435:
3436:         with path.open("r", encoding="utf-8") as handle:
3437:             return json.load(handle)
3438:
3439:     def _extract_method_from_entry(entry: object) -> str | None:
3440:         """Return the method string encoded in `entry` if present."""
3441:
3442:         if isinstance(entry, str) and METHOD_PATTERN.match(entry):
3443:             return entry
3444:
3445:         if isinstance(entry, Mapping):
3446:             # Architecture steps are stored as {"Class.method": "description"}
3447:             for key in entry:
3448:                 if isinstance(key, str) and METHOD_PATTERN.match(key):
3449:                     return key
3450:
3451:             # Some entries are dictionaries using {"name": "method"}. These lack
3452:             # class information and therefore cannot be enforced reliably.
3453:             name = entry.get("name") if isinstance(entry.get("name"), str) else None
3454:             if name and METHOD_PATTERN.match(name):
3455:                 return name
3456:
3457:         return None
3458:
3459:     def _extract_methods_from_string(value: str) -> Iterable[str]:
3460:         """Extract additional method references embedded in textual descriptions."""
3461:
3462:         for candidate in re.findall(r"[A-Za-z_][A-Za-z0-9_]*\.[A-Za-z_][A-Za-z0-9_]*", value):
3463:             if METHOD_PATTERN.match(candidate):
3464:                 yield candidate
3465:
```

```
3466: def extract_architecture_methods(spec: Mapping[str, object]) -> tuple[dict[str, dict[str, list[str]]], list[str]]:
3467:     """Extract method sequences per dimension and global method references."""
3468:
3469:     policy_spec = spec.get("policy_analysis_architecture", {})
3470:     if not isinstance(policy_spec, Mapping):
3471:         raise ValueError("Malformed architecture specification: missing 'policy_analysis_architecture'.")
3472:
3473:     per_dimension: dict[str, dict[str, list[str]]] = {}
3474:     global_methods: list[str] = []
3475:
3476:     # --- Component level methods -----
3477:     orchestration = policy_spec.get("orchestration_flow", {})
3478:     if isinstance(orchestration, Mapping):
3479:         for component in orchestration.get("components", []):
3480:             if not isinstance(component, Mapping):
3481:                 continue
3482:             for entry in component.get("key_methods", []):
3483:                 method = _extract_method_from_entry(entry)
3484:                 if method:
3485:                     global_methods.append(method.strip())
3486:                 if isinstance(entry, Mapping):
3487:                     description = entry.get("description")
3488:                     if isinstance(description, str):
3489:                         global_methods.extend(list(_extract_methods_from_string(description)))
3490:
3491:     # --- Phase 0 (initialisation) -----
3492:     phase_zero = policy_spec.get("phase_0_inicializacion_y_carga", {})
3493:     if isinstance(phase_zero, Mapping):
3494:         for step in phase_zero.get("steps", []):
3495:             if not isinstance(step, Mapping):
3496:                 continue
3497:             for action in step.get("actions", []):
3498:                 method = _extract_method_from_entry(action)
3499:                 if method:
3500:                     global_methods.append(method.strip())
3501:                 if isinstance(action, Mapping):
3502:                     for value in action.values():
3503:                         if isinstance(value, str):
3504:                             global_methods.extend(list(_extract_methods_from_string(value)))
3505:
3506:     # --- Analytical dimensions -----
3507:     for dimension in policy_spec.get("dimensiones", []):
3508:         if not isinstance(dimension, Mapping):
3509:             continue
3510:         dim_id = str(dimension.get("id", "UNKNOWN"))
3511:         dimension_methods: dict[str, list[str]] = {}
3512:         for subdimension in dimension.get("subdimension", []):
3513:             if not isinstance(subdimension, Mapping):
3514:                 continue
3515:             pregunta_id = str(subdimension.get("pregunta", "UNKNOWN"))
3516:             methods: list[str] = []
3517:             for step in subdimension.get("cadena_metodos", []):
3518:                 method = _extract_method_from_entry(step)
3519:                 if method:
3520:                     methods.append(method.strip())
3521:             if isinstance(step, Mapping):
```

```
3522:             for value in step.values():
3523:                 if isinstance(value, str):
3524:                     methods.extend(list(_extract_methods_from_string(value)))
3525:             dimension_methods[question_id] = methods
3526:         if dimension_methods:
3527:             per_dimension[dim_id] = dimension_methods
3528:
3529: # --- Transversal modules -----
3530: transversal = policy_spec.get("modulos_transversales", {})
3531: if isinstance(transversal, Mapping):
3532:     metricas = transversal.get("metricas_rendimiento", {})
3533:     if isinstance(metricas, Mapping):
3534:         for component in metricas.get("componentes", []):
3535:             if not isinstance(component, Mapping):
3536:                 continue
3537:             method = _extract_method_from_entry(component)
3538:             if method:
3539:                 global_methods.append(method.strip())
3540:             description = component.get("descripcion") or component.get("description")
3541:             if isinstance(description, str):
3542:                 global_methods.extend(list(_extract_methods_from_string(description)))
3543:
3544:     return per_dimension, global_methods
3545:
3546: def load_method_inventory(path: Path) -> tuple[set[str], set[str]]:
3547:     """Load available class methods and module functions from the inventory."""
3548:
3549:     with path.open("r", encoding="utf-8") as handle:
3550:         inventory = json.load(handle)
3551:
3552:     available_methods: set[str] = set()
3553:     functions: set[str] = set()
3554:
3555:     candidate_files = inventory.get("files", {})
3556:     for file_name in candidate_files:
3557:         file_path = ROOT_DIR / file_name
3558:         if not file_path.exists():
3559:             continue
3560:         try:
3561:             tree = ast.parse(file_path.read_text(encoding="utf-8"))
3562:         except SyntaxError:
3563:             continue
3564:
3565:         for node in tree.body:
3566:             if isinstance(node, (ast.FunctionDef, ast.AsyncFunctionDef)):
3567:                 functions.add(node.name)
3568:             elif isinstance(node, ast.ClassDef):
3569:                 for item in node.body:
3570:                     if isinstance(item, (ast.FunctionDef, ast.AsyncFunctionDef)):
3571:                         available_methods.add(f"{node.name}.{item.name}")
3572:
3573:     return available_methods, functions
3574:
3575: def _resolve_method_reference(
3576:     reference: str,
3577:     available_methods: set[str],
```

```
3578:     available_functions: set[str],  
3579: ) -> str | None:  
3580:     """Resolve a method reference using the available inventory."""  
3581:  
3582:     reference = reference.strip()  
3583:  
3584:     alias_target = ALIAS_MAP.get(reference)  
3585:     if alias_target is _EXTERNAL_REFERENCE:  
3586:         return reference  
3587:     if isinstance(alias_target, str):  
3588:         if alias_target == reference:  
3589:             return reference if reference in available_methods else None  
3590:         resolved_alias = _resolve_method_reference(alias_target, available_methods, available_functions)  
3591:         if resolved_alias:  
3592:             return resolved_alias  
3593:  
3594:     if METHOD_PATTERN.match(reference):  
3595:         if reference in available_methods:  
3596:             return reference  
3597:         # Allow ``Class.init`` aliases for ``Class.__init__``  
3598:         if reference.endswith(".init"):  
3599:             init_alias = reference[:-4] + "__init__"  
3600:             if init_alias in available_methods:  
3601:                 return init_alias  
3602:             return None  
3603:  
3604:     # Plain function reference  
3605:     if reference in available_functions:  
3606:         return reference  
3607:     return None  
3608:  
3609: def validate_architecture(spec_path: Path, inventory_path: Path) -> ArchitectureValidationResult:  
3610:     """Validate that every method described in the architecture exists."""  
3611:  
3612:     spec = load_architecture_spec(spec_path)  
3613:     per_dimension, global_methods = extract_architecture_methods(spec)  
3614:  
3615:     available_methods, available_functions = load_method_inventory(inventory_path)  
3616:  
3617:     resolved_methods: set[str] = set()  
3618:     missing_methods: dict[str, dict[str, list[str]]] = {}  
3619:  
3620:     # Validate global references  
3621:     for method in global_methods:  
3622:         resolved = _resolve_method_reference(method, available_methods, available_functions)  
3623:         if resolved:  
3624:             resolved_methods.add(resolved)  
3625:         else:  
3626:             missing_methods.setdefault("__global__", {}).setdefault("__global__", []).append(method)  
3627:  
3628:     # Validate per-dimension references  
3629:     for dimension, question_map in per_dimension.items():  
3630:         for question, methods in question_map.items():  
3631:             for method in methods:  
3632:                 resolved = _resolve_method_reference(method, available_methods, available_functions)  
3633:                 if resolved:
```

```
3634:             resolved_methods.add(resolved)
3635:         else:
3636:             missing_methods.setdefault(dimension, {}).setdefault(question, []).append(method)
3637:
3638:     total_references = len(global_methods)
3639:     total_references += sum(len(methods) for question_map in per_dimension.values() for methods in question_map.values())
3640:     total_references = max(total_references, 1)
3641:
3642:     coverage = len(resolved_methods) / total_references
3643:
3644:     return ArchitectureValidationResult(
3645:         resolved_methods=resolved_methods,
3646:         missing_methods=missing_methods,
3647:         coverage=coverage,
3648:         total_spec_methods=total_references,
3649:         total_available_methods=len(available_methods),
3650:         per_dimension=per_dimension,
3651:         global_methods=tuple(global_methods),
3652:     )
3653:
3654: def write_validation_report(result: ArchitectureValidationResult, output_path: Path) -> None:
3655:     """Write the validation report to `output_path` in JSON format."""
3656:
3657:     output_path.parent.mkdir(parents=True, exist_ok=True)
3658:     with output_path.open("w", encoding="utf-8") as handle:
3659:         json.dump(result.to_dict(), handle, indent=2, ensure_ascii=False)
3660:
3661: def main() -> None:
3662:     """Entry point for CLI usage."""
3663:
3664:     spec_path = ROOT_DIR / "policy_analysis_architecture.json"
3665:     inventory_path = ROOT_DIR / "COMPLETE_METHOD_CLASS_MAP.json"
3666:     report_path = ROOT_DIR / "validation" / "architecture_validation_report.json"
3667:
3668:     result = validate_architecture(spec_path, inventory_path)
3669:     write_validation_report(result, report_path)
3670:
3671:     if result.missing_methods:
3672:         missing_total = sum(len(methods) for dimension in result.missing_methods.values() for methods in dimension.values())
3673:         print(
3674:             f"Architecture validation completed with {missing_total} missing methods. "
3675:             f"Report saved to {report_path}."
3676:         )
3677:     else:
3678:         print(f"Architecture validation successful. Report saved to {report_path}.")
3679:
3680: # Note: Main entry point removed to maintain I/O boundary separation.
3681: # For usage, call main() function from a script or see examples/ directory.
3682:
3683:
3684:
3685: =====
3686: FILE: src/farfan_pipeline/utils/validation/contract_logger.py
3687: =====
3688:
3689: #!/usr/bin/env python3
```

```
3690: """
3691: Contract Error Logger - Structured logging for contract validation errors.
3692:
3693: Provides a unified interface for logging contract violations in a machine-readable
3694: format conforming to schemas/contract_error_log.schema.json.
3695:
3696: Usage:
3697:     from farfan_pipeline.utils.validation.contract_logger import ContractErrorLogger
3698:
3699:     logger = ContractErrorLogger(module_name="scoring")
3700:     logger.log_contract_mismatch(
3701:         function="apply_scoring",
3702:         key="confidence",
3703:         needed="float",
3704:         got=evidence.get("confidence"),
3705:         file=__file__,
3706:         line=234,
3707:         remediation="Convert confidence to float between 0.0 and 1.0"
3708:     )
3709: """
3710:
3711: import json
3712: import sys
3713: import traceback
3714: import uuid
3715: from datetime import datetime, timezone
3716: from typing import Any
3717: from farfan_pipeline.core.calibration.decorators import calibrated_method
3718:
3719:
3720: class ContractErrorLogger:
3721:     """Structured logger for contract validation errors."""
3722:
3723:     # Standard error codes
3724:     ERR_CONTRACT_MISMATCH = "ERR_CONTRACT_MISMATCH"
3725:     ERR_TYPE_VIOLATION = "ERR_TYPE_VIOLATION"
3726:     ERR_SCHEMA_VALIDATION = "ERR_SCHEMA_VALIDATION"
3727:     ERR_MISSING_REQUIRED_FIELD = "ERR_MISSING_REQUIRED_FIELD"
3728:     ERR_INVALID_MODALITY = "ERR_INVALID_MODALITY"
3729:     ERR_DETERMINISM_VIOLATION = "ERR_DETERMINISM_VIOLATION"
3730:
3731:     # Severity levels
3732:     CRITICAL = "CRITICAL"
3733:     ERROR = "ERROR"
3734:     WARNING = "WARNING"
3735:     INFO = "INFO"
3736:
3737:     def __init__(self, module_name: str, enable_stdout: bool = True) -> None:
3738:         """
3739:             Initialize contract error logger.
3740:
3741:             Args:
3742:                 module_name: Name of the module using this logger
3743:                 enable_stdout: Whether to output to stdout (default: True)
3744:             """
3745:             self.module_name = module_name
```

```
3746:         self.enable_stdout = enable_stdout
3747:         self.request_id = str(uuid.uuid4())
3748:
3749:     def _log(
3750:         self,
3751:         error_code: str,
3752:         function: str,
3753:         message: str,
3754:         severity: str,
3755:         context: dict,
3756:         remediation: str | None = None,
3757:         stack_trace: list[str] | None = None
3758:     ) -> None:
3759:         """
3760:             Internal method to log structured error.
3761:
3762:             Args:
3763:                 error_code: Standard error code
3764:                 function: Function name where error occurred
3765:                 message: Human-readable error message
3766:                 severity: Error severity level
3767:                 context: Structured context dictionary
3768:                 remediation: Optional remediation steps
3769:                 stack_trace: Optional stack trace
3770:
3771:             log_entry = {
3772:                 "error_code": error_code,
3773:                 "timestamp": datetime.now(timezone.utc).isoformat(),
3774:                 "severity": severity,
3775:                 "function": f"{self.module_name}.{function}",
3776:                 "message": message,
3777:                 "context": context,
3778:                 "request_id": self.request_id
3779:             }
3780:
3781:             if remediation:
3782:                 log_entry["remediation"] = remediation
3783:
3784:             if stack_trace:
3785:                 log_entry["stack_trace"] = stack_trace
3786:
3787:             # Output as single-line JSON
3788:             log_line = json.dumps(log_entry, separators=(',', ':'))
3789:
3790:             if self.enable_stdout:
3791:                 print(log_line, file=sys.stderr)
3792:
3793:     def log_contract_mismatch(
3794:         self,
3795:         function: str,
3796:         key: str,
3797:         needed: Any,
3798:         got: Any,
3799:         index: int | None = None,
3800:         file: str | None = None,
3801:         line: int | None = None,
```

```
3802:     remediation: str | None = None
3803: ) -> None:
3804: """
3805:     Log a contract mismatch error (ERR_CONTRACT_MISMATCH).
3806:
3807: Args:
3808:     function: Function name where error occurred
3809:     key: Parameter/field name that failed
3810:     needed: Expected type/value
3811:     got: Actual value received
3812:     index: Optional index in collection
3813:     file: Optional source file
3814:     line: Optional line number
3815:     remediation: Optional remediation steps
3816: """
3817: context = {
3818:     "key": key,
3819:     "needed": needed,
3820:     "got": got
3821: }
3822:
3823: if index is not None:
3824:     context["index"] = index
3825: if file:
3826:     context["file"] = file
3827: if line:
3828:     context["line"] = line
3829:
3830: message = f"Contract violation: required parameter '{key}' is missing or invalid"
3831: if got is None:
3832:     message = f"Contract violation: required parameter '{key}' is missing"
3833:
3834: self._log(
3835:     error_code=self.ERR_CONTRACT_MISMATCH,
3836:     function=function,
3837:     message=message,
3838:     severity=self.ERROR,
3839:     context=context,
3840:     remediation=remediation
3841: )
3842:
3843: def log_typeViolation(
3844:     self,
3845:     function: str,
3846:     key: str,
3847:     expected_type: str,
3848:     got: Any,
3849:     file: str | None = None,
3850:     line: int | None = None,
3851:     remediation: str | None = None
3852: ) -> None:
3853: """
3854:     Log a type violation error (ERR_TYPE_VIOLATION).
3855:
3856: Args:
3857:     function: Function name where error occurred
```

```
3858:             key: Parameter/field name with wrong type
3859:             expected_type: Expected type name
3860:             got: Actual value received
3861:             file: Optional source file
3862:             line: Optional line number
3863:             remediation: Optional remediation steps
3864:         """
3865:         actual_type = type(got).__name__
3866:
3867:         context = {
3868:             "key": key,
3869:             "needed": expected_type,
3870:             "got": str(got) if got is not None else None
3871:         }
3872:
3873:         if file:
3874:             context["file"] = file
3875:         if line:
3876:             context["line"] = line
3877:
3878:         message = f"Type violation: expected {expected_type} for '{key}', got {actual_type}"
3879:
3880:         self._log(
3881:             error_code=self.ERR_TYPE_VIOLATION,
3882:             function=function,
3883:             message=message,
3884:             severity=self.ERROR,
3885:             context=context,
3886:             remediation=remediation
3887:         )
3888:
3889:     def log_invalid_modality(
3890:         self,
3891:         function: str,
3892:         modality: str,
3893:         allowed_modalities: list[str],
3894:         file: str | None = None,
3895:         line: int | None = None
3896:     ) -> None:
3897:         """
3898:             Log an invalid modality error (ERR_INVALID_MODALITY).
3899:
3900:         Args:
3901:             function: Function name where error occurred
3902:             modality: Invalid modality value
3903:             allowed_modalities: List of allowed modalities
3904:             file: Optional source file
3905:             line: Optional line number
3906:         """
3907:         context = {
3908:             "key": "modality",
3909:             "needed": "|".join(allowed_modalities),
3910:             "got": modality
3911:         }
3912:
3913:         if file:
```

```
3914:         context["file"] = file
3915:         if line:
3916:             context["line"] = line
3917:
3918:             message = f"Invalid modality: {modality} is not in allowed modalities"
3919:             remediation = f"Use one of the allowed modality types: {', '.join(allowed_modalities)}"
3920:
3921:             self._log(
3922:                 error_code=self.ERR_INVALID_MODALITY,
3923:                 function=function,
3924:                 message=message,
3925:                 severity=self.ERROR,
3926:                 context=context,
3927:                 remediation=remediation
3928:             )
3929:
3930:     def log_determinismViolation(
3931:         self,
3932:         function: str,
3933:         description: str,
3934:         expected_hash: str,
3935:         actual_hash: str,
3936:         file: str | None = None,
3937:         line: int | None = None
3938:     ) -> None:
3939:         """
3940:             Log a determinism violation (ERR_DETERMINISM_VIOLATION).
3941:
3942:             Args:
3943:                 function: Function name where error occurred
3944:                 description: Description of what failed determinism check
3945:                 expected_hash: Expected hash value
3946:                 actual_hash: Actual hash value
3947:                 file: Optional source file
3948:                 line: Optional line number
3949:         """
3950:
3951:         context = {
3952:             "key": "determinism_check",
3953:             "needed": expected_hash,
3954:             "got": actual_hash
3955:         }
3956:
3957:         if file:
3958:             context["file"] = file
3959:         if line:
3960:             context["line"] = line
3961:
3962:         message = f"Determinism violation: {description}"
3963:         remediation = "Check for non-deterministic operations (random, time, concurrency)"
3964:
3965:         # Include stack trace for determinism violations
3966:         stack_trace = traceback.format_stack()
3967:
3968:         self._log(
3969:             error_code=self.ERR_DETERMINISM_VIOLATION,
3970:             function=function,
```

```
3970:         message=message,
3971:         severity=self.CRITICAL,
3972:         context=context,
3973:         remediation=remediation,
3974:         stack_trace=stack_trace
3975:     )
3976:
3977: # Example usage
3978: # Note: Example usage removed to maintain I/O boundary separation.
3979: # For usage examples, see examples/ directory.
3980:
3981:
3982:
3983: =====
3984: FILE: src/farfan_pipeline/utils/validation/golden_rule.py
3985: =====
3986:
3987: """Golden Rule enforcement utilities."""
3988:
3989: from __future__ import annotations
3990:
3991: import hashlib
3992: from typing import TYPE_CHECKING
3993: from farfan_pipeline.core.calibration.decorators import calibrated_method
3994:
3995: if TYPE_CHECKING:
3996:     from collections.abc import Iterable
3997:
3998: class GoldenRuleViolation(Exception):
3999:     """Raised when a Golden Rule assertion is violated."""
4000:
4001: class GoldenRuleValidator:
4002:     """Enforces the Golden Rules across orchestrated execution phases."""
4003:
4004:     def __init__(self, questionnaire_hash: str, step_catalog: Iterable[str]) -> None:
4005:         self._baseline_questionnaire_hash = questionnaire_hash
4006:         self._baseline_step_signature = self._hash_sequence(step_catalog)
4007:         self._baseline_step_catalog = list(step_catalog)
4008:         self._state_ids: set[int] = set()
4009:         self._predicate_signature: set[str] | None = None
4010:
4011:     @staticmethod
4012:     def _hash_sequence(sequence: Iterable[str]) -> str:
4013:         canonical = "|".join(str(item) for item in sequence)
4014:         return hashlib.sha256(canonical.encode("utf-8")).hexdigest()
4015:
4016:     def assert_immutable_metadata(
4017:         self,
4018:         questionnaire_hash: str,
4019:         step_catalog: Iterable[str]
4020:     ) -> None:
4021:         """Ensure canonical questionnaire and step catalog remain unchanged."""
4022:
4023:         if self._baseline_questionnaire_hash:
4024:             if questionnaire_hash and questionnaire_hash != self._baseline_questionnaire_hash:
4025:                 raise GoldenRuleViolation("Questionnaire metadata hash mismatch")
```

```
4026:  
4027:     if self._hash_sequence(step_catalog) != self._baseline_step_signature:  
4028:         raise GoldenRuleViolation("Execution step catalog mutated")  
4029:  
4030:     @calibrated_method("farfan_core.utils.validation.golden_rule.GoldenRuleValidator.reset_atomic_state")  
4031:     def reset_atomic_state(self) -> None:  
4032:         """Reset atomic state tracking between phases."""  
4033:  
4034:         self._state_ids.clear()  
4035:  
4036:     @calibrated_method("farfan_core.utils.validation.golden_rule.GoldenRuleValidator.assert_atomic_context")  
4037:     def assert_atomic_context(self, state_obj: object) -> None:  
4038:         """Ensure copy-on-write semantics for per-step state."""  
4039:  
4040:         obj_id = id(state_obj)  
4041:         if obj_id in self._state_ids:  
4042:             raise GoldenRuleViolation("State object reused across steps")  
4043:  
4044:         self._state_ids.add(obj_id)  
4045:  
4046:     @calibrated_method("farfan_core.utils.validation.golden_rule.GoldenRuleValidator.assert_deterministic_dag")  
4047:     def assert_deterministic_dag(self, step_ids: list[str]) -> None:  
4048:         """Validate deterministic ordering and absence of cycles."""  
4049:  
4050:         if len(step_ids) != len(set(step_ids)):  
4051:             raise GoldenRuleViolation("Duplicate step identifiers detected")  
4052:  
4053:         # Validate that step_ids is a subsequence of the canonical step catalog and in the same order  
4054:         canonical = self._baseline_step_catalog  
4055:         canonical_indices = {step_id: idx for idx, step_id in enumerate(canonical)}  
4056:         try:  
4057:             indices = [canonical_indices[step_id] for step_id in step_ids]  
4058:         except KeyError as e:  
4059:             raise GoldenRuleViolation(f"Step ID '{e.args[0]}' not found in canonical step catalog")  
4060:         if indices != sorted(indices):  
4061:             raise GoldenRuleViolation("Execution chain deviates from canonical order")  
4062:     @calibrated_method("farfan_core.utils.validation.golden_rule.GoldenRuleValidator.assert_homogeneous_treatment")  
4063:     def assert_homogeneous_treatment(self, predicate_set: Iterable[str]) -> None:  
4064:         """Ensure identical predicate set is applied across all questions."""  
4065:  
4066:         fingerprint = {str(item) for item in predicate_set}  
4067:  
4068:         if self._predicate_signature is None:  
4069:             self._predicate_signature = fingerprint  
4070:         return  
4071:  
4072:         if fingerprint != self._predicate_signature:  
4073:             raise GoldenRuleViolation("Predicate set mismatch detected")  
4074:  
4075:     @property  
4076:     @calibrated_method("farfan_core.utils.validation.golden_rule.GoldenRuleValidator.baseline_step_catalog")  
4077:     def baseline_step_catalog(self) -> list[str]:  
4078:         """Expose the baseline step catalog for downstream validation."""  
4079:  
4080:         return list(self._baseline_step_catalog)  
4081:
```

```
4082:
4083:
4084: =====
4085: FILE: src/farfan_pipeline/utils/validation/predicates.py
4086: =====
4087:
4088: #!/usr/bin/env python3
4089: """
4090: Validation Predicates - Precondition Checks for Execution
4091: =====
4092:
4093: Provides reusable predicates for validating preconditions before
4094: executing analysis steps.
4095:
4096: Author: Integration Team - Agent 3
4097: Version: 1.0.0
4098: Python: 3.10+
4099: """
4100:
4101: from dataclasses import dataclass
4102: from typing import Any
4103:
4104:
4105: @dataclass
4106: class ValidationResult:
4107:     """Result of a validation check."""
4108:
4109:     is_valid: bool
4110:     severity: str # ERROR, WARNING, INFO
4111:     message: str
4112:     context: dict[str, Any]
4113:
4114:
4115: class ValidationPredicates:
4116:     """
4117:     Collection of validation predicates for precondition checking.
4118:
4119:     These predicates verify that all required preconditions are met
4120:     before executing specific analysis steps.
4121:     """
4122:
4123:     @staticmethod
4124:     def verify_scoring_preconditions(
4125:         question_spec: dict[str, Any], execution_results: dict[str, Any], plan_text: str
4126:     ) -> ValidationResult:
4127:         """
4128:             Verify preconditions for TYPE_A scoring modality.
4129:
4130:             PRECONDITIONS for TYPE_A (Binary presence/absence):
4131:             - question_spec must have expected_elements list
4132:             - execution_results must be non-empty dict
4133:             - plan_text must be non-empty string
4134:
4135:             Args:
4136:                 question_spec: Question specification from rubric
4137:                 execution_results: Results from execution pipeline
```

```
4138:         plan_text: Full plan document text
4139:
4140:     Returns:
4141:         ValidationResult indicating if preconditions are met
4142:
4143:     """
4144:     errors = []
4145:
4146:     # Check question_spec has expected_elements
4147:     if not isinstance(question_spec, dict):
4148:         errors.append("question_spec must be a dictionary")
4149:     elif "expected_elements" not in question_spec:
4150:         errors.append("question_spec must have 'expected_elements' field")
4151:     elif not isinstance(question_spec.get("expected_elements"), list):
4152:         errors.append("expected_elements must be a list")
4153:     elif len(question_spec.get("expected_elements", [])) == 0:
4154:         errors.append("expected_elements cannot be empty")
4155:
4156:     # Check execution_results
4157:     if not isinstance(execution_results, dict):
4158:         errors.append("execution_results must be a dictionary")
4159:     elif len(execution_results) == 0:
4160:         errors.append("execution_results cannot be empty")
4161:
4162:     # Check plan_text
4163:     if not isinstance(plan_text, str):
4164:         errors.append("plan_text must be a string")
4165:     elif len(plan_text.strip()) == 0:
4166:         errors.append("plan_text cannot be empty")
4167:
4168:     if errors:
4169:         return ValidationResult(
4170:             is_valid=False,
4171:             severity="ERROR",
4172:             message="; ".join(errors),
4173:             context={
4174:                 "question_id": question_spec.get("id", "UNKNOWN"),
4175:                 "errors": errors,
4176:             },
4177:         )
4178:
4179:     return ValidationResult(
4180:         is_valid=True,
4181:         severity="INFO",
4182:         message="All scoring preconditions met",
4183:         context={
4184:             "question_id": question_spec.get("id"),
4185:             "expected_elements_count": len(
4186:                 question_spec.get("expected_elements", [])
4187:             ),
4188:             "execution_results_keys": list(execution_results.keys()),
4189:         },
4190:     )
4191:     @staticmethod
4192:     def _validate_expected_elements_types(
4193:         question_schema: list[Any] | dict[str, Any] | None,
```

```
4194:     chunk_schema: list[Any] | dict[str, Any] | None,
4195:     question_id: str,
4196: ) -> ValidationResult:
4197:     """
4198:         Validate that question_schema and chunk_schema have valid types.
4199:
4200:     Args:
4201:         question_schema: Schema object (None, list, or dict)
4202:         chunk_schema: Schema object (None, list, or dict)
4203:         question_id: Question identifier for error context
4204:
4205:     Returns:
4206:         ValidationResult indicating if schema types are valid
4207:
4208:     Raises:
4209:         TypeError: If either schema has invalid type (not None, list, or dict)
4210:         ValueError: If schemas are both lists with different lengths or both dicts with different keys
4211:     """
4212:     if question_schema is None and chunk_schema is None:
4213:         return ValidationResult(
4214:             is_valid=True,
4215:             severity="INFO",
4216:             message=f"Question {question_id} has valid schema types",
4217:             context={
4218:                 "question_id": question_id,
4219:                 "question_schema_type": "None",
4220:                 "chunk_schema_type": "None",
4221:             },
4222:         )
4223:
4224:     errors = []
4225:
4226:     if question_schema is not None and not isinstance(question_schema, list | dict):
4227:         errors.append("question_schema has invalid type (not None, list, or dict)")
4228:
4229:     if chunk_schema is not None and not isinstance(chunk_schema, list | dict):
4230:         errors.append("chunk_schema has invalid type (not None, list, or dict)")
4231:
4232:     if errors:
4233:         raise TypeError(
4234:             f"Question {question_id} has invalid schema types: {'; '.join(errors)}"
4235:         )
4236:
4237:     if (
4238:         isinstance(question_schema, list)
4239:         and isinstance(chunk_schema, list)
4240:         and len(question_schema) != len(chunk_schema)
4241:     ):
4242:         raise ValueError(
4243:             f"Question {question_id} schema length mismatch: "
4244:             f"question_schema has {len(question_schema)} elements, "
4245:             f"chunk_schema has {len(chunk_schema)} elements"
4246:         )
4247:
4248:     if isinstance(question_schema, dict) and isinstance(chunk_schema, dict):
4249:         question_keys = set(question_schema.keys())
```

```
4250:         chunk_keys = set(chunk_schema.keys())
4251:         key_diff = question_keys.symmetric_difference(chunk_keys)
4252:         if key_diff:
4253:             sorted_diff = sorted(key_diff)
4254:             raise ValueError(
4255:                 f"Question {question_id} schema key mismatch: "
4256:                 f"symmetric difference in keys: {sorted_diff}")
4257:         )
4258:
4259:         question_type = (
4260:             "None"
4261:             if question_schema is None
4262:             else "list" if isinstance(question_schema, list) else "dict"
4263:         )
4264:         chunk_type = (
4265:             "None"
4266:             if chunk_schema is None
4267:             else "list" if isinstance(chunk_schema, list) else "dict"
4268:         )
4269:
4270:         return ValidationResult(
4271:             is_valid=True,
4272:             severity="INFO",
4273:             message=f"Question {question_id} has valid schema types",
4274:             context={
4275:                 "question_id": question_id,
4276:                 "question_schema_type": question_type,
4277:                 "chunk_schema_type": chunk_type,
4278:             },
4279:         )
4280:
4281:     @staticmethod
4282:     def _validate_element_compatibility(
4283:         question_elements: list[dict[str, Any]],
4284:         chunk_elements: list[dict[str, Any]],
4285:         question_id: str,
4286:     ) -> int:
4287:         """
4288:             Validate compatibility between question elements and chunk elements.
4289:
4290:             Validates type, required, and minimum fields for all element pairs.
4291:
4292:             Args:
4293:                 question_elements: List of element dictionaries from question schema
4294:                 chunk_elements: List of element dictionaries from chunk schema
4295:                 question_id: Question identifier for error context
4296:
4297:             Returns:
4298:                 Count of successfully validated element pairs
4299:
4300:             Raises:
4301:                 ValueError: If element validation fails (type mismatch, required field missing,
4302:                             or minimum threshold not met)
4303:             """
4304:             if not question_elements and not chunk_elements:
4305:                 return 0
```

```
4306:
4307:     if len(question_elements) != len(chunk_elements):
4308:         raise ValueError(
4309:             f"Question {question_id}: element count mismatch - "
4310:             f"question has {len(question_elements)} elements, "
4311:             f"chunk has {len(chunk_elements)} elements"
4312:         )
4313:
4314:     validated_count = 0
4315:
4316:     for idx, (q_elem, c_elem) in enumerate(
4317:         zip(question_elements, chunk_elements, strict=True)
4318:     ):
4319:         if not isinstance(q_elem, dict) or not isinstance(c_elem, dict):
4320:             raise ValueError(
4321:                 f"Question {question_id}: element at index {idx} must be a dictionary"
4322:             )
4323:
4324:         # Type validation
4325:         q_type = q_elem.get("type")
4326:         c_type = c_elem.get("type")
4327:         if q_type and c_type and q_type != c_type:
4328:             raise ValueError(
4329:                 f"Question {question_id}: element type mismatch at index {idx} - "
4330:                 f"question expects type '{q_type}', chunk has type '{c_type}'"
4331:             )
4332:
4333:         # Required field validation
4334:         q_required = q_elem.get("required", False)
4335:         c_required = c_elem.get("required", False)
4336:         if q_required and not c_required:
4337:             raise ValueError(
4338:                 f"Question {question_id}: required field mismatch at index {idx} - "
4339:                 f"question requires element but chunk does not"
4340:             )
4341:
4342:         # Minimum field validation
4343:         q_minimum = q_elem.get("minimum", 0)
4344:         c_minimum = c_elem.get("minimum", 0)
4345:         if c_minimum < q_minimum:
4346:             element_type = q_elem.get("type", "unknown")
4347:             raise ValueError(
4348:                 f"Question {question_id}: minimum threshold not met at index {idx} - "
4349:                 f"element type '{element_type}' requires minimum {q_minimum}, "
4350:                 f"but chunk has minimum {c_minimum}"
4351:             )
4352:
4353:             validated_count += 1
4354:
4355:     return validated_count
4356:
4357:     @staticmethod
4358:     def verify_expected_elements(
4359:         question_spec: dict[str, Any], cuestionario_data: dict[str, Any] | None = None
4360:     ) -> ValidationResult:
4361:         """
```

```
4362:     Verify that expected_elements are defined correctly.
4363:
4364:     Args:
4365:         question_spec: Question specification from rubric
4366:         cuestionario_data: Full cuestionario metadata
4367:
4368:     Returns:
4369:         ValidationResult indicating if expected_elements are valid
4370:         """
4371:         question_id = question_spec.get("id", "UNKNOWN")
4372:
4373:         expected_elements = question_spec.get("expected_elements")
4374:         question_schema_raw = expected_elements
4375:         chunk_schema_raw = question_spec.get("chunk_schema")
4376:
4377:         if question_schema_raw is None:
4378:             question_type = "None"
4379:         elif isinstance(question_schema_raw, list):
4380:             question_type = "list"
4381:         elif isinstance(question_schema_raw, dict):
4382:             question_type = "dict"
4383:         else:
4384:             question_type = "invalid"
4385:
4386:         if chunk_schema_raw is None:
4387:             chunk_type = "None"
4388:         elif isinstance(chunk_schema_raw, list):
4389:             chunk_type = "list"
4390:         elif isinstance(chunk_schema_raw, dict):
4391:             chunk_type = "dict"
4392:         else:
4393:             chunk_type = "invalid"
4394:
4395:         schema_types = (question_type, chunk_type)
4396:
4397:         ValidationPredicates._validate_expected_elements_types(
4398:             question_schema_raw, chunk_schema_raw, question_id
4399:         )
4400:
4401:         if question_type == "None":
4402:             return ValidationResult(
4403:                 is_valid=False,
4404:                 severity="WARNING",
4405:                 message=f"Question {question_id} has no expected_elements defined",
4406:                 context={"question_id": question_id, "schema_types": schema_types},
4407:             )
4408:
4409:         if (
4410:             question_type == "list"
4411:             and expected_elements is not None
4412:             and len(expected_elements) == 0
4413:         ):
4414:             return ValidationResult(
4415:                 is_valid=False,
4416:                 severity="WARNING",
4417:                 message=f"Question {question_id} has empty expected_elements",
```

```
4418:             context={"question_id": question_id, "schema_types": schema_types},
4419:         )
4420:
4421:     # Validate element compatibility if both schemas are lists of dicts
4422:     validated_element_count = 0
4423:     if (
4424:         question_type == "list"
4425:         and chunk_type == "list"
4426:         and expected_elements
4427:         and all(isinstance(e, dict) for e in expected_elements)
4428:     ):
4429:         chunk_elements = (
4430:             chunk_schema_raw if isinstance(chunk_schema_raw, list) else []
4431:         )
4432:         if chunk_elements and all(isinstance(e, dict) for e in chunk_elements):
4433:             try:
4434:                 validated_element_count = (
4435:                     ValidationPredicates._validate_element_compatibility(
4436:                         expected_elements, chunk_elements, question_id
4437:                     )
4438:                 )
4439:             except ValueError as e:
4440:                 return ValidationResult(
4441:                     is_valid=False,
4442:                     severity="ERROR",
4443:                     message=str(e),
4444:                     context={
4445:                         "question_id": question_id,
4446:                         "schema_types": schema_types,
4447:                     },
4448:                 )
4449:
4450:     element_count = None
4451:     if question_type == "list" and expected_elements is not None:
4452:         element_count = len(expected_elements)
4453:
4454:     return ValidationResult(
4455:         is_valid=True,
4456:         severity="INFO",
4457:         message=f"Question {question_id} has valid expected_elements",
4458:         context={
4459:             "question_id": question_id,
4460:             "expected_elements": expected_elements,
4461:             "count": element_count,
4462:             "schema_types": schema_types,
4463:             "validated_element_count": validated_element_count,
4464:         },
4465:     )
4466:
4467:     @staticmethod
4468:     def verify_execution_context(
4469:         question_id: str, policy_area: str, dimension: str
4470:     ) -> ValidationResult:
4471:         """
4472:             Verify execution context parameters are valid.
4473:
```

```
4474:     Args:
4475:         question_id: Canonical question ID (P#-D#-Q#)
4476:         policy_area: Policy area (P1-P10)
4477:         dimension: Dimension (D1-D6)
4478:
4479:     Returns:
4480:         ValidationResult indicating if context is valid
4481:     """
4482:     errors = []
4483:
4484:     # Validate question_id format
4485:     if not question_id or not isinstance(question_id, str):
4486:         errors.append("question_id must be a non-empty string")
4487:     elif not question_id.startswith("P"):
4488:         errors.append(f"question_id '{question_id}' must start with 'P'")
4489:
4490:     # Validate policy_area
4491:     if not policy_area or not isinstance(policy_area, str):
4492:         errors.append("policy_area must be a non-empty string")
4493:     elif not policy_area.startswith("P"):
4494:         errors.append(f"policy_area '{policy_area}' must start with 'P'")
4495:     else:
4496:         try:
4497:             area_num = int(policy_area[1:])
4498:             if not (1 <= area_num <= 10):
4499:                 errors.append(f"policy_area '{policy_area}' must be P1-P10")
4500:             except ValueError:
4501:                 errors.append(f"Invalid policy_area format: '{policy_area}'")
4502:
4503:     # Validate dimension
4504:     if not dimension or not isinstance(dimension, str):
4505:         errors.append("dimension must be a non-empty string")
4506:     elif not dimension.startswith("D"):
4507:         errors.append(f"dimension '{dimension}' must start with 'D'")
4508:     else:
4509:         try:
4510:             dim_num = int(dimension[1:])
4511:             if not (1 <= dim_num <= 6):
4512:                 errors.append(f"dimension '{dimension}' must be D1-D6")
4513:             except ValueError:
4514:                 errors.append(f"Invalid dimension format: '{dimension}'")
4515:
4516:     if errors:
4517:         return ValidationResult(
4518:             is_valid=False,
4519:             severity="ERROR",
4520:             message="; ".join(errors),
4521:             context={
4522:                 "question_id": question_id,
4523:                 "policy_area": policy_area,
4524:                 "dimension": dimension,
4525:                 "errors": errors,
4526:             },
4527:         )
4528:
4529:     return ValidationResult(
```

```
4530:             is_valid=True,
4531:             severity="INFO",
4532:             message="Execution context is valid",
4533:             context={
4534:                 "question_id": question_id,
4535:                 "policy_area": policy_area,
4536:                 "dimension": dimension,
4537:             },
4538:         )
4539:
4540:     @staticmethod
4541:     def verify_producer_availability(
4542:         producer_name: str, producers_dict: dict[str, Any]
4543:     ) -> ValidationResult:
4544:
4545:         """
4546:             Verify that a producer module is available and initialized.
4547:
4548:             Args:
4549:                 producer_name: Name of the producer (e.g., 'derek_beach')
4550:                 producers_dict: Dictionary of initialized producers
4551:
4552:             Returns:
4553:                 ValidationResult indicating if producer is available
4554:
4555:             if producer_name not in producers_dict:
4556:                 return ValidationResult(
4557:                     is_valid=False,
4558:                     severity="ERROR",
4559:                     message=f"Producer '{producer_name}' not found in initialized producers",
4560:                     context={
4561:                         "producer_name": producer_name,
4562:                         "available_producers": list(producers_dict.keys()),
4563:                     },
4564:                 )
4565:
4566:             producer = producers_dict[producer_name]
4567:
4568:             # Check if producer is initialized
4569:             if isinstance(producer, dict):
4570:                 status = producer.get("status")
4571:                 if status != "initialized":
4572:                     return ValidationResult(
4573:                         is_valid=False,
4574:                         severity="ERROR",
4575:                         message=f"Producer '{producer_name}' status is '{status}'",
4576:                         context={
4577:                             "producer_name": producer_name,
4578:                             "status": status,
4579:                             "error": producer.get("error"),
4580:                         },
4581:                     )
4582:
4583:             return ValidationResult(
4584:                 is_valid=True,
4585:                 severity="INFO",
4586:                 message=f"Producer '{producer_name}' is available and initialized",
```

```
4586:         context={"producer_name": producer_name, "status": "initialized"},  
4587:     )  
4588:  
4589:
```