

```
src/farfan_pipeline/orchestration/verification_manifest.py
```

```
"""
Verification Manifest Generation with Cryptographic Integrity

Generates verification manifests for pipeline executions with HMAC signatures
for tamper detection and comprehensive execution environment tracking.
"""

from __future__ import annotations

import hashlib
import hmac
import json
import logging
import os
import platform
import sys
from datetime import datetime
from typing import Any

logger = logging.getLogger(__name__)

# Manifest schema version
MANIFEST_VERSION = "1.0.0"

class VerificationManifest:
    """
    Builder for verification manifests with cryptographic integrity.

    Features:
    - JSON Schema validation
    - HMAC-SHA256 integrity signatures
    - Execution environment tracking
    - Determinism metadata
    - Phase and artifact tracking
    """

    def __init__(self, hmac_secret: str | None = None) -> None:
        """
        Initialize manifest builder.

        Args:
            hmac_secret: Secret key for HMAC generation. If None, uses
                         environment variable VERIFICATION_HMAC_SECRET.
                         If not set, generates warning (integrity disabled).
        """
        self.hmac_secret = hmac_secret or os.getenv("VERIFICATION_HMAC_SECRET")
        if not self.hmac_secret:
            logger.warning(
                "No HMAC secret provided. Integrity verification disabled. "
                "Set VERIFICATION_HMAC_SECRET environment variable."
            )
"""
```

```

self.manifest_data: dict[str, Any] = {
    "version": MANIFEST_VERSION,
    "timestamp": datetime.utcnow().isoformat() + "Z",
    "success": False, # Default to false, set explicitly
}

def set_success(self, success: bool):
    """Set overall pipeline success flag."""
    self.manifest_data["success"] = success
    return self

def set_pipeline_hash(self, pipeline_hash: str):
    """Set SHA256 hash of pipeline execution."""
    self.manifest_data["pipeline_hash"] = pipeline_hash
    return self

def set_environment(self):
    """
    Capture execution environment information.

    Automatically captures:
    - Python version
    - Platform (OS)
    - CPU count
    - Available memory (if psutil available)
    - UTC timestamp
    """
    env_data = {
        "python_version": sys.version,
        "platform": platform.platform(),
        "cpu_count": os.cpu_count() or 1,
        "timestamp_utc": datetime.utcnow().isoformat() + "Z",
    }

    # Try to get memory info
    try:
        import psutil
        mem = psutil.virtual_memory()
        env_data["memory_gb"] = round(mem.total / (1024**3), 2)
    except ImportError:
        logger.debug("psutil not available, skipping memory info")
    except Exception as e:
        logger.debug(f"Failed to get memory info: {e}")

    self.manifest_data["environment"] = env_data
    return self

def add_environment_info(self, environment: dict[str, Any] | None = None):
    """
    Merge extra environment attributes into the manifest.

    Args:
        environment: Optional mapping of additional environment data.
    """

```

```

"""
if environment:
    current = self.manifest_data.get("environment", {})
    current.update(environment)
    self.manifest_data["environment"] = current
elif "environment" not in self.manifest_data:
    self.set_environment()
return self

def set_determinism(
    self,
    seed_version: str,
    base_seed: int | None = None,
    policy_unit_id: str | None = None,
    correlation_id: str | None = None,
    seeds_by_component: dict[str, int] | None = None
):
    """
    Set determinism tracking information.

    Args:
        seed_version: Seed derivation algorithm version
        base_seed: Base seed used
        policy_unit_id: Policy unit identifier
        correlation_id: Execution correlation ID
        seeds_by_component: Dict mapping component names to seeds
    """

    determinism_data = {
        "seed_version": seed_version
    }

    if base_seed is not None:
        determinism_data["base_seed"] = base_seed
    if policy_unit_id:
        determinism_data["policy_unit_id"] = policy_unit_id
    if correlation_id:
        determinism_data["correlation_id"] = correlation_id
    if seeds_by_component:
        determinism_data["seeds_by_component"] = seeds_by_component

    self.manifest_data["determinism"] = determinism_data
    return self

def set_determinism_info(self, determinism_info: dict[str, Any]):
    """Alias for setting determinism metadata directly."""
    if determinism_info:
        self.manifest_data["determinism"] = determinism_info
    return self

def set_calibrations(
    self,
    version: str,
    calibration_hash: str,
    methods_calibrated: int,

```

```

methods_missing: list[str]
):
"""
Set calibration information.

Args:
    version: Calibration registry version
    calibration_hash: SHA256 hash of calibration data
    methods_calibrated: Number of calibrated methods
    methods_missing: List of methods without calibration
"""

self.manifest_data["calibrations"] = {
    "version": version,
    "hash": calibration_hash,
    "methods_calibrated": methods_calibrated,
    "methods_missing": methods_missing
}
return self

def set_calibration_info(self, calibration_info: dict[str, Any]):
    """Set calibration metadata using a snapshot dictionary."""
    if calibration_info:
        self.manifest_data["calibration"] = calibration_info
    return self

def set_ingestion(
    self,
    method: str,
    chunk_count: int,
    text_length: int,
    sentence_count: int,
    chunk_strategy: str | None = None,
    chunk_overlap: int | None = None
):
    """
Set ingestion information.

Args:
    method: Ingestion method ("SPC" or "CPP")
    chunk_count: Number of chunks
    text_length: Total text length
    sentence_count: Number of sentences
    chunk_strategy: Chunking strategy used
    chunk_overlap: Chunk overlap in characters
"""

    ingestion_data = {
        "method": method,
        "chunk_count": chunk_count,
        "text_length": text_length,
        "sentence_count": sentence_count
    }

    if chunk_strategy:
        ingestion_data["chunk_strategy"] = chunk_strategy

```

```

    if chunk_overlap is not None:
        ingestion_data["chunk_overlap"] = chunk_overlap

    self.manifest_data["ingestion"] = ingestion_data
    return self

def set_spc_utilization(self, spc_utilization: dict[str, Any]):
    """
    Set SPC utilization metrics (Phase 2).

    Args:
        spc_utilization: Dictionary containing SPC metrics

    """
    if spc_utilization:
        self.manifest_data["spc_utilization"] = spc_utilization
    return self

def set_path_import_verification(self, report):
    """
    Set path and import verification results.

    Args:
        report: PolicyReport object from observability.path_import_policy

    Returns:
        self for chaining
    """
    # Use PolicyReport.to_dict() as canonical serialization
    self.manifest_data["path_import_verification"] = report.to_dict()
    return self

def set_parametrization(self, parametrization: dict[str, Any]):
    """
    Record executor/config parameterization data.

    Args:
        parametrization: dict[str, Any]
    """
    if parametrization:
        self.manifest_data["parametrization"] = parametrization
    return self

def add_phase(
    self,
    phase_id: int,
    phase_name: str,
    success: bool,
    duration_ms: float | None = None,
    items_processed: int | None = None,
    error: str | None = None
):
    """
    Add phase execution information.

    Args:
        phase_id: Phase numeric identifier
        phase_name: Phase human-readable name
        success: Phase execution success
    """

```

```

duration_ms: Duration in milliseconds
items_processed: Number of items processed
error: Error message if failed
"""
if "phases" not in self.manifest_data:
    self.manifest_data[ "phases" ] = [ ]

phase_data = {
    "phase_id": phase_id,
    "phase_name": phase_name,
    "success": success
}

if duration_ms is not None:
    phase_data[ "duration_ms" ] = duration_ms
if items_processed is not None:
    phase_data[ "items_processed" ] = items_processed
if error:
    phase_data[ "error" ] = error

container = self.manifest_data.get( "phases" )
if isinstance(container, dict):
    entries = container.setdefault("entries", [])
    entries.append(phase_data)
else:
    container.append(phase_data)
return self

def add_artifact(
    self,
    artifact_id: str,
    path: str,
    artifact_hash: str,
    size_bytes: int | None = None
):
    """
    Add artifact information.

    Args:
        artifact_id: Artifact identifier
        path: Artifact file path
        artifact_hash: SHA256 hash of artifact
        size_bytes: Artifact size in bytes
    """
    if "artifacts" not in self.manifest_data:
        self.manifest_data[ "artifacts" ] = {}

    artifact_data = {
        "path": path,
        "hash": artifact_hash
    }

    if size_bytes is not None:
        artifact_data[ "size_bytes" ] = size_bytes

```

```

self.manifest_data[ "artifacts" ][artifact_id] = artifact_data
return self

def _compute_hmac(self, content: str) -> str:
    """
    Compute HMAC-SHA256 of manifest content.

    Args:
        content: JSON string of manifest (without HMAC field)

    Returns:
        Hex-encoded HMAC signature
    """
    if not self.hmac_secret:
        return "00" * 32 # Placeholder if no secret

    signature = hmac.new(
        self.hmac_secret.encode("utf-8"),
        content.encode("utf-8"),
        hashlib.sha256
    )
    return signature.hexdigest()

def build(self) -> dict[str, Any]:
    """
    Build final manifest with HMAC signature.

    Returns:
        Complete manifest dictionary with integrity_hmac
    """
    # Create canonical JSON (without HMAC)
    canonical = json.dumps(
        self.manifest_data,
        sort_keys=True,
        indent=None,
        separators=( ' ', ' : ' )
    )

    # Compute HMAC
    hmac_signature = self._compute_hmac(canonical)

    # Add HMAC to manifest
    final_manifest = dict(self.manifest_data)
    final_manifest[ "integrity_hmac" ] = hmac_signature

    return final_manifest

def build_json(self, indent: int = 2) -> str:
    """
    Build manifest as JSON string.

    Args:
        indent: JSON indentation level
    """

```

```

>Returns:
    Pretty-printed JSON string
"""

manifest = self.build()
return json.dumps(manifest, indent=indent)

def write(self, filepath: str) -> None:
    """
    Write manifest to file.

    Args:
        filepath: Path to write manifest JSON
    """

    manifest_json = self.build_json()

    with open(filepath, 'w', encoding='utf-8') as f:
        f.write(manifest_json)

    logger.info(f"Verification manifest written to: {filepath}")

def verify_manifest_integrity(
    manifest: dict[str, Any],
    hmac_secret: str
) -> bool:
    """
    Verify HMAC integrity of a manifest.

    Args:
        manifest: Manifest dictionary (with integrity_hmac)
        hmac_secret: HMAC secret key

    Returns:
        True if HMAC is valid, False otherwise
    """

    if "integrity_hmac" not in manifest:
        logger.error("Manifest missing integrity_hmac field")
        return False

    # Extract HMAC
    provided_hmac = manifest["integrity_hmac"]

    # Rebuild manifest without HMAC
    manifest_without_hmac = {k: v for k, v in manifest.items() if k != "integrity_hmac"}

    # Compute canonical JSON
    canonical = json.dumps(
        manifest_without_hmac,
        sort_keys=True,
        indent=None,
        separators=(',', ',')
    )

```

```
# Compute expected HMAC
expected_hmac = hmac.new(
    hmac_secret.encode("utf-8"),
    canonical.encode("utf-8"),
    hashlib.sha256
).hexdigest()

# Constant-time comparison
is_valid = hmac.compare_digest(provided_hmac, expected_hmac)

if not is_valid:
    logger.error("HMAC verification failed - manifest may be tampered")

return is_valid
```

```

src/farfan_pipeline/orchestration/versions.py

"""
Version Tracking for Contract Enforcement

Centralized version management for all contract enforcement components.
Enables compatibility checking and rollback safety.
"""

# Pipeline version
PIPELINE_VERSION = "2.0.0"

# Calibration version (from calibration_registry.py)
CALIBRATION_VERSION = "2.0.0" # Should match calibration_registry.CALIBRATION_VERSION

# Signal version
SIGNAL_VERSION = "1.0.0"

# Advanced module version
ADVANCED_MODULE_VERSION = "1.0.0"

# Seed registry version
SEED_VERSION = "sha256_v1" # Should match seed_registry.SEED_VERSION

# Verification manifest version
MANIFEST_VERSION = "1.0.0" # Should match verification_manifest.MANIFEST_VERSION

# Minimum supported versions for backward compatibility
MIN_CALIBRATION_VERSION = "2.0.0"
MIN_SIGNAL_VERSION = "1.0.0"
MIN_PIPELINE_VERSION = "2.0.0"

def check_version_compatibility(component: str, version: str, min_version: str) -> bool:
    """
    Check if a version meets minimum requirements.

    Args:
        component: Component name (for error messages)
        version: Current version string (e.g., "2.1.0")
        min_version: Minimum required version (e.g., "2.0.0")

    Returns:
        True if version >= min_version

    Raises:
        ValueError: If version is incompatible
    """
    try:
        v_parts = [int(x) for x in version.split(".")]
        min_parts = [int(x) for x in min_version.split(".")]

        # Pad to same length
        while len(v_parts) < len(min_parts):
            v_parts.insert(0, 0)
    
```

```

        v_parts.append(0)
    while len(min_parts) < len(v_parts):
        min_parts.append(0)

    # Compare tuple
    if tuple(v_parts) < tuple(min_parts):
        raise ValueError(
            f"{component} version {version} is below minimum required {min_version}.

            "Please upgrade or regenerate calibration data."
        )

    return True
except (ValueError, AttributeError) as e:
    if "below minimum" in str(e):
        raise
    raise ValueError(
        f"Invalid version format for {component}: {version}. "
        "Expected semantic version like '1.0.0'"
    ) from e

def get_all_versions() -> dict[str, str]:
    """
    Get all component versions for manifest inclusion.

    Returns:
        Dictionary mapping component names to version strings
    """
    return {
        "pipeline": PIPELINE_VERSION,
        "calibration": CALIBRATION_VERSION,
        "signal": SIGNAL_VERSION,
        "advanced_module": ADVANCED_MODULE_VERSION,
        "seed": SEED_VERSION,
        "manifest": MANIFEST_VERSION,
    }

```

```
src/farfan_pipeline/orchestration/wiring/__init__.py
```

```
"""Wiring System - Fine-Grained Module Connection and Contract Validation.
```

```
This package implements the complete wiring architecture for SAAAAAA,  
providing deterministic initialization, contract validation, and observability  
for all module connections.
```

Architecture:

- Ports and adapters (hexagonal architecture)
- Dependency injection via constructors
- Feature flags for conditional wiring
- Contract validation between all links
- OpenTelemetry observability
- Deterministic initialization order

Key Modules:

- errors: Typed error classes for wiring failures
- contracts: Pydantic models for link contracts
- feature\_flags: Typed feature flags
- bootstrap: Deterministic initialization engine
- validation: Contract validation between links
- observability: Tracing and metrics
- analysis\_factory: Factory for analysis module dependency injection

```
"""
```

```
from orchestration.wiring.analysis_factory import (  
    create_analysis_components,  
    create_bayesian_confidence_calculator,  
    create_contradiction_detector,  
    create_document_loader,  
    create_municipal_analyzer,  
    create_municipal_ontology,  
    create_performance_analyzer,  
    create_semantic_analyzer,  
    create_temporal_logic_verifier,  
)
```

```
__all__ = [  
    'create_analysis_components',  
    'create_bayesian_confidence_calculator',  
    'create_contradiction_detector',  
    'create_document_loader',  
    'create_municipal_analyzer',  
    'create_municipal_ontology',  
    'create_performance_analyzer',  
    'create_semantic_analyzer',  
    'create_temporal_logic_verifier',  
]
```

```
src/farfan_pipeline/phases/Phase_eight/__init__.py
```

```
"""
```

```
Phase Eight: Strategic Integration
```

```
=====
```

```
Phase 8 of the F.A.R.F.A.N pipeline responsible for strategic-level integration  
and synthesis of processed policy documents.
```

```
"""
```

```
__all__ = [ ]
```

```
src/farfan_pipeline/phases/Phase_eight/recommendation_engine.py
```

```
# recommendation_engine.py - Rule-Based Recommendation Engine
```

```
"""
```

```
Recommendation Engine - Multi-Level Rule-Based Recommendations
```

```
=====
```

```
This module implements a rule-based recommendation engine that:
```

1. Loads and validates recommendation rules from JSON files
2. Evaluates conditions against score data at MICRO, MESO, and MACRO levels
3. Generates actionable recommendations with specific interventions
4. Renders templates with context-specific variable substitution

```
Supports three levels of recommendations:
```

- MICRO: Question-level recommendations (PA-DIM combinations)
- MESO: Cluster-level recommendations (CL01-CL04)
- MACRO: Plan-level strategic recommendations

```
Author: Integration Team
```

```
Version: 2.0.0
```

```
Python: 3.10+
```

```
"""
```

```
import logging
import re
from dataclasses import asdict, dataclass, field
from datetime import datetime, timezone
from pathlib import Path
from typing import Any

import jsonschema
from farfan_pipeline.core.parameters import ParameterLoaderV2
from cross_cutting_infrastructure.capaz_calibration_parametrization.calibration.decorators
import calibrated_method

logger = logging.getLogger(__name__)

_REQUIRED_ENHANCED_FEATURES = {
    "template_parameterization",
    "execution_logic",
    "measurable_indicators",
    "unambiguous_time_horizons",
    "testable_verification",
    "cost_tracking",
    "authority_mapping",
}
# =====
# DATA STRUCTURES FOR RECOMMENDATIONS
# =====

@dataclass
class Recommendation:
```

```

"""
Structured recommendation with full intervention details.

Supports both v1.0 (simple) and v2.0 (enhanced with 7 advanced features):
1. Template parameterization
2. Execution logic
3. Measurable indicators
4. Unambiguous time horizons
5. Testable verification
6. Cost tracking
7. Authority mapping
"""

rule_id: str
level: str # MICRO, MESO, or MACRO
problem: str
intervention: str
indicator: dict[str, Any]
responsible: dict[str, Any]
horizon: dict[str, Any] # Changed from Dict[str, str] to support enhanced fields
verification: list[Any] # Changed from List[str] to support structured verification
metadata: dict[str, Any] = field(default_factory=dict)

# Enhanced fields (v2.0) - optional for backward compatibility
execution: dict[str, Any] | None = None
budget: dict[str, Any] | None = None
template_id: str | None = None
template_params: dict[str, Any] | None = None

@calibrated_method("farfan_core.analysis.recommendation_engine.Recommendation.to_dict")
def to_dict(self) -> dict[str, Any]:
    """Convert to dictionary for JSON serialization"""
    result = asdict(self)
    # Remove None values for cleaner output
    return {k: v for k, v in result.items() if v is not None}

@dataclass
class RecommendationSet:
    """
Collection of recommendations with metadata
    """

    level: str
    recommendations: list[Recommendation]
    generated_at: str
    total_rules_evaluated: int
    rules_matched: int
    metadata: dict[str, Any] = field(default_factory=dict)

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationSet.to_dict")
def to_dict(self) -> dict[str, Any]:
    """Convert to dictionary for JSON serialization"""
    return {

```

```

        'level': self.level,
        'recommendations': [r.to_dict() for r in self.recommendations],
        'generated_at': self.generated_at,
        'total_rules_evaluated': self.total_rules_evaluated,
        'rules_matched': self.rules_matched,
        'metadata': self.metadata
    }

# =====
# RECOMMENDATION ENGINE
# =====

class RecommendationEngine:
    """
    Core recommendation engine that evaluates rules and generates recommendations.

    Uses canonical notation for dimension and policy area validation.
    """

    def __init__(
        self,
        rules_path: str = "config/recommendation_rules_enhanced.json",
        schema_path: str = "rules/recommendation_rules.schema.json",
        questionnaire_provider=None,
        orchestrator=None
    ) -> None:
        """
        Initialize recommendation engine

        Args:
            rules_path: Path to recommendation rules JSON file
            schema_path: Path to JSON schema for validation
            questionnaire_provider: QuestionnaireResourceProvider instance (injected via
DI)
            orchestrator: Orchestrator instance for accessing thresholds and patterns

        ARCHITECTURAL NOTE: Thresholds should come from questionnaire monolith
        via QuestionnaireResourceProvider, not from hardcoded values.
        """

        self.rules_path = Path(rules_path)
        self.schema_path = Path(schema_path)
        self.questionnaire_provider = questionnaire_provider
        self.orchestrator = orchestrator
        self.rules: dict[str, Any] = {}
        self.schema: dict[str, Any] = {}
        self.rules_by_level: dict[str, list[dict[str, Any]]] = {
            'MICRO': [],
            'MESO': [],
            'MACRO': []
        }

        # Load canonical notation for validation
        self._load_canonical_notation()

```

```

# Load rules and schema
self._load_schema()
self._load_rules()

logger.info(
    f"Recommendation engine initialized with "
    f"{len(self.rules_by_level['MICRO'])} MICRO, "
    f"{len(self.rules_by_level['MESO'])} MESO, "
    f"{len(self.rules_by_level['MACRO'])} MACRO rules"
)

```

@calibrated\_method("farfan\_core.analysis.recommendation\_engine.RecommendationEngine.\_load\_canonical\_notation")

```

def _load_canonical_notation(self) -> None:
    """Load canonical notation for validation"""
    try:
        from farfan_pipeline.core.canonical_notation import get_all_dimensions,
get_all_policy_areas
        self.canonical_dimensions = get_all_dimensions()
        self.canonical_policy_areas = get_all_policy_areas()
        logger.info(
            f"Canonical notation loaded: {len(self.canonical_dimensions)} dimensions, "
            f"{len(self.canonical_policy_areas)} policy areas"
        )
    except Exception as e:
        logger.warning(f"Could not load canonical notation: {e}")
        self.canonical_dimensions = {}
        self.canonical_policy_areas = {}

```

@calibrated\_method("farfan\_core.analysis.recommendation\_engine.RecommendationEngine.\_load\_schema")

```

def _load_schema(self) -> None:
    """Load JSON schema for rule validation"""
    # Delegate to factory for I/O operation
    from farfan_pipeline.analysis.factory import load_json

    try:
        self.schema = load_json(self.schema_path)
        logger.info(f"Loaded recommendation rules schema from {self.schema_path}")
    except Exception as e:
        logger.error(f"Failed to load schema: {e}")
        raise

```

@calibrated\_method("farfan\_core.analysis.recommendation\_engine.RecommendationEngine.\_load\_rules")

```

def _load_rules(self) -> None:
    """Load and validate recommendation rules"""
    # Delegate to factory for I/O operation
    from farfan_pipeline.analysis.factory import load_json

```

```

try:
    self.rules = load_json(self.rules_path)

    # Validate against schema
    jsonschema.validate(instance=self.rules, schema=self.schema)
    self._validate_ruleset_metadata()

    # Organize rules by level
    for rule in self.rules.get('rules', []):
        self._validate_rule(rule)
        level = rule.get('level')
        if level in self.rules_by_level:
            self.rules_by_level[level].append(rule)

    logger.info(f"Loaded and validated {len(self.rules.get('rules', []))} rules
from {self.rules_path}")
except jsonschema.ValidationError as e:
    logger.error(f"Rule validation failed: {e.message}")
    raise
except Exception as e:
    logger.error(f"Failed to load rules: {e}")
    raise

```

```

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine.reload_rules")
def reload_rules(self) -> None:
    """Reload rules from disk (useful for hot-reloading)"""
    self.rules_by_level = {'MICRO': [], 'MESO': [], 'MACRO': []}
    self._load_rules()

```

```

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine.get_thresholds_from_monolith")
def get_thresholds_from_monolith(self) -> dict[str, Any]:
    """
    Get scoring thresholds from questionnaire monolith.
    """

```

Returns:

Dictionary of thresholds by question\_id or default thresholds

ARCHITECTURAL NOTE: This method demonstrates proper access to questionnaire data via QuestionnaireResourceProvider, not direct I/O.

```

    """
if self.questionnaire_provider is None:
    logger.warning("No questionnaire provider attached, using default
thresholds")
    return {
        'default_micro_threshold': 2.0,
        'default_meso_threshold': 55.0,
        'default_macro_threshold': 65.0
    }

```

# Get questionnaire data via provider

```

questionnaire_data = self.questionnaire_provider.get_data()

# Extract thresholds from monolith structure
thresholds = {}
blocks = questionnaire_data.get('blocks', {})
micro_questions = blocks.get('micro_questions', [])

for question in micro_questions:
    question_id = question.get('question_id')
    scoring_info = question.get('scoring', {})
    threshold = scoring_info.get('threshold')

    if question_id and threshold is not None:
        thresholds[question_id] = threshold

logger.info(f"Loaded {len(thresholds)} thresholds from questionnaire monolith")
return thresholds

# =====
# MICRO LEVEL RECOMMENDATIONS
# =====

def generate_micro_recommendations(
    self,
    scores: dict[str, float],
    context: dict[str, Any] | None = None
) -> RecommendationSet:
    """
    Generate MICRO-level recommendations based on PA-DIM scores

    Args:
        scores: Dictionary mapping "PA##-DIM##" to scores
    (ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.
    get_thresholds_from_monolith", "auto_param_L279_63", 0.0)-3.0)
        context: Additional context for template rendering

    Returns:
        RecommendationSet with matched recommendations
    """
    recommendations = []
    rules_evaluated = 0

    for rule in self.rules_by_level['MICRO']:
        rules_evaluated += 1

        # Extract condition
        when = rule.get('when', {})
        pa_id = when.get('pa_id')
        dim_id = when.get('dim_id')
        score_lt = when.get('score_lt')

        # Build score key
        score_key = f"{pa_id}-{dim_id}"

```

```

# Check if condition matches
if score_key in scores and scores[score_key] < score_lt:
    # Render template
    template = rule.get('template', {})
    rendered = self._render_micro_template(template, pa_id, dim_id, context)

    # Create recommendation with enhanced fields (v2.0) if available
    rec = Recommendation(
        rule_id=rule.get('rule_id'),
        level='MICRO',
        problem=rendered['problem'],
        intervention=rendered['intervention'],
        indicator=rendered['indicator'],
        responsible=rendered['responsible'],
        horizon=rendered['horizon'],
        verification=rendered['verification'],
        metadata={
            'score_key': score_key,
            'actual_score': scores[score_key],
            'threshold': score_lt,
            'gap': score_lt - scores[score_key]
        },
        # Enhanced fields (v2.0)
        execution=rule.get('execution'),
        budget=rule.get('budget'),
        template_id=rendered.get('template_id'),
        template_params=rendered.get('template_params')
    )
    recommendations.append(rec)

return RecommendationSet(
    level='MICRO',
    recommendations=recommendations,
    generated_at=datetime.now(timezone.utc).isoformat(),
    total_rules_evaluated=rules_evaluated,
    rules_matched=len(recommendations)
)

def _render_micro_template(
    self,
    template: dict[str, Any],
    pa_id: str,
    dim_id: str,
    context: dict[str, Any] | None = None
) -> dict[str, Any]:
    """
    Render MICRO template with variable substitution

    Variables supported:
    - {{PAxx}}: Policy area (e.g., PA01)
    - {{DIMxx}}: Dimension (e.g., DIM01)
    - {{Q###}}: Question number (from context)
    """
    ctx = context or {}

```

```

substitutions = {
    'PAxx': pa_id,
    'DIMxx': dim_id,
    'pa_id': pa_id,
    'dim_id': dim_id,
}

question_hint = ctx.get('question_id')
template_params = template.get('template_params', {}) if isinstance(template, dict) else {}
if isinstance(template_params, dict):
    for key, value in template_params.items():
        if isinstance(value, str):
            substitutions.setdefault(key, value)
            substitutions.setdefault(key.upper(), value)
        if key == 'question_id':
            question_hint = value

if isinstance(question_hint, str):
    substitutions.setdefault(question_hint, question_hint)
    substitutions.setdefault('question_id', question_hint)
    substitutions.setdefault('Q001', question_hint)

for key, value in ctx.items():
    if isinstance(value, str):
        substitutions.setdefault(key, value)

return self._render_template(template, substitutions)

# =====
# MESO LEVEL RECOMMENDATIONS
# =====

def generate_meso_recommendations(
    self,
    cluster_data: dict[str, Any],
    context: dict[str, Any] | None = None
) -> RecommendationSet:
    """
    Generate MESO-level recommendations based on cluster performance
    """

    Args:
        cluster_data: Dictionary with cluster metrics:
        {
            'CL01': {'score': 75.0, 'variance': ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.get_thresholds_from_monolith", "auto_param_L398_56", 0.15), 'weak_pa': 'PA02'},
            'CL02': {'score': 62.0, 'variance': ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.get_thresholds_from_monolith", "auto_param_L399_56", 0.22), 'weak_pa': 'PA05'},
            ...
        }
        context: Additional context for template rendering
    """

```

```

>Returns:
    RecommendationSet with matched recommendations
"""

recommendations = []
rules_evaluated = 0

for rule in self.rules_by_level['MESO']:
    rules_evaluated += 1

    # Extract condition
    when = rule.get('when', {})
    cluster_id = when.get('cluster_id')
    score_band = when.get('score_band')
    variance_level = when.get('variance_level')
    variance_threshold = when.get('variance_threshold')
    weak_pa_id = when.get('weak_pa_id')

    # Get cluster data
    cluster = cluster_data.get(cluster_id, {})
    cluster_score = cluster.get('score', 0)
    cluster_variance = cluster.get('variance', 0)
    cluster_weak_pa = cluster.get('weak_pa')

    # Check conditions
    if not self._check_meso_conditions(
        cluster_score, cluster_variance, cluster_weak_pa,
        score_band, variance_level, variance_threshold, weak_pa_id
    ):
        continue

    # Render template
    template = rule.get('template', {})
    rendered = self._render_meso_template(template, cluster_id, context)

    # Create recommendation with enhanced fields (v2.0) if available
    rec = Recommendation(
        rule_id=rule.get('rule_id'),
        level='MESO',
        problem=rendered['problem'],
        intervention=rendered['intervention'],
        indicator=rendered['indicator'],
        responsible=rendered['responsible'],
        horizon=rendered['horizon'],
        verification=rendered['verification'],
        metadata={
            'cluster_id': cluster_id,
            'score': cluster_score,
            'score_band': score_band,
            'variance': cluster_variance,
            'variance_level': variance_level,
            'weak_pa': cluster_weak_pa
        },
        # Enhanced fields (v2.0)

```

```

        execution=rule.get('execution'),
        budget=rule.get('budget'),
        template_id=rendered.get('template_id'),
        template_params=rendered.get('template_params')
    )
    recommendations.append(rec)

    return RecommendationSet(
        level='MESO',
        recommendations=recommendations,
        generated_at=datetime.now(timezone.utc).isoformat(),
        total_rules_evaluated=rules_evaluated,
        rules_matched=len(recommendations)
    )
}

def _check_meso_conditions(
    self,
    score: float,
    variance: float,
    weak_pa: str | None,
    score_band: str,
    variance_level: str,
    variance_threshold: float | None,
    weak_pa_id: str | None
) -> bool:
    """Check if MESO conditions are met"""
    # Check score band
    if score_band == 'BAJO' and score >= 55 or score_band == 'MEDIO' and (score < 55
or score >= 75) or score_band == 'ALTO' and score < 75:
        return False

    # Check variance level
    if variance_level == 'BAJA' and variance >=
ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.g
et_thresholds_from_monolith", "auto_param_L488_52", 0.08) or variance_level == 'MEDIA'
and
                                         (variance <
ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.g
et_thresholds_from_monolith", "auto_param_L488_102", 0.08) or variance >=
ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.g
et_thresholds_from_monolith", "auto_param_L488_122", 0.18)):
        return False
    elif variance_level == 'ALTA':
        if variance_threshold and variance < variance_threshold / 100 or not
variance_threshold
                                         and
                                         variance <
ParameterLoaderV2.get("farfan_core.analysis.recommendation_engine.RecommendationEngine.g
et_thresholds_from_monolith", "auto_param_L491_115", 0.18):
        return False

    # Check weak PA if specified
    return not (weak_pa_id and weak_pa != weak_pa_id)

def _render_meso_template(
    self,
    template: dict[str, Any],

```

```

cluster_id: str,
context: dict[str, Any] | None = None
) -> dict[str, Any]:
    """Render MESO template with variable substitution"""

    substitutions = {
        'cluster_id': cluster_id,
    }

    if isinstance(template, dict):
        params = template.get('template_params', {})
        if isinstance(params, dict):
            for key, value in params.items():
                if isinstance(value, str):
                    substitutions.setdefault(key, value)
                    substitutions.setdefault(key.upper(), value)

    if context:
        for key, value in context.items():
            if isinstance(value, str):
                substitutions.setdefault(key, value)

    return self._render_template(template, substitutions)

# =====
# MACRO LEVEL RECOMMENDATIONS
# =====

def generate_macro_recommendations(
    self,
    macro_data: dict[str, Any],
    context: dict[str, Any] | None = None
) -> RecommendationSet:
    """
    Generate MACRO-level strategic recommendations
    """

    Args:
        macro_data: Dictionary with plan-level metrics:
        {
            'macro_band': 'SATISFACTORIO',
            'clusters_below_target': ['CL02', 'CL03'],
            'variance_alert': 'MODERADA',
            'priority_micro_gaps': ['PA01-DIM05', 'PA04-DIM04']
        }
        context: Additional context for template rendering

    Returns:
        RecommendationSet with matched recommendations
    """

    recommendations = []
    rules_evaluated = 0

    for rule in self.rules_by_level['MACRO']:
        rules_evaluated += 1

```

```

# Extract condition
when = rule.get('when', {})
macro_band = when.get('macro_band')
clusters_below = set(when.get('clusters_below_target', []))
variance_alert = when.get('variance_alert')
priority_gaps = set(when.get('priority_micro_gaps', []))

# Get macro data
actual_band = macro_data.get('macro_band')
actual_clusters = set(macro_data.get('clusters_below_target', []))
actual_variance = macro_data.get('variance_alert')
actual_gaps = set(macro_data.get('priority_micro_gaps', []))

# Check conditions
if macro_band and macro_band != actual_band:
    continue
if variance_alert and variance_alert != actual_variance:
    continue

# Check if clusters match (subset or exact match)
if clusters_below and not clusters_below.issubset(actual_clusters):
    # For MACRO, we want exact match or the rule's clusters to be present
    if clusters_below != actual_clusters and not
actual_clusters.issubset(clusters_below):
        continue

# Check if priority gaps match (subset)
if priority_gaps and not priority_gaps.issubset(actual_gaps):
    continue

# Render template
template = rule.get('template', {})
rendered = self._render_macro_template(template, context)

# Create recommendation with enhanced fields (v2.0) if available
rec = Recommendation(
    rule_id=rule.get('rule_id'),
    level='MACRO',
    problem=rendered['problem'],
    intervention=rendered['intervention'],
    indicator=rendered['indicator'],
    responsible=rendered['responsible'],
    horizon=rendered['horizon'],
    verification=rendered['verification'],
    metadata={
        'macro_band': actual_band,
        'clusters_below_target': list(actual_clusters),
        'variance_alert': actual_variance,
        'priority_micro_gaps': list(actual_gaps)
    },
    # Enhanced fields (v2.0)
    execution=rule.get('execution'),
    budget=rule.get('budget'),

```

```

        template_id=rendered.get('template_id'),
        template_params=rendered.get('template_params')
    )
    recommendations.append(rec)

    return RecommendationSet(
        level='MACRO',
        recommendations=recommendations,
        generated_at=datetime.now(timezone.utc).isoformat(),
        total_rules_evaluated=rules_evaluated,
        rules_matched=len(recommendations)
    )

def _render_macro_template(
    self,
    template: dict[str, Any],
    context: dict[str, Any] | None = None
) -> dict[str, Any]:
    """Render MACRO template with variable substitution"""

    substitutions = {}

    if context:
        for key, value in context.items():
            if isinstance(value, str):
                substitutions.setdefault(key, value)

    if isinstance(template, dict):
        params = template.get('template_params', {})
        if isinstance(params, dict):
            for key, value in params.items():
                if isinstance(value, str):
                    substitutions.setdefault(key, value)
                    substitutions.setdefault(key.upper(), value)

    return self._render_template(template, substitutions)

# =====
# UTILITY METHODS
# =====

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._substitute_variables")
def _substitute_variables(self, text: str, substitutions: dict[str, str]) -> str:
    """
    Substitute variables in text using {{variable}} syntax
    """

    Args:
        text: Text with variables
        substitutions: Dictionary of variable_name -> value

    Returns:
        Text with variables substituted

```

```

"""
result = text
for var, value in substitutions.items():
    pattern = r'\{\' + re.escape(var) + r'\}\}'
    result = re.sub(pattern, value, result)
return result

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._render_template")
def _render_template(self, template: dict[str, Any], substitutions: dict[str, str]) -> dict[str, Any]:
    """Recursively render a template applying substitutions to nested structures."""

    def render_value(value: Any) -> Any:
        if isinstance(value, str):
            return self._substitute_variables(value, substitutions)
        if isinstance(value, list):
            return [render_value(item) for item in value]
        if isinstance(value, dict):
            return {k: render_value(v) for k, v in value.items()}
        return value

    return render_value(template)

# =====
# VALIDATION UTILITIES
# =====

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._validate_rule")
def _validate_rule(self, rule: dict[str, Any]) -> None:
    """Apply structural validation to guarantee rigorous recommendations."""
    rule_id = rule.get('rule_id')
    if not isinstance(rule_id, str) or not rule_id.strip():
        raise ValueError("Recommendation rule missing rule_id")

    level = rule.get('level')
    if level not in self.rules_by_level:
        raise ValueError(f"Rule {rule_id} declares unsupported level: {level}")

    when = rule.get('when', {})
    if not isinstance(when, dict):
        raise ValueError(f"Rule {rule_id} has invalid 'when' definition")

    if level == 'MICRO':
        self._validate_micro_when(rule_id, when)
    elif level == 'MESO':
        self._validate_meso_when(rule_id, when)
    elif level == 'MACRO':
        self._validate_macro_when(rule_id, when)

    template = rule.get('template')

```

```

if not isinstance(template, dict):
    raise ValueError(f"Rule {rule_id} lacks a structured template")

self._validate_template(rule_id, template, level)

execution = rule.get('execution')
if execution is None:
    raise ValueError(f"Rule {rule_id} is missing execution block required for
enhanced rules")
self._validate_execution(rule_id, execution)

budget = rule.get('budget')
if budget is None:
    raise ValueError(f"Rule {rule_id} is missing budget block required for
enhanced rules")
self._validate_budget(rule_id, budget)

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._val
ide_micro_when")
def _validate_micro_when(self, rule_id: str, when: dict[str, Any]) -> None:
    required_keys = ('pa_id', 'dim_id', 'score_lt')
    for key in required_keys:
        if key not in when:
            raise ValueError(f"Rule {rule_id} missing '{key}' in MICRO condition")

    pa_id = when['pa_id']
    dim_id = when['dim_id']
    if not isinstance(pa_id, str) or not pa_id.strip():
        raise ValueError(f"Rule {rule_id} has invalid pa_id")
    if not isinstance(dim_id, str) or not dim_id.strip():
        raise ValueError(f"Rule {rule_id} has invalid dim_id")

    score_lt = when['score_lt']
    if not self._is_number(score_lt):
        raise ValueError(f"Rule {rule_id} has non-numeric MICRO threshold")
    if not 0 <= float(score_lt) <= 3:
        raise ValueError(f"Rule {rule_id} MICRO threshold must be between 0 and 3")

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._val
ide_meso_when")
def _validate_meso_when(self, rule_id: str, when: dict[str, Any]) -> None:
    cluster_id = when.get('cluster_id')
    if not isinstance(cluster_id, str) or not cluster_id.strip():
        raise ValueError(f"Rule {rule_id} missing cluster_id for MESO condition")

    condition_counter = 0

    score_band = when.get('score_band')
    if score_band is not None:
        if score_band not in {'BAJO', 'MEDIO', 'ALTO'}:
            raise ValueError(f"Rule {rule_id} has invalid MESO score_band")
        condition_counter += 1

```

```

variance_level = when.get('variance_level')
if variance_level is not None:
    if variance_level not in {'BAJA', 'MEDIA', 'ALTA'}:
        raise ValueError(f"Rule {rule_id} has invalid MESO variance_level")
    condition_counter += 1

variance_threshold = when.get('variance_threshold')
if variance_threshold is not None and not self._is_number(variance_threshold):
    raise ValueError(f"Rule {rule_id} has non-numeric variance_threshold")

weak_pa_id = when.get('weak_pa_id')
if weak_pa_id is not None:
    if not isinstance(weak_pa_id, str) or not weak_pa_id.strip():
        raise ValueError(f"Rule {rule_id} has invalid weak_pa_id")
    condition_counter += 1

if condition_counter == 0:
    raise ValueError(
        f"Rule {rule_id} must specify at least one discriminant condition for
MESO"
    )

```

`@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._validate_macro_when")`

```

def _validate_macro_when(self, rule_id: str, when: dict[str, Any]) -> None:
    discriminants = 0

    macro_band = when.get('macro_band')
    if macro_band is not None:
        if not isinstance(macro_band, str) or not macro_band.strip():
            raise ValueError(f"Rule {rule_id} has invalid macro_band")
        discriminants += 1

    clusters = when.get('clusters_below_target')
    if clusters is not None:
        if not isinstance(clusters, list) or not clusters:
            raise ValueError(f"Rule {rule_id} must declare non-empty
clusters_below_target")
        if not all(isinstance(item, str) and item.strip() for item in clusters):
            raise ValueError(f"Rule {rule_id} has invalid cluster identifiers")
        discriminants += 1

    variance_alert = when.get('variance_alert')
    if variance_alert is not None:
        if not isinstance(variance_alert, str) or not variance_alert.strip():
            raise ValueError(f"Rule {rule_id} has invalid variance_alert")
        discriminants += 1

    priority_gaps = when.get('priority_micro_gaps')
    if priority_gaps is not None:
        if not isinstance(priority_gaps, list) or not priority_gaps:
            raise ValueError(f"Rule {rule_id} must declare non-empty
priority_gaps")

```

```

priority_micro_gaps")
        if not all(isinstance(item, str) and item.strip() for item in
priority_gaps):
            raise ValueError(f"Rule {rule_id} has invalid priority_micro_gaps
entries")
    discriminants += 1

    if discriminants == 0:
        raise ValueError(
            f"Rule {rule_id} must specify at least one MACRO discriminant condition"
        )

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._val
ideate_template")
def _validate_template(self, rule_id: str, template: dict[str, Any], level: str) ->
None:
    required_fields = ['problem', 'intervention', 'indicator', 'responsible',
'horizon', 'verification', 'template_id', 'template_params']
    for field in required_fields:
        if field not in template:
            raise ValueError(f"Rule {rule_id} template missing '{field}'")

    for text_field in ('problem', 'intervention'):
        value = template[text_field]
        if not isinstance(value, str):
            raise ValueError(f"Rule {rule_id} template field '{text_field}' must be
text")
        stripped = value.strip()
        if len(stripped) < 40 or len(stripped.split()) < 12:
            raise ValueError(
                f"Rule {rule_id} template field '{text_field}' lacks actionable
detail"
            )

    indicator = template['indicator']
    if not isinstance(indicator, dict):
        raise ValueError(f"Rule {rule_id} indicator must be an object")
    for key in ('name', 'target', 'unit'):
        if key not in indicator:
            raise ValueError(f"Rule {rule_id} indicator missing '{key}' field")

    if not isinstance(indicator['name'], str) or len(indicator['name'].strip()) < 5:
        raise ValueError(f"Rule {rule_id} indicator name too short")

    target = indicator['target']
    if not self._is_number(target):
        raise ValueError(f"Rule {rule_id} indicator target must be numeric")

    unit = indicator['unit']
    if not isinstance(unit, str) or not unit.strip():
        raise ValueError(f"Rule {rule_id} indicator unit missing or empty")

    acceptable_range = indicator.get('acceptable_range')

```

```

if acceptable_range is not None:
    if not isinstance(acceptable_range, list) or len(acceptable_range) != 2:
        raise ValueError(f"Rule {rule_id} acceptable_range must have two numeric
bounds")
    if not all(self._is_number(bound) for bound in acceptable_range):
        raise ValueError(f"Rule {rule_id} acceptable_range values must be
numeric")
    lower, upper = acceptable_range
    if float(lower) >= float(upper):
        raise ValueError(f"Rule {rule_id} acceptable_range lower bound must be <
upper bound")

template_id = template['template_id']
if not isinstance(template_id, str) or not template_id.strip():
    raise ValueError(f"Rule {rule_id} template_id must be a non-empty string")

template_params = template['template_params']
if not isinstance(template_params, dict):
    raise ValueError(f"Rule {rule_id} template_params must be an object")
allowed_param_keys = {'pa_id', 'dim_id', 'cluster_id', 'question_id'}
unknown_params = set(template_params) - allowed_param_keys
if unknown_params:
    raise ValueError(f"Rule {rule_id} template_params contains unsupported keys:
{sorted(unknown_params)}")

required_params: set[str] = set()
if level == 'MICRO':
    required_params = {'pa_id', 'dim_id', 'question_id'}
elif level == 'MESO':
    required_params = {'cluster_id'}

missing_params = required_params - set(template_params)
if missing_params:
    raise ValueError(
        f"Rule {rule_id} template_params missing required keys for {level}:
{sorted(missing_params)}"
    )

if level != 'MACRO' and not template_params:
    raise ValueError(f"Rule {rule_id} template_params cannot be empty for
{level} level")

responsible = template['responsible']
if not isinstance(responsible, dict):
    raise ValueError(f"Rule {rule_id} responsible must be an object")
for key in ('entity', 'role'):
    value = responsible.get(key)
    if not isinstance(value, str) or not value.strip():
        raise ValueError(f"Rule {rule_id} responsible missing '{key}'")

partners = responsible.get('partners')
if partners is None or not isinstance(partners, list) or not partners:
    raise ValueError(f"Rule {rule_id} responsible must enumerate partners")
    if any(not isinstance(partner, str) or not partner.strip() for partner in

```

```

partners):
    raise ValueError(f"Rule {rule_id} responsible partners must be non-empty
strings")

horizon = template['horizon']
if not isinstance(horizon, dict):
    raise ValueError(f"Rule {rule_id} horizon must be an object")
for key in ('start', 'end'):
    value = horizon.get(key)
    if not isinstance(value, str) or not value.strip():
        raise ValueError(f"Rule {rule_id} horizon missing '{key}'")

verification = template['verification']
if not isinstance(verification, list) or not verification:
    raise ValueError(f"Rule {rule_id} must define verification artifacts")
for artifact in verification:
    if not isinstance(artifact, dict):
        raise ValueError(
            f"Rule {rule_id} verification entries must be structured
dictionaries"
    )
    required_artifact_fields = (
        'id',
        'type',
        'artifact',
        'format',
        'approval_required',
        'approver',
        'due_date',
        'required_sections',
        'automated_check',
    )
    for key in required_artifact_fields:
        if key not in artifact:
            raise ValueError(
                f"Rule {rule_id} verification artifact missing required field
'{key}'"
            )
    # Special handling for boolean fields - they can be False
    if key in ('approval_required', 'automated_check'):
        if not isinstance(artifact[key], bool):
            raise ValueError(
                f"Rule {rule_id} verification artifact field '{key}' must be
a boolean"
            )
    # Special handling for required_sections - must be a list
    elif key == 'required_sections':
        if not isinstance(artifact[key], list) or not all(isinstance(s, str)
and s.strip() for s in artifact[key]):
            raise ValueError(
                f"Rule {rule_id} verification required_sections must be a
list of strings (may be empty)"
            )
    # For other non-boolean fields, check for empty values

```

```

        elif not artifact[key]:
            raise ValueError(
                f"Rule {rule_id} verification artifact field '{key}' cannot be
empty"
        )
    )

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._val
ide_execution")
def _validate_execution(self, rule_id: str, execution: dict[str, Any]) -> None:
    if not isinstance(execution, dict):
        raise ValueError(f"Rule {rule_id} execution block must be an object")

    required_keys = [
        'trigger_condition',
        'blocking',
        'auto_apply',
        'requires_approval',
        'approval_roles',
    ]
    missing = required_keys - execution.keys()
    if missing:
        raise ValueError(f"Rule {rule_id} execution block missing keys:
{sorted(missing)}")

        if not isinstance(execution['trigger_condition'], str) or not
execution['trigger_condition'].strip():
            raise ValueError(f"Rule {rule_id} execution trigger_condition must be a
non-empty string")
        for flag in ('blocking', 'auto_apply', 'requires_approval'):
            if not isinstance(execution[flag], bool):
                raise ValueError(f"Rule {rule_id} execution field '{flag}' must be
boolean")

    roles = execution['approval_roles']
    if not isinstance(roles, list) or not roles:
        raise ValueError(f"Rule {rule_id} execution approval_roles must be a
non-empty list")
    if any(not isinstance(role, str) or not role.strip() for role in roles):
        raise ValueError(f"Rule {rule_id} execution approval_roles must contain
non-empty strings")

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._val
ide_budget")
def _validate_budget(self, rule_id: str, budget: dict[str, Any]) -> None:
    if not isinstance(budget, dict):
        raise ValueError(f"Rule {rule_id} budget block must be an object")

    required_keys = {'estimated_cost_cop', 'cost_breakdown', 'funding_sources',
'fiscal_year'}
    missing = required_keys - budget.keys()
    if missing:
        raise ValueError(f"Rule {rule_id} budget block missing keys:

```

```

{sorted(missing)}")

    if not self._is_number(budget['estimated_cost_cop']):
        raise ValueError(f"Rule {rule_id} budget estimated_cost_cop must be numeric")

    cost_breakdown = budget['cost_breakdown']
    if not isinstance(cost_breakdown, dict) or not cost_breakdown:
        raise ValueError(f"Rule {rule_id} cost_breakdown must be a non-empty object")
    for key, value in cost_breakdown.items():
        if not isinstance(key, str) or not key.strip():
            raise ValueError(f"Rule {rule_id} cost_breakdown keys must be non-empty strings")
        if not self._is_number(value):
            raise ValueError(f"Rule {rule_id} cost_breakdown values must be numeric")

    funding_sources = budget['funding_sources']
    if not isinstance(funding_sources, list) or not funding_sources:
        raise ValueError(f"Rule {rule_id} funding_sources must be a non-empty list")
    for source in funding_sources:
        if not isinstance(source, dict):
            raise ValueError(f"Rule {rule_id} funding source entries must be objects")
        for key in ('source', 'amount', 'confirmed'):
            if key not in source:
                raise ValueError(f"Rule {rule_id} funding source missing '{key}'")
            if not isinstance(source['source'], str) or not source['source'].strip():
                raise ValueError(f"Rule {rule_id} funding source name must be a non-empty string")
            if not self._is_number(source['amount']):
                raise ValueError(f"Rule {rule_id} funding source amount must be numeric")
            if not isinstance(source['confirmed'], bool):
                raise ValueError(f"Rule {rule_id} funding source confirmed flag must be boolean")

    fiscal_year = budget['fiscal_year']
    if not isinstance(fiscal_year, int):
        raise ValueError(f"Rule {rule_id} fiscal_year must be an integer")

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._validate_ruleset_metadata")
def _validate_ruleset_metadata(self) -> None:
    version = self.rules.get('version')
    if not isinstance(version, str) or not version.startswith('2.0'):
        raise ValueError(
            "Enhanced recommendation engine requires ruleset version 2.0"
        )

    features = self.rules.get('enhanced_features')
    if not isinstance(features, list) or not features:

```

```

        raise ValueError("Enhanced recommendation engine requires enhanced_features
list")

    feature_set = {feature for feature in features if isinstance(feature, str)}
    missing = _REQUIRED_ENHANCED_FEATURES - feature_set
    if missing:
        raise ValueError(
            f"Enhanced recommendation rules missing required features:
{sorted(missing)}"
        )

    @staticmethod
    def _is_number(value: Any) -> bool:
        return isinstance(value, (int, float)) and not isinstance(value, bool)

    def generate_all_recommendations(
        self,
        micro_scores: dict[str, float],
        cluster_data: dict[str, Any],
        macro_data: dict[str, Any],
        context: dict[str, Any] | None = None
    ) -> dict[str, RecommendationSet]:
        """
        Generate recommendations at all three levels
        """

        Args:
            micro_scores: PA-DIM scores for MICRO recommendations
            cluster_data: Cluster metrics for MESO recommendations
            macro_data: Plan-level metrics for MACRO recommendations
            context: Additional context

        Returns:
            Dictionary with 'MICRO', 'MESO', and 'MACRO' recommendation sets
        """
        return {
            'MICRO': self.generate_micro_recommendations(micro_scores, context),
            'MESO': self.generate_meso_recommendations(cluster_data, context),
            'MACRO': self.generate_macro_recommendations(macro_data, context)
        }

    def export_recommendations(
        self,
        recommendations: dict[str, RecommendationSet],
        output_path: str,
        format: str = 'json'
    ) -> None:
        """
        Export recommendations to file
        """

        Args:
            recommendations: Dictionary of recommendation sets
            output_path: Path to output file
            format: Output format ('json' or 'markdown')
        """

```

```

# Delegate to factory for I/O operation
from farfan_pipeline.analysis.factory import save_json, write_text_file

if format == 'json':
    save_json(
        {level: rec_set.to_dict() for level, rec_set in
recommendations.items()},
        output_path
    )
elif format == 'markdown':
    write_text_file(
        self._format_as_markdown(recommendations),
        output_path
    )
else:
    raise ValueError(f"Unsupported format: {format}")

logger.info(f"Exported recommendations to {output_path} in {format} format")

@calibrated_method("farfan_core.analysis.recommendation_engine.RecommendationEngine._for
mat_as_markdown")
def _format_as_markdown(self, recommendations: dict[str, RecommendationSet]) -> str:
    """Format recommendations as Markdown"""
    lines = ["# Recomendaciones del Plan de Desarrollo\n"]

    for level in ['MICRO', 'MESO', 'MACRO']:
        rec_set = recommendations.get(level)
        if not rec_set:
            continue

        lines.append(f"\n## Nivel {level}\n")
        lines.append(f"**Generado:** {rec_set.generated_at}\n")
        lines.append(f"**Reglas evaluadas:** {rec_set.total_rules_evaluated}\n")
        lines.append(f"**Recomendaciones:** {rec_set.rules_matched}\n")

        for i, rec in enumerate(rec_set.recommendations, 1):
            lines.append(f"\n### {i}. {rec.rule_id}\n")
            lines.append(f"**Problema:** {rec.problem}\n")
            lines.append(f"**Intervención:** {rec.intervention}\n")
            lines.append("\n**Indicador:**\n")
            lines.append(f"- Nombre: {rec.indicator.get('name')}\n")
            lines.append(f"- Meta: {rec.indicator.get('target')}\n")
            lines.append(f"- Unidad: {rec.indicator.get('unit')}\n")
            lines.append(f"**Responsable:** {rec.responsible.get('entity')}\n")
            lines.append(f"**Role:** {rec.responsible.get('role')}\n")
            lines.append(f"**Socios:** {', '.join(rec.responsible.get('partners', []))}\n")
            lines.append(f"**Horizonte:** {rec.horizon.get('start')} ?\n{rec.horizon.get('end')}\n")
            lines.append("\n**Verificación:**\n")
            for v in rec.verification:
                if isinstance(v, dict):
                    descriptor = f"[{v.get('type', 'ARTIFACT')}] {v.get('artifact',"

```

```

'Sin artefacto')}"
    due = v.get('due_date')
    approver = v.get('approver')
    suffix_parts: list[str] = []
    if due:
        suffix_parts.append(f"entrega: {due}")
    if approver:
        suffix_parts.append(f"aprueba: {approver}")
    suffix = f" ({'; '.join(suffix_parts)})" if suffix_parts else ""
    lines.append(f"- {descriptor}{suffix}")
    sections = v.get('required_sections') or []
    if sections:
        lines.append(
            " - Secciones requeridas: " + ", ".join(str(section)
for section in sections)
        )
    else:
        lines.append(f"- {v}")
    lines.append("")

    return "\n".join(lines)

# =====
# CONVENIENCE FUNCTIONS
# =====

def load_recommendation_engine(
    rules_path: str = "config/recommendation_rules_enhanced.json",
    schema_path: str = "rules/recommendation_rules.schema.json"
) -> RecommendationEngine:
    """
    Convenience function to load recommendation engine

    Args:
        rules_path: Path to rules JSON
        schema_path: Path to schema JSON

    Returns:
        Initialized RecommendationEngine
    """
    return RecommendationEngine(rules_path=rules_path, schema_path=schema_path)

# Note: Main entry point removed to maintain I/O boundary separation.
# For usage examples, see examples/ directory.

```

```
src/farfan_pipeline/phases/Phase_eight/recommendation_engine_adapter.py
```

```
"""
```

```
Recommendation engine adapter for infrastructure layer.
```

```
Implements RecommendationEnginePort using the concrete RecommendationEngine from the analysis module. This adapter follows the Ports and Adapters pattern, allowing the orchestrator to depend on abstractions rather than concrete implementations.
```

```
Version: 1.0.0
```

```
"""
```

```
import logging
from pathlib import Path
from typing import Any
```

```
logger = logging.getLogger(__name__)
```

```
class RecommendationEngineAdapter:
```

```
    """Adapter implementing RecommendationEnginePort.
```

```
This adapter wraps the concrete RecommendationEngine from the analysis module, providing a clean boundary between the orchestrator (core) and analysis (domain).
```

```
"""
```

```
def __init__(
    self,
    rules_path: str | Path,
    schema_path: str | Path,
    questionnaire_provider: Any = None,
    orchestrator: Any = None,
) -> None:
    """Initialize recommendation engine adapter.
```

```
Args:
```

```
    rules_path: Path to recommendation rules JSON file
    schema_path: Path to JSON schema for validation
    questionnaire_provider: QuestionnaireResourceProvider instance
    orchestrator: Orchestrator instance for accessing thresholds (can be set later)
```

```
Raises:
```

```
    ImportError: If RecommendationEngine cannot be imported
    Exception: If engine initialization fails
```

```
"""
```

```
    self._rules_path = Path(rules_path)
    self._schema_path = Path(schema_path)
    self._questionnaire_provider = questionnaire_provider
    self._orchestrator = orchestrator
    self._engine: Any = None
```

```
    self._initialize_engine()
```

```

def set_orchestrator(self, orchestrator: Any) -> None:
    """Set orchestrator reference after construction.

    This handles circular dependency between orchestrator and recommendation engine.

    Args:
        orchestrator: Orchestrator instance
    """
    self._orchestrator = orchestrator
    if self._engine is not None:
        self._engine.orchestrator = orchestrator

def __initialize_engine(self) -> None:
    """Initialize the concrete RecommendationEngine.

    Raises:
        ImportError: If RecommendationEngine cannot be imported
        Exception: If engine initialization fails
    """
    try:
        from farfan_pipeline.analysis.recommendation_engine import (
            RecommendationEngine,
        )

        self._engine = RecommendationEngine(
            rules_path=str(self._rules_path),
            schema_path=str(self._schema_path),
            questionnaire_provider=self._questionnaire_provider,
            orchestrator=self._orchestrator,
        )
        logger.info(
            f"RecommendationEngine initialized via adapter: "
            f"{len(self._engine.rules_by_level.get('MICRO', []))} MICRO, "
            f"{len(self._engine.rules_by_level.get('MESO', []))} MESO, "
            f"{len(self._engine.rules_by_level.get('MACRO', []))} MACRO rules"
        )
    except ImportError as e:
        logger.error(f"Failed to import RecommendationEngine: {e}")
        raise ImportError(
            "RecommendationEngine not available. "
            "Ensure farfan_pipeline.analysis.recommendation_engine is installed."
        ) from e
    except Exception as e:
        logger.error(f"Failed to initialize RecommendationEngine: {e}")
        raise

def generate_all_recommendations(
    self,
    micro_scores: dict[str, float],
    cluster_data: dict[str, Any],
    macro_data: dict[str, Any],
    context: dict[str, Any] | None = None,
) -> dict[str, Any]:

```

```

    """Generate recommendations at all three levels.

    Delegates to the concrete RecommendationEngine implementation.

    Args:
        micro_scores: Dictionary mapping "PA##-DIM##" to scores
        cluster_data: Dictionary with cluster metrics
        macro_data: Dictionary with macro-level metrics
        context: Optional context for template rendering

    Returns:
        Dictionary mapping level to RecommendationSet

    Raises:
        RuntimeError: If engine is not initialized
    """
    if self._engine is None:
        raise RuntimeError("RecommendationEngine not initialized")

    return self._engine.generate_all_recommendations(
        micro_scores=micro_scores,
        cluster_data=cluster_data,
        macro_data=macro_data,
        context=context,
    )

def generate_micro_recommendations(
    self, scores: dict[str, float], context: dict[str, Any] | None = None
) -> Any:
    """Generate MICRO-level recommendations.

    Args:
        scores: Dictionary mapping "PA##-DIM##" to scores
        context: Optional context for template rendering

    Returns:
        RecommendationSet with MICRO recommendations

    Raises:
        RuntimeError: If engine is not initialized
    """
    if self._engine is None:
        raise RuntimeError("RecommendationEngine not initialized")

    return self._engine.generate_micro_recommendations(
        scores=scores, context=context
    )

def generate_meso_recommendations(
    self, cluster_data: dict[str, Any], context: dict[str, Any] | None = None
) -> Any:
    """Generate MESO-level recommendations.

    Args:

```

```
cluster_data: Dictionary with cluster metrics
context: Optional context for template rendering

>Returns:
    RecommendationSet with MESO recommendations

>Raises:
    RuntimeError: If engine is not initialized
"""

if self._engine is None:
    raise RuntimeError("RecommendationEngine not initialized")

return self._engine.generate_meso_recommendations(
    cluster_data=cluster_data, context=context
)

def generate_macro_recommendations(
    self, macro_data: dict[str, Any], context: dict[str, Any] | None = None
) -> Any:
    """Generate MACRO-level recommendations.

>Args:
    macro_data: Dictionary with macro-level metrics
    context: Optional context for template rendering

>Returns:
    RecommendationSet with MACRO recommendations

>Raises:
    RuntimeError: If engine is not initialized
"""

if self._engine is None:
    raise RuntimeError("RecommendationEngine not initialized")

return self._engine.generate_macro_recommendations(
    macro_data=macro_data, context=context
)

def reload_rules(self) -> None:
    """Reload recommendation rules from disk.

>Raises:
    RuntimeError: If engine is not initialized
"""

if self._engine is None:
    raise RuntimeError("RecommendationEngine not initialized")

self._engine.reload_rules()
logger.info("Recommendation rules reloaded via adapter")

def create_recommendation_engine_adapter(
    rules_path: str | Path,
    schema_path: str | Path,
```

```
questionnaire_provider: Any = None,
orchestrator: Any = None,
) -> RecommendationEngineAdapter:
    """Factory function to create RecommendationEngineAdapter.

    This is the primary entry point for creating recommendation engine instances
    in the infrastructure layer. It handles initialization and error handling.

    Args:
        rules_path: Path to recommendation rules JSON file
        schema_path: Path to JSON schema for validation
        questionnaire_provider: QuestionnaireResourceProvider instance
        orchestrator: Orchestrator instance for accessing thresholds

    Returns:
        RecommendationEngineAdapter instance

    Raises:
        ImportError: If RecommendationEngine cannot be imported
        Exception: If engine initialization fails

    Example:
        >>> from pathlib import Path
        >>> adapter = create_recommendation_engine_adapter(
        ...     rules_path=Path("config/recommendation_rules_enhanced.json"),
        ...     schema_path=Path("rules/recommendation_rules_enhanced.schema.json"),
        ...     questionnaire_provider=provider,
        ...     orchestrator=orch
        ... )
        """
        return RecommendationEngineAdapter(
            rules_path=rules_path,
            schema_path=schema_path,
            questionnaire_provider=questionnaire_provider,
            orchestrator=orchestrator,
        )

    __all__ = [
        "RecommendationEngineAdapter",
        "create_recommendation_engine_adapter",
    ]
```

```
src/farfan_pipeline/phases/Phase_eight/signal_enriched_recommendations.py
```

```
"""Phase 8 Signal-Enriched Recommendation Module
```

```
Extends Phase 8 recommendation engine with signal-based enhancements for context-aware, data-driven recommendations with enhanced precision and determinism.
```

Enhancement Value:

- Signal-based rule matching using questionnaire patterns
- Signal-driven priority scoring for interventions
- Pattern-based intervention template selection
- Enhanced recommendation provenance with signal metadata

Integration: Used by RecommendationEngine to enhance rule evaluation with signal intelligence.

Author: F.A.R.F.A.N Pipeline Team

Version: 1.0.0

```
"""
```

```
from __future__ import annotations

import logging
from typing import Any, TYPE_CHECKING

if TYPE_CHECKING:
    try:
        from cross_cutting_infrastructure.irrigation_using_signals.SISAS.signal_registry
    import (
        QuestionnaireSignalRegistry,
    )
    except ImportError:
        QuestionnaireSignalRegistry = Any # type: ignore

logger = logging.getLogger(__name__)

# Default fallback values (replaced with actual values in production)
DEFAULT_POLICY_AREA = "PA01"
DEFAULT_QUESTION_FORMAT = "{:03d}" # Format for question IDs

# Pattern/indicator thresholds for signal support
STRONG_PATTERN_THRESHOLD = 5 # Pattern count for strong support
STRONG_INDICATOR_THRESHOLD = 3 # Indicator count for strong support

# Priority scoring thresholds
CRITICAL_SCORE_THRESHOLD = 0.3 # Scores below this are critical
LOW_SCORE_THRESHOLD = 0.5 # Scores below this are low
CRITICAL_PRIORITY_BOOST = 0.3 # Priority boost for critical scores
LOW_PRIORITY_BOOST = 0.2 # Priority boost for low scores
INSUFFICIENT_QUALITY_BOOST = 0.2 # Boost for insufficient quality
ACTIONABILITY_PATTERN_THRESHOLD = 10 # Pattern count for actionability
ACTIONABILITY_INDICATOR_THRESHOLD = 5 # Indicator count for actionability
ACTIONABILITY_BOOST = 0.15 # Boost for high actionability
```

```

__all__ = [
    "SignalEnrichedRecommender",
    "enhance_rule_matching",
    "prioritize_interventions",
]

class SignalEnrichedRecommender:
    """Signal-enriched recommender for Phase 8 with pattern-based matching.

    Enhances recommendation engine with signal intelligence for better
    rule matching, priority scoring, and intervention selection.

    Attributes:
        signal_registry: Optional signal registry for signal access
        enable_pattern_matching: Enable signal-based pattern matching
        enable_priority_scoring: Enable signal-driven priority scoring
    """

    def __init__(
        self,
        signal_registry: QuestionnaireSignalRegistry | None = None,
        enable_pattern_matching: bool = True,
        enable_priority_scoring: bool = True,
    ) -> None:
        """Initialize signal-enriched recommender.

        Args:
            signal_registry: Optional signal registry for signal access
            enable_pattern_matching: Enable pattern matching feature
            enable_priority_scoring: Enable priority scoring feature
        """
        self.signal_registry = signal_registry
        self.enable_pattern_matching = enable_pattern_matching
        self.enable_priority_scoring = enable_priority_scoring

        logger.info(
            f"SignalEnrichedRecommender initialized: "
            f"registry={'enabled' if signal_registry else 'disabled'}, "
            f"pattern_match={enable_pattern_matching}, "
            f"priority_score={enable_priority_scoring}"
        )

    def enhance_rule_condition(
        self,
        rule_id: str,
        condition: dict[str, Any],
        score_data: dict[str, Any],
    ) -> tuple[bool, dict[str, Any]]:
        """Enhance rule condition evaluation with signal-based pattern matching.

        Uses signal patterns to improve rule matching accuracy and provide
        additional context for condition evaluation.
        """

```

```

Args:
    rule_id: Rule identifier
    condition: Rule condition dict
    score_data: Score data to evaluate against

Returns:
    Tuple of (condition_met, evaluation_details)
"""

if not self.enable_pattern_matching:
    return False, {"enhancement": "disabled"}

evaluation_details: dict[str, Any] = {
    "rule_id": rule_id,
    "base_condition": condition,
    "enhancements": [],
}

condition_met = False

try:
    # Extract condition parameters
    field = condition.get("field", "")
    operator = condition.get("operator", "")
    threshold = condition.get("value", 0.0)

    # Get actual value from score_data
    actual_value = score_data.get(field, 0.0)

    # Basic evaluation
    if operator == "lt":
        base_met = actual_value < threshold
    elif operator == "lte":
        base_met = actual_value <= threshold
    elif operator == "gt":
        base_met = actual_value > threshold
    elif operator == "gte":
        base_met = actual_value >= threshold
    elif operator == "eq":
        base_met = actual_value == threshold
    else:
        base_met = False

    evaluation_details["base_evaluation"] = {
        "met": base_met,
        "actual_value": actual_value,
        "threshold": threshold,
        "operator": operator,
    }

    # Signal-based enhancement: Check if patterns support this condition
    if self.signal_registry and base_met:
        try:
            # Try to get related signals

```

```

        # In production, extract actual policy_area from score_data
        signal_pack = self.signal_registry.get_micro_answering_signals(
            f"Q{int(score_data.get('question_global', 1)):03d}"
        )

        # Check if condition field matches signal patterns
        pattern_count = len(signal_pack.patterns) if hasattr(signal_pack,
        'patterns') else 0
        indicator_count = len(signal_pack.indicators) if
hasattr(signal_pack, 'indicators') else 0

        # Add confidence boost if signals support the condition
        if pattern_count > 5: # Strong pattern support
            evaluation_details["enhancements"].append({
                "type": "pattern_support",
                "pattern_count": pattern_count,
                "confidence_boost": 0.1,
            })

        if indicator_count > 3: # Strong indicator support
            evaluation_details["enhancements"].append({
                "type": "indicator_support",
                "indicator_count": indicator_count,
                "confidence_boost": 0.05,
            })

    except Exception as e:
        logger.debug(f"Signal enhancement failed for rule {rule_id}: {e}")

    condition_met = base_met
    evaluation_details["final_result"] = condition_met

except Exception as e:
    logger.warning(f"Failed to enhance rule condition for {rule_id}: {e}")
    evaluation_details["error"] = str(e)
    condition_met = False

return condition_met, evaluation_details

def compute_intervention_priority(
    self,
    recommendation: dict[str, Any],
    score_data: dict[str, Any],
) -> tuple[float, dict[str, Any]]:
    """Compute intervention priority using signal-driven scoring.

    Analyzes recommendation and score data with signal intelligence to
    determine intervention priority for resource allocation.

```

Args:

recommendation: Recommendation dict with intervention details  
score\_data: Score data for context

Returns:

```

        Tuple of (priority_score, priority_details)
"""

if not self.enable_priority_scoring:
    return 0.5, {"priority": "disabled"}

priority_details: dict[str, Any] = {
    "factors": [],
    "base_priority": 0.5,
}

priority_score = 0.5 # Neutral base

try:
    # Factor 1: Score severity (lower scores = higher priority)
    actual_score = score_data.get("score", 0.5)
    if actual_score < 0.3:
        severity_boost = 0.3 # Critical priority
        priority_details["factors"].append({
            "type": "critical_score",
            "score": actual_score,
            "boost": severity_boost,
        })
        priority_score += severity_boost
    elif actual_score < 0.5:
        severity_boost = 0.2 # High priority
        priority_details["factors"].append({
            "type": "low_score",
            "score": actual_score,
            "boost": severity_boost,
        })
        priority_score += severity_boost

    # Factor 2: Quality level
    quality_level = score_data.get("quality_level", "")
    if quality_level == "INSUFICIENTE":
        quality_boost = 0.2
        priority_details["factors"].append({
            "type": "insufficient_quality",
            "quality": quality_level,
            "boost": quality_boost,
        })
        priority_score += quality_boost

    # Factor 3: Signal-based pattern density (more patterns = more actionable)
    if self.signal_registry:
        try:
            question_id = f"Q{int(score_data.get('question_global', 1)):03d}"
            signal_pack = self.signal_registry.get_micro_answering_signals(question_id)

            pattern_count = len(signal_pack.patterns) if hasattr(signal_pack,
            'patterns') else 0
            indicator_count = len(signal_pack.indicators) if
            hasattr(signal_pack, 'indicators') else 0
        except Exception as e:
            print(f"Error processing signal for question {question_id}: {e}")
            priority_score -= 0.1

```

```

        # High pattern/indicator density suggests actionability
        if pattern_count > 10 and indicator_count > 5:
            actionability_boost = 0.15
            priority_details["factors"].append({
                "type": "high_actionability",
                "pattern_count": pattern_count,
                "indicator_count": indicator_count,
                "boost": actionability_boost,
            })
            priority_score += actionability_boost

    except Exception as e:
        logger.debug(f"Signal-based priority scoring failed: {e}")

    # Clamp priority to [0.0, 1.0]
    priority_score = max(0.0, min(1.0, priority_score))
    priority_details["final_priority"] = priority_score

except Exception as e:
    logger.warning(f"Failed to compute intervention priority: {e}")
    priority_details["error"] = str(e)
    priority_score = 0.5

return priority_score, priority_details

```

```

def select_intervention_template(
    self,
    problem_type: str,
    score_data: dict[str, Any],
) -> tuple[str | None, dict[str, Any]]:
    """Select intervention template using signal-based pattern analysis.

```

Uses signal patterns to identify the most appropriate intervention template for the identified problem.

Args:

```

    problem_type: Type of problem identified
    score_data: Score data for context

```

Returns:

```

    Tuple of (template_id, selection_details)
"""

```

```

selection_details: dict[str, Any] = {
    "problem_type": problem_type,
    "candidates": [],
}

```

```

template_id = None

```

try:

```

    # Map problem types to template patterns
    template_mapping = {
        "insufficient_baseline": "intervention_baseline",

```

```

    "weak_causal_links": "intervention_causality",
    "missing_indicators": "intervention_measurement",
    "low_coverage": "intervention_coverage",
    "inconsistent_data": "intervention_quality",
}

# Base template selection
template_id = template_mapping.get(problem_type, "intervention_generic")
selection_details["base_template"] = template_id

# Signal-based refinement
if self.signal_registry:
    try:
        question_id = f"Q{int(score_data.get('question_global', 1)):03d}"
        signal_pack =
self.signal_registry.get_micro_answering_signals(question_id)

        # Analyze signal patterns to refine template choice
        patterns = signal_pack.patterns if hasattr(signal_pack, 'patterns')
else []

        # Check for specific pattern types that suggest template refinement
        temporal_patterns = [p for p in patterns if 'temporal' in
str(p).lower()]
        causal_patterns = [p for p in patterns if 'caus' in str(p).lower()
or 'efect' in str(p).lower()]

        if len(temporal_patterns) > 3:
            selection_details["candidates"].append({
                "template": "intervention_temporal",
                "reason": "high_temporal_pattern_density",
                "count": len(temporal_patterns),
            })
            # Refine template for temporal focus
            if problem_type in ["insufficient_baseline",
"missing_indicators"]:
                template_id = "intervention_temporal"

            if len(causal_patterns) > 3:
                selection_details["candidates"].append({
                    "template": "intervention_causality",
                    "reason": "high_causal_pattern_density",
                    "count": len(causal_patterns),
                })
                # Refine template for causal focus
                if problem_type in ["weak_causal_links", "inconsistent_data"]:
                    template_id = "intervention_causality"

            except Exception as e:
                logger.debug(f"Signal-based template selection failed: {e}")

            selection_details["selected_template"] = template_id

        except Exception as e:

```

```

        logger.warning(f"Failed to select intervention template: {e}")
        selection_details["error"] = str(e)
        template_id = "intervention_generic"

    return template_id, selection_details

def enrich_recommendation(
    self,
    recommendation: dict[str, Any],
    evaluation_details: dict[str, Any],
    priority_details: dict[str, Any],
    template_details: dict[str, Any],
) -> dict[str, Any]:
    """Enrich recommendation with signal-based metadata.

    Adds signal provenance and enhancement details to recommendation
    for full transparency and reproducibility.

    Args:
        recommendation: Base recommendation dict
        evaluation_details: Rule evaluation details
        priority_details: Priority scoring details
        template_details: Template selection details

    Returns:
        Enriched recommendation dict
    """

    enriched = {
        **recommendation,
        "signal_enrichment": {
            "enabled": True,
            "registry_available": self.signal_registry is not None,
            "evaluation": evaluation_details,
            "priority": priority_details,
            "template_selection": template_details,
        },
    }

    return enriched

```

```

def enhance_rule_matching(
    signal_registry: QuestionnaireSignalRegistry | None,
    rule_id: str,
    condition: dict[str, Any],
    score_data: dict[str, Any],
) -> tuple[bool, dict[str, Any]]:
    """Convenience function for signal-enhanced rule matching.

    Creates a temporary SignalEnrichedRecommender and evaluates condition.

    Args:
        signal_registry: Signal registry instance (optional)
        rule_id: Rule identifier

```

```
    condition: Rule condition to evaluate
    score_data: Score data for evaluation

>Returns:
    Tuple of (condition_met, evaluation_details)
"""

recommender = SignalEnrichedRecommender(signal_registry=signal_registry)
return recommender.enhance_rule_condition(
    rule_id=rule_id,
    condition=condition,
    score_data=score_data,
)
```

```
def prioritize_interventions(
    signal_registry: QuestionnaireSignalRegistry | None,
    recommendations: list[dict[str, Any]],
    score_data: dict[str, Any],
) -> list[tuple[dict[str, Any], float, dict[str, Any]]]:
    """Prioritize interventions using signal-driven scoring.
```

Args:

```
    signal_registry: Signal registry instance (optional)
    recommendations: List of recommendation dicts
    score_data: Score data for context
```

Returns:

```
    List of tuples: (recommendation, priority_score, priority_details)
"""

recommender = SignalEnrichedRecommender(signal_registry=signal_registry)

prioritized = []
for rec in recommendations:
    priority_score, priority_details = recommender.compute_intervention_priority(
        recommendation=rec,
        score_data=score_data,
    )
    prioritized.append((rec, priority_score, priority_details))

# Sort by priority score (descending)
prioritized.sort(key=lambda x: x[1], reverse=True)

return prioritized
```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/__init__.py
```

```
"""
```

```
Phase Four Through Seven: Advanced Processing
```

```
=====
```

```
Combined phases 4-7 of the F.A.R.F.A.N pipeline responsible for advanced  
processing and analysis of policy documents.
```

```
Phase 4-7: Aggregation Pipeline
```

```
-----
```

```
This module implements the complete aggregation pipeline for the policy analysis system:
```

- FASE 4: Dimension aggregation (60 dimensions:  $6 \times 10$  policy areas)
- FASE 5: Policy area aggregation (10 areas)
- FASE 6: Cluster aggregation (4 MESO questions)
- FASE 7: Macro evaluation (1 holistic question)

```
New additions:
```

- Choquet aggregator for multi-layer calibration with interaction terms

```
"""
```

```
from .aggregation import (  
    DimensionAggregator,  
    AreaPolicyAggregator,  
    ClusterAggregator,  
    AggregationSettings,  
    DimensionScore,  
    AreaScore,  
    ClusterScore,  
    MacroScore,  
)  
  
from .choquet_aggregator import (  
    ChoquetAggregator,  
    ChoquetConfig,  
    CalibrationResult,  
    CalibrationBreakdown,  
    CalibrationConfigError,  
)  
  
from .aggregation_enhancements import (  
    EnhancedDimensionAggregator,  
    EnhancedAreaAggregator,  
    EnhancedClusterAggregator,  
    EnhancedMacroAggregator,  
    ConfidenceInterval,  
    DispersionMetrics,  
    HermeticityDiagnosis,  
    StrategicAlignmentMetrics,  
    enhance_aggregator,  
)  
  
from .aggregation_validation import (
```

```
validate_phase4_output,
validate_phase5_output,
validate_phase6_output,
validate_phase7_output,
validate_full_aggregation_pipeline,
enforce_validation_or_fail,
ValidationResult,
AggregationValidationError,
```

```
)
```

```
__all__ = [
    # Existing aggregation
    "DimensionAggregator",
    "AreaPolicyAggregator",
    "ClusterAggregator",
    "AggregationSettings",
    "DimensionScore",
    "AreaScore",
    "ClusterScore",
    "MacroScore",
    # Choquet aggregator
    "ChoquetAggregator",
    "ChoquetConfig",
    "CalibrationResult",
    "CalibrationBreakdown",
    "CalibrationConfigError",
    # Enhancements and contracts
    "EnhancedDimensionAggregator",
    "EnhancedAreaAggregator",
    "EnhancedClusterAggregator",
    "EnhancedMacroAggregator",
    "ConfidenceInterval",
    "DispersionMetrics",
    "HermeticityDiagnosis",
    "StrategicAlignmentMetrics",
    "enhance_aggregator",
    # Validation
    "validate_phase4_output",
    "validate_phase5_output",
    "validate_phase6_output",
    "validate_phase7_output",
    "validate_full_aggregation_pipeline",
    "enforce_validation_or_fail",
    "ValidationResult",
    "AggregationValidationError",
```

```
]
```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/adaptive_meso_scoring.py
```

```
"""
```

## Adaptive Meso-Level Scoring Module

This module implements enhanced, adaptive scoring mechanisms for cluster (meso-level) aggregation that address the audit findings:

1. Adaptive penalty weights based on scenario characteristics (dispersion vs convergence)
2. Non-linear penalty scaling for extreme dispersion cases
3. Strengthened penalties for high-dispersion scenarios
4. Maintained mathematical correctness (100%)

### Mathematical Foundation:

```
-----
```

The adaptive scoring mechanism uses a context-sensitive penalty function that adjusts based on:

- Coefficient of Variation (CV): Relative dispersion measure
- Dispersion Index (DI): Normalized range measure
- Score distribution shape: Convergence vs dispersion pattern

### Formula:

```
adjusted_score = weighted_score × penalty_factor
```

Where `penalty_factor` is computed adaptively using:

```
penalty_factor = 1.0 - penalty_strength  
penalty_strength = base_penalty × sensitivity_multiplier × shape_factor
```

`base_penalty`: Derived from normalized standard deviation

`sensitivity_multiplier`: Adaptive [0.8-2.0] based on CV and DI

`shape_factor`: Non-linear scaling [1.0-2.5] for extreme cases

```
"""
```

```
from __future__ import annotations
```

```
import logging
```

```
from dataclasses import dataclass
```

```
from typing import Any
```

```
logger = logging.getLogger(__name__)
```

```
@dataclass
```

```
class AdaptiveScoringConfig:
```

```
    """Configuration for adaptive scoring behavior."""
```

```
    max_score: float = 3.0
```

```
    # Base penalty weight (replaces fixed PENALTY_WEIGHT=0.3)  
    base_penalty_weight: float = 0.35
```

```
    # Convergence thresholds (low dispersion)  
    convergence_cv_threshold: float = 0.15
```

```

convergence_di_threshold: float = 0.20

# Dispersion thresholds (high dispersion)
high_dispersion_cv_threshold: float = 0.40
extreme_dispersion_cv_threshold: float = 0.60

# Sensitivity multipliers
convergence_multiplier: float = 0.5 # Reduce penalty for convergence
moderate_multiplier: float = 1.0      # Normal penalty
high_dispersion_multiplier: float = 1.5 # Increase penalty
extreme_dispersion_multiplier: float = 2.0 # Strong penalty

# Non-linear scaling factors for extreme cases
extreme_shape_factor: float = 1.8 # Exponential scaling
bimodal_penalty_boost: float = 1.3 # Additional penalty for bimodal

@dataclass
class ScoringMetrics:
    """Computed scoring metrics for analysis."""

    mean: float
    variance: float
    std_dev: float
    coefficient_variation: float
    dispersion_index: float
    normalized_std: float
        scenario_type: str      # "convergence", "moderate", "high_dispersion",
"extreme_dispersion"
    shape_classification: str # "uniform", "clustered", "bimodal", "dispersed"

class AdaptiveMesoScoring:
    """
    Adaptive scoring mechanism for meso-level cluster aggregation.

    This class implements mathematically-sound, adaptive penalty mechanisms
    that are sensitive to different dispersion and convergence scenarios.
    """

    def __init__(self, config: AdaptiveScoringConfig | None = None) -> None:
        """
        Initialize adaptive scoring.

        Args:
            config: Optional configuration (uses defaults if not provided)
        """
        self.config = config or AdaptiveScoringConfig()

    def compute_metrics(self, scores: list[float]) -> ScoringMetrics:
        """
        Compute comprehensive scoring metrics.

        Args:

```

```

scores: List of policy area scores [0-3]

Returns:
    ScoringMetrics with all computed values
"""

if not scores:
    raise ValueError("Empty scores list")

n = len(scores)

# Basic statistics
mean = sum(scores) / n
variance = sum((s - mean) ** 2 for s in scores) / n
std_dev = variance ** 0.5

# Coefficient of variation (relative dispersion)
cv = std_dev / mean if mean > 0 else 0.0

# Dispersion index (normalized range)
score_range = max(scores) - min(scores) if scores else 0.0
dispersion_index = score_range / self.config.max_score

# Normalized standard deviation
normalized_std = std_dev / self.config.max_score

# Classify scenario type based on CV and DI
if cv < self.config.convergence_cv_threshold and dispersion_index <
self.config.convergence_di_threshold:
    scenario_type = "convergence"
elif cv < self.config.high_dispersion_cv_threshold:
    scenario_type = "moderate"
elif cv < self.config.extreme_dispersion_cv_threshold:
    scenario_type = "high_dispersion"
else:
    scenario_type = "extreme_dispersion"

# Classify shape
shape_classification = self._classify_distribution_shape(scores, mean, std_dev)

return ScoringMetrics(
    mean=mean,
    variance=variance,
    std_dev=std_dev,
    coefficient_variation=cv,
    dispersion_index=dispersion_index,
    normalized_std=normalized_std,
    scenario_type=scenario_type,
    shape_classification=shape_classification
)

def _classify_distribution_shape(
    self,
    scores: list[float],
    mean: float,

```

```

    std_dev: float
) -> str:
"""
Classify the distribution shape of scores.

Classifications:
- uniform: Scores evenly distributed
- clustered: Scores tightly grouped
- bimodal: Two distinct clusters
- dispersed: Wide spread of scores

Args:
    scores: List of scores
    mean: Mean of scores
    std_dev: Standard deviation

Returns:
    Shape classification string
"""

if len(scores) < 3:
    return "insufficient_data"

# Check for clustering
within_1std = sum(1 for s in scores if abs(s - mean) <= std_dev)
clustering_ratio = within_1std / len(scores)

if clustering_ratio > 0.85:
    return "clustered"

# Check for bimodal pattern
sorted_scores = sorted(scores)
n = len(sorted_scores)

# Detect gap in middle
if n >= 4:
    midpoint = n // 2
    gap_size = sorted_scores[midpoint] - sorted_scores[midpoint - 1]
    avg_gap = (max(scores) - min(scores)) / (n - 1) if n > 1 else 0

    if gap_size > 2 * avg_gap:
        return "bimodal"

# Check for uniform distribution
gaps = [sorted_scores[i+1] - sorted_scores[i] for i in range(n-1)]
if gaps:
    gap_variance = sum((g - sum(gaps)/len(gaps)) ** 2 for g in gaps) / len(gaps)
    if gap_variance < 0.1:
        return "uniform"

return "dispersed"

def compute_adaptive_penalty_factor(
    self,
    metrics: ScoringMetrics
)

```

```

) -> tuple[float, dict[str, Any]]:
    """
    Compute adaptive penalty factor based on scenario characteristics.

    This is the core improvement over the fixed PENALTY_WEIGHT=0.3 approach.

    Algorithm:
    1. Determine sensitivity multiplier based on scenario type
    2. Apply shape-based adjustments for extreme cases
    3. Compute non-linear penalty for high dispersion
    4. Return penalty factor bounded to [0, 1]

    Args:
        metrics: ScoringMetrics from compute_metrics

    Returns:
        Tuple of (penalty_factor, computation_details)
    """
    # Step 1: Determine base penalty from normalized std
    base_penalty = metrics.normalized_std * self.config.base_penalty_weight

    # Step 2: Select sensitivity multiplier based on scenario
    if metrics.scenario_type == "convergence":
        sensitivity_multiplier = self.config.convergence_multiplier
    elif metrics.scenario_type == "moderate":
        sensitivity_multiplier = self.config.moderate_multiplier
    elif metrics.scenario_type == "high_dispersion":
        sensitivity_multiplier = self.config.high_dispersion_multiplier
    else: # extreme_dispersion
        sensitivity_multiplier = self.config.extreme_dispersion_multiplier

    # Step 3: Apply shape-based adjustments
    shape_factor = 1.0

    if metrics.shape_classification == "bimodal":
        # Bimodal distributions indicate policy inconsistency - add penalty
        shape_factor = self.config.bimodal_penalty_boost
    elif metrics.scenario_type == "extreme_dispersion":
        # Apply non-linear scaling for extreme cases
        # Use exponential scaling: factor = 1 + (DI ^ extreme_shape_factor)
        shape_factor = 1.0 + (metrics.dispersion_index ** self.config.extreme_shape_factor)
    elif metrics.dispersion_index > 0.7:
        # Moderate non-linear scaling for high dispersion
        shape_factor = 1.0 + (0.3 * metrics.dispersion_index)

    # Step 4: Compute final penalty strength
    penalty_strength = base_penalty * sensitivity_multiplier * shape_factor

    # Step 5: Compute penalty factor (bounded to [0, 1])
    # Ensure we don't go below 0.5 (max 50% penalty) for stability
    penalty_factor = max(0.5, min(1.0, 1.0 - penalty_strength))

    # Computation details for transparency

```

```

details = {
    "base_penalty": base_penalty,
    "sensitivity_multiplier": sensitivity_multiplier,
    "shape_factor": shape_factor,
    "penalty_strength": penalty_strength,
    "penalty_factor": penalty_factor,
    "scenario_type": metrics.scenario_type,
    "shape_classification": metrics.shape_classification,
    "bounded_to_min": penalty_factor == 0.5,
    "bounded_to_max": penalty_factor == 1.0
}

return penalty_factor, details

```

**def compute\_adjusted\_score(**

- self,**
- scores: list[float],**
- weights: list[float] | None = None**

**) -> tuple[float, dict[str, Any]]:**

"""

Compute adjusted cluster score with adaptive penalties.

This is the main entry point that replaces the current ClusterAggregator scoring logic with adaptive mechanisms.

**Args:**

- scores: List of policy area scores [0-3]**
- weights: Optional weights (defaults to equal weights)**

**Returns:**

- Tuple of (adjusted\_score, computation\_details)**

**Raises:**

- ValueError: If inputs are invalid**

"""

**if not scores:**

**raise ValueError("Empty scores list")**

**# Validate and set weights**

**if weights is None:**

**weights = [1.0 / len(scores)] \* len(scores)**

**if len(weights) != len(scores):**

**raise ValueError(f"Weight length mismatch: {len(weights)} != {len(scores)}")**

**weight\_sum = sum(weights)**

**if abs(weight\_sum - 1.0) > 1e-6:**

**raise ValueError(f"Weights don't sum to 1.0: {weight\_sum:.6f}")**

**# Compute weighted score**

**weighted\_score = sum(s \* w for s, w in zip(scores, weights, strict=True))**

**# Compute metrics**

**metrics = self.compute\_metrics(scores)**

```

# Compute adaptive penalty factor
penalty_factor, penalty_details = self.compute_adaptive_penalty_factor(metrics)

# Apply penalty
adjusted_score = weighted_score * penalty_factor

# Compute coherence (for compatibility with existing code)
if len(scores) <= 1:
    coherence = 1.0
else:
    coherence = max(0.0, 1.0 - metrics.normalized_std)

# Full computation details
details = {
    "scores": scores,
    "weights": weights,
    "weighted_score": weighted_score,
    "metrics": {
        "mean": metrics.mean,
        "variance": metrics.variance,
        "std_dev": metrics.std_dev,
        "cv": metrics.coefficient_variation,
        "dispersion_index": metrics.dispersion_index,
        "normalized_std": metrics.normalized_std,
        "scenario_type": metrics.scenario_type,
        "shape_classification": metrics.shape_classification
    },
    "penalty_computation": penalty_details,
    "adjusted_score": adjusted_score,
    "coherence": coherence,
    "improvement_over_fixed": self._compute_improvement_metric(
        weighted_score, adjusted_score, metrics
    )
}

logger.debug(
    f"Adaptive scoring: weighted={weighted_score:.4f}, "
    f"penalty_factor={penalty_factor:.4f}, adjusted={adjusted_score:.4f}, "
    f"scenario={metrics.scenario_type}"
)

return adjusted_score, details

def _compute_improvement_metric(
    self,
    weighted_score: float,
    adjusted_score: float,
    metrics: ScoringMetrics
) -> dict[str, Any]:
    """
    Compute improvement metric vs fixed penalty approach.
    """

Args:

```

```

        weighted_score: Base weighted score
        adjusted_score: Adjusted score with adaptive penalty
        metrics: Scoring metrics

    Returns:
        Dictionary with improvement analysis
    """
    # Compute what the old fixed approach would have done
    old_penalty_factor = 1.0 - (metrics.normalized_std * 0.3)
    old_adjusted_score = weighted_score * old_penalty_factor

    # Compare
    score_difference = adjusted_score - old_adjusted_score
    penalty_difference = (1.0 - old_penalty_factor) - (1.0 - (adjusted_score / weighted_score if weighted_score > 0 else 0))

    return {
        "old_fixed_approach": {
            "penalty_factor": old_penalty_factor,
            "adjusted_score": old_adjusted_score
        },
        "new_adaptive_approach": {
            "penalty_factor": adjusted_score / weighted_score if weighted_score > 0
else 0,
            "adjusted_score": adjusted_score
        },
        "score_difference": score_difference,
        "penalty_difference": penalty_difference,
        "is_more_lenient": score_difference > 0.01,
        "is_stricter": score_difference < -0.01
    }

def get_sensitivity_analysis(
    self,
    scores: list[float],
    weights: list[float] | None = None
) -> dict[str, Any]:
    """
    Perform sensitivity analysis for a given score set.

    This method provides detailed analysis of how the adaptive mechanism
    responds to the specific scenario characteristics.

    Args:
        scores: List of policy area scores
        weights: Optional weights

    Returns:
        Dictionary with comprehensive sensitivity analysis
    """
    adjusted_score, details = self.compute_adjusted_score(scores, weights)

    # Test variations to show sensitivity
    metrics = self.compute_metrics(scores)

```

```

# Analyze sensitivity to small perturbations
perturbations = []
for i, score in enumerate(scores):
    for delta in [-0.2, -0.1, 0.1, 0.2]:
        if 0 <= score + delta <= self.config.max_score:
            perturbed = scores.copy()
            perturbed[i] = score + delta
            perturbed_score, _ = self.compute_adjusted_score(perturbed, weights)
            perturbations.append({
                "position": i,
                "delta": delta,
                "original_score": score,
                "perturbed_score": perturbed_score,
                "impact": perturbed_score - adjusted_score
            })
return {
    "base_analysis": details,
    "sensitivity_to_perturbations": {
        "max_positive_impact": max((p["impact"] for p in perturbations),
default=0),
        "max_negative_impact": min((p["impact"] for p in perturbations),
default=0),
        "avg_abs_impact": sum(abs(p["impact"])) for p in perturbations) /
len(perturbations) if perturbations else 0,
        "perturbations": perturbations[:5] # Sample
    },
    "robustness_metrics": {
        "scenario_stability": "stable" if metrics.coefficient_variation < 0.3
else "unstable",
        "penalty_proportionality": details["penalty_computation"]["penalty_strength"],
        "adaptive_advantage": details["improvement_over_fixed"]["score_difference"]
    }
}

```

```

def create_adaptive_scorer(
    base_penalty_weight: float = 0.35,
    extreme_shape_factor: float = 1.8
) -> AdaptiveMesoScoring:
    """
    Factory function to create AdaptiveMesoScoring instance.

    Args:
        base_penalty_weight: Base penalty weight (default 0.35, slightly higher than old
0.3)
        extreme_shape_factor: Non-linear scaling factor for extreme cases

    Returns:
        Configured AdaptiveMesoScoring instance
    """

```

```
config = AdaptiveScoringConfig(  
    base_penalty_weight=base_penalty_weight,  
    extreme_shape_factor=extreme_shape_factor  
)  
return AdaptiveMesoScoring(config)
```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/aggregation.py
```

```
"""
```

## Aggregation Module - Hierarchical Score Aggregation System

This module implements the complete aggregation pipeline for the policy analysis system:

- FASE 4: Dimension aggregation (60 dimensions: 6 × 10 policy areas)
- FASE 5: Policy area aggregation (10 areas)
- FASE 6: Cluster aggregation (4 MESO questions)
- FASE 7: Macro evaluation (1 holistic question)

### Requirements:

- Validation of weights, thresholds, and hermeticity
- Comprehensive logging and abortability at each level
- No strategic simplification
- Full alignment with monolith specifications
- Uses canonical notation for dimension and policy area validation

### Architecture:

- DimensionAggregator: Aggregates 5 micro questions ? 1 dimension score
- AreaPolicyAggregator: Aggregates 6 dimension scores ? 1 area score
- ClusterAggregator: Aggregates multiple area scores ? 1 cluster score
- MacroAggregator: Aggregates all cluster scores ? 1 holistic evaluation

```
"""
```

```
from __future__ import annotations

import logging
from collections import defaultdict
from dataclasses import dataclass, field
from typing import TYPE_CHECKING, Any, TypeVar

# Calibration removed - stub placeholders
def get_parameter_loader() -> Any:
    """Stub: Calibration parameter loader removed."""
    return None

def calibrated_method(method_path: str) -> Any:
    """Stub decorator: Calibration removed."""
    def decorator(func: Any) -> Any:
        return func
    return decorator

from farfan_pipeline.core.parameters import ParameterLoaderV2

# SOTA imports
from farfan_pipeline.processing.aggregation_provenance import (
    AggregationDAG,
    ProvenanceNode,
)
from farfan_pipeline.processing.uncertainty_quantification import (
    BootstrapAggregator,
    UncertaintyMetrics,
    aggregate_with_uncertainty,
```

```

)
from farfan_pipeline.processing.choquet_adapter import (
    ChoquetProcessingAdapter,
    create_default_choquet_adapter,
)
if TYPE_CHECKING:
    from collections.abc import Callable, Iterable
        from cross_cutting_infrastructure.irrigation_using_signals.SISAS.signal_registry
import QuestionnaireSignalRegistry

T = TypeVar('T')

@dataclass(frozen=True)
class AggregationSettings:
    """Resolved aggregation settings derived from the questionnaire monolith.

    SISAS: Now supports construction from signal registry for deterministic irrigation.
    """

    dimension_group_by_keys: list[str]
    area_group_by_keys: list[str]
    cluster_group_by_keys: list[str]
    dimension_question_weights: dict[str, dict[str, float]]
    policy_area_dimension_weights: dict[str, dict[str, float]]
    cluster_policy_area_weights: dict[str, dict[str, float]]
    macro_cluster_weights: dict[str, float]
    dimension_expected_counts: dict[tuple[str, str], int]
    area_expected_dimension_counts: dict[str, int]
    # SISAS: Signal provenance for byte-reproducibility
    source_hash: str | None = None
    sisas_source: str = "legacy" # "legacy" or "sisas_registry"

    @classmethod
    def from_signal_registry(
        cls,
        registry: "QuestionnaireSignalRegistry",
        level: str = "MACRO_1",
    ) -> "AggregationSettings":
        """SISAS: Build aggregation settings from signal registry.

        This method provides deterministic, signal-driven aggregation by:
        1. Using AssemblySignalPack for canonical weights
        2. Tracking source_hash for byte-reproducibility
        3. Ensuring single source of truth from registry
        """

        Args:
            registry: SISAS signal registry with assembly signals
            level: Assembly level (MESO_1, MESO_2, MACRO_1)

        Returns:
            AggregationSettings with signal-derived weights
    
```

```

Raises:
    SignalExtractionError: If assembly signals cannot be retrieved
"""

        logger.info(f"SISAS: Building AggregationSettings from registry
(level={level})")

try:
    # Get assembly signals from registry
    # Note: Using "meso" as the canonical level for aggregation assembly signals
    assembly_pack = registry.get_assembly_signals("meso")
    source_hash = getattr(assembly_pack, 'source_hash', None)

    # Extract weights from assembly pack
    cluster_policy_areas = getattr(assembly_pack, 'cluster_policy_areas', {})
    dimension_weights = getattr(assembly_pack, 'dimension_weights', {})
    aggregation_methods = getattr(assembly_pack, 'aggregation_methods', {})

    # Build cluster weights from cluster_policy_areas mapping
    cluster_policy_area_weights: dict[str, dict[str, float]] = {}
    for cluster_id, area_ids in cluster_policy_areas.items():
        if area_ids:
            equal_weight = 1.0 / len(area_ids)
            cluster_policy_area_weights[cluster_id] = {
                area_id: equal_weight for area_id in area_ids
            }

    # Build macro weights (equal across clusters)
    cluster_ids = list(cluster_policy_areas.keys())
    macro_cluster_weights: dict[str, float] = {}
    if cluster_ids:
        equal_weight = 1.0 / len(cluster_ids)
        macro_cluster_weights = {cid: equal_weight for cid in cluster_ids}

    settings = cls(
        dimension_group_by_keys=["policy_area", "dimension"],
        area_group_by_keys=["area_id"],
        cluster_group_by_keys=["cluster_id"],
        dimension_question_weights=dimension_weights,
        policy_area_dimension_weights={}, # Populated from registry if
available
        cluster_policy_area_weights=cluster_policy_area_weights,
        macro_cluster_weights=macro_cluster_weights,
        dimension_expected_counts={},
        area_expected_dimension_counts={},
        source_hash=source_hash,
        sisas_source="sisas_registry",
    )

    logger.info(
        f"SISAS: AggregationSettings built from registry - "
        f"clusters={len(cluster_policy_area_weights)}",
        source_hash={source_hash[:16] if source_hash else 'N/A'}...
    )

```

```

    return settings

except Exception as e:
    logger.warning(
        "SISAS: Failed to build from registry (%s); falling back to legacy empty
weights "
        "(callers will use equal-weight defaults.)",
        e,
    )
# Return empty settings as fallback
return cls(
    dimension_group_by_keys=["policy_area", "dimension"],
    area_group_by_keys=["area_id"],
    cluster_group_by_keys=["cluster_id"],
    dimension_question_weights={},
    policy_area_dimension_weights={},
    cluster_policy_area_weights={},
    macro_cluster_weights={},
    dimension_expected_counts={},
    area_expected_dimension_counts={},
    source_hash=None,
    sisas_source="legacy_fallback",
)

```

`@classmethod`

```

def from_monolith(cls, monolith: dict[str, Any] | None) -> "AggregationSettings":
    """Build aggregation settings from canonical questionnaire data."""
    if not monolith:
        return cls(
            dimension_group_by_keys=["policy_area", "dimension"],
            area_group_by_keys=["area_id"],
            cluster_group_by_keys=["cluster_id"],
            dimension_question_weights={},
            policy_area_dimension_weights={},
            cluster_policy_area_weights={},
            macro_cluster_weights={},
            dimension_expected_counts={},
            area_expected_dimension_counts={},
            source_hash=None, # SISAS: No hash for empty monolith
            sisas_source="legacy_monolith", # SISAS: Track source
        )

    blocks = monolith.get("blocks", {})
    niveles = blocks.get("niveles_abstraccion", {})
    policy_areas = niveles.get("policy_areas", [])
    clusters = niveles.get("clusters", [])
    micro_questions = blocks.get("micro_questions", [])

    aggregation_block = (
        monolith.get("aggregation")
        or blocks.get("aggregation")
        or monolith.get("rubric", {}).get("aggregation")
        or {}
    )

```

```

# Map question_id ? base_slot for later normalization
question_slot_lookup: dict[str, str] = {}
dimension_slot_map: dict[str, set[str]] = defaultdict(set)
dimension_expected_counts: dict[tuple[str, str], int] = defaultdict(int)

for question in micro_questions:
    qid = question.get("question_id")
    dim_id = question.get("dimension_id") or question.get("dimension")
    area_id = question.get("policy_area_id") or question.get("policy_area")
    base_slot = question.get("base_slot")

    if dim_id and qid and not base_slot:
        base_slot = f"{dim_id}-{qid}"

    if qid and base_slot:
        question_slot_lookup[qid] = base_slot
        dimension_slot_map[dim_id].add(base_slot)

    if area_id and dim_id:
        dimension_expected_counts[(area_id, dim_id)] += 1

area_expected_dimension_counts: dict[str, int] = {}
for area in policy_areas:
    area_id = area.get("policy_area_id") or area.get("id")
    if not area_id:
        continue
    dims = area.get("dimension_ids") or []
    area_expected_dimension_counts[area_id] = len(dims)

group_by_block = aggregation_block.get("group_by_keys") or {}
dimension_group_by_keys = cls._coerce_str_list(
    group_by_block.get("dimension"),
    fallback=["policy_area", "dimension"],
)
area_group_by_keys = cls._coerce_str_list(
    group_by_block.get("area"),
    fallback=["area_id"],
)
cluster_group_by_keys = cls._coerce_str_list(
    group_by_block.get("cluster"),
    fallback=["cluster_id"],
)

dimension_question_weights = cls._build_dimension_weights(
    aggregation_block.get("dimension_question_weights") or {},
    question_slot_lookup,
    dimension_slot_map,
)
policy_area_dimension_weights = cls._build_area_dimension_weights(
    aggregation_block.get("policy_area_dimension_weights") or {},
    policy_areas,
)
cluster_policy_area_weights = cls._build_cluster_weights(

```

```

        aggregation_block.get("cluster_policy_area_weights") or {},
        clusters,
    )

    macro_cluster_weights = cls._build_macro_weights(
        aggregation_block.get("macro_cluster_weights") or {},
        clusters,
    )

    return cls(
        dimension_group_by_keys=dimension_group_by_keys,
        area_group_by_keys=area_group_by_keys,
        cluster_group_by_keys=cluster_group_by_keys,
        dimension_question_weights=dimension_question_weights,
        policy_area_dimension_weights=policy_area_dimension_weights,
        cluster_policy_area_weights=cluster_policy_area_weights,
        macro_cluster_weights=macro_cluster_weights,
        dimension_expected_counts=dict(dimension_expected_counts),
        area_expected_dimension_counts=area_expected_dimension_counts,
        source_hash=None, # SISAS: No hash from raw monolith
        sisas_source="legacy_monolith", # SISAS: Track source
    )

```

`@classmethod`

```

def from_monolith_or_registry(
    cls,
    monolith: dict[str, Any] | None = None,
    registry: "QuestionnaireSignalRegistry" | None" = None,
    level: str = "MACRO_1",
) -> "AggregationSettings":
    """SISAS: Transition method - prefer registry, fallback to monolith.

This method enables gradual migration from legacy monolith-based
configuration to SISAS signal-driven configuration.

```

`Args:`

```

    monolith: Legacy questionnaire monolith (fallback)
    registry: SISAS signal registry (preferred)
    level: Assembly level for registry (MESO_1, MESO_2, MACRO_1)

```

`Returns:`

```
    AggregationSettings from registry if available, else from monolith
```

`Raises:`

```
    ValueError: If neither registry nor monolith provided
"""

```

```
if registry is not None:
    logger.info("SISAS: Building AggregationSettings from registry (preferred
```

```
path)")

    return cls.from_signal_registry(registry, level)
```

```
elif monolith is not None:
    logger.info("SISAS: Building AggregationSettings from monolith (fallback
```

```
path)")

    return cls.from_monolith(monolith)

else:
```

```

        raise ValueError("Must provide either registry or monolith")

@staticmethod
def _coerce_str_list(value: Any, *, fallback: list[str]) -> list[str]:
    if isinstance(value, list) and all(isinstance(item, str) for item in value):
        return value or fallback
    return fallback

@staticmethod
def _normalize_weights(weight_map: dict[str, float]) -> dict[str, float]:
    if not weight_map:
        return {}
    positive_map = {
        key: float(value)
        for key, value in weight_map.items()
        if isinstance(value, (float, int)) and float(value) >= 0.0
    }
    if not positive_map:
        return {}
    total = sum(positive_map.values())
    if total <= 0:
        return {}
    return {k: value / total for k, value in positive_map.items()}

@classmethod
def _build_dimension_weights(
    cls,
    raw_weights: dict[str, dict[str, Any]],
    question_slot_lookup: dict[str, str],
    dimension_slot_map: dict[str, set[str]],
) -> dict[str, dict[str, float]]:
    dimension_weights: dict[str, dict[str, float]] = {}
    if raw_weights:
        for dim_id, weights in raw_weights.items():
            resolved: dict[str, float] = {}
            for qid, weight in weights.items():
                slot = question_slot_lookup.get(qid, qid)
                try:
                    resolved[slot] = float(weight)
                except (TypeError, ValueError):
                    continue
            if resolved:
                normalized = cls._normalize_weights(resolved)
                if normalized:
                    dimension_weights[dim_id] = normalized

    if not dimension_weights:
        for dim_id, slots in dimension_slot_map.items():
            if not slots:
                continue
            equal = 1.0 / len(slots)
            dimension_weights[dim_id] = {slot: equal for slot in slots}

    return dimension_weights

```

```

@classmethod
def _build_area_dimension_weights(
    cls,
    raw_weights: dict[str, dict[str, Any]],
    policy_areas: list[dict[str, Any]],
) -> dict[str, dict[str, float]]:
    area_weights: dict[str, dict[str, float]] = {}
    if raw_weights:
        for area_id, weights in raw_weights.items():
            resolved: dict[str, float] = {}
            for dim_id, value in weights.items():
                try:
                    resolved[dim_id] = float(value)
                except (TypeError, ValueError):
                    continue
            if resolved:
                normalized = cls._normalize_weights(resolved)
                if normalized:
                    area_weights[area_id] = normalized

    if not area_weights:
        for area in policy_areas:
            area_id = area.get("policy_area_id") or area.get("id")
            dims = area.get("dimension_ids") or []
            if not area_id or not dims:
                continue
            equal = 1.0 / len(dims)
            area_weights[area_id] = {dim: equal for dim in dims}

    return area_weights


@classmethod
def _build_cluster_weights(
    cls,
    raw_weights: dict[str, dict[str, Any]],
    clusters: list[dict[str, Any]],
) -> dict[str, dict[str, float]]:
    cluster_weights: dict[str, dict[str, float]] = {}
    if raw_weights:
        for cluster_id, weights in raw_weights.items():
            resolved: dict[str, float] = {}
            for area_id, value in weights.items():
                try:
                    resolved[area_id] = float(value)
                except (TypeError, ValueError):
                    continue
            if resolved:
                normalized = cls._normalize_weights(resolved)
                if normalized:
                    cluster_weights[cluster_id] = normalized

    if not cluster_weights:
        for cluster in clusters:

```

```

        cluster_id = cluster.get("cluster_id")
        area_ids = cluster.get("policy_area_ids") or []
        if not cluster_id or not area_ids:
            continue
        equal = 1.0 / len(area_ids)
        cluster_weights[cluster_id] = {area_id: equal for area_id in area_ids}

    return cluster_weights

@classmethod
def _build_macro_weights(
    cls,
    raw_weights: dict[str, Any],
    clusters: list[dict[str, Any]],
) -> dict[str, float]:
    if raw_weights:
        resolved = {}
        for cluster_id, weight in raw_weights.items():
            try:
                resolved[cluster_id] = float(weight)
            except (TypeError, ValueError):
                continue
        normalized = cls._normalize_weights(resolved)
        if normalized:
            return normalized

        cluster_ids = [cluster.get("cluster_id") for cluster in clusters if
cluster.get("cluster_id")]
        if not cluster_ids:
            return {}
        equal = 1.0 / len(cluster_ids)
        return {cluster_id: equal for cluster_id in cluster_ids}

```

```
def group_by(items: Iterable[T], key_func: Callable[[T], tuple]) -> dict[tuple,
list[T]]:
    """

```

Groups a sequence of items into a dictionary based on a key function.

This utility function iterates over a collection, applies a key function to each item, and collects items into lists, keyed by the result of the key function.

The key function must return a tuple. This is because dictionary keys must be hashable, and tuples are hashable whereas lists are not. Using a tuple allows for grouping by multiple attributes.

If the input iterable `items` is empty, this function will return an empty dictionary.

Example:

```

>>> from dataclasses import dataclass
>>> @dataclass
... class Record:
...     category: str
...     value: int

```

```
...
>>> data = [Record("A", 1), Record("B", 2), Record("A", 3)]
>>> group_by(data, key_func=lambda r: (r.category,))
{('A',): [Record(category='A', value=1), Record(category='A', value=3)],
 ('B',): [Record(category='B', value=2)]}
```

Args:

```
    items: An iterable of items to be grouped.
    key_func: A callable that accepts an item and returns a tuple to be
              used as the grouping key.
```

Returns:

```
    A dictionary where keys are the result of the key function and values are
    lists of items belonging to that group.
```

"""

```
grouped = defaultdict(list)
for item in items:
    grouped[key_func(item)].append(item)
return dict(grouped)
```

```
def validate_scored_results(results: list[dict[str, Any]]) -> list[ScoredResult]:
```

"""

```
    Validates a list of dictionaries and converts them to ScoredResult objects.
```

Args:

```
    results: A list of dictionaries representing scored results.
```

Returns:

```
    A list of ScoredResult objects.
```

Raises:

```
    ValidationError: If any of the dictionaries are invalid.
```

"""

```
validated_results: list[ScoredResult] = []
required_keys: dict[str, type[Any] | tuple[type[Any], ...]] = {
    "question_global": (int, float, str),
    "base_slot": str,
    "policy_area": str,
    "dimension": str,
    "score": (float, int),
    "quality_level": str,
    "evidence": dict,
    "raw_results": dict,
}
for i, res_dict in enumerate(results):
    missing_keys = set(required_keys) - set(res_dict)
    if missing_keys:
        raise ValidationError(
            f"Invalid ScoredResult at index {i}: missing keys {missing_keys}"
        )

    normalized = dict(res_dict)
    qid = normalized["question_global"]
    if isinstance(qid, bool):
```

```

        raise ValidationError(
            f"Invalid type for key 'question_global' at index {i}. "
            f"Expected int|float|str, got bool."
        )
    if isinstance(qid, float):
        if not qid.is_integer():
            raise ValidationError(
                f"Invalid value for key 'question_global' at index {i}. "
                f"Expected integer-like float, got {qid}.")
    else:
        normalized["question_global"] = int(qid)
    elif isinstance(qid, str):
        qid_str = qid.strip()
        if qid_str.isdigit():
            normalized["question_global"] = int(qid_str)
        else:
            normalized["question_global"] = qid_str

    score_value = normalized["score"]
    if isinstance(score_value, bool):
        raise ValidationError(
            f"Invalid type for key 'score' at index {i}. Expected float|int, got "
            f"bool."
        )
    normalized["score"] = float(score_value)

    for key, expected_type in required_keys.items():
        if key in {"question_global", "score"}:
            continue
        if not isinstance(normalized[key], expected_type):
            raise ValidationError(
                f"Invalid type for key '{key}' at index {i}. "
                f"Expected {expected_type}, got {type(normalized[key])}.")
    try:
        validated_results.append(ScoredResult(**normalized))
    except TypeError as e:
        raise ValidationError(f"Invalid ScoredResult at index {i}: {e}") from e
    return validated_results

def _normalize_question_node_id(question_id: int | str) -> str:
    if isinstance(question_id, int):
        return f"Q{question_id:03d}"

    question_id_str = question_id.strip()
    suffix = question_id_str.lstrip("Qq")
    if suffix.isdigit():
        return f"Q{int(suffix):03d}"
    return question_id_str

# Import canonical notation for validation
try:
    from farfan_pipeline.core.canonical_notation import get_all_dimensions,
    get_all_policy_areas

```

```

HAS_CANONICAL_NOTATION = True
except ImportError:
    HAS_CANONICAL_NOTATION = False

logger = logging.getLogger(__name__)

@dataclass
class ScoredResult:
    """Represents a single, scored micro-question, forming the input for aggregation."""
    question_global: int | str
    base_slot: str
    policy_area: str
    dimension: str
    score: float
    quality_level: str
    evidence: dict[str, Any]
    raw_results: dict[str, Any]

@dataclass
class DimensionScore:
    """
    Aggregated score for a single dimension within a policy area.

    SOTA Extensions:
    - Uncertainty quantification (mean, std, CI)
    - Provenance tracking (DAG node ID)
    - Aggregation method recording
    """

    dimension_id: str
    area_id: str
    score: float
    quality_level: str
    contributing_questions: list[int | str]
    validation_passed: bool = True
    validation_details: dict[str, Any] = field(default_factory=dict)

    # SOTA: Uncertainty quantification
    score_std: float = 0.0
    confidence_interval_95: tuple[float, float] = field(default_factory=lambda: (0.0,
0.0))
    epistemic_uncertainty: float = 0.0
    aleatoric_uncertainty: float = 0.0

    # SOTA: Provenance tracking
    provenance_node_id: str = ""
    aggregation_method: str = "weighted_average"

@dataclass
class AreaScore:
    """
    Aggregated score for a policy area, based on its constituent dimensions.

    SOTA Extensions:
    - Uncertainty quantification

```

```

- Provenance tracking
"""

area_id: str
area_name: str
score: float
quality_level: str
dimension_scores: list[DimensionScore]
validation_passed: bool = True
validation_details: dict[str, Any] = field(default_factory=dict)
cluster_id: str | None = None

# SOTA: Uncertainty quantification
score_std: float = 0.0
confidence_interval_95: tuple[float, float] = field(default_factory=lambda: (0.0,
0.0))

# SOTA: Provenance tracking
provenance_node_id: str = ""
aggregation_method: str = "weighted_average"

@dataclass
class ClusterScore:
"""

Aggregated score for a MESO cluster, based on its policy areas.

SOTA Extensions:
- Uncertainty quantification
- Provenance tracking
"""

cluster_id: str
cluster_name: str
areas: list[str]
score: float
coherence: float
variance: float
weakest_area: str | None
area_scores: list[AreaScore]
validation_passed: bool = True
validation_details: dict[str, Any] = field(default_factory=dict)

# SOTA: Uncertainty quantification
score_std: float = 0.0
confidence_interval_95: tuple[float, float] = field(default_factory=lambda: (0.0,
0.0))

# SOTA: Provenance tracking
provenance_node_id: str = ""
aggregation_method: str = "weighted_average"

@dataclass
class MacroScore:
"""Represents the final, holistic macro evaluation score for the entire system."""
score: float
quality_level: str

```

```
cross_cutting_coherence: float # Coherence across all clusters
systemic_gaps: list[str]
strategic_alignment: float
cluster_scores: list[ClusterScore]
validation_passed: bool = True
validation_details: dict[str, Any] = field(default_factory=dict)

class AggregationError(Exception):
    """Base exception for aggregation errors."""
    pass

class ValidationError(AggregationError):
    """Raised when validation fails."""
    pass

class WeightValidationError(ValidationError):
    """Raised when weight validation fails."""
    pass

class ThresholdValidationError(ValidationError):
    """Raised when threshold validation fails."""
    pass

class HermeticityValidationError(ValidationError):
    """Raised when hermeticity validation fails."""
    pass

class CoverageError(AggregationError):
    """Raised when coverage requirements are not met."""
    pass

def calculate_weighted_average(scores: list[float], weights: list[float] | None = None) \
-> float:
    if not scores:
        return 0.0

    if weights is None:
        weights = [1.0 / len(scores)] * len(scores)

    if len(weights) != len(scores):
        msg = f"Weight length mismatch: {len(weights)} weights for {len(scores)} scores"
        logger.error(msg)
        raise WeightValidationError(msg)

    expected_sum = 1.0
    weight_sum = sum(weights)
    tolerance = 1e-6
    if abs(weight_sum - expected_sum) > tolerance:
        msg = f"Weight sum validation failed: sum={weight_sum:.6f}, "
    expected=f"expected={expected_sum}"
    logger.error(msg)
    raise WeightValidationError(msg)

    return sum(score * weight for score, weight in zip(scores, weights, strict=False))
```

```

class DimensionAggregator:
    """
    Aggregates micro question scores into dimension scores.

    Responsibilities:
    - Aggregate 5 micro questions (Q1-Q5) per dimension
    - Validate weights sum to 1.0
    - Apply rubric thresholds
    - Ensure coverage (abort if insufficient)
    - Provide detailed logging
    """

    def __init__(
        self,
        monolith: dict[str, Any] | None = None,
        abort_on_insufficient: bool = True,
        aggregation_settings: AggregationSettings | None = None,
        enable_sota_features: bool = True,
        signal_registry: "QuestionnaireSignalRegistry" | None = None, # SISAS injection
    ) -> None:
        """
        Initialize dimension aggregator.

        Args:
            monolith: Questionnaire monolith configuration (optional, required for
run())
            abort_on_insufficient: Whether to abort on insufficient coverage
            aggregation_settings: Resolved aggregation settings
            enable_sota_features: Enable SOTA features (Choquet, UQ, provenance)
            signal_registry: SISAS signal registry for signal-driven aggregation

        Raises:
            ValueError: If monolith is None and required for operations
        """
        self.monolith = monolith
        self.abort_on_insufficient = abort_on_insufficient
        self._signal_registry = signal_registry # SISAS: Cache for signal-driven config

        # SISAS: Use transition method for automatic detection
        # Handle case where both are None gracefully
        if aggregation_settings is not None:
            self.aggregation_settings = aggregation_settings
        elif signal_registry is not None or monolith is not None:
            self.aggregation_settings = AggregationSettings.from_monolith_or_registry(
                monolith=monolith,
                registry=signal_registry,
                level="MACRO_1"
            )
        else:
            # Both None - use empty settings (legacy fallback)
            self.aggregation_settings = AggregationSettings.from_monolith(None)
        self.dimension_group_by_keys = (
            self.aggregation_settings.dimension_group_by_keys or ["policy_area",

```

```

"dimension" ]
)
self.enable_sota_features = enable_sota_features

# Extract configuration if monolith provided
if monolith is not None:
    self.scoring_config = monolith["blocks"]["scoring"]
    self.niveles = monolith["blocks"]["niveles_abstraccion"]
else:
    self.scoring_config = None
    self.niveles = None

# SOTA: Initialize provenance DAG
if self.enable_sota_features:
    self.provenance_dag = AggregationDAG()
    self.bootstrap_aggregator = BootstrapAggregator(n_samples=1000,
random_seed=42)
    logger.info("DimensionAggregator initialized with SOTA features enabled")
else:
    self.provenance_dag = None
    self.bootstrap_aggregator = None
    logger.info("DimensionAggregator initialized (legacy mode)")

# Validate canonical notation if available
if HAS_CANONICAL_NOTATION:
    try:
        canonical_dims = get_all_dimensions()
        canonical_areas = get_all_policy_areas()
        logger.info(
            f"Canonical notation loaded: {len(canonical_dims)} dimensions, "
            f"{len(canonical_areas)} policy areas"
        )
    except Exception as e:
        logger.warning(f"Could not load canonical notation: {e}")

@calibrated_method("farfan_core.processing.aggregation.DimensionAggregator.validate_dimension_id")
def validate_dimension_id(self, dimension_id: str) -> bool:
    """
    Validate dimension ID against canonical notation.

    Args:
        dimension_id: Dimension ID to validate (e.g., "DIM01")

    Returns:
        True if dimension ID is valid

    Raises:
        ValidationError: If dimension ID is invalid and abort_on_insufficient is
True
    """
    if not HAS_CANONICAL_NOTATION:
        logger.debug("Canonical notation not available, skipping validation")

```

```

        return True

    try:
        canonical_dims = get_all_dimensions()
        # Check if dimension_id is a valid code
        valid_codes = {info.code for info in canonical_dims.values()}
        if dimension_id in valid_codes:
            return True

            msg = f"Invalid dimension ID: {dimension_id}. Valid codes:
{sorted(valid_codes)}"
            logger.error(msg)
            if self.abort_on_insufficient:
                raise ValidationError(msg)
            return False
        except Exception as e:
            logger.warning(f"Could not validate dimension ID: {e}")
            return True # Don't fail if validation can't be performed

@calibrated_method("farfan_core.processing.aggregation.DimensionAggregator.validate_policy_area_id")
def validate_policy_area_id(self, area_id: str) -> bool:
    """
    Validate policy area ID against canonical notation.

    Args:
        area_id: Policy area ID to validate (e.g., "PA01")

    Returns:
        True if policy area ID is valid

    Raises:
        ValidationError: If policy area ID is invalid and abort_on_insufficient is
True
    """
    if not HAS_CANONICAL_NOTATION:
        logger.debug("Canonical notation not available, skipping validation")
        return True

    try:
        canonical_areas = get_all_policy_areas()
        if area_id in canonical_areas:
            return True

            msg = f"Invalid policy area ID: {area_id}. Valid codes:
{sorted(canonical_areas.keys())}"
            logger.error(msg)
            if self.abort_on_insufficient:
                raise ValidationError(msg)
            return False
        except Exception as e:
            logger.warning(f"Could not validate policy area ID: {e}")
            return True # Don't fail if validation can't be performed

```

```

@calibrated_method("farfan_core.processing.aggregation.DimensionAggregator.validate_weights")
def validate_weights(self, weights: list[float]) -> tuple[bool, str]:
    """
        Ensures that a list of weights sums to
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L582_47", 1.0) within a small tolerance.

Args:
    weights: A list of floating-point weights.

Returns:
    A tuple containing a boolean indicating validity and a descriptive message.

Raises:
    WeightValidationError: If `abort_on_insufficient` is True and validation fails.

    """
    if not weights:
        msg = "No weights provided"
        logger.error(msg)
        if self.abort_on_insufficient:
            raise WeightValidationError(msg)
        return False, msg

    weight_sum = sum(weights)
    tolerance = 1e-6

        if abs(weight_sum -
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L603_28", 1.0)) > tolerance:
            expected_weight =
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L604_81", 1.0)
            msg = f"Weight sum validation failed: sum={weight_sum:.6f}, "
expected={expected_weight}"
            logger.error(msg)
            if self.abort_on_insufficient:
                raise WeightValidationError(msg)
            return False, msg

    logger.debug(f"Weight validation passed: sum={weight_sum:.6f}")
    return True, "Weights valid"

def validate_coverage(
    self,
    results: list[ScoredResult],
    expected_count: int = 5
) -> tuple[bool, str]:
    """
        Checks if the number of results meets a minimum expectation.

```

```

Args:
    results: A list of ScoredResult objects.
    expected_count: The minimum number of results required.

Returns:
    A tuple containing a boolean indicating validity and a descriptive message.

Raises:
    CoverageError: If `abort_on_insufficient` is True and coverage is
insufficient.

"""
actual_count = len(results)

if actual_count < expected_count:
    msg = (
        f"Coverage validation failed: "
        f"expected {expected_count} questions, got {actual_count}"
    )
    logger.error(msg)
    if self.abort_on_insufficient:
        raise CoverageError(msg)
    return False, msg

    logger.debug(f"Coverage validation passed: {actual_count}/{expected_count}
questions")
    return True, "Coverage sufficient"

def calculate_weighted_average(
    self,
    scores: list[float],
    weights: list[float] | None = None
) -> float:
    """
    Calculates a weighted average, defaulting to an equal weighting if none
provided.

Args:
    scores: A list of scores to be averaged.
    weights: An optional list of weights. If None, equal weights are assumed.

Returns:
    The calculated weighted average.

Raises:
    WeightValidationError: If the weights are invalid (e.g., mismatched length).

"""
    weighted_sum = calculate_weighted_average(scores, weights)
    logger.debug(
        "Weighted average calculated: scores=%s, weights=%s, result=%.4f",
        scores,
        weights,
        weighted_sum,
    )
    return weighted_sum

```

```

def aggregate_with_sota(
    self,
    scores: list[float],
    weights: list[float] | None = None,
    method: str = "choquet",
    compute_uncertainty: bool = True,
) -> tuple[float, UncertaintyMetrics | None]:
    """
    SOTA aggregation with Choquet integral and uncertainty quantification.

    This method provides:
    1. Non-linear aggregation via Choquet integral (captures synergies)
    2. Bayesian uncertainty quantification via bootstrap
    3. Full reproducibility with fixed random seed
    """

    Args:
        scores: Input scores to aggregate
        weights: Optional weights (default: uniform)
        method: Aggregation method ("choquet" or "weighted_average")
        compute_uncertainty: Whether to compute uncertainty metrics

    Returns:
        Tuple of (aggregated_score, uncertainty_metrics)
        If compute_uncertainty=False, uncertainty_metrics is None

    Raises:
        ValueError: If scores is empty or method invalid
    """
    if not scores:
        raise ValueError("Cannot aggregate empty score list")

    if method == "choquet":
        # Use Choquet integral for non-linear aggregation
        choquet_adapter = create_default_choquet_adapter(len(scores))
        score = choquet_adapter.aggregate(scores, weights)
        logger.info(f"Choquet aggregation: {len(scores)} inputs ? {score:.4f}")
    elif method == "weighted_average":
        # Fall back to standard weighted average
        score = self.calculate_weighted_average(scores, weights)
    else:
        raise ValueError(f"Unknown aggregation method: {method}")

    # Compute uncertainty if requested
    uncertainty = None
    if compute_uncertainty and self.bootstrap_aggregator:
        _, uncertainty = aggregate_with_uncertainty(
            scores, weights, n_bootstrap=1000, random_seed=42
        )
        logger.debug(
            f"Uncertainty: mean={uncertainty.mean:.4f}, "
            f"std={uncertainty.std:.4f}, "
            f"CI95={uncertainty.confidence_interval_95}"
        )

```

```

        return score, uncertainty

    def apply_rubric_thresholds(
        self,
        score: float,
        thresholds: dict[str, float] | None = None
    ) -> str:
        """
        Apply rubric thresholds to determine quality level.

        Args:
            score: Aggregated score (0-3 range)
            thresholds: Optional threshold definitions (dict with keys: EXCELENTE,
        BUENO, ACEPTABLE)
                Each value should be a normalized threshold (0-1 range)

        Returns:
            Quality level (EXCELENTE, BUENO, ACEPTABLE, INSUFICIENTE)
        """
        # Clamp score to valid range [0, 3]
        clamped_score = max(ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L714_28", 0.0), min(3.0, score))

        # Normalize to 0-1 range
        normalized_score = clamped_score / 3.0

        # Use provided thresholds or defaults
        if thresholds:
            excellent_threshold = thresholds.get('EXCELENTE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L721_62", 0.85))
            good_threshold = thresholds.get('BUENO',
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L722_53", 0.70))
            acceptable_threshold = thresholds.get('ACEPTABLE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "auto_param_L723_63", 0.55))
        else:
            excellent_threshold = ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "excellent_threshold", 0.85) # Refactored
            good_threshold = ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "good_threshold", 0.7) # Refactored
            acceptable_threshold = ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights", "acceptable_threshold", 0.55) # Refactored

        # Apply thresholds
        if normalized_score >= excellent_threshold:
            quality = "EXCELENTE"
        elif normalized_score >= good_threshold:

```

```

        quality = "BUENO"
    elif normalized_score >= acceptable_threshold:
        quality = "ACCEPTABLE"
    else:
        quality = "INSUFICIENTE"

    logger.debug(
        f"Rubric applied: score={score:.4f}, "
        f"normalized={normalized_score:.4f}, quality={quality}"
    )

    return quality

def aggregate_dimension(
    self,
    scored_results: list[ScoredResult],
    group_by_values: dict[str, Any],
    weights: list[float] | None = None,
) -> DimensionScore:
    """
    Aggregate a single dimension from micro question results.

    Args:
        scored_results: List of scored results for this dimension/area.
        group_by_values: Dictionary of grouping keys and their values.
        weights: Optional weights for questions (defaults to equal weights).

    Returns:
        DimensionScore with aggregated score and quality level.

    Raises:
        ValidationError: If validation fails.
        CoverageError: If coverage is insufficient.
    """
    dimension_id = group_by_values.get("dimension", "UNKNOWN")
    area_id = group_by_values.get("policy_area", "UNKNOWN")
    logger.info(f"Aggregating dimension {dimension_id} for area {area_id}")

    validation_details = {}

    # In this context, scored_results are already grouped, so we can use them
    # directly.
    dim_results = scored_results

    expected_count = self._expected_question_count(area_id, dimension_id)

    # Validate coverage
    try:
        coverage_valid, coverage_msg = self.validate_coverage(
            dim_results,
            expected_count=expected_count or 5,
        )
        validation_details["coverage"] = {
            "valid": coverage_valid,
        }
    except ValidationError as e:
        validation_details["coverage"] = {
            "valid": False,
            "msg": str(e),
        }

    return DimensionScore(score=dim_results[0].score, quality=quality)

```

```

        "message": coverage_msg,
        "count": len(dim_results)
    }
except CoverageError as e:
    logger.error(f"Covrera validation failed for {dimension_id}/{area_id}:
{e}")
    # Return minimal score if aborted
    return DimensionScore(
        dimension_id=dimension_id,
        area_id=area_id,
        score=ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights",
        "auto_param_L795_22", 0.0),
        quality_level="INSUFICIENTE",
        contributing_questions=[],
        validation_passed=False,
        validation_details={"error": str(e), "type": "coverage"})
)

if not dim_results:
    logger.warning(f"No results for dimension {dimension_id}/{area_id}")
    return DimensionScore(
        dimension_id=dimension_id,
        area_id=area_id,
        score=ParameterLoaderV2.get("farfan_core.processing.aggregation.DimensionAggregator.validate_weights",
        "auto_param_L807_22", 0.0),
        quality_level="INSUFICIENTE",
        contributing_questions=[],
        validation_passed=False,
        validation_details={"error": "No results", "type": "empty"})
)

raw_scores = [r.score for r in dim_results]
scores = [min(3.0, max(0.0, score)) for score in raw_scores]
n_clamped = sum(
    1 for raw, clamped in zip(raw_scores, scores, strict=False) if raw != clamped
)
if n_clamped:
    validation_details["clamping"] = {
        "applied": True,
        "n_clamped": n_clamped,
        "original_min": min(raw_scores),
        "original_max": max(raw_scores),
        "clamped_min": min(scores),
        "clamped_max": max(scores),
    }
question_ids = [r.question_global for r in dim_results]

# Calculate weighted average with SOTA features
resolved_weights = weights or self._resolve_dimension_weights(dimension_id,
dim_results)

```

```

# SOTA: Use Choquet + uncertainty if enabled
if self.enable_sota_features and len(scores) >= 3:
    try:
        avg_score, uncertainty = self.aggregate_with_sota(
            scores,
            resolved_weights,
            method="choquet",
            compute_uncertainty=True,
        )
        validation_details["aggregation"] = {
            "method": "choquet",
            "uncertainty": uncertainty.to_dict() if uncertainty else None,
        }
    except Exception as e:
        logger.warning(f"SOTA aggregation failed, falling back to standard: {e}")
        avg_score = self.calculate_weighted_average(scores, resolved_weights)
        uncertainty = None
        validation_details["aggregation"] = {"method": "weighted_average",
                                             "fallback": True}
    else:
        # Standard aggregation
        avg_score = self.calculate_weighted_average(scores, resolved_weights)
        uncertainty = None
        validation_details["aggregation"] = {"method": "weighted_average"}

validation_details["weights"] = {
    "valid": True,
    "weights": resolved_weights if resolved_weights else "equal",
    "score": avg_score
}

# Apply rubric thresholds
quality_level = self.apply_rubric_thresholds(avg_score)
validation_details["rubric"] = {
    "score": avg_score,
    "quality_level": quality_level
}
validation_details["score_max"] = 3.0

# SOTA: Add provenance tracking
provenance_node_id = f"DIM_{dimension_id}_{area_id}"
if self.enable_sota_features and self.provenance_dag:
    # Add dimension node
    dim_node = ProvenanceNode(
        node_id=provenance_node_id,
        level="dimension",
        score=avg_score,
        quality_level=quality_level,
        metadata={
            "dimension_id": dimension_id,
            "area_id": area_id,
            "n_questions": len(question_ids),
        }
)

```

```

        },
    )
    self.provenance_dag.add_node(dim_node)

    # Add aggregation edges from questions to dimension
    question_node_ids = [_normalize_question_node_id(qid) for qid in
question_ids]
    for qid_str, score_value in zip(question_node_ids, scores, strict=False):
        # Add question node if not exists
        if qid_str not in self.provenance_dag.nodes:
            q_node = ProvenanceNode(
                node_id=qid_str,
                level="micro",
                score=score_value,
                quality_level="UNKNOWN",
            )
            self.provenance_dag.add_node(q_node)

    # Record aggregation operation
    self.provenance_dag.add_aggregation_edge(
        source_ids=question_node_ids,
        target_id=provenance_node_id,
        operation="choquet" if self.enable_sota_features else
"weighted_average",
        weights=resolved_weights or [1.0 / len(question_ids)] *
len(question_ids),
        metadata={"dimension": dimension_id, "area": area_id},
    )

    logger.debug(f"Provenance recorded: {len(question_node_ids)} questions ?
{provenance_node_id}")

    logger.info(
        f"? Dimension {dimension_id}/{area_id}: "
        f"score={avg_score:.4f}, quality={quality_level}"
        + (f", std={uncertainty.std:.4f}" if uncertainty else ""))
    )

    return DimensionScore(
        dimension_id=dimension_id,
        area_id=area_id,
        score=avg_score,
        quality_level=quality_level,
        contributing_questions=question_ids,
        validation_passed=True,
        validation_details=validation_details,
        # SOTA fields
        score_std=uncertainty.std if uncertainty else 0.0,
        confidence_interval_95=uncertainty.confidence_interval_95 if uncertainty
else (0.0, 0.0),
        epistemic_uncertainty=uncertainty.epistemic_uncertainty if uncertainty else
0.0,
        aleatoric_uncertainty=uncertainty.aleatoric_uncertainty if uncertainty else
0.0,
    )

```

```

provenance_node_id=provenance_node_id if self.enable_sota_features else "",
        aggregation_method="choquet" if (self.enable_sota_features and len(scores)
>= 3) else "weighted_average",
    )

def run(
    self,
    scored_results: list[ScoredResult],
    group_by_keys: list[str]
) -> list[DimensionScore]:
    """
    Run the dimension aggregation process.

    Args:
        scored_results: List of all scored results.
        group_by_keys: List of keys to group by.

    Returns:
        A list of DimensionScore objects.
    """
    def key_func(r):
        return tuple(getattr(r, key) for key in group_by_keys)
    grouped_results = group_by(scored_results, key_func)

    dimension_scores = []
    for group_key, results in grouped_results.items():
        group_by_values = dict(zip(group_by_keys, group_key, strict=False))
        score = self.aggregate_dimension(results, group_by_values)
        dimension_scores.append(score)

    return dimension_scores

@calibrated_method("farfan_core.processing.aggregation.DimensionAggregator._expected_question_count")
def _expected_question_count(self, area_id: str, dimension_id: str) -> int | None:
    if not self.aggregation_settings.dimension_expected_counts:
        return None
    return self.aggregation_settings.dimension_expected_counts.get((area_id,
dimension_id))

def _resolve_dimension_weights(
    self,
    dimension_id: str,
    dim_results: list[ScoredResult],
) -> list[float] | None:
    mapping = self.aggregation_settings.dimension_question_weights.get(dimension_id)
    if not mapping:
        return None

    weights: list[float] = []
    for result in dim_results:
        slot = result.base_slot
        weight = mapping.get(slot)

```

```

        if weight is None:
            logger.debug(
                "Missing weight for slot %s in dimension %s ? falling back to equal
weights",
                slot,
                dimension_id,
            )
            return None
        weights.append(weight)

    total = sum(weights)
    if total <= 0:
        return None
    return [w / total for w in weights]

def run_aggregation_pipeline(
    scored_results: list[dict[str, Any]],
    monolith: dict[str, Any],
    abort_on_insufficient: bool = True
) -> list[ClusterScore]:
    """
    Orchestrates the end-to-end aggregation pipeline.
    """

```

This function provides a high-level entry point to the aggregation system, demonstrating the sequential wiring of the aggregator components. It ensures that data flows from raw scored results through dimension, area, and finally cluster aggregation in a controlled and validated manner.

**Note on Parallelization:** This implementation is sequential. For very large datasets, the `group\_by` operations in each aggregator's `run` method could be parallelized (e.g., using `concurrent.futures`) to process independent groups concurrently.

**Args:**

- scored\_results: A list of dictionaries, each representing a raw scored result.
- monolith: The central monolith configuration object.
- abort\_on\_insufficient: If True, the pipeline will stop on validation errors.

**Returns:**

- A list of aggregated ClusterScore objects.

```

    """
# 1. Input Validation (Pre-flight check)
validated_scored_results = validate_scored_results(scored_results)

aggregation_settings = AggregationSettings.from_monolith(monolith)

# 2. FASE 4: Dimension Aggregation
dim_aggregator = DimensionAggregator(
    monolith,
    abort_on_insufficient,
    aggregation_settings=aggregation_settings,
)
dimension_scores = dim_aggregator.run(
    validated_scored_results,
)
```

```

        group_by_keys=dim_aggregator.dimension_group_by_keys,
    )

# 3. FASE 5: Area Policy Aggregation
area_aggregator = AreaPolicyAggregator(
    monolith,
    abort_on_insufficient,
    aggregation_settings=aggregation_settings,
)
area_scores = area_aggregator.run(
    dimension_scores,
    group_by_keys=area_aggregator.area_group_by_keys,
)

# 4. FASE 6: Cluster Aggregation
cluster_aggregator = ClusterAggregator(
    monolith,
    abort_on_insufficient,
    aggregation_settings=aggregation_settings,
)
cluster_definitions = monolith["blocks"]["niveles_abstraccion"]["clusters"]
cluster_scores = cluster_aggregator.run(
    area_scores,
    cluster_definitions
)

return cluster_scores

class AreaPolicyAggregator:

    """
    Aggregates dimension scores into policy area scores.

    Responsibilities:
    - Aggregate 6 dimension scores per policy area
    - Validate dimension completeness
    - Apply area-level rubric thresholds
    - Ensure hermeticity (no dimension overlap)
    """

    def __init__(
        self,
        monolith: dict[str, Any] | None = None,
        abort_on_insufficient: bool = True,
        aggregation_settings: AggregationSettings | None = None,
    ) -> None:
        """
        Initialize area aggregator.

        Args:
            monolith: Questionnaire monolith configuration (optional, required for run())
            abort_on_insufficient: Whether to abort on insufficient coverage

        Raises:
        """

```

```

        ValueError: If monolith is None and required for operations
    """
    self.monolith = monolith
    self.abort_on_insufficient = abort_on_insufficient
            self.aggregation_settings = aggregation_settings or
AggregationSettings.from_monolith(monolith)
            self.area_group_by_keys = self.aggregation_settings.area_group_by_keys or
[ "area_id" ]

# Extract configuration if monolith provided
if monolith is not None:
    self.scoring_config = monolith["blocks"]["scoring"]
    self.niveles = monolith["blocks"]["niveles_abstraccion"]
    self.policy_areas = self.niveles["policy_areas"]
    self.dimensions = self.niveles["dimensions"]
else:
    self.scoring_config = None
    self.niveles = None
    self.policy_areas = None
    self.dimensions = None

logger.info("AreaPolicyAggregator initialized")

def validate_hermeticity(
    self,
    dimension_scores: list[DimensionScore],
    area_id: str
) -> tuple[bool, str]:
    """
    Validate hermeticity (no dimension overlap/gaps).
    Uses scoped validation based on policy_area.dimension_ids from monolith.

Args:
    dimension_scores: List of dimension scores for the area
    area_id: Policy area ID

Returns:
    Tuple of (is_valid, message)

Raises:
    HermeticityValidationError: If hermeticity is violated
    """
    # Get expected dimensions for this specific policy area
    area_def = next(
        (a for a in self.policy_areas if a["policy_area_id"] == area_id),
        None
    )

    if area_def and "dimension_ids" in area_def:
        expected_dimension_ids = set(area_def["dimension_ids"])
    else:
        # Fallback to all global dimensions if not specified
        expected_dimension_ids = {d["dimension_id"] for d in self.dimensions}

```

```

actual_dimension_ids = {d.dimension_id for d in dimension_scores}
len(expected_dimension_ids)
len(dimension_scores)

# Check for missing dimensions
missing_dims = expected_dimension_ids - actual_dimension_ids
if missing_dims:
    msg = (
        f"Hermeticity violation for area {area_id}: "
        f"missing dimensions {missing_dims}"
    )
    logger.error(msg)
    if self.abort_on_insufficient:
        raise HermeticityValidationError(msg)
    return False, msg

# Check for unexpected dimensions
extra_dims = actual_dimension_ids - expected_dimension_ids
if extra_dims:
    msg = (
        f"Hermeticity violation for area {area_id}: "
        f"unexpected dimensions {extra_dims}"
    )
    logger.error(msg)
    if self.abort_on_insufficient:
        raise HermeticityValidationError(msg)
    return False, msg

# Check for duplicate dimensions
dimension_ids = [d.dimension_id for d in dimension_scores]
if len(dimension_ids) != len(set(dimension_ids)):
    msg = f"Hermeticity violation for area {area_id}: duplicate dimensions found"
    logger.error(msg)
    if self.abort_on_insufficient:
        raise HermeticityValidationError(msg)
    return False, msg

logger.debug(f"Hermeticity validation passed for area {area_id}")
return True, "Hermeticity validated"

@calibrated_method("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores")
def normalize_scores(self, dimension_scores: list[DimensionScore]) -> list[float]:
    """
    Normalize dimension scores to 0-1 range.

    Args:
        dimension_scores: List of dimension scores

    Returns:
        List of normalized scores
    """

```

```

normalized = []
for d in dimension_scores:
    # Extract max_expected from validation_details or default to 3.0
    max_expected = d.validation_details.get('score_max', 3.0) if
d.validation_details else 3.0

normalized.append(max(ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1148_34", 0.0), min(max_expected, d.score)) /
max_expected)

logger.debug(f"Scores normalized: {normalized}")
return normalized

def apply_rubric_thresholds(
    self,
    score: float,
    thresholds: dict[str, float] | None = None
) -> str:
    """
    Apply area-level rubric thresholds.

    Args:
        score: Aggregated score (0-3 range)
        thresholds: Optional threshold definitions (dict with keys: EXCELENTE, BUENO, ACEPTABLE)
            Each value should be a normalized threshold (0-1 range)

    Returns:
        Quality level (EXCELENTE, BUENO, ACEPTABLE, INSUFICIENTE)
    """
    # Clamp score to valid range [0, 3]
    clamped_score = max(ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1170_28", 0.0), min(3.0, score))

    # Normalize to 0-1 range
    normalized_score = clamped_score / 3.0

    # Use provided thresholds or defaults
    if thresholds:
        excellent_threshold = thresholds.get('EXCELENTE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1177_62", 0.85))
        good_threshold = thresholds.get('BUENO',
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1178_53", 0.70))
        acceptable_threshold = thresholds.get('ACEPTABLE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1179_63", 0.55))
    else:
        excellent_threshold = ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "excellent_threshold", 0.85) # Refactored
        good_threshold =

```

```

ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize
_scores", "good_threshold", 0.7) # Refactored
                                                 acceptable_threshold      =
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize
_scores", "acceptable_threshold", 0.55) # Refactored

    # Apply thresholds
    if normalized_score >= excellent_threshold:
        quality = "EXCELENTE"
    elif normalized_score >= good_threshold:
        quality = "BUENO"
    elif normalized_score >= acceptable_threshold:
        quality = "ACEPTABLE"
    else:
        quality = "INSUFICIENTE"

    logger.debug(
        f"Area rubric applied: score={score:.4f}, "
        f"normalized={normalized_score:.4f}, quality={quality}"
    )

    return quality

def aggregate_area(
    self,
    dimension_scores: list[DimensionScore],
    group_by_values: dict[str, Any],
    weights: list[float] | None = None,
) -> AreaScore:
    """
    Aggregate a single policy area from dimension scores.

    Args:
        dimension_scores: List of dimension scores for this area.
        group_by_values: Dictionary of grouping keys and their values.
        weights: Optional list of weights for dimension scores.

    Returns:
        AreaScore with aggregated score and quality level.

    Raises:
        ValidationError: If validation fails.
    """
    area_id = group_by_values.get("area_id", "UNKNOWN")
    logger.info(f"Aggregating policy area {area_id}")

    validation_details = {}

    # The dimension_scores are already grouped.
    area_dim_scores = dimension_scores

    # Validate hermeticity
    try:
        hermetic_valid, hermetic_msg = self.validate_hermeticity(area_dim_scores,

```

```

area_id)
    validation_details["hermeticity"] = {
        "valid": hermetic_valid,
        "message": hermetic_msg,
        "dimension_count": len(area_dim_scores)
    }
except HermeticityValidationError as e:
    logger.error(f"Hermeticity validation failed for area {area_id}: {e}")
    # Get area name
    area_name = next(
        (a["i18n"]["keys"]["label_es"] for a in self.policy_areas
         if a["policy_area_id"] == area_id),
        area_id
    )
    return AreaScore(
        area_id=area_id,
        area_name=area_name,
        score=ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1249_22", 0.0),
        quality_level="INSUFICIENTE",
        dimension_scores=[],
        validation_passed=False,
        validation_details={"error": str(e), "type": "hermeticity"}
    )

if not area_dim_scores:
    logger.warning(f"No dimension scores for area {area_id}")
    area_name = next(
        (a["i18n"]["keys"]["label_es"] for a in self.policy_areas
         if a["policy_area_id"] == area_id),
        area_id
    )
    return AreaScore(
        area_id=area_id,
        area_name=area_name,
        score=ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_scores", "auto_param_L1266_22", 0.0),
        quality_level="INSUFICIENTE",
        dimension_scores=[],
        validation_passed=False,
        validation_details={"error": "No dimensions", "type": "empty"}
    )

# Normalize scores
normalized = self.normalize_scores(area_dim_scores)
validation_details["normalization"] = {
    "original": [d.score for d in area_dim_scores],
    "normalized": normalized
}

# Calculate weighted average score
scores = [d.score for d in area_dim_scores]

```

```

        resolved_weights = weights or self._resolve_area_weights(area_id,
area_dim_scores)
    avg_score = calculate_weighted_average(scores, weights=resolved_weights)

    # Apply rubric thresholds
    quality_level = self.apply_rubric_thresholds(avg_score)
    validation_details["rubric"] = {
        "score": avg_score,
        "quality_level": quality_level
    }

    # Get area name
    area_name = next(
        (a["i18n"]["keys"]["label_es"] for a in self.policy_areas
         if a["policy_area_id"] == area_id),
        area_id
    )

    logger.info(
        f"? Policy area {area_id} ({area_name}): "
        f"score={avg_score:.4f}, quality={quality_level}"
    )

    return AreaScore(
        area_id=area_id,
        area_name=area_name,
        score=avg_score,
        quality_level=quality_level,
        dimension_scores=area_dim_scores,
        validation_passed=True,
        validation_details=validation_details
    )
}

def run(
    self,
    dimension_scores: list[DimensionScore],
    group_by_keys: list[str]
) -> list[AreaScore]:
    """
    Run the area aggregation process.

    Args:
        dimension_scores: List of all dimension scores.
        group_by_keys: List of keys to group by.

    Returns:
        A list of AreaScore objects.
    """
    def key_func(d):
        return tuple(getattr(d, key) for key in group_by_keys)
    grouped_scores = group_by(dimension_scores, key_func)

    area_scores = []
    for group_key, scores in grouped_scores.items():

```

```

group_by_values = dict(zip(group_by_keys, group_key, strict=False))
score = self.aggregate_area(scores, group_by_values, weights=None)
area_scores.append(score)

return area_scores

def _resolve_area_weights(
    self,
    area_id: str,
    dimension_scores: list[DimensionScore],
) -> list[float] | None:
    mapping = self.aggregation_settings.policy_area_dimension_weights.get(area_id)
    if not mapping:
        return None

    weights: list[float] = []
    for dim_score in dimension_scores:
        weight = mapping.get(dim_score.dimension_id)
        if weight is None:
            logger.debug(
                "Missing weight for dimension %s in area %s ? falling back to equal
weights",
                dim_score.dimension_id,
                area_id,
            )
        return None
    weights.append(weight)

    total = sum(weights)
    if total <= 0:
        return None
    return [w / total for w in weights]

class ClusterAggregator:

    """
    Aggregates policy area scores into cluster scores (MESO level).

    Responsibilities:
    - Aggregate multiple area scores per cluster
    - Apply cluster-specific weights
    - Calculate coherence metrics
    - Validate cluster hermeticity
    """

    PENALTY_WEIGHT = ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize
_scores", "PENALTY_WEIGHT", 0.3)
    MAX_SCORE = 3.0

    def __init__(
        self,
        monolith: dict[str, Any] | None = None,
        abort_on_insufficient: bool = True,
        aggregation_settings: AggregationSettings | None = None,
    ):

```

```

) -> None:
    """
    Initialize cluster aggregator.

Args:
    monolith: Questionnaire monolith configuration (optional, required for
run())
    abort_on_insufficient: Whether to abort on insufficient coverage

Raises:
    ValueError: If monolith is None and required for operations
    """
    self.monolith = monolith
    self.abort_on_insufficient = abort_on_insufficient
        self.aggregation_settings = aggregation_settings or
AggregationSettings.from_monolith(monolith)
    self.cluster_group_by_keys = self.aggregation_settings.cluster_group_by_keys or
["cluster_id"]

    # Extract configuration if monolith provided
    if monolith is not None:
        self.scoring_config = monolith["blocks"]["scoring"]
        self.niveles = monolith["blocks"]["niveles_abstraccion"]
        self.clusters = self.niveles["clusters"]
    else:
        self.scoring_config = None
        self.niveles = None
        self.clusters = None

    logger.info("ClusterAggregator initialized")

def validate_cluster_hermeticity(
    self,
    cluster_def: dict[str, Any],
    area_scores: list[AreaScore]
) -> tuple[bool, str]:
    """
    Validate cluster hermeticity.

Args:
    cluster_def: Cluster definition from monolith
    area_scores: List of area scores for this cluster

Returns:
    Tuple of (is_valid, message)

Raises:
    HermeticityValidationError: If hermeticity is violated
    """
    expected_areas = cluster_def.get("policy_area_ids", [])
    actual_areas = [a.area_id for a in area_scores]

    # Check for duplicate areas
    if len(actual_areas) != len(set(actual_areas)):

```

```

        msg = (
            f"Cluster hermeticity violation: "
            f"duplicate areas found for cluster {cluster_def['cluster_id']}"
        )
        logger.error(msg)
        if self.abort_on_insufficient:
            raise HermeticityValidationError(msg)
        return False, msg

    # Check that all expected areas are present
    missing_areas = set(expected_areas) - set(actual_areas)
    if missing_areas:
        msg = (
            f"Cluster hermeticity violation: "
            f"missing areas {missing_areas} for cluster {cluster_def['cluster_id']}"
        )
        logger.error(msg)
        if self.abort_on_insufficient:
            raise HermeticityValidationError(msg)
        return False, msg

    # Check for unexpected areas
    extra_areas = set(actual_areas) - set(expected_areas)
    if extra_areas:
        msg = (
            f"Cluster hermeticity violation: "
            f"unexpected areas {extra_areas} for cluster "
            f"{cluster_def['cluster_id']}"
        )
        logger.error(msg)
        if self.abort_on_insufficient:
            raise HermeticityValidationError(msg)
        return False, msg

    logger.debug(f"Cluster hermeticity validated for {cluster_def['cluster_id']}")
    return True, "Cluster hermeticity validated"

```

```

def apply_cluster_weights(
    self,
    area_scores: list[AreaScore],
    weights: list[float] | None = None
) -> float:
    """
    Apply cluster-specific weights to area scores.

```

Args:

```

    area_scores: List of area scores
    weights: Optional weights (defaults to equal weights)

```

Returns:

```

    Weighted average score

```

Raises:

```

    WeightValidationError: If weights validation fails

```

```

"""
scores = [a.score for a in area_scores]

if weights is None:
    # Equal weights
    weights      =
[ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_
e_scores", "auto_param_L1495_23", 1.0) / len(scores)] * len(scores)

# Validate weights length matches scores length
if len(weights) != len(scores):
    msg = (
        f"Cluster weight length mismatch: "
        f"{len(weights)} weights for {len(scores)} area scores"
    )
    logger.error(msg)
    if self.abort_on_insufficient:
        raise WeightValidationError(msg)

        # Validate weights sum to
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_
_scores", "auto_param_L1507_34", 1.0)
    weight_sum = sum(weights)
    tolerance = 1e-6
        if abs(weight_sum -
ParameterLoaderV2.get("farfan_core.processing.aggregation.AreaPolicyAggregator.normalize_
_scores", "auto_param_L1510_28", 1.0)) > tolerance:
            msg = f"Cluster weight validation failed: sum={weight_sum:.6f}"
            logger.error(msg)
            if self.abort_on_insufficient:
                raise WeightValidationError(msg)

# Calculate weighted average
weighted_avg = sum(s * w for s, w in zip(scores, weights, strict=False))

logger.debug(
    f"Cluster weights applied: scores={scores}, "
    f"weights={weights}, result={weighted_avg:.4f}"
)

return weighted_avg

@calibrated_method("farfan_core.processing.aggregation.ClusterAggregator.analyze_coheren_
ce")
def analyze_coherence(self, area_scores: list[AreaScore]) -> float:
    """
    Analyze cluster coherence.

    Coherence is measured as the inverse of standard deviation.
    Higher coherence means scores are more consistent.

    Args:
        area_scores: List of area scores

```

```

>Returns:
    Coherence value (0-1, where 1 is perfect coherence)
"""

scores = [a.score for a in area_scores]

if len(scores) <= 1:
    return ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1543_19", 1.0)

    # Calculate mean
    mean = sum(scores) / len(scores)

    # Calculate standard deviation
    variance = sum((s - mean) ** 2 for s in scores) / len(scores)
        std_dev      =      variance      **
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1550_30", 0.5)

    # Convert to coherence (inverse relationship)
    # Normalize by max possible std dev (3.0 for 0-3 range)
    max_std = 3.0
        coherence      =
max(ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence",
                           "auto_param_L1555_24", 0.0),
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1555_29", 1.0) - (std_dev / max_std))

logger.debug(
    f"Coherence analysis: mean={mean:.4f}, "
    f"std_dev={std_dev:.4f}, coherence={coherence:.4f}"
)

return coherence

def aggregate_cluster(
    self,
    area_scores: list[AreaScore],
    group_by_values: dict[str, Any],
    weights: list[float] | None = None,
) -> ClusterScore:
    """
Aggregate a single MESO cluster from area scores.

Args:
    area_scores: List of area scores for this cluster.
    group_by_values: Dictionary of grouping keys and their values.
    weights: Optional cluster-specific weights.

>Returns:
    ClusterScore with aggregated score and coherence.

Raises:

```

```

        ValidationError: If validation fails.

cluster_id = group_by_values.get("cluster_id", "UNKNOWN")
logger.info(f"Aggregating cluster {cluster_id}")

validation_details = {}

# Get cluster definition
cluster_def = next(
    (c for c in self.clusters if c["cluster_id"] == cluster_id), None
)

if not cluster_def:
    logger.error(f"Cluster definition not found: {cluster_id}")
    return ClusterScore(
        cluster_id=cluster_id,
        cluster_name=cluster_id,
        areas=[],
        score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1600_22", 0.0),
        coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1601_26", 0.0),
        variance=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1602_25", 0.0),
        weakest_area=None,
        area_scores=[],
        validation_passed=False,
        validation_details={"error": "Definition not found", "type": "config"},
    )

    cluster_name = cluster_def["i18n"]["keys"]["label_es"]
    expected_areas = cluster_def["policy_area_ids"]

    # The area_scores are already grouped.
    cluster_area_scores = area_scores

    # Validate hermeticity
    try:
        hermetic_valid, hermetic_msg = self.validate_cluster_hermeticity(
            cluster_def,
            cluster_area_scores
        )
        validation_details["hermeticity"] = {
            "valid": hermetic_valid,
            "message": hermetic_msg
        }
    except HermeticityValidationError as e:
        logger.error(f"Cluster hermeticity validation failed: {e}")
        return ClusterScore(
            cluster_id=cluster_id,
            cluster_name=cluster_name,
            score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1600_22", 0.0),
            coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1601_26", 0.0),
            variance=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1602_25", 0.0),
            weakest_area=None,
            area_scores=[],
            validation_passed=False,
            validation_details={"error": "Hermeticity validation failed", "type": "config"}
        )

```

```

        areas=expected_areas,

score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1631_22", 0.0),

coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1632_26", 0.0),

variance=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1633_25", 0.0),
    weakest_area=None,
    area_scores=[],
    validation_passed=False,
    validation_details={"error": str(e), "type": "hermeticity"}
)

if not cluster_area_scores:
    logger.warning(f"No area scores for cluster {cluster_id}")
    return ClusterScore(
        cluster_id=cluster_id,
        cluster_name=cluster_name,
        areas=expected_areas,
    )

score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1646_22", 0.0),

coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1647_26", 0.0),

variance=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1648_25", 0.0),
    weakest_area=None,
    area_scores=[],
    validation_passed=False,
    validation_details={"error": "No areas", "type": "empty"}
)

# Apply cluster weights
resolved_weights = weights or self._resolve_cluster_weights(cluster_id,
cluster_area_scores)
try:
    weighted_score = self.apply_cluster_weights(cluster_area_scores,
resolved_weights)
    validation_details["weights"] = {
        "valid": True,
        "weights": resolved_weights if resolved_weights else "equal",
        "score": weighted_score
    }
except WeightValidationError as e:
    logger.error(f"Cluster weight validation failed: {e}")
    return ClusterScore(
        cluster_id=cluster_id,
        cluster_name=cluster_name,
        areas=expected_areas,
    )

```

```

score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1670_22", 0.0),

coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1671_26", 0.0),

variance=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1672_25", 0.0),
    weakest_area=None,
    area_scores=cluster_area_scores,
    validation_passed=False,
    validation_details={"error": str(e), "type": "weights"}
)

# Analyze coherence and variance metrics
coherence = self.analyze_coherence(cluster_area_scores)
scores_array = [a.score for a in cluster_area_scores]
if scores_array:
    mean_score = sum(scores_array) / len(scores_array)
    variance = sum((score - mean_score) ** 2 for score in scores_array) / len(scores_array)
else:
    variance = ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "variance", 0.0) # Refactored
    weakest_area = min(cluster_area_scores, key=lambda a: a.score, default=None)

    std_dev = variance **
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1689_30", 0.5)
    normalized_std = min(std_dev / self.MAX_SCORE,
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1690_55", 1.0)) if std_dev > 0 else
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1690_80", 0.0)
    penalty_factor = ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1691_25", 1.0) - (normalized_std * self.PENALTY_WEIGHT)
    adjusted_score = weighted_score * penalty_factor

    validation_details["coherence"] = {
        "value": coherence,
        "interpretation": "high" if coherence >
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1696_52", 0.8) else "medium" if coherence >
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1696_85", 0.6) else "low"
    }
    validation_details["variance"] = variance
    if weakest_area:
        validation_details["weakest_area"] = weakest_area.area_id
    validation_details["imbalance_penalty"] = {
        "std_dev": std_dev,

```

```

        "penalty_factor": penalty_factor,
        "raw_score": weighted_score,
        "adjusted_score": adjusted_score,
    }

    logger.info(
        f"? Cluster {cluster_id} ({cluster_name}): "
        f"score={adjusted_score:.4f}, coherence={coherence:.4f}"
    )
}

return ClusterScore(
    cluster_id=cluster_id,
    cluster_name=cluster_name,
    areas=expected_areas,
    score=adjusted_score,
    coherence=coherence,
    variance=variance,
    weakest_area=weakest_area.area_id if weakest_area else None,
    area_scores=cluster_area_scores,
    validation_passed=True,
    validation_details=validation_details
)
)

def run(
    self,
    area_scores: list[AreaScore],
    cluster_definitions: list[dict[str, Any]]
) -> list[ClusterScore]:
    """
    Run the cluster aggregation process.

    Args:
        area_scores: List of all area scores.
        cluster_definitions: List of cluster definitions from the monolith.

    Returns:
        A list of ClusterScore objects.
    """

    # Create a mapping from area_id to cluster_id
    area_to_cluster = {}
    for cluster in cluster_definitions:
        for area_id in cluster["policy_area_ids"]:
            area_to_cluster[area_id] = cluster["cluster_id"]

    # Assign cluster_id to each area score
    for score in area_scores:
        score.cluster_id = area_to_cluster.get(score.area_id)

    def key_func(area_score: AreaScore) -> tuple:
        return tuple(getattr(area_score, key) for key in self.cluster_group_by_keys)

    grouped_scores = group_by([s for s in area_scores if hasattr(s, 'cluster_id')], key_func)

```

```

cluster_scores = []
for group_key, scores in grouped_scores.items():
    group_by_values = dict(zip(self.cluster_group_by_keys, group_key,
strict=False))
    score = self.aggregate_cluster(scores, group_by_values)
    cluster_scores.append(score)

return cluster_scores

def _resolve_cluster_weights(
    self,
    cluster_id: str,
    area_scores: list[AreaScore],
) -> list[float] | None:
    mapping = self.aggregation_settings.cluster_policy_area_weights.get(cluster_id)
    if not mapping:
        return None

    weights: list[float] = []
    for area_score in area_scores:
        weight = mapping.get(area_score.area_id)
        if weight is None:
            logger.debug(
                "Missing weight for area %s in cluster %s ? falling back to equal
weights",
                area_score.area_id,
                cluster_id,
            )
            return None
        weights.append(weight)

    total = sum(weights)
    if total <= 0:
        return None
    return [w / total for w in weights]

class MacroAggregator:
    """
    Performs holistic macro evaluation (Q305).

    Responsibilities:
    - Aggregate all cluster scores
    - Calculate cross-cutting coherence
    - Identify systemic gaps
    - Assess strategic alignment
    """
    def __init__(
        self,
        monolith: dict[str, Any] | None = None,
        abort_on_insufficient: bool = True,
        aggregation_settings: AggregationSettings | None = None,
    ) -> None:
        """
        """

```

```

Initialize macro aggregator.

Args:
    monolith: Questionnaire monolith configuration (optional, required for
run())
    abort_on_insufficient: Whether to abort on insufficient coverage

Raises:
    ValueError: If monolith is None and required for operations
"""

self.monolith = monolith
self.abort_on_insufficient = abort_on_insufficient
self.aggregation_settings = aggregation_settings or
AggregationSettings.from_monolith(monolith)

# Extract configuration if monolith provided
if monolith is not None:
    self.scoring_config = monolith["blocks"]["scoring"]
    self.niveles = monolith["blocks"]["niveles_abstraccion"]
else:
    self.scoring_config = None
    self.niveles = None

logger.info("MacroAggregator initialized")

def calculate_cross_cutting_coherence(
    self,
    cluster_scores: list[ClusterScore]
) -> float:
    """
    Calculate cross-cutting coherence across all clusters.

Args:
    cluster_scores: List of cluster scores

Returns:
    Cross-cutting coherence value (0-1)
"""

scores = [c.score for c in cluster_scores]

if len(scores) <= 1:
    return
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1847_19", 1.0)

# Calculate mean
mean = sum(scores) / len(scores)

# Calculate standard deviation
variance = sum((s - mean) ** 2 for s in scores) / len(scores)
std_dev = variance **
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1854_30", 0.5)

```

```

# Convert to coherence
max_std = 3.0

coherence = max(ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence",
                                         "auto_param_L1858_24", 0.0),
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence",
                                         "auto_param_L1858_29", 1.0) - (std_dev / max_std))

logger.debug(
    f"Cross-cutting coherence: mean={mean:.4f}, "
    f"std_dev={std_dev:.4f}, coherence={coherence:.4f}"
)

return coherence


def identify_systemic_gaps(
    self,
    area_scores: list[AreaScore]
) -> list[str]:
    """
    Identify systemic gaps (areas with INSUFICIENTE quality).

    Args:
        area_scores: List of area scores

    Returns:
        List of area names with systemic gaps
    """
    gaps = []
    for area in area_scores:
        if area.quality_level == "INSUFICIENTE":
            gaps.append(area.area_name)
            logger.warning(f"Systemic gap identified: {area.area_name}")

    logger.info(f"Systemic gaps identified: {len(gaps)}")
    return gaps


def assess_strategic_alignment(
    self,
    cluster_scores: list[ClusterScore],
    dimension_scores: list[DimensionScore]
) -> float:
    """
    Assess strategic alignment across all levels.

    Args:
        cluster_scores: List of cluster scores
        dimension_scores: List of dimension scores

    Returns:
        Strategic alignment score (0-1)
    """
    # Calculate average cluster coherence
    cluster_coherence = (

```

```

        sum(c.coherence for c in cluster_scores) / len(cluster_scores)
                                            if      cluster_scores      else
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "auto_param_L1907_35", 0.0)
    )

    # Calculate dimension validation rate
    validated_dims = sum(1 for d in dimension_scores if d.validation_passed)
    validation_rate = validated_dims / len(dimension_scores) if dimension_scores
else
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "auto_param_L1912_90", 0.0)

    # Strategic alignment is weighted combination
                                            alignment      =
(ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coh
erence", "auto_param_L1915_21", 0.6) * cluster_coherence) +
(ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coh
erence", "auto_param_L1915_49", 0.4) * validation_rate)

    logger.debug(
        f"Strategic alignment: cluster_coherence={cluster_coherence:.4f}, "
        f"validation_rate={validation_rate:.4f}, alignment={alignment:.4f}"
    )

    return alignment

def apply_rubric_thresholds(
    self,
    score: float,
    thresholds: dict[str, float] | None = None
) -> str:
    """
    Apply macro-level rubric thresholds.

    Args:
        score: Aggregated macro score (0-3 range)
        thresholds: Optional threshold definitions (dict with keys: EXCELENTE,
BUENO, ACEPTABLE)
            Each value should be a normalized threshold (0-1 range)

    Returns:
        Quality level (EXCELENTE, BUENO, ACEPTABLE, INSUFICIENTE)
    """
    # Clamp score to valid range [0, 3]
                                            clamped_score      =
max(ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_
coherence", "auto_param_L1941_28", 0.0), min(3.0, score))

    # Normalize to 0-1 range
    normalized_score = clamped_score / 3.0

    # Use provided thresholds or defaults
    if thresholds:

```

```

                excellent_threshold = thresholds.get('EXCELENTE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "auto_param_L1948_62", 0.85))
                good_threshold = thresholds.get('BUENO',
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "auto_param_L1949_53", 0.70))
                acceptable_threshold = thresholds.get('ACCEPTABLE',
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "auto_param_L1950_63", 0.55))
            else:
                excellent_threshold =
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "excellent_threshold", 0.85) # Refactored
                good_threshold =
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "good_threshold", 0.7) # Refactored
                acceptable_threshold =
ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_cohe
rence", "acceptable_threshold", 0.55) # Refactored

# Apply thresholds
if normalized_score >= excellent_threshold:
    quality = "EXCELENTE"
elif normalized_score >= good_threshold:
    quality = "BUENO"
elif normalized_score >= acceptable_threshold:
    quality = "ACCEPTABLE"
else:
    quality = "INSUFICIENTE"

logger.debug(
    f"Macro rubric applied: score={score:.4f}, "
    f"normalized={normalized:.4f}, quality={quality}"
)
return quality

def evaluate_macro(
    self,
    cluster_scores: list[ClusterScore],
    area_scores: list[AreaScore],
    dimension_scores: list[DimensionScore]
) -> MacroScore:
    """
    Perform holistic macro evaluation (Q305).
    """

    Args:
        cluster_scores: List of cluster scores (MESO level)
        area_scores: List of area scores
        dimension_scores: List of dimension scores

    Returns:
        MacroScore with holistic evaluation
    """

```

```

logger.info("Performing macro holistic evaluation (Q305)")

validation_details = {}

if not cluster_scores:
    logger.error("No cluster scores available for macro evaluation")
    return MacroScore()

score=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1997_22", 0.0),
    quality_level="INSUFICIENTE",

cross_cutting_coherence=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L1999_40", 0.0),
    systemic_gaps=[ ],

strategic_alignment=ParameterLoaderV2.get("farfan_core.processing.aggregation.ClusterAggregator.analyze_coherence", "auto_param_L2001_36", 0.0),
    cluster_scores=[ ],
    validation_passed=False,
    validation_details={"error": "No clusters", "type": "empty"}
)

# Calculate cross-cutting coherence
cross_cutting_coherence = self.calculate_cross_cutting_coherence(cluster_scores)
validation_details["coherence"] = {
    "value": cross_cutting_coherence,
    "clusters": len(cluster_scores)
}

# Identify systemic gaps
systemic_gaps = self.identify_systemic_gaps(area_scores)
validation_details["gaps"] = {
    "count": len(systemic_gaps),
    "areas": systemic_gaps
}

# Assess strategic alignment
strategic_alignment = self.assess_strategic_alignment(
    cluster_scores,
    dimension_scores
)
validation_details["alignment"] = {
    "value": strategic_alignment
}

# Calculate overall macro score (weighted average of clusters)
macro_score = self._calculate_macro_score(cluster_scores)

# Apply quality rubric
quality_level = self.apply_rubric_thresholds(macro_score)
validation_details["rubric"] = {
    "score": macro_score,
    "quality_level": quality_level
}

```

```

        }

    logger.info(
        f"? Macro evaluation (Q305): score={macro_score:.4f}, "
        f"quality={quality_level}, coherence={cross_cutting_coherence:.4f}, "
        f"alignment={strategic_alignment:.4f}, gaps={len(systemic_gaps)}"
    )

    return MacroScore(
        score=macro_score,
        quality_level=quality_level,
        cross_cutting_coherence=cross_cutting_coherence,
        systemic_gaps=systemic_gaps,
        strategic_alignment=strategic_alignment,
        cluster_scores=cluster_scores,
        validation_passed=True,
        validation_details=validation_details
    )
}

@calibrated_method("farfan_core.processing.aggregation.MacroAggregator._calculate_macro_
score")
def _calculate_macro_score(self, cluster_scores: list[ClusterScore]) -> float:
    weights = self.aggregation_settings.macro_cluster_weights
    if not cluster_scores:
        return ParameterLoaderV2.get("farfan_core.processing.aggregation.MacroAggregator._calculate_ma
cro_score", "auto_param_L2061_19", 0.0)
    if not weights:
        return sum(c.score for c in cluster_scores) / len(cluster_scores)

    resolved_weights: list[float] = []
    for cluster in cluster_scores:
        weight = weights.get(cluster.cluster_id)
        if weight is None:
            logger.debug(
                "Missing macro weight for cluster %s ? falling back to equal
weights",
                cluster.cluster_id,
            )
        return sum(c.score for c in cluster_scores) / len(cluster_scores)
    resolved_weights.append(weight)

    total = sum(resolved_weights)
    if total <= 0:
        return sum(c.score for c in cluster_scores) / len(cluster_scores)

    normalized = [w / total for w in resolved_weights]
    return sum(
        cluster.score * weight
        for cluster, weight in zip(cluster_scores, normalized, strict=False)
    )

```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/aggregation_enhancements.py
```

```
"""
```

## Aggregation Enhancements - Surgical Performance Improvements

This module provides targeted enhancements to the aggregation pipeline:

1. Enhanced provenance tracking with DAG visualization
2. Improved uncertainty quantification with confidence intervals
3. Advanced coherence metrics with dispersion analysis
4. Contract-enforced validation at all levels
5. Performance monitoring and telemetry

Enhancement Windows Identified:

- [EW-001] DimensionAggregator: Add confidence interval tracking
- [EW-002] AreaPolicyAggregator: Enhanced hermeticity with diagnosis
- [EW-003] ClusterAggregator: Adaptive coherence thresholds
- [EW-004] MacroAggregator: Strategic alignment with PAxDIM matrix
- [EW-005] All: Contract integration for dura lex enforcement

```
"""
```

```
from __future__ import annotations

import logging
from dataclasses import dataclass, field
from typing import TYPE_CHECKING, Any

from cross_cutting_infrastructure.contractual.dura_lex.aggregation_contract import (
    create_aggregation_contract,
    AggregationContractViolation,
)

if TYPE_CHECKING:
    from farfan_pipeline.processing.aggregation_provenance import ProvenanceNode

logger = logging.getLogger(__name__)

@dataclass
class ConfidenceInterval:
    """Enhanced confidence interval with provenance."""

    lower_bound: float
    upper_bound: float
    confidence_level: float # e.g., 0.95 for 95% CI
    method: str # "bootstrap", "wilson", "analytical"
    n_samples: int | None = None
    provenance: dict[str, Any] = field(default_factory=dict)

@dataclass
class DispersionMetrics:
    """
    Enhanced dispersion metrics for coherence analysis.

```

```

Metrics:
- coefficient_of_variation: CV = std_dev / mean (relative dispersion)
- dispersion_index: DI = range / MAX_SCORE (normalized spread)
- quartile_coefficient: QC = (Q3 - Q1) / (Q3 + Q1)
- gini_coefficient: Gini for inequality measurement
"""

coefficient_of_variation: float
dispersion_index: float
quartile_coefficient: float
mean: float
std_dev: float
min_val: float
max_val: float
q1: float
median: float
q3: float
scenario: str # "convergence", "moderate", "high_dispersion", "extreme_dispersion"

@dataclass
class HermeticityDiagnosis:
    """Enhanced hermeticity diagnosis with remediation hints."""

    is_hermetic: bool
    missing_ids: set[str]
    extra_ids: set[str]
    duplicate_ids: list[str]
    severity: str # "CRITICAL", "HIGH", "MEDIUM", "LOW"
    remediation_hint: str
    validation_details: dict[str, Any] = field(default_factory=dict)

@dataclass
class StrategicAlignmentMetrics:
"""

Strategic alignment metrics for macro evaluation.

Tracks:
- PA×DIM coverage (60 cells: 10 areas × 6 dimensions)
- Cross-cutting coherence
- Systemic gap identification
- Cluster balance assessment
"""

pa_dim_coverage: dict[tuple[str, str], float]
coverage_rate: float # Percentage of PA×DIM cells covered
cluster_coherence_mean: float
cluster_coherence_std: float
systemic_gaps: list[str]
weakest_dimensions: list[tuple[str, float]]
strongest_dimensions: list[tuple[str, float]]
balance_score: float # 0-1, measures evenness across clusters

```

```

class EnhancedDimensionAggregator:
    """
    [EW-001] Enhanced dimension aggregator with confidence tracking.

    Enhancements:
    - Contract-enforced validation
    - Confidence interval computation
    - Enhanced provenance tracking
    """

    def __init__(self, base_aggregator: Any, enable_contracts: bool = True):
        """
        Initialize enhanced aggregator.

        Args:
            base_aggregator: Base DimensionAggregator instance
            enable_contracts: Whether to enable contract enforcement
        """
        self.base = base_aggregator
        self.enable_contracts = enable_contracts
        self.contract = create_aggregation_contract("dimension") if enable_contracts
        else None

    def aggregate_with_confidence(
        self,
        scores: list[float],
        weights: list[float] | None = None,
        confidence_level: float = 0.95
    ) -> tuple[float, ConfidenceInterval]:
        """
        Aggregate with confidence interval.

        Args:
            scores: Input scores
            weights: Optional weights
            confidence_level: Confidence level (default 0.95)

        Returns:
            Tuple of (aggregated_score, confidence_interval)
        """
        # Contract validation
        if self.contract and weights:
            self.contract.validate_weight_normalization(
                weights,
                context={"method": "aggregate_with_confidence", "n_scores": len(scores)}
            )

        # Use base aggregator for weighted average
        aggregated = self.base.calculate_weighted_average(scores, weights)

        # Contract validation of result
        if self.contract:
            self.contract.validate_score_bounds(

```

```

        aggregated,
        context={ "method": "aggregate_with_confidence" }
    )
    self.contract.validate_convexity(
        aggregated,
        scores,
        context={ "method": "aggregate_with_confidence" }
    )

# Compute confidence interval using Wilson score
# For aggregated scores, use bootstrap approximation
        if      hasattr(self.base,      'bootstrap_aggregator')      and
self.base.bootstrap_aggregator:
    # Use existing bootstrap aggregator
    try:
        from farfan_pipeline.processing.uncertainty_quantification import
aggregate_with_uncertainty
        _, uq_metrics = aggregate_with_uncertainty(
            scores,
            weights or [1.0 / len(scores)] * len(scores),
            method="weighted_average",
            bootstrap_samples=1000
        )
        ci = ConfidenceInterval(
            lower_bound=uq_metrics.confidence_interval_lower,
            upper_bound=uq_metrics.confidence_interval_upper,
            confidence_level=confidence_level,
            method="bootstrap",
            n_samples=1000,
            provenance={"uq_metrics": uq_metrics}
        )
    except Exception as e:
        logger.warning(f"Bootstrap CI failed, using analytical: {e}")
        ci = self._analytical_ci(aggregated, scores, confidence_level)
    else:
        ci = self._analytical_ci(aggregated, scores, confidence_level)

    return aggregated, ci

def _analytical_ci(
    self,
    mean: float,
    scores: list[float],
    confidence_level: float
) -> ConfidenceInterval:
    """Compute analytical confidence interval."""
    import math

    n = len(scores)
    if n <= 1:
        # No variance with single score
        return ConfidenceInterval(
            lower_bound=mean,
            upper_bound=mean,

```

```

        confidence_level=confidence_level,
        method="analytical",
        n_samples=n
    )

    # Calculate standard error
    variance = sum((s - mean) ** 2 for s in scores) / (n - 1)
    std_error = math.sqrt(variance / n)

    # Use t-distribution critical value (approximation: ~2 for 95% CI)
    z_score = 1.96 if confidence_level == 0.95 else 2.576 if confidence_level ==
0.99 else 1.645

    margin = z_score * std_error

    return ConfidenceInterval(
        lower_bound=max(0.0, mean - margin),
        upper_bound=min(3.0, mean + margin),
        confidence_level=confidence_level,
        method="analytical",
        n_samples=n,
        provenance={"std_error": std_error, "z_score": z_score}
    )

```

```
class EnhancedAreaAggregator:
```

```
"""
```

```
[EW-002] Enhanced area aggregator with improved hermeticity diagnosis.
```

**Enhancements:**

- Detailed hermeticity diagnosis with remediation hints
- Contract-enforced validation
- Enhanced provenance

```
"""
```

```
def __init__(self, base_aggregator: Any, enable_contracts: bool = True):
    self.base = base_aggregator
    self.enable_contracts = enable_contracts
    self.contract = create_aggregation_contract("area") if enable_contracts else
```

None

```
def diagnose_hermeticity(
    self,
    actual_dimension_ids: set[str],
    expected_dimension_ids: set[str],
    area_id: str
) -> HermeticityDiagnosis:
    """
    Enhanced hermeticity diagnosis with remediation hints.

```

**Args:**

```
    actual_dimension_ids: Dimensions actually present
    expected_dimension_ids: Dimensions expected per monolith
    area_id: Policy area ID for context
```

```

>Returns:
    HermeticityDiagnosis with detailed analysis
"""

missing = expected_dimension_ids - actual_dimension_ids
extra = actual_dimension_ids - expected_dimension_ids

# Check for duplicates
# (In practice, actual_dimension_ids is a set, so no duplicates possible here)
# This is a placeholder for list-based inputs
duplicates = []

is_hermetic = not (missing or extra or duplicates)

# Determine severity
if missing:
    severity = "CRITICAL"
elif extra:
    severity = "HIGH"
elif duplicates:
    severity = "MEDIUM"
else:
    severity = "LOW"

# Generate remediation hint
if missing:
    remediation = f"Add missing dimensions: {', '.join(sorted(missing))}"
elif extra:
    remediation = f"Remove unexpected dimensions: {', '.join(sorted(extra))}"
elif duplicates:
    remediation = f"Remove duplicate dimensions: {', '.join(duplicates)}"
else:
    remediation = "No action needed - hermeticity validated"

diagnosis = HermeticityDiagnosis(
    is_hermetic=is_hermetic,
    missing_ids=missing,
    extra_ids=extra,
    duplicate_ids=duplicates,
    severity=severity,
    remediation_hint=remediation,
    validation_details={
        "area_id": area_id,
        "expected_count": len(expected_dimension_ids),
        "actual_count": len(actual_dimension_ids)
    }
)

# Contract validation
if self.contract:
    try:
        self.contract.validate_hermeticity(
            actual_dimension_ids,
            expected_dimension_ids,

```

```

        context={"area_id": area_id}
    )
except ValueError:
    # Contract will have logged the violation
    pass

return diagnosis

class EnhancedClusterAggregator:
"""
[EW-003] Enhanced cluster aggregator with adaptive coherence.

Enhancements:
- Adaptive coherence thresholds based on dispersion
- Enhanced dispersion metrics
- Contract-enforced validation
"""

def __init__(self, base_aggregator: Any, enable_contracts: bool = True):
    self.base = base_aggregator
    self.enable_contracts = enable_contracts
    self.contract = create_aggregation_contract("cluster") if enable_contracts else
None

def compute_dispersion_metrics(self, scores: list[float]) -> DispersionMetrics:
"""
Compute enhanced dispersion metrics.

Args:
    scores: List of scores to analyze

Returns:
    DispersionMetrics with comprehensive analysis
"""

import statistics

if not scores:
    return DispersionMetrics(
        coefficient_of_variation=0.0,
        dispersion_index=0.0,
        quartile_coefficient=0.0,
        mean=0.0,
        std_dev=0.0,
        min_val=0.0,
        max_val=0.0,
        q1=0.0,
        median=0.0,
        q3=0.0,
        scenario="convergence"
    )

sorted_scores = sorted(scores)
n = len(scores)

```

```

mean = sum(scores) / n
variance = sum((s - mean) ** 2 for s in scores) / n if n > 1 else 0.0
std_dev = variance ** 0.5

# Coefficient of Variation
cv = std_dev / mean if mean > 0 else 0.0

# Dispersion Index
min_val = min(scores)
max_val = max(scores)
dispersion_index = (max_val - min_val) / 3.0 if 3.0 > 0 else 0.0

# Quartiles
q1 = statistics.quantiles(sorted_scores, n=4)[0] if n >= 4 else sorted_scores[0]
median = statistics.median(sorted_scores)
q3 = statistics.quantiles(sorted_scores, n=4)[2] if n >= 4 else sorted_scores[-1]

# Quartile Coefficient of Dispersion
qc = (q3 - q1) / (q3 + q1) if (q3 + q1) > 0 else 0.0

# Classify scenario
if cv < 0.15:
    scenario = "convergence"
elif cv < 0.40:
    scenario = "moderate"
elif cv < 0.60:
    scenario = "high_dispersion"
else:
    scenario = "extreme_dispersion"

return DispersionMetrics(
    coefficient_of_variation=cv,
    dispersion_index=dispersion_index,
    quartile_coefficient=qc,
    mean=mean,
    std_dev=std_dev,
    min_val=min_val,
    max_val=max_val,
    q1=q1,
    median=median,
    q3=q3,
    scenario=scenario
)

def adaptive_penalty(self, dispersion: DispersionMetrics) -> float:
    """
    Calculate adaptive penalty based on dispersion scenario.

    Enhancement from audit: Replaces fixed PENALTY_WEIGHT=0.3 with
    adaptive mechanism that responds to dispersion patterns.
    """

```

Args:

```

        dispersion: Dispersion metrics

    Returns:
        Adaptive penalty multiplier (0.0-1.0)
    """
base_penalty = 0.3

    if dispersion.scenario == "convergence":
        multiplier = 0.5
    elif dispersion.scenario == "moderate":
        multiplier = 1.0
    elif dispersion.scenario == "high_dispersion":
        multiplier = 1.5
    else: # extreme_dispersion
        multiplier = 2.0

    # Round for deterministic/contract-friendly comparisons (tests assert exact
decimals).
    return round(base_penalty * multiplier, 2)

```

```

class EnhancedMacroAggregator:
    """
    [EW-004] Enhanced macro aggregator with strategic alignment.

    Enhancements:
    - PAxDIM matrix coverage tracking (60 cells)
    - Strategic alignment metrics
    - Cluster balance assessment
    - Contract-enforced validation
    """
    def __init__(self, base_aggregator: Any, enable_contracts: bool = True):
        self.base = base_aggregator
        self.enable_contracts = enable_contracts
        self.contract = create_aggregation_contract("macro") if enable_contracts else
None

```

```

def compute_strategic_alignment(
    self,
    cluster_scores: list[Any],
    area_scores: list[Any],
    dimension_scores: list[Any]
) -> StrategicAlignmentMetrics:
    """
    Compute strategic alignment metrics with PAxDIM coverage.

```

Args:

- cluster\_scores: List of ClusterScore objects
- area\_scores: List of AreaScore objects
- dimension\_scores: List of DimensionScore objects

Returns:

StrategicAlignmentMetrics with comprehensive analysis

```

"""
# Build PAxDIM coverage matrix
pa_dim_coverage: dict[tuple[str, str], float] = {}

for dim in dimension_scores:
    area_id = getattr(dim, 'policy_area_id', getattr(dim, 'area_id', 'UNKNOWN'))
    dim_id = getattr(dim, 'dimension_id', 'UNKNOWN')
    score = getattr(dim, 'score', 0.0)

    pa_dim_coverage[(area_id, dim_id)] = score

# Calculate coverage rate (expected: 60 cells)
expected_cells = 60 # 10 areas x 6 dimensions
actual_cells = len(pa_dim_coverage)
coverage_rate = actual_cells / expected_cells if expected_cells > 0 else 0.0

# Cluster coherence statistics
cluster_coherences = [getattr(c, 'coherence', 0.0) for c in cluster_scores]
cluster_coherence_mean = sum(cluster_coherences) / len(cluster_coherences) if cluster_coherences else 0.0

variance = sum((c - cluster_coherence_mean) ** 2 for c in cluster_coherences) / len(cluster_coherences) if len(cluster_coherences) > 1 else 0.0
cluster_coherence_std = variance ** 0.5

# Identify systemic gaps (areas with insufficient quality)
systemic_gaps = []
for area in area_scores:
    quality = getattr(area, 'quality_level', 'UNKNOWN')
    if quality == "INSUFICIENTE":
        area_name = getattr(area, 'area_name', getattr(area, 'area_id', 'UNKNOWN'))
    systemic_gaps.append(area_name)

# Identify weakest and strongest dimensions (aggregate across areas)
dimension_aggregates: dict[str, list[float]] = {}
for dim in dimension_scores:
    dim_id = getattr(dim, 'dimension_id', 'UNKNOWN')
    score = getattr(dim, 'score', 0.0)
    if dim_id not in dimension_aggregates:
        dimension_aggregates[dim_id] = []
    dimension_aggregates[dim_id].append(score)

dimension_means = {
    dim_id: sum(scores) / len(scores) if scores else 0.0
    for dim_id, scores in dimension_aggregates.items()
}

sorted_dims = sorted(dimension_means.items(), key=lambda x: x[1])
weakest = sorted_dims[:3] # Bottom 3
strongest = sorted_dims[-3:] # Top 3

# Calculate balance score (inverse of cluster score standard deviation)
cluster_scores_values = [getattr(c, 'score', 0.0) for c in cluster_scores]

```

```

        if len(cluster_scores_values) > 1:
            cluster_mean = sum(cluster_scores_values) / len(cluster_scores_values)
            cluster_variance = sum((s - cluster_mean) ** 2 for s in
cluster_scores_values) / len(cluster_scores_values)
            cluster_std = cluster_variance ** 0.5
            # Normalize: 1.0 = perfect balance (std=0), 0.0 = max imbalance (std=3.0)
            balance_score = max(0.0, 1.0 - (cluster_std / 3.0))
        else:
            balance_score = 1.0

    return StrategicAlignmentMetrics(
        pa_dim_coverage=pa_dim_coverage,
        coverage_rate=coverage_rate,
        cluster_coherence_mean=cluster_coherence_mean,
        cluster_coherence_std=cluster_coherence_std,
        systemic_gaps=systemic_gaps,
        weakest_dimensions=weakest,
        strongest_dimensions=strongest,
        balance_score=balance_score
    )

# Utility function to enhance existing aggregators
def enhance_aggregator(aggregator: Any, level: str, enable_contracts: bool = True) ->
Any:
    """
    Factory function to wrap existing aggregators with enhancements.

    Args:
        aggregator: Base aggregator instance
        level: Aggregation level (dimension, area, cluster, macro)
        enable_contracts: Whether to enable contract enforcement

    Returns:
        Enhanced aggregator instance

    Raises:
        ValueError: If level is invalid
    """
    enhancers = {
        "dimension": EnhancedDimensionAggregator,
        "area": EnhancedAreaAggregator,
        "cluster": EnhancedClusterAggregator,
        "macro": EnhancedMacroAggregator,
    }

    if level.lower() not in enhancers:
        raise ValueError(f"Invalid aggregation level: {level}. Must be one of
{list(enhancers.keys())}")

    return enhancers[level.lower()](aggregator, enable_contracts=enable_contracts)

```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/aggregation_provenance.py
```

```
"""
```

## Aggregation Provenance System - DAG-based Lineage Tracking

This module provides full provenance tracking for the hierarchical aggregation pipeline. It implements W3C PROV-compliant directed acyclic graphs (DAGs) to capture:

- Data lineage: Which micro-questions contributed to which macro-scores
- Operation tracking: How aggregation operations transformed data
- Sensitivity analysis: Which inputs have highest impact on outputs
- Counterfactual reasoning: What-if analysis for policy decisions

### Architecture:

- ProvenanceNode: Immutable record of a score at any aggregation level
- AggregationDAG: NetworkX-based graph with attribution methods
- SHAPAttribution: Shapley value computation for feature importance

### References:

- W3C PROV: <https://www.w3.org/TR/prov-overview/>
- Shapley values: Shapley, L.S. (1953). "A value for n-person games"
- NetworkX: Hagberg et al. (2008). "Exploring network structure, dynamics, and function"

```
"""
```

```
from __future__ import annotations
```

```
import hashlib
import json
import logging
from collections import defaultdict
from dataclasses import asdict, dataclass, field
from datetime import datetime, timezone
from typing import Any
```

```
import networkx as nx
import numpy as np
```

```
logger = logging.getLogger(__name__)
```

```
@dataclass(frozen=True)
```

```
class ProvenanceNode:
```

```
"""
```

```
    """
```

```
    Immutable provenance record for a score node in the aggregation DAG.
```

#### Attributes:

```
    node_id: Unique identifier (e.g., "Q001", "DIM01_PA05", "CLUSTER_MESO_1")
    level: Abstraction level in hierarchy
    score: Numeric score value
    quality_level: Quality classification (EXCELENTE, BUENO, etc.)
    timestamp: ISO timestamp of computation
    metadata: Additional context (weights, confidence, etc.)
```

```
"""
```

```
node_id: str
```

```
level: str # "micro", "dimension", "area", "cluster", "macro"
```

```

score: float
quality_level: str
    timestamp: str      = field(default_factory=lambda:
datetime.now(timezone.utc).isoformat())
metadata: dict[str, Any] = field(default_factory=dict)

def to_dict(self) -> dict[str, Any]:
    """Convert to dictionary for serialization."""
    return asdict(self)

def compute_hash(self) -> str:
    """Compute deterministic hash for reproducibility."""
    content = json.dumps(
        {
            "node_id": self.node_id,
            "level": self.level,
            "score": self.score,
            "quality_level": self.quality_level,
        },
        sort_keys=True,
    )
    return hashlib.sha256(content.encode("utf-8")).hexdigest()[:16]

```

```

@dataclass
class AggregationEdge:
    """
    Edge in the provenance DAG representing an aggregation operation.

    Attributes:
        source_id: Parent node contributing to aggregation
        target_id: Child node receiving aggregated score
        operation: Type of aggregation (weighted_average, choquet, max, etc.)
        weight: Contribution weight (for weighted operations)
        timestamp: When this edge was created
    """

    source_id: str
    target_id: str
    operation: str
    weight: float
        timestamp: str      = field(default_factory=lambda:
datetime.now(timezone.utc).isoformat())
    metadata: dict[str, Any] = field(default_factory=dict)

```

```

class AggregationDAG:
    """
    Directed Acyclic Graph for full provenance tracking of aggregation pipeline.

```

This class maintains the complete lineage of how micro-question scores propagate through dimension ? area ? cluster ? macro aggregation.

Features:

- Cycle detection (enforces DAG property)

```

- Topological sorting for dependency resolution
- Shapley value attribution for sensitivity analysis
- GraphML export for visualization in Gephi/Cytoscape
- W3C PROV-compliant serialization
"""

def __init__(self):
    """Initialize empty DAG."""
    self.graph: nx.DiGraph = nx.DiGraph()
    self.nodes: dict[str, ProvenanceNode] = {}
    self.edges: list[AggregationEdge] = []
    logger.info("AggregationDAG initialized")

def add_node(self, node: ProvenanceNode) -> None:
    """
    Add a provenance node to the DAG.

    Args:
        node: ProvenanceNode to add

    Raises:
        ValueError: If node_id already exists
    """
    if node.node_id in self.nodes:
        logger.warning(f"Node {node.node_id} already exists, skipping")
        return

    self.nodes[node.node_id] = node
    self.graph.add_node(
        node.node_id,
        level=node.level,
        score=node.score,
        quality=node.quality_level,
        timestamp=node.timestamp,
        metadata=node.metadata,
    )
    logger.debug(f"Added node {node.node_id} (level={node.level}, score={node.score:.2f})")

def add_aggregation_edge(
    self,
    source_ids: list[str],
    target_id: str,
    operation: str,
    weights: list[float],
    metadata: dict[str, Any] | None = None,
) -> None:
    """
    Record an aggregation operation: sources ? target.

    Args:
        source_ids: List of source node IDs (e.g., micro-questions)
        target_id: Target node ID (e.g., dimension score)
        operation: Aggregation type (weighted_average, choquet, etc.)
    """

```

```

weights: Contribution weights for each source
metadata: Additional operation metadata

Raises:
    ValueError: If weights don't match sources or cycle detected
"""

if len(source_ids) != len(weights):
    raise ValueError(
        f"Mismatch: {len(source_ids)} sources but {len(weights)} weights"
    )

if metadata is None:
    metadata = {}

for source_id, weight in zip(source_ids, weights):
    if source_id not in self.graph:
        logger.warning(f"Source node {source_id} not found, adding placeholder")
        self.graph.add_node(source_id, level="unknown", score=0.0)

    if target_id not in self.graph:
        logger.warning(f"Target node {target_id} not found, adding placeholder")
        self.graph.add_node(target_id, level="unknown", score=0.0)

    edge = AggregationEdge(
        source_id=source_id,
        target_id=target_id,
        operation=operation,
        weight=weight,
        metadata=metadata,
    )
    self.edges.append(edge)

    self.graph.add_edge(
        source_id,
        target_id,
        operation=operation,
        weight=weight,
        timestamp=edge.timestamp,
        metadata=metadata,
    )

# Verify DAG property (no cycles)
if not nx.is_directed_acyclic_graph(self.graph):
    # Rollback last edges
    for source_id in source_ids:
        if self.graph.has_edge(source_id, target_id):
            self.graph.remove_edge(source_id, target_id)
    raise ValueError(f"Adding edges {source_ids} ? {target_id} would create a
cycle")

logger.info(
    f"Added aggregation: {len(source_ids)} sources ? {target_id} "
    f"(operation={operation})"
)

```

```

def trace_lineage(self, target_id: str) -> dict[str, Any]:
    """
    Trace complete lineage of a target node.

    Returns all ancestor nodes, the aggregation path, and sensitivity metrics.

    Args:
        target_id: Node ID to trace

    Returns:
        Dictionary with:
        - ancestors: Set of all ancestor node IDs
        - path: Topologically sorted path from sources to target
        - depth: Maximum path length from any micro-question
        - breadth: Number of unique micro-questions contributing

    Raises:
        ValueError: If target_id not in graph
    """
    if target_id not in self.graph:
        raise ValueError(f"Node {target_id} not found in DAG")

    ancestors = nx.ancestors(self.graph, target_id)
    subgraph = self.graph.subgraph(ancestors | {target_id})

    # Compute metrics
    depth = nx.dag_longest_path_length(subgraph) if subgraph.nodes else 0

    # Count micro-questions (level="micro")
    micro_nodes = [
        n for n in ancestors
        if self.graph.nodes[n].get("level") == "micro"
    ]

    # Get topological path
    topo_path = list(nx.topological_sort(subgraph))

    return {
        "target_id": target_id,
        "ancestor_count": len(ancestors),
        "ancestors": sorted(ancestors),
        "topological_path": topo_path,
        "depth": depth,
        "micro_question_count": len(micro_nodes),
        "micro_questions": sorted(micro_nodes),
    }

def compute_shapley_attribution(self, target_id: str) -> dict[str, float]:
    """
    Compute Shapley values for feature attribution.

    Shapley values represent the marginal contribution of each source node
    to the target score, accounting for all possible coalitions.

```

```

Args:
    target_id: Node to attribute

Returns:
    Dictionary mapping source node IDs to Shapley values (sum = target score)

Note:
    This is an exact computation for weighted averages. For non-linear
    aggregations (Choquet), this uses kernel SHAP approximation.

"""
if target_id not in self.graph:
    raise ValueError(f"Node {target_id} not found in DAG")

# Get direct predecessors (sources)
predecessors = list(self.graph.predecessors(target_id))

if not predecessors:
    logger.warning(f"Node {target_id} has no predecessors")
    return {}

# Extract weights and scores
weights = []
scores = []
for pred in predecessors:
    edge_data = self.graph.get_edge_data(pred, target_id)
    weights.append(edge_data.get("weight", 0.0))
    scores.append(self.graph.nodes[pred].get("score", 0.0))

weights = np.array(weights)
scores = np.array(scores)

# For weighted average, Shapley value = weight × score
# This is exact because weighted average is a linear function
shapley_values = weights * scores

# Normalize to sum to target score
target_score = self.graph.nodes[target_id].get("score", 0.0)
if np.sum(shapley_values) > 0:
    shapley_values = shapley_values * (target_score / np.sum(shapley_values))

attribution = {
    pred: float(shap_val)
    for pred, shap_val in zip(predecessors, shapley_values)
}

logger.debug(
    f"Shapley attribution for {target_id}: "
    f"{len(attribution)} sources, sum={sum(attribution.values()):.4f}"
)

return attribution

def get_critical_path(self, target_id: str, top_k: int = 5) -> list[tuple[str,

```

```

float]]:

"""
Identify the most critical source nodes for a target.

Uses Shapley values to rank sources by importance.

Args:
    target_id: Target node
    top_k: Number of top sources to return

Returns:
    List of (node_id, shapley_value) tuples, sorted by importance
"""

attribution = self.compute_shapley_attribution(target_id)

# Sort by absolute Shapley value (handle negative contributions)
sorted_attribution = sorted(
    attribution.items(),
    key=lambda x: abs(x[1]),
    reverse=True,
)

return sorted_attribution[:top_k]

def export_graphml(self, path: str) -> None:
    """
    Export DAG to GraphML format for visualization.

    Compatible with:
    - Gephi: https://gephi.org/
    - Cytoscape: https://cytoscape.org/
    - yEd: https://www.yworks.com/products/yed
    """

    Args:
        path: Output file path (e.g., "aggregation_dag.graphml")
    """

    # Create a copy with serialized dict attributes (GraphML doesn't support dicts)
    g_copy = self.graph.copy()
    for node, data in g_copy.nodes(data=True):
        if "metadata" in data and isinstance(data["metadata"], dict):
            data["metadata_json"] = json.dumps(data["metadata"])
            del data["metadata"]
    for u, v, data in g_copy.edges(data=True):
        if "metadata" in data and isinstance(data["metadata"], dict):
            data["metadata_json"] = json.dumps(data["metadata"])
            del data["metadata"]
        if "weights" in data and isinstance(data["weights"], list):
            data["weights_json"] = json.dumps(data["weights"])
            del data["weights"]
    nx.write_graphml(g_copy, path)
    logger.info(
        f"Exported DAG to {path}: "
        f"{self.graph.number_of_nodes()} nodes, "
        f"{self.graph.number_of_edges()} edges"
    )

```

```

    )

def export_prov_json(self, path: str) -> None:
    """
    Export to W3C PROV-JSON format.

    Spec: https://www.w3.org/Submission/prov-json/

    Args:
        path: Output file path (e.g., "provenance.json")
    """
    prov_doc = {
        "prefix": {
            "prov": "http://www.w3.org/ns/prov#",
            "farfan": "http://farfan.org/ns/aggregation#",
        },
        "entity": {},
        "activity": {},
        "wasGeneratedBy": {},
        "used": {},
    }

    # Entities: All nodes
    for node_id, node_data in self.graph.nodes(data=True):
        prov_doc["entity"][node_id] = {
            "prov:type": "farfan:ScoreEntity",
            "farfan:level": node_data.get("level"),
            "farfan:score": node_data.get("score"),
            "farfan:quality": node_data.get("quality"),
            "prov:generatedAtTime": node_data.get("timestamp"),
        }

    # Activities: Aggregation operations
    for idx, edge in enumerate(self.edges):
        activity_id = f"agg_{idx}"
        prov_doc["activity"][activity_id] = {
            "prov:type": f"farfan:{edge.operation}",
            "farfan:weight": edge.weight,
            "prov:startedAtTime": edge.timestamp,
        }

        # wasGeneratedBy: target was generated by this activity
        if edge.target_id not in prov_doc["wasGeneratedBy"]:
            prov_doc["wasGeneratedBy"][edge.target_id] = []
        prov_doc["wasGeneratedBy"][edge.target_id].append(activity_id)

        # used: activity used source
        if activity_id not in prov_doc["used"]:
            prov_doc["used"][activity_id] = []
        prov_doc["used"][activity_id].append(edge.source_id)

    with open(path, "w", encoding="utf-8") as f:
        json.dump(prov_doc, f, indent=2, ensure_ascii=False)

```

```

logger.info(f"Exported PROV-JSON to {path}")

def get_statistics(self) -> dict[str, Any]:
    """
    Get DAG statistics for monitoring and validation.

    Returns:
        Dictionary with graph metrics
    """
    return {
        "node_count": self.graph.number_of_nodes(),
        "edge_count": self.graph.number_of_edges(),
        "max_depth": nx.dag_longest_path_length(self.graph) if self.graph.nodes else
0,
        "is_dag": nx.is_directed_acyclic_graph(self.graph),
        "weakly_connected_components": nx.number_weakly_connected_components(self.graph),
        "nodes_by_level": self._count_by_level(),
    }

def _count_by_level(self) -> dict[str, int]:
    """Count nodes by abstraction level."""
    counts = defaultdict(int)
    for node_id in self.graph.nodes:
        level = self.graph.nodes[node_id].get("level", "unknown")
        counts[level] += 1
    return dict(counts)

__all__ = [
    "ProvenanceNode",
    "AggregationEdge",
    "AggregationDAG",
]

```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/aggregation_validation.py
```

```
"""
```

```
Aggregation Validation Module - Hard Validation for Phases 4-7
```

```
This module provides strict validation to ensure aggregation pipeline  
produces non-trivial, traceable results. Prevents silent failures.
```

```
Requirements from Issue #[P0]:
```

- Fail hard if any phase returns empty or cannot be traced to source micro-questions
- Non-zero macro score is required for proper inputs
- All phases must chain outputs to inputs

```
"""
```

```
from __future__ import annotations
```

```
from dataclasses import dataclass  
from typing import TYPE_CHECKING
```

```
if TYPE_CHECKING:
```

```
    from .aggregation import (  
        DimensionScore,  
        AreaScore,  
        ClusterScore,  
        MacroScore,  
        ScoredResult,  
)
```

```
@dataclass
```

```
class ValidationResult:
```

```
    """Result of aggregation validation."""
```

```
    passed: bool  
    phase: str  
    error_message: str = ""  
    details: dict[str, any] = None
```

```
class AggregationValidationError(Exception):
```

```
    """Raised when aggregation validation fails."""
```

```
pass
```

```
def validate_phase4_output(
```

```
    dimension_scores: list["DimensionScore"],  
    input_scored_results: list["ScoredResult"],
```

```
) -> ValidationResult:
```

```
    """
```

```
    Validate Phase 4 (Dimension Aggregation) output.
```

```
Requirements:
```

- Must produce non-empty dimension scores

- Each dimension score must trace to source micro questions
- Score values must be non-negative and within valid range [0, 3]

Args:

```
dimension_scores: Output from Phase 4
input_scored_results: Input to Phase 4 (for traceability check)
```

Returns:

```
ValidationResult with pass/fail status
```

Raises:

```
AggregationValidationError: If validation fails in strict mode
```

"""

```
if not dimension_scores:
    return ValidationResult(
        passed=False,
        phase="Phase 4 (Dimension Aggregation)",
        error_message="Dimension aggregation returned EMPTY list. No dimensions were
aggregated.",
        details={"input_count": len(input_scored_results), "output_count": 0}
    )
```

# Check traceability

```
non_traceable = []
```

```
invalid_scores = []
```

for dim\_score in dimension\_scores:

# Check traceability

```
if not dim_score.contributing_questions or len(dim_score.contributing_questions)
```

== 0:

```
    non_traceable.append(f"{dim_score.dimension_id}/{dim_score.area_id}")
```

# Check score validity

```
if dim_score.score < 0 or dim_score.score > 3:
```

```
    invalid_scores.append((f"{dim_score.dimension_id}/{dim_score.area_id}", dim_score.score))
```

if non\_traceable:

```
    return ValidationResult(
```

passed=False,

phase="Phase 4 (Dimension Aggregation)",

```
        error_message=f"Dimension scores not traceable to source micro questions:
{non_traceable[:5]}",
        details={"non_traceable_count": len(non_traceable), "examples":
```

```
non_traceable[:5]}
```

)

if invalid\_scores:

```
    return ValidationResult(
```

passed=False,

phase="Phase 4 (Dimension Aggregation)",

```
        error_message=f"Invalid dimension scores (outside [0, 3] range):
{invalid_scores[:5]}",
        details={"invalid_count": len(invalid_scores), "examples":
```

```
invalid_scores[:5]}
```

```

invalid_scores[:5]
)

return ValidationResult(
    passed=True,
    phase="Phase 4 (Dimension Aggregation)",
    details={
        "dimension_count": len(dimension_scores),
        "input_count": len(input_scored_results),
        "traceable": True
    }
)
)

def validate_phase5_output(
    area_scores: list["AreaScore"],
    input_dimension_scores: list["DimensionScore"],
) -> ValidationResult:
    """
    Validate Phase 5 (Area Policy Aggregation) output.

    Requirements:
    - Must produce non-empty area scores
    - Each area score must trace to dimension scores
    - Score values must be non-negative and within valid range [0, 3]

    Args:
        area_scores: Output from Phase 5
        input_dimension_scores: Input to Phase 5

    Returns:
        ValidationResult with pass/fail status
    """
    if not area_scores:
        return ValidationResult(
            passed=False,
            phase="Phase 5 (Area Policy Aggregation)",
            error_message="Area aggregation returned EMPTY list. No policy areas were aggregated.",
            details={"input_count": len(input_dimension_scores), "output_count": 0}
        )

    # Check traceability
    non_traceable = []
    invalid_scores = []

    for area_score in area_scores:
        # Check traceability
        if not area_score.dimension_scores or len(area_score.dimension_scores) == 0:
            non_traceable.append(area_score.area_id)

        # Check score validity
        if area_score.score < 0 or area_score.score > 3:
            invalid_scores.append((area_score.area_id, area_score.score))

```

```

if non_traceable:
    return ValidationResult(
        passed=False,
        phase="Phase 5 (Area Policy Aggregation)",
        error_message=f"Area scores not traceable to dimension scores: {non_traceable[:5]}",
        details={"non_traceable_count": len(non_traceable), "examples": non_traceable[:5]}
    )

if invalid_scores:
    return ValidationResult(
        passed=False,
        phase="Phase 5 (Area Policy Aggregation)",
        error_message=f"Invalid area scores (outside [0, 3] range): {invalid_scores[:5]}",
        details={"invalid_count": len(invalid_scores), "examples": invalid_scores[:5]}
    )

return ValidationResult(
    passed=True,
    phase="Phase 5 (Area Policy Aggregation)",
    details={
        "area_count": len(area_scores),
        "input_count": len(input_dimension_scores),
        "traceable": True
    }
)

```

```

def validate_phase6_output(
    cluster_scores: list["ClusterScore"],
    input_area_scores: list["AreaScore"],
) -> ValidationResult:
    """
    Validate Phase 6 (Cluster Aggregation) output.
    """

```

Requirements:

- Must produce non-empty cluster scores
- Each cluster score must trace to area scores
- Score values must be non-negative and within valid range [0, 3]

Args:

```

    cluster_scores: Output from Phase 6
    input_area_scores: Input to Phase 6

```

Returns:

```
    ValidationResult with pass/fail status
    """

```

```

if not cluster_scores:
    return ValidationResult(
        passed=False,

```

```

        phase="Phase 6 (Cluster Aggregation)",
        error_message="Cluster aggregation returned EMPTY list. No clusters were
aggregated.",
        details={"input_count": len(input_area_scores), "output_count": 0}
    )

# Check traceability
non_traceable = []
invalid_scores = []

for cluster_score in cluster_scores:
    # Check traceability
    if not cluster_score.area_scores or len(cluster_score.area_scores) == 0:
        non_traceable.append(cluster_score.cluster_id)

    # Check score validity
    if cluster_score.score < 0 or cluster_score.score > 3:
        invalid_scores.append((cluster_score.cluster_id, cluster_score.score))

if non_traceable:
    return ValidationResult(
        passed=False,
        phase="Phase 6 (Cluster Aggregation)",
        error_message=f"Cluster scores not traceable to area scores:
{non_traceable[:5]}",
        details={"non_traceable_count": len(non_traceable), "examples":
non_traceable[:5]}
    )

if invalid_scores:
    return ValidationResult(
        passed=False,
        phase="Phase 6 (Cluster Aggregation)",
        error_message=f"Invalid cluster scores (outside [0, 3] range):
{invalid_scores[:5]}",
        details={"invalid_count": len(invalid_scores), "examples":
invalid_scores[:5]}
    )

return ValidationResult(
    passed=True,
    phase="Phase 6 (Cluster Aggregation)",
    details={
        "cluster_count": len(cluster_scores),
        "input_count": len(input_area_scores),
        "traceable": True
    }
)

```

```

def validate_phase7_output(
    macro_score: "MacroScore",
    input_cluster_scores: list["ClusterScore"],
    input_area_scores: list["AreaScore"],

```

```

    input_dimension_scores: list["DimensionScore"],
) -> ValidationResult:
"""
Validate Phase 7 (Macro Evaluation) output.

Requirements:
- Macro score must be non-zero for valid inputs
- Must trace to cluster scores
- Score values must be non-negative and within valid range [0, 3]
- Cross-cutting coherence must be in [0, 1]
- Strategic alignment must be in [0, 1]

Args:
    macro_score: Output from Phase 7
    input_cluster_scores: Input cluster scores
    input_area_scores: Input area scores
    input_dimension_scores: Input dimension scores

Returns:
    ValidationResult with pass/fail status
"""

# Check if macro score is zero when inputs are non-empty and non-zero
has_valid_inputs = (
    input_cluster_scores
    and input_area_scores
    and input_dimension_scores
    and any(cs.score > 0 for cs in input_cluster_scores)
)

if has_valid_inputs and macro_score.score == 0:
    return ValidationResult(
        passed=False,
        phase="Phase 7 (Macro Evaluation)",
        error_message=(
            "Macro score is ZERO despite valid non-zero inputs. "
            "This indicates a broken aggregation chain."
        ),
        details={
            "macro_score": macro_score.score,
            "cluster_count": len(input_cluster_scores),
            "area_count": len(input_area_scores),
            "dimension_count": len(input_dimension_scores),
            "non_zero_clusters": sum(1 for cs in input_cluster_scores if cs.score >
0)
        }
    )

# Check traceability
if not macro_score.cluster_scores or len(macro_score.cluster_scores) == 0:
    return ValidationResult(
        passed=False,
        phase="Phase 7 (Macro Evaluation)",
        error_message="Macro score not traceable to cluster scores (empty
cluster_scores list).",
    )

```

```

        details={"macro_score": macro_score.score, "cluster_scores_count": 0}
    )

# Check score validity
if macro_score.score < 0 or macro_score.score > 3:
    return ValidationResult(
        passed=False,
        phase="Phase 7 (Macro Evaluation)",
        error_message=f"Invalid macro score (outside [0, 3] range): {macro_score.score}",
        details={"macro_score": macro_score.score}
    )

# Check coherence and alignment ranges
if not (0 <= macro_score.cross_cutting_coherence <= 1):
    return ValidationResult(
        passed=False,
        phase="Phase 7 (Macro Evaluation)",
        error_message=f"Invalid cross-cutting coherence (outside [0, 1] range): {macro_score.cross_cutting_coherence}",
        details={"coherence": macro_score.cross_cutting_coherence}
    )

if not (0 <= macro_score.strategic_alignment <= 1):
    return ValidationResult(
        passed=False,
        phase="Phase 7 (Macro Evaluation)",
        error_message=f"Invalid strategic alignment (outside [0, 1] range): {macro_score.strategic_alignment}",
        details={"alignment": macro_score.strategic_alignment}
    )

return ValidationResult(
    passed=True,
    phase="Phase 7 (Macro Evaluation)",
    details={
        "macro_score": macro_score.score,
        "quality_level": macro_score.quality_level,
        "coherence": macro_score.cross_cutting_coherence,
        "alignment": macro_score.strategic_alignment,
        "systemic_gaps_count": len(macro_score.systemic_gaps),
        "traceable": True
    }
)

def validate_full_aggregation_pipeline(
    dimension_scores: list["DimensionScore"],
    area_scores: list["AreaScore"],
    cluster_scores: list["ClusterScore"],
    macro_score: "MacroScore",
    input_scored_results: list["ScoredResult"],
) -> tuple[bool, list[ValidationResult]]:
    """

```

```
Validate the entire aggregation pipeline (Phases 4-7).
```

```
Performs comprehensive validation across all phases to ensure:
```

- No empty results at any phase
- Traceability from macro down to micro questions
- Valid score ranges
- Non-zero macro score for valid inputs

```
Args:
```

```
    dimension_scores: Phase 4 output  
    area_scores: Phase 5 output  
    cluster_scores: Phase 6 output  
    macro_score: Phase 7 output  
    input_scored_results: Phase 3 output (input to Phase 4)
```

```
Returns:
```

```
    Tuple of (all_passed, validation_results_list)
```

```
Raises:
```

```
    AggregationValidationError: If any validation fails in strict mode
```

```
"""
```

```
validation_results = []
```

```
# Validate Phase 4
```

```
phase4_result = validate_phase4_output(dimension_scores, input_scored_results)  
validation_results.append(phase4_result)
```

```
# Validate Phase 5
```

```
phase5_result = validate_phase5_output(area_scores, dimension_scores)  
validation_results.append(phase5_result)
```

```
# Validate Phase 6
```

```
phase6_result = validate_phase6_output(cluster_scores, area_scores)  
validation_results.append(phase6_result)
```

```
# Validate Phase 7
```

```
phase7_result = validate_phase7_output(  
    macro_score, cluster_scores, area_scores, dimension_scores  
)  
validation_results.append(phase7_result)
```

```
# Check if all passed
```

```
all_passed = all(result.passed for result in validation_results)
```

```
return all_passed, validation_results
```

```
def enforce_validation_or_fail(  
    validation_results: list[ValidationResult],  
    allow_failure: bool = False,  
) -> None:
```

```
"""
```

```
Enforce validation results or raise exception.
```

Args:

```
validation_results: List of validation results from pipeline
allow_failure: If False, raises exception on any failure
```

Raises:

```
AggregationValidationError: If validation fails and allow_failure=False
```

"""

```
failed_results = [r for r in validation_results if not r.passed]
```

```
if failed_results and not allow_failure:
```

```
    error_messages = []
```

```
    for result in failed_results:
```

```
        error_messages.append(
```

```
            f"\n{result.phase}: {result.error_message}\n"
```

```
            f"\n    Details: {result.details}"
```

```
)
```

```
full_error = (
```

```
    f"Aggregation validation failed at {len(failed_results)} phase(s):\n\n"
```

```
    + "\n\n".join(error_messages)
```

```
)
```

```
raise AggregationValidationError(full_error)
```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/choquet_aggregator.py
```

```
"""
```

### Choquet Aggregator - Non-linear Multi-Layer Calibration Aggregation

This module implements the Choquet integral for aggregating multi-layer calibration scores with interaction terms. The Choquet aggregator captures synergies and complementarities between layers that simple weighted averages cannot represent.

Formula:

$$\text{Cal}(I) = \sum_{k=1}^K a_k x_k + \sum_{1 \leq i < j \leq K} a_{ij} \min(x_i, x_j)$$

Where:

- $x_k$ : Score for layer  $k$  (normalized to  $[0,1]$ )
- $a_k$ : Linear weight for layer  $k$
- $a_{ij}$ : Interaction weight for layer pair  $(i,j)$
- $\text{Cal}(I)$ : Choquet-aggregated calibration score  $\in [0,1]$

Architecture:

- ChoquetConfig: Configuration with linear and interaction weights
- ChoquetAggregator: Main aggregation engine
- CalibrationResult: Output with score breakdown and rationales

Requirements:

- Boundedness:  $0.0 \leq \text{Cal}(I) \leq 1.0$  (enforced via validation)
- Monotonicity: Higher layer scores  $\rightarrow$  higher aggregate score
- Normalization: Weights normalized to ensure boundedness
- Determinism: Fixed random seeds for reproducible results

```
"""
```

```
from __future__ import annotations

import logging
from dataclasses import dataclass, field
from typing import Any

logger = logging.getLogger(__name__)

class CalibrationConfigError(Exception):
    """Raised when calibration configuration validation fails."""
    pass

@dataclass(frozen=True)
class ChoquetConfig:
    """
    Configuration for Choquet aggregation with interaction terms.

    Attributes:
        linear_weights: Dictionary mapping layer_id -> linear weight (a_k)
        interaction_weights: Dictionary mapping (layer_i, layer_j) -> interaction weight (a_ij)
        validate_boundedness: Whether to validate that Cal(I) in [0,1]
    """
    linear_weights: dict[int, float]
    interaction_weights: dict[tuple[int, int], float]
    validate_boundedness: bool
```

```
normalize_weights: Whether to normalize weights automatically
```

Example:

```
>>> config = ChoquetConfig(
...     linear_weights={"@b": 0.4, "@chain": 0.3, "@q": 0.2},
...     interaction_weights={("@b", "@chain"): 0.1},
...     validate_boundedness=True
... )
"""

linear_weights: dict[str, float]
interaction_weights: dict[tuple[str, str], float] = field(default_factory=dict)
validate_boundedness: bool = True
normalize_weights: bool = True

def __post_init__(self) -> None:
    """Validate configuration on construction."""
    if not self.linear_weights:
        raise CalibrationConfigError("linear_weights cannot be empty")

    for layer_id, weight in self.linear_weights.items():
        if not isinstance(layer_id, str):
            raise CalibrationConfigError(f"Layer ID must be string, got {type(layer_id)}")
        if not isinstance(weight, (int, float)):
            raise CalibrationConfigError(f"Weight must be numeric, got {type(weight)}")
        if weight < 0.0:
            raise CalibrationConfigError(f"Negative weight not allowed: {layer_id}={weight}")

    for (layer_i, layer_j), weight in self.interaction_weights.items():
        if not isinstance(layer_i, str) or not isinstance(layer_j, str):
            raise CalibrationConfigError(f"Layer IDs must be strings: ({layer_i}, {layer_j})")
        if not isinstance(weight, (int, float)):
            raise CalibrationConfigError(f"Interaction weight must be numeric: {weight}")
        if weight < 0.0:
            raise CalibrationConfigError(f"Negative interaction weight: ({layer_i}, {layer_j})={weight}")

    if layer_i not in self.linear_weights:
        raise CalibrationConfigError(f"Interaction layer {layer_i} not in linear_weights")
    if layer_j not in self.linear_weights:
        raise CalibrationConfigError(f"Interaction layer {layer_j} not in linear_weights")

@dataclass
class CalibrationBreakdown:
"""

Detailed breakdown of Choquet aggregation computation.
```

Attributes:

```

linear_contribution: Total contribution from linear terms ?(a?·x?)
interaction_contribution: Total contribution from interaction terms
?(a??·min(x?,x?))
per_layer_contributions: Dictionary mapping layer_id -> a?·x?
per_interaction_contributions: Dictionary mapping (layer_i, layer_j) ->
a??·min(x?,x?)
per_layer_rationales: Dictionary mapping layer_id -> human-readable rationale
per_interaction_rationales: Dictionary mapping (layer_i, layer_j) -> rationale
"""

linear_contribution: float
interaction_contribution: float
per_layer_contributions: dict[str, float]
per_interaction_contributions: dict[tuple[str, str], float]
per_layer_rationales: dict[str, str]
per_interaction_rationales: dict[tuple[str, str], str]

"""

@dataclass
class CalibrationResult:
"""

Result of Choquet aggregation with full breakdown and metadata.

Attributes:
    subject: Identifier for the aggregated subject (e.g., method name)
    calibration_score: Final Cal(I) score ? [0,1]
    breakdown: Detailed CalibrationBreakdown
    layer_scores: Input layer scores dictionary
    metadata: Additional metadata (config hash, timestamp, etc.)
    validation_passed: Whether boundedness validation passed
    validation_details: Details of validation checks

"""

subject: str
calibration_score: float
breakdown: CalibrationBreakdown
layer_scores: dict[str, float]
metadata: dict[str, Any] = field(default_factory=dict)
validation_passed: bool = True
validation_details: dict[str, Any] = field(default_factory=dict)

```

```

class ChoquetAggregator:
"""

Choquet integral aggregator with interaction terms.

```

This aggregator computes:  
 $\text{Cal}(I) = \text{linear\_sum} + \text{interaction\_sum}$

Where:  
 $\text{linear\_sum} = ?(a?·x?)$  over all layers  
 $\text{interaction\_sum} = ?(a??·\min(x?,x?))$  over all interaction pairs

The aggregator enforces boundedness ( $0 \leq \text{Cal}(I) \leq 1$ ) and provides detailed breakdowns showing contribution of each layer and interaction.

Example:

```
>>> config = ChoquetConfig(  
...     linear_weights={"@b": 0.4, "@chain": 0.3},  
...     interaction_weights={( "@b", "@chain"): 0.2}  
... )  
>>> aggregator = ChoquetAggregator(config)  
>>> result = aggregator.aggregate(  
...     subject="method_X",  
...     layer_scores={"@b": 0.8, "@chain": 0.7}  
... )  
>>> print(f"Cal(I) = {result.calibration_score:.4f}")  
"""
```

```
def __init__(self, config: ChoquetConfig) -> None:  
    """
```

Initialize Choquet aggregator with configuration.

Args:

config: ChoquetConfig with linear and interaction weights

Raises:

CalibrationConfigError: If configuration is invalid

```
"""
```

self.config = config

self.\_normalized\_linear\_weights = self.\_normalize\_linear\_weights()

self.\_normalized\_interaction\_weights = self.\_normalize\_interaction\_weights()

logger.info(

f"ChoquetAggregator initialized: "

f"{len(self.config.linear\_weights)} layers, "

f"{len(self.config.interaction\_weights)} interactions"

)

```
def _normalize_linear_weights(self) -> dict[str, float]:  
    """
```

Normalize linear weights to ensure boundedness.

Returns:

Normalized linear weights dictionary

```
"""
```

if not self.config.normalize\_weights:

return dict(self.config.linear\_weights)

total = sum(self.config.linear\_weights.values())

if total <= 0:

raise CalibrationConfigError("Total linear weight sum must be positive")

normalized = {

layer: weight / total

for layer, weight in self.config.linear\_weights.items()

}

logger.debug(f"Linear weights normalized: total={total:.4f}")

return normalized

```

def _normalize_interaction_weights(self) -> dict[tuple[str, str], float]:
    """
    Normalize interaction weights relative to linear weights.

    Interaction weights are constrained to ensure boundedness:
    ?(a??) ? min?(a?), 1.0
    """

    Returns:
        Normalized interaction weights dictionary
    """
    if not self.config.interaction_weights:
        return {}

    if not self.config.normalize_weights:
        return dict(self.config.interaction_weights)

    total_interaction = sum(self.config.interaction_weights.values())
    max_allowed = min(sum(self._normalized_linear_weights.values()), 1.0) * 0.5

    if total_interaction > max_allowed:
        scale_factor = max_allowed / total_interaction
        normalized = {
            pair: weight * scale_factor
            for pair, weight in self.config.interaction_weights.items()
        }
        logger.warning(
            f"Interaction weights scaled by {scale_factor:.4f} to ensure
boundedness"
        )
    else:
        normalized = dict(self.config.interaction_weights)

    return normalized

def _compute_linear_sum(
    self,
    layer_scores: dict[str, float]
) -> tuple[float, dict[str, float], dict[str, str]]:
    """
    Compute linear contribution: ?(a?·x?)

    Args:
        layer_scores: Dictionary mapping layer_id -> score ? [0,1]

    Returns:
        Tuple of (total_sum, per_layer_contributions, per_layer_rationales)
    """
    linear_sum = 0.0
    per_layer_contributions: dict[str, float] = {}
    per_layer_rationales: dict[str, str] = {}

    for layer_id, weight in self._normalized_linear_weights.items():
        score = layer_scores.get(layer_id, 0.0)

```

```

        if score < 0.0 or score > 1.0:
            logger.warning(f"Layer {layer_id} score {score} outside [0,1],
clamping")
            score = max(0.0, min(1.0, score))

        contribution = weight * score
        linear_sum += contribution
        per_layer_contributions[layer_id] = contribution

        per_layer_rationales[layer_id] = (
            f"Layer {layer_id}: weight={weight:.4f} x score={score:.4f} "
            f"= {contribution:.4f}"
        )

    logger.debug(f"Linear sum computed: {linear_sum:.4f}")
    return linear_sum, per_layer_contributions, per_layer_rationales

def _compute_interaction_sum(
    self,
    layer_scores: dict[str, float]
) -> tuple[float, dict[tuple[str, str], float], dict[tuple[str, str], str]]:
    """
    Compute interaction contribution: ?(a??·min(x?,x?))
    """

    Args:
        layer_scores: Dictionary mapping layer_id -> score ? [0,1]

    Returns:
        Tuple of (total_sum, per_interaction_contributions,
per_interaction_rationales)
    """

    interaction_sum = 0.0
    per_interaction_contributions: dict[tuple[str, str], float] = {}
    per_interaction_rationales: dict[tuple[str, str], str] = {}

    for (layer_i, layer_j), weight in self._normalized_interaction_weights.items():
        score_i = layer_scores.get(layer_i, 0.0)
        score_j = layer_scores.get(layer_j, 0.0)

        if score_i < 0.0 or score_i > 1.0:
            logger.warning(f"Layer {layer_i} score {score_i} outside [0,1],
clamping")
            score_i = max(0.0, min(1.0, score_i))

        if score_j < 0.0 or score_j > 1.0:
            logger.warning(f"Layer {layer_j} score {score_j} outside [0,1],
clamping")
            score_j = max(0.0, min(1.0, score_j))

        min_score = min(score_i, score_j)
        contribution = weight * min_score
        interaction_sum += contribution
        per_interaction_contributions[(layer_i, layer_j)] = contribution

```

```

per_interaction_rationales[(layer_i, layer_j)] = (
    f"Interaction ({layer_i}, {layer_j}): "
    f"weight={weight:.4f} × min({score_i:.4f}, {score_j:.4f}) "
    f"= {contribution:.4f}"
)

logger.debug(f"Interaction sum computed: {interaction_sum:.4f}")
return interaction_sum, per_interaction_contributions,
per_interaction_rationales

def _validate_boundedness(self, calibration_score: float) -> tuple[bool, dict[str, Any]]:
    """
    Validate that calibration score is bounded in [0,1].
    """

    Args:
        calibration_score: Computed Cal(I) score

    Returns:
        Tuple of (is_valid, validation_details)

    Raises:
        CalibrationConfigError: If boundedness is violated and validation enabled
    """
    validation_details = {
        "score": calibration_score,
        "lower_bound": 0.0,
        "upper_bound": 1.0,
        "bounded": True,
        "message": "Boundedness validated"
    }

    if calibration_score < 0.0 or calibration_score > 1.0:
        validation_details["bounded"] = False
        validation_details["message"] = (
            f"Boundedness violation: Cal(I)={calibration_score:.6f} not in [0,1]"
        )

    if self.config.validate_boundedness:
        logger.error(validation_details["message"])
        raise CalibrationConfigError(validation_details["message"])
    else:
        logger.warning(validation_details["message"])
        return False, validation_details

    logger.debug(f"Boundedness validated: {calibration_score:.4f} ? [0,1]")
    return True, validation_details

def aggregate(
    self,
    subject: str,
    layer_scores: dict[str, float],
    metadata: dict[str, Any] | None = None
)

```

```

) -> CalibrationResult:
    """
    Aggregate layer scores using Choquet integral with interaction terms.

    Args:
        subject: Identifier for aggregated subject (e.g., method name)
        layer_scores: Dictionary mapping layer_id -> score ? [0,1]
        metadata: Optional metadata to include in result

    Returns:
        CalibrationResult with score, breakdown, and validation details

    Raises:
        CalibrationConfigError: If boundedness validation fails
        ValueError: If required layers are missing from layer_scores

    Example:
        >>> result = aggregator.aggregate(
        ...     subject="BayesianAnalyzer",
        ...     layer_scores={"@b": 0.85, "@chain": 0.75, "@q": 0.90}
        ... )
        >>> print(f"Cal(I) = {result.calibration_score:.4f}")
        >>> print(f"Linear: {result.breakdown.linear_contribution:.4f}")
        >>> print(f"Interaction: {result.breakdown.interaction_contribution:.4f}")
    """
    logger.info(f"Aggregating calibration for subject: {subject}")

    missing_layers = set(self.config.linear_weights.keys()) -
    set(layer_scores.keys())
    if missing_layers:
        raise ValueError(
            f"Missing required layers in layer_scores: {missing_layers}. "
            f"Expected: {set(self.config.linear_weights.keys())}"
        )

    linear_sum, per_layer_contrib, per_layer_rationale =
    self._compute_linear_sum(layer_scores)

    interaction_sum, per_interaction_contrib, per_interaction_rationale = (
        self._compute_interaction_sum(layer_scores)
    )

    calibration_score = linear_sum + interaction_sum

    validation_passed, validation_details =
    self._validate_boundedness(calibration_score)

    calibration_score = max(0.0, min(1.0, calibration_score))

    breakdown = CalibrationBreakdown(
        linear_contribution=linear_sum,
        interaction_contribution=interaction_sum,
        per_layer_contributions=per_layer_contrib,
        per_interaction_contributions=per_interaction_contrib,

```

```
        per_layer_rationales=per_layer_rationale,
        per_interaction_rationales=per_interaction_rationale
    )

result_metadata = metadata or {}
result_metadata.update({
    "n_layers": len(layer_scores),
    "n_interactions": len(self._normalized_interaction_weights),
    "normalized_weights": self.config.normalize_weights,
})

result = CalibrationResult(
    subject=subject,
    calibration_score=calibration_score,
    breakdown=breakdown,
    layer_scores=layer_scores,
    metadata=result_metadata,
    validation_passed=validation_passed,
    validation_details=validation_details
)

logger.info(
    f"? Aggregation complete: subject={subject}, "
    f"Cal(I)={calibration_score:.4f}, "
    f"linear={linear_sum:.4f}, interaction={interaction_sum:.4f}"
)

return result
```

```
src/farfan_pipeline/phases/Phase_four_five_six_seven/choquet_examples/run_examples.py
```

```
"""
```

### Executable Examples for Choquet Aggregator

This script runs all examples from the markdown documentation with actual code, allowing verification of calculations and interactive exploration.

Usage:

```
python run_examples.py                      # Run all examples
python run_examples.py --example 1           # Run specific example
python run_examples.py --example 2 --verbose  # Verbose output
```

```
"""
```

```
import sys
from pathlib import Path

sys.path.insert(0, str(Path(__file__).parent.parent.parent.parent))

from canonic_phases.Phase_four_five_six_seven.choquet_aggregator import (
    ChoquetAggregator,
    ChoquetConfig,
)

def print_separator(title: str = "") -> None:
    """Print a visual separator."""
    width = 80
    if title:
        padding = (width - len(title) - 2) // 2
        print(f"\n{'=' * padding} {title} {'=' * padding}")
    else:
        print(f"\n{'=' * width}")

def print_result(result, verbose: bool = False) -> None:
    """Print aggregation result in formatted way."""
    print(f"\n{ 'Subject':<30} {result.subject}")
    print(f"{ 'Calibration Score':<30} {result.calibration_score:.4f}")
    print(f"{ 'Linear Contribution':<30} {result.breakdown.linear_contribution:.4f}")
    print(f"{'Interaction' :<30} {result.breakdown.interaction_contribution:.4f} { 'Contribution':<30} {result.breakdown.interaction_contribution:.4f}")
    print(f"{'Validation Passed':<30} {result.validation_passed}")

    if verbose:
        print("\nPer-Layer Contributions:")
        for layer, contrib in result.breakdown.per_layer_contributions.items():
            print(f"  {layer:<10} {contrib:.4f}")

        if result.breakdown.per_interaction_contributions:
            print("\nPer-Interaction Contributions:")
            for pair, contrib in result.breakdown.per_interaction_contributions.items():
                print(f"  {pair[0]:<10} x {pair[1]:<10} {contrib:.4f}")
```

```

    if verbose and result.breakdown.per_layer_rationales:
        print("\nRationales:")
        for layer, rationale in result.breakdown.per_layer_rationales.items():
            print(f"  {rationale}")

def example_1_basic_calculation(verbose: bool = False) -> None:
    """Example 1: Basic calculation without interactions."""
    print_separator("EXAMPLE 1: Basic Calculation")

    print("\nConfiguration: 3 layers, no interactions")
    config = ChoquetConfig(
        linear_weights={
            "@b": 0.4,
            "@chain": 0.3,
            "@q": 0.3
        },
        interaction_weights={},
        normalize_weights=False,
        validate_boundedness=True
    )

    print("\nLayer Scores:")
    layer_scores = {
        "@b": 0.85,
        "@chain": 0.75,
        "@q": 0.90
    }
    for layer, score in layer_scores.items():
        print(f"  {layer:<10} {score:.2f}")

    aggregator = ChoquetAggregator(config)
    result = aggregator.aggregate(
        subject="ExampleMethod",
        layer_scores=layer_scores
    )

    print_result(result, verbose)

    print("\n? Linear-only aggregation produces weighted average")
    print("  Expected: 0.4×0.85 + 0.3×0.75 + 0.3×0.90 = 0.835")
    print(f"  Actual:   {result.calibration_score:.4f}")

def example_2_with_interactions(verbose: bool = False) -> None:
    """Example 2: Full Choquet with interaction terms."""
    print_separator("EXAMPLE 2: With Interactions")

    print("\nConfiguration: 3 layers, 2 interaction pairs")
    config = ChoquetConfig(
        linear_weights={
            "@b": 0.35,
            "@chain": 0.30,
            "@q": 0.25
        },
        interaction_weights={
            ("@b", "@chain"): 0.1,
            ("@b", "@q"): 0.05
        }
    )

```

```

        },
        interaction_weights={
            ("@b", "@chain"): 0.10,
            ("@chain", "@q"): 0.05
        },
        normalize_weights=False,
        validate_boundedness=True
    )

print("\nLayer Scores:")
layer_scores = {
    "@b": 0.80,
    "@chain": 0.70,
    "@q": 0.85
}
for layer, score in layer_scores.items():
    print(f" {layer:<10} {score:.2f}")

aggregator = ChoquetAggregator(config)
result = aggregator.aggregate(
    subject="ExampleMethodWithSynergies",
    layer_scores=layer_scores
)

print_result(result, verbose)

linear_only = 0.35 * 0.80 + 0.30 * 0.70 + 0.25 * 0.85
synergy_bonus = result.calibration_score - linear_only

print("\n? Synergy Analysis:")
print(f" Linear-only score: {linear_only:.4f}")
print(f" With interactions: {result.calibration_score:.4f}")
    print(f" Synergy bonus: {synergy_bonus:.4f}")
({synergy_bonus/linear_only*100:.1f}%)")
        print(f"\n           Interaction (@b, @chain):"
{result.breakdown.per_interaction_contributions[('@b', '@chain')]:.4f}"))
        print(f"           Interaction (@chain, @q):"
{result.breakdown.per_interaction_contributions[('@chain', '@q')]:.4f}"))

def example_3_normalization(verbose: bool = False) -> None:
    """Example 3: Automatic weight normalization."""
    print_separator("EXAMPLE 3: Normalization")

    print("\nConfiguration: Unnormalized weights (raw values)")
    print(" Raw linear weights: @b=4.0, @chain=3.0, @q=2.0, @d=1.0")
    print(" Raw interaction weights: (@b,@chain)=1.5, (@q,@d)=0.5")

    config = ChoquetConfig(
        linear_weights={
            "@b": 4.0,
            "@chain": 3.0,
            "@q": 2.0,
            "@d": 1.0
        }
    )

```

```

        },
        interaction_weights={
            ("@b", "@chain"): 1.5,
            ("@q", "@d"): 0.5
        },
        normalize_weights=True,
        validate_boundedness=True
    )

print("\nLayer Scores:")
layer_scores = {
    "@b": 0.90,
    "@chain": 0.85,
    "@q": 0.75,
    "@d": 0.80
}
for layer, score in layer_scores.items():
    print(f" {layer:<10} {score:.2f}")

aggregator = ChoquetAggregator(config)

print("\n? Normalized Weights:")
print(" Linear weights (sum = 1.0):")
for layer, weight in aggregator._normalized_linear_weights.items():
    print(f" {layer:<10} {weight:.4f}")

print("\n Interaction weights (scaled for boundedness):")
for pair, weight in aggregator._normalized_interaction_weights.items():
    print(f" {pair[0]:<10} x {pair[1]:<10} {weight:.4f}")

result = aggregator.aggregate(
    subject="HighPerformingMethod",
    layer_scores=layer_scores
)

print_result(result, verbose)

```

```

def example_4_boundary_cases(verbose: bool = False) -> None:
    """Example 4: Boundary cases and edge conditions."""
    print_separator("EXAMPLE 4: Boundary Cases")

    # Case 1: All zeros
    print("\n--- Case 1: All Scores = 0 (Worst Case) ---")
    config = ChoquetConfig(
        linear_weights={"@b": 0.5, "@chain": 0.3, "@q": 0.2},
        interaction_weights={("@b", "@chain"): 0.1},
        normalize_weights=False
    )
    aggregator = ChoquetAggregator(config)
    result = aggregator.aggregate(
        "WorstCase",
        {"@b": 0.0, "@chain": 0.0, "@q": 0.0}
    )

```

```

print(f"  Cal(I) = {result.calibration_score:.4f}")
print("  ? Perfect lower bound: 0.0")

# Case 2: All ones
print("\n--- Case 2: All Scores = 1 (Perfect Case) ---")
config = ChoquetConfig(
    linear_weights={"@b": 0.4, "@chain": 0.3, "@q": 0.3},
    interaction_weights={},
    normalize_weights=False
)
aggregator = ChoquetAggregator(config)
result = aggregator.aggregate(
    "PerfectCase",
    {"@b": 1.0, "@chain": 1.0, "@q": 1.0}
)
print(f"  Cal(I) = {result.calibration_score:.4f}")
print("  ? Perfect upper bound: 1.0")

# Case 3: Single layer dominates
print("\n--- Case 3: Single Layer Dominates ---")
config = ChoquetConfig(
    linear_weights={"@b": 0.95, "@chain": 0.03, "@q": 0.02},
    interaction_weights={},
    normalize_weights=False
)
aggregator = ChoquetAggregator(config)
result = aggregator.aggregate(
    "DominantLayer",
    {"@b": 0.50, "@chain": 1.00, "@q": 1.00}
)
print(f"  Cal(I) = {result.calibration_score:.4f}")
print("  ? Dominated by @b (95% weight, 0.50 score)")
print("  Despite @chain and @q being perfect (1.0)")

# Case 4: Asymmetric interaction
print("\n--- Case 4: Extreme Asymmetry in Interactions ---")
config = ChoquetConfig(
    linear_weights={"@b": 0.5, "@chain": 0.5},
    interaction_weights={("@b", "@chain"): 0.3},
    normalize_weights=False,
    validate_boundedness=False
)
aggregator = ChoquetAggregator(config)
result = aggregator.aggregate(
    "AsymmetricInteraction",
    {"@b": 1.0, "@chain": 0.0}
)
print(f"  Cal(I) = {result.calibration_score:.4f}")
print(f"  Linear: {result.breakdown.linear_contribution:.4f}")
print(f"  Interaction: {result.breakdown.interaction_contribution:.4f}")
print("  ? Interaction vanishes: min(1.0, 0.0) = 0.0")

# Case 5: Empty interactions (weighted average)
print("\n--- Case 5: Empty Interactions (Degenerates to WA) ---")

```

```

config = ChoquetConfig(
    linear_weights={"@b": 0.3, "@chain": 0.3, "@q": 0.4},
    interaction_weights={},
    normalize_weights=False
)
aggregator = ChoquetAggregator(config)
result = aggregator.aggregate(
    "WeightedAverage",
    {"@b": 0.7, "@chain": 0.6, "@q": 0.8}
)
wa_expected = 0.3 * 0.7 + 0.3 * 0.6 + 0.4 * 0.8
print(f"Cal(I) = {result.calibration_score:.4f}")
print(f"Expected WA = {wa_expected:.4f}")
print(" ? Choquet = WA when no interactions")

def example_5_real_world_scenario(verbose: bool = False) -> None:
    """Example 5: Realistic calibration scenario."""
    print_separator("EXAMPLE 5: Real-World Scenario")

    print("\nScenario: Calibrating a Bayesian analysis method")
    print(" Layers: @b (base), @chain (dependencies), @q (question), @d (data)")
    print(" Interactions: Base-chain synergy, Question-data synergy")

    config = ChoquetConfig(
        linear_weights={
            "@b": 0.35,          # Base quality (implementation)
            "@chain": 0.25,       # Chain quality (dependencies)
            "@q": 0.25,          # Question relevance
            "@d": 0.15           # Data quality
        },
        interaction_weights={
            ("@b", "@chain"): 0.08,   # Good code + good deps = multiplicative
            ("@q", "@d"): 0.05       # Relevant questions + good data = synergy
        },
        normalize_weights=True,
        validate_boundedness=True
    )

    print("\nLayer Scores (from actual calibration):")
    layer_scores = {
        "@b": 0.82,          # Strong implementation (test coverage, typing, docs)
        "@chain": 0.68,        # Moderate dependencies (some technical debt)
        "@q": 0.91,          # Excellent question alignment
        "@d": 0.74           # Good data quality
    }
    for layer, score in layer_scores.items():
        print(f" {layer:<10} {score:.2f}")

    aggregator = ChoquetAggregator(config)
    result = aggregator.aggregate(
        subject="BayesianCounterfactualAuditor",
        layer_scores=layer_scores,
        metadata={}
    )

```

```

        "cohort": "COHORT_2024",
        "calibration_date": "2024-12-15",
        "method_role": "SCORE_Q"
    }
)

print_result(result, verbose)

print("\n? Interpretation:")
print(f"  Final Calibration: {result.calibration_score:.4f}")

if result.calibration_score >= 0.80:
    quality_tier = "EXCELLENT"
elif result.calibration_score >= 0.70:
    quality_tier = "GOOD"
elif result.calibration_score >= 0.60:
    quality_tier = "ACCEPTABLE"
else:
    quality_tier = "NEEDS_IMPROVEMENT"

print(f"  Quality Tier: {quality_tier}")
print("\n  Strengths:")
print("      - Excellent question alignment (0.91)")
print("      - Strong base implementation (0.82)")
print("\n  Opportunities:")
print("      - Improve chain dependencies (0.68) - bottleneck for synergy")
print("      - Enhance data quality (0.74)")

synergy_total = result.breakdown.interaction_contribution
print(f"\n  Synergy Contribution: {synergy_total:.4f}")
print("  This represents the value captured by interaction terms")
print("  that a simple weighted average would miss.")

def main() -> None:
    """Run examples based on command-line arguments."""
    import argparse

    parser = argparse.ArgumentParser(
        description="Run Choquet Aggregator examples"
    )
    parser.add_argument(
        "--example",
        type=int,
        choices=[1, 2, 3, 4, 5],
        help="Run specific example (1-5)"
    )
    parser.add_argument(
        "--verbose",
        action="store_true",
        help="Enable verbose output with detailed breakdowns"
    )

    args = parser.parse_args()

```

```
examples = {
    1: example_1_basic_calculation,
    2: example_2_with_interactions,
    3: example_3_normalization,
    4: example_4_boundary_cases,
    5: example_5_real_world_scenario,
}

if args.example:
    examples[args.example](args.verbose)
else:
    print("Running all examples...")
    for i, func in examples.items():
        func(args.verbose)
        print()

print_separator()
print("\n? All examples completed successfully!")
print("\nFor detailed documentation, see:")
print(" - example_1_basic_calculation.md")
print(" - example_2_with_interactions.md")
print(" - example_3_normalization.md")
print(" - example_4_boundary_cases.md")
print(" - README.md")

if __name__ == "__main__":
    main()
```